

COMPSC 760

Advanced Neural Networks

Deep Learning - Lecture 3

Learning outcomes

Lecture 1: Deep Neural Networks Review

- ▶ Review what a deep neural network is and the differences classical ML approaches.
- ▶ Review the different steps involved in training a deep NN.
- ▶ Review network initialisation and the different activation functions.
- ▶ Review what the hyperparameters of a DNN are.
- ▶ Review the different strategies to improve the performance of a deep NN.
- ▶ Review the different strategies to tune a deep NN.

Lectures 2 & 3: Learning with sequences (RNNs, Transformers, LLMs)

- ▶ Understand how recurrent neural networks work.
- ▶ Recognise commonly used neural network architectures based on RNN (LSTM, GRU).
- ▶ Understand how transformers work.
- ▶ Understand the principles of Large Language models.

Transformers and LLMs

- ▶ Understand how transformers work.
- ▶ Understand the principles of Large Language Models (LLMs).
- ▶ Be aware of models/tools using LLMs.

If you want to go further:

- ▶ Deep Learning, Part II Deep Networks, chap. 10.

<https://www.deeplearningbook.org/>,

- ▶ Stanford Deep Learning courses:

<https://stanford.edu/~shervine/teaching/cs-230/>

<https://cs230.stanford.edu/>

<http://cs231n.stanford.edu/>

Transformer - Motivation

- ▶ Traditional RNNs suffer some drawbacks
 1. Sequential processing leading to low training times (especially for long sequences) + hard to parallelise.
 2. Difficulty to model long term dependencies as information from earlier steps becomes increasingly diluted.
- ▶ The Transformer architecture addresses these issues by using the self-attention mechanism.

Self-attention mechanism

Self-attention allows the model to weigh the importance of different parts of the input sequence when making predictions.

E.g., with language models, the model “focuses” on different words of the sequence depending on their relevance to the task at hand:

Predicting the next word in: “The rabbits are eating”

What are the important words to focus on for the prediction?

Individual words are represented as vector of numerical values in a lower dimension space. Such representations are called word embeddings.

Word embeddings aim at capturing the meaning of words and their relationship to other words (semantic and syntax).

- ▶ Bag of words (BOW) are one of the simplest word embedding, but they can be intensive to compute, and they fail to capture the relationship between words (i.e., do not consider the order).
- ▶ Modern word embeddings are learned through ML and take in considering the local or global context of the words. A few popular embeddings: Word2Vec, GloVe, FastText and ELMO.

Self-attention – basic version

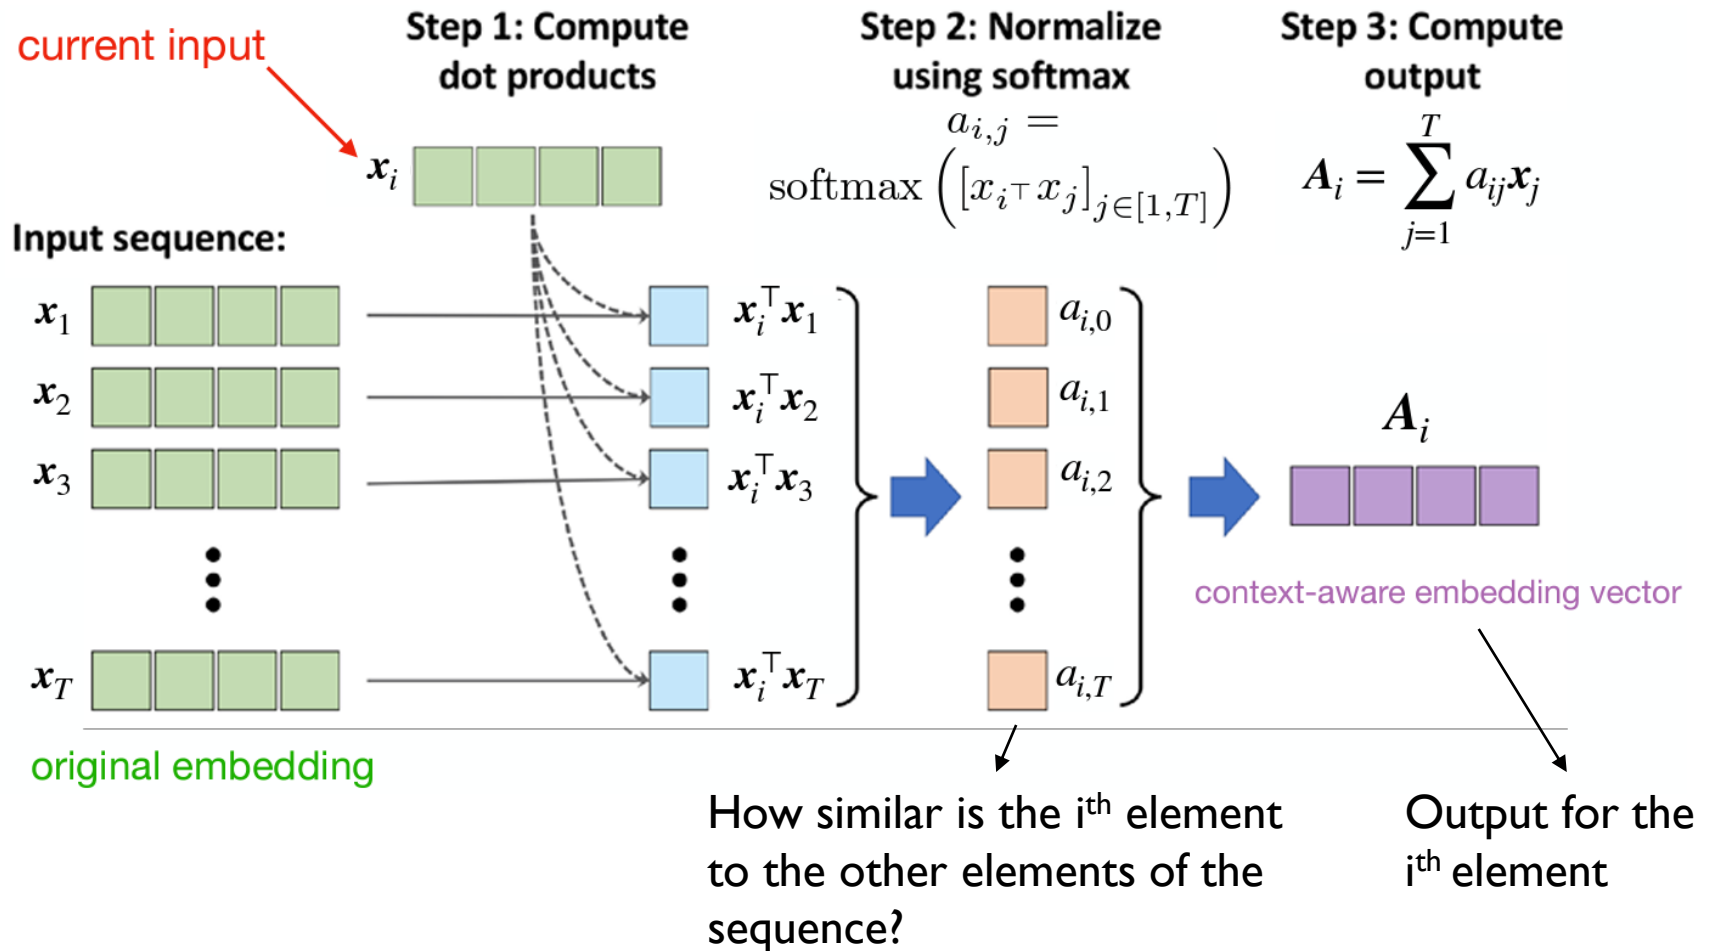
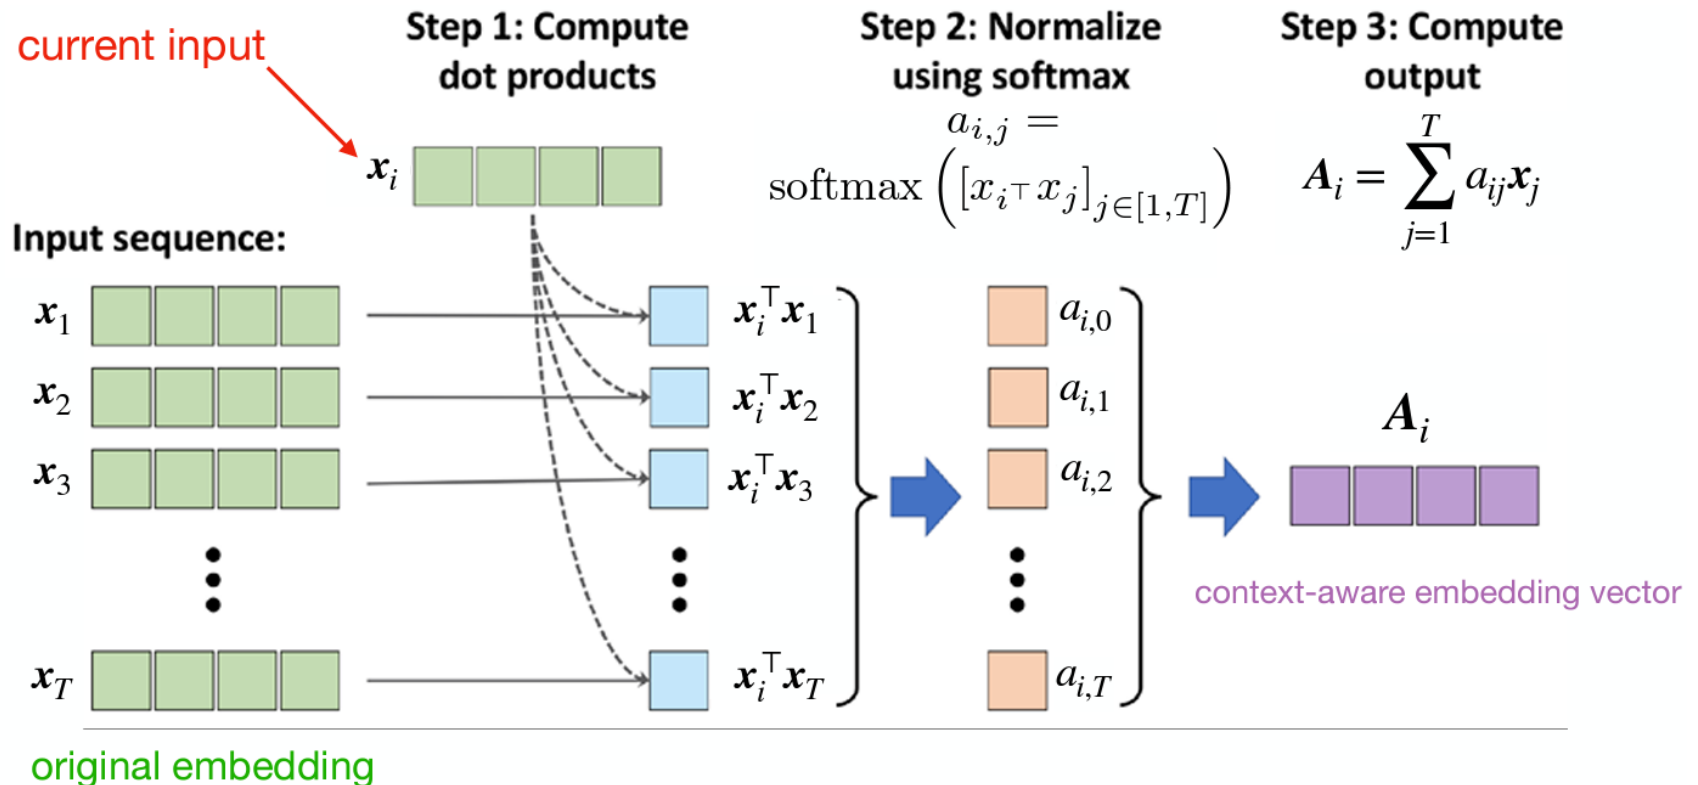


Image source: Raschka & Mirjalili 2019. Python Machine Learning, 3rd edition

Self-attention – basic version



Limitation of this basic version: no learnable parameters!

Image source: Raschka & Mirjalili 2019. Python Machine Learning, 3rd edition


Scaled Dot-Product attention

The Transformer architecture uses an implementation of self-attention called « Scaled Dot-Product attention ».

1. Transform the input matrix $X = [x_1, x_2, \dots, x_n]$ into 3 matrices:
 - ▶ Query: $Q = W^Q X$
 - ▶ Key: $K = W^K X$
 - ▶ Value: $V = W^V X$

W^Q, W^K, W^V are learnable weight matrices that transform the input matrix into query, key and value.

Scaled Dot-Product Attention



[Attention Is All You Need, Vaswani et al., 2017]

Scaled Dot-Product attention


2. Calculate the attention « score »:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

d_k is the size of the keys and queries
($d_k = d_q = d_v$ in the original paper).

Scaling prevents the dot products to grow large, thus avoiding vanishing gradients (application of softmax to large values would yield small gradients).

Scaled Dot-Product Attention




[Attention Is All You Need, Vaswani et al., 2017]

Multi-Head attention

The Transformer architecture actually stacks several scaled dot-product attention layers in parallel.

- ▶ h parallel layers, also called “heads”.
- ▶ Values, keys and queries are projected linearly h times with different learned linear projections.
- ▶ Each projection is the input of an head, which has its own different W^Q , W^K , W^V matrices.



[Attention Is All You Need, Vaswani et al., 2017]


Multi-Head attention

The Transformer architecture actually stacks several scaled dot-product attention layers in parallel.

- ▶ Results are concatenated and projected again to obtain the output:

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W^O$$

where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$




[Attention Is All You Need, Vaswani et al., 2017]

Multi-Head attention

Multi-head attention is beneficial because it allows the Transformer to focus on different aspects of the input sequence.

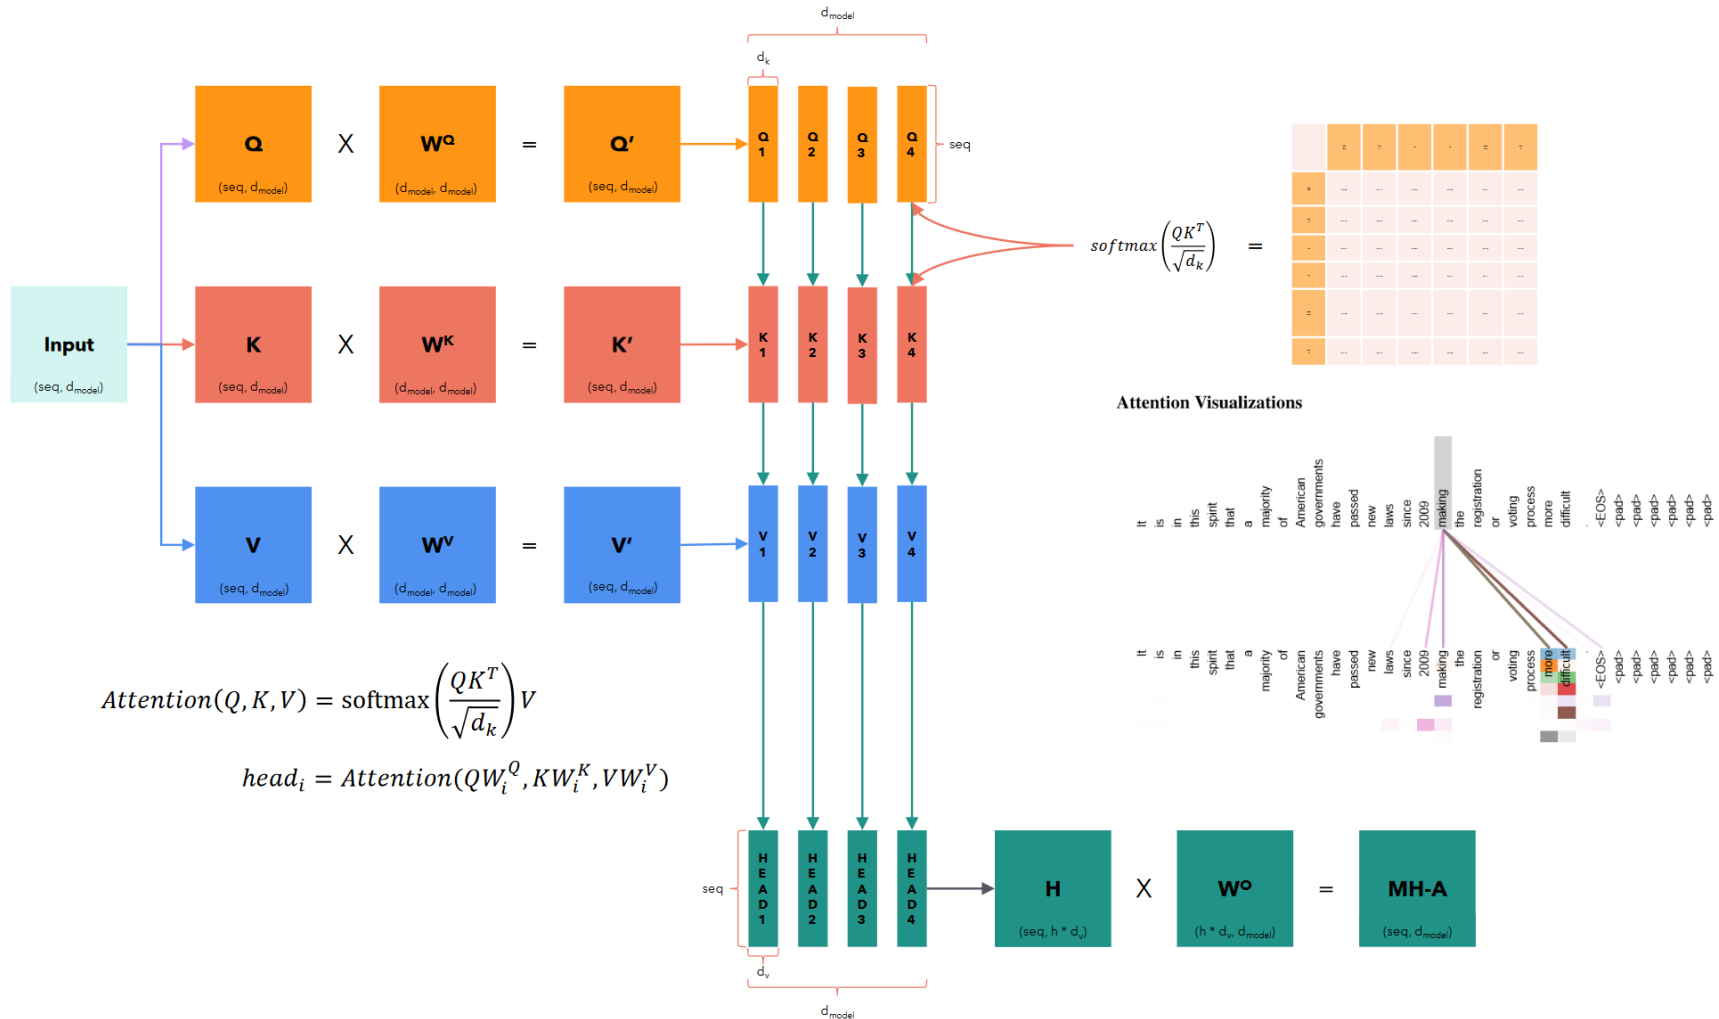
Using a single attention head would have a limiting effect.

Intuitively, the multi-head attention allows the Transformer to spread its attention on different parts of the sequence, instead of averaging it over the full sequence.



[Attention Is All You Need, Vaswani et al., 2017]

Multi-Head attention



Transformers – Attention Visualisation


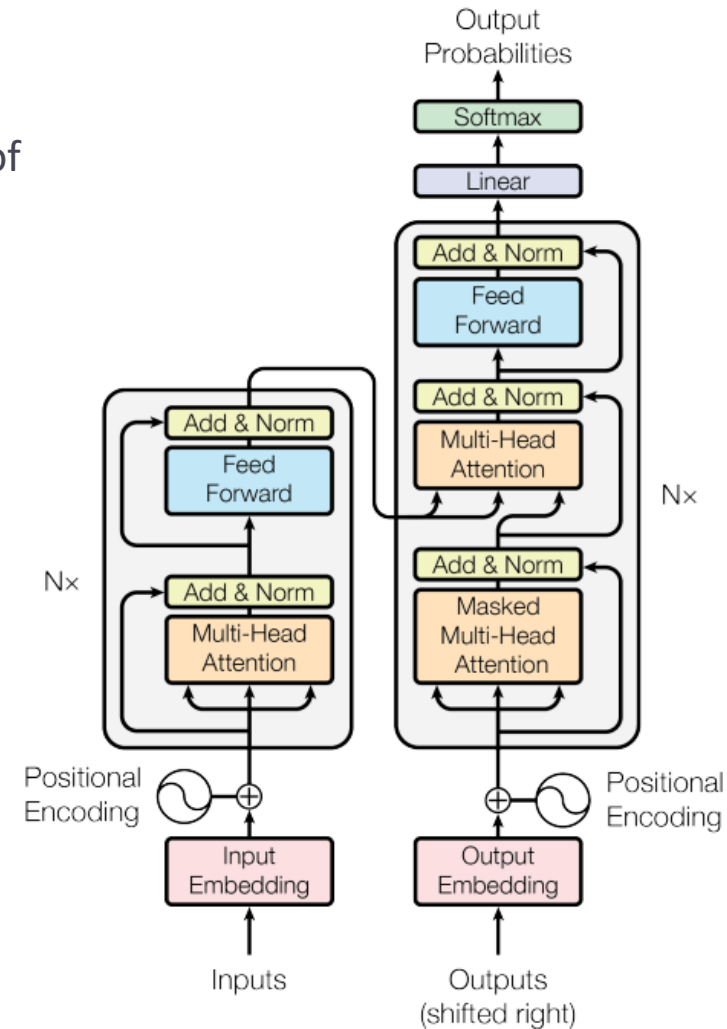


Figure 5: Many of the attention heads exhibit behaviour that seems related to the structure of the sentence. We give two such examples above, from two different heads from the encoder self-attention at layer 5 of 6. The heads clearly learned to perform different tasks.

Transformer – Full architecture

- ▶ Seq2Seq architecture (encoder/decoder)
 - ▶ Encoder: takes an input sequence and produces a set of hidden representations, also known as context vector.
 - ▶ Decoder: takes the context vector and generates the output sequence.
- ▶ Not a RNN!
 - ✚ No sequential processing, it uses embedded representations to **encode positions in the sequence**.
- ▶ Uses the **attention mechanism** to retain information about which parts of the sequence are important.

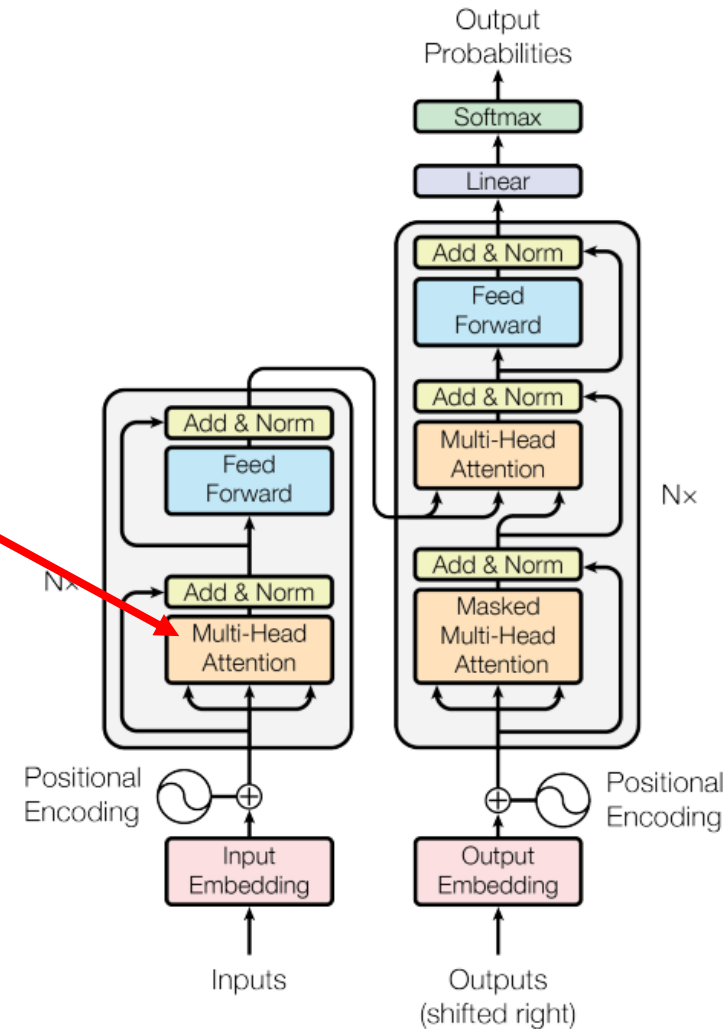


[Attention Is All You Need, Vaswani 2017]

Transformer – Full architecture

Encoder block components:

1. Multi-head attention




[Attention Is All You Need, Vaswani 2017]

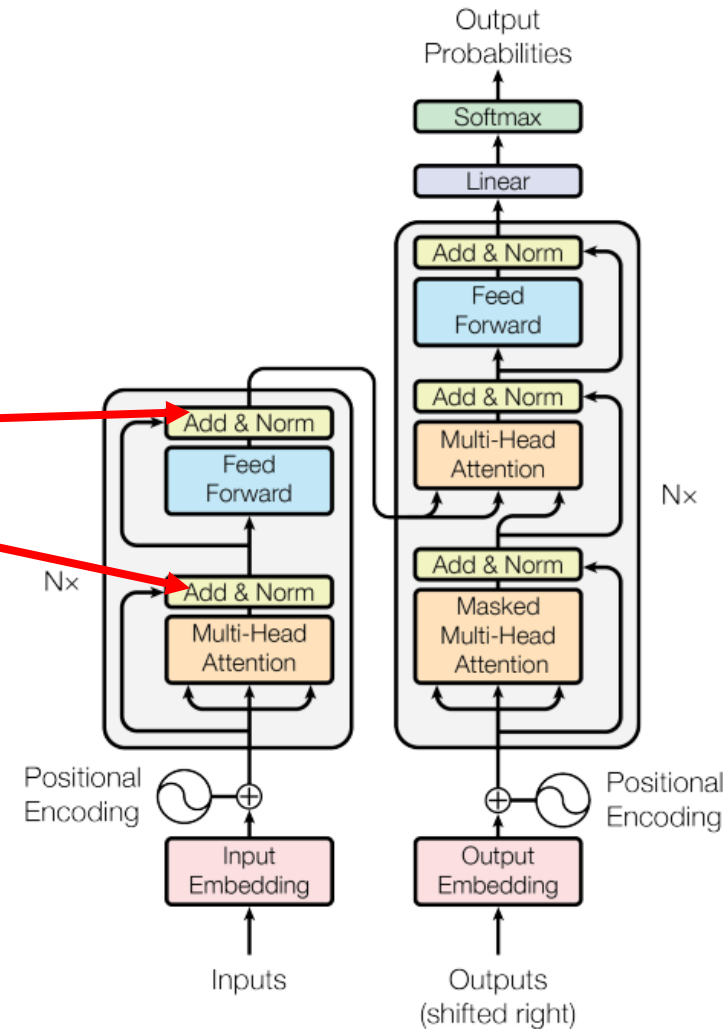
Transformer – Full architecture

Encoder block components:

1. Multi-head attention
2. Layer normalisation



[Group normalization. ECCV, Wu, Y., & He, K., 2018]



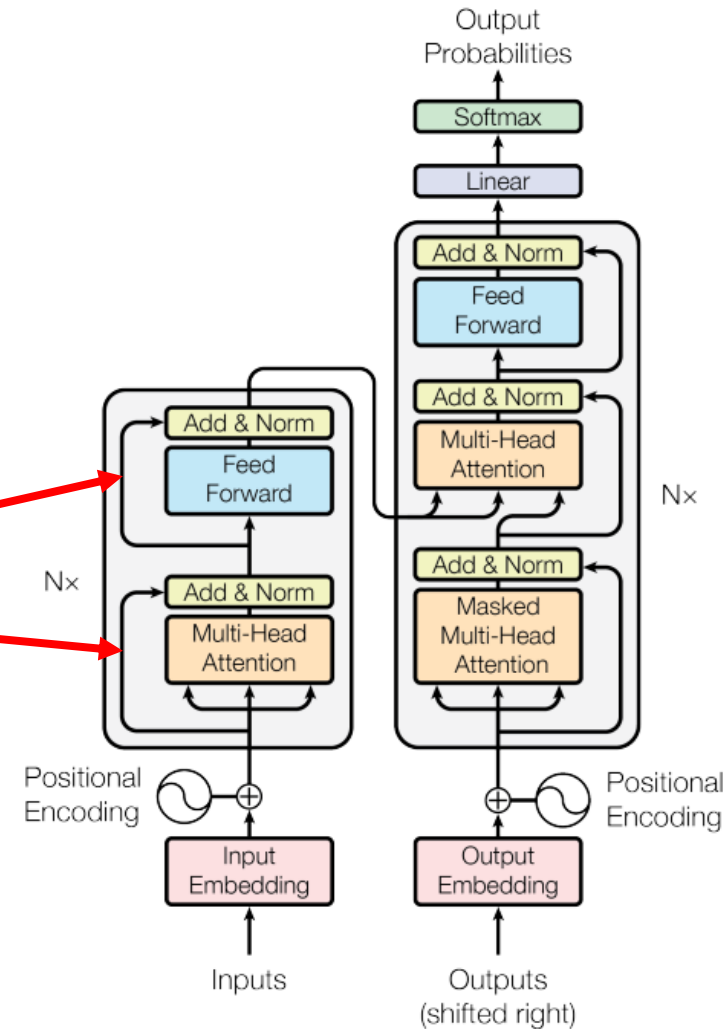
[Attention Is All You Need, Vaswani 2017]

Transformer – Full architecture

Encoder block components:

1. Multi-head attention
2. Layer normalisation
3. Residual/skip connection

$$\text{out}_{\text{norm}} = \text{LayerNorm}(x + \text{Sublayer}(x))$$

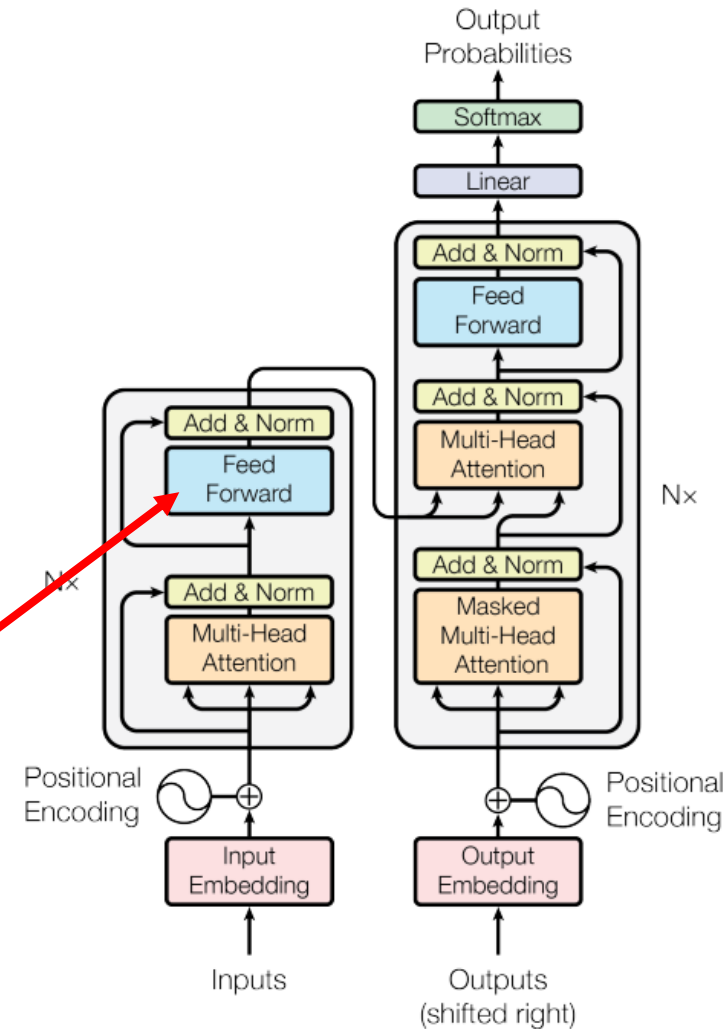


[Attention Is All You Need, Vaswani 2017]

Transformer – Full architecture

Encoder block components:

1. Multi-head attention
2. Layer normalisation
3. Residual/skip connection
4. Fully-connected feed forward NN (multilayer perceptron)
 - ▶ Transforms each attention vector into a form adapted for the next block → can be easily parallelised (treat all words at the same time).

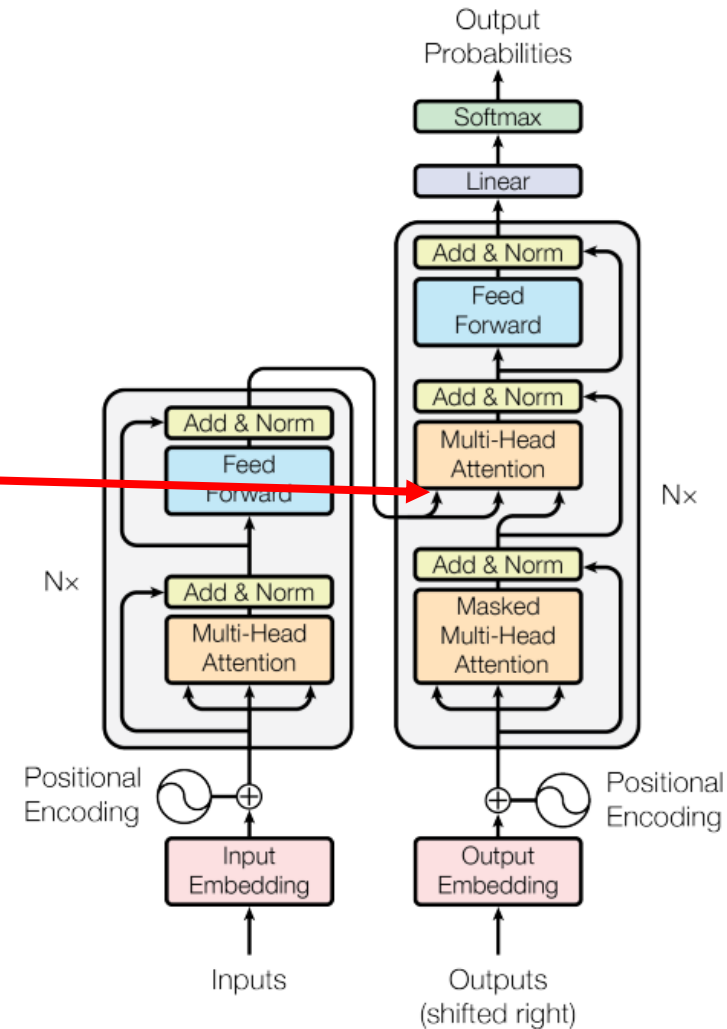


[Attention Is All You Need, Vaswani 2017]

Transformer – Full architecture

Decoder block components:

- ▶ Similar components as the encoder.
- ▶ Encoder output passed to multi-head attention.

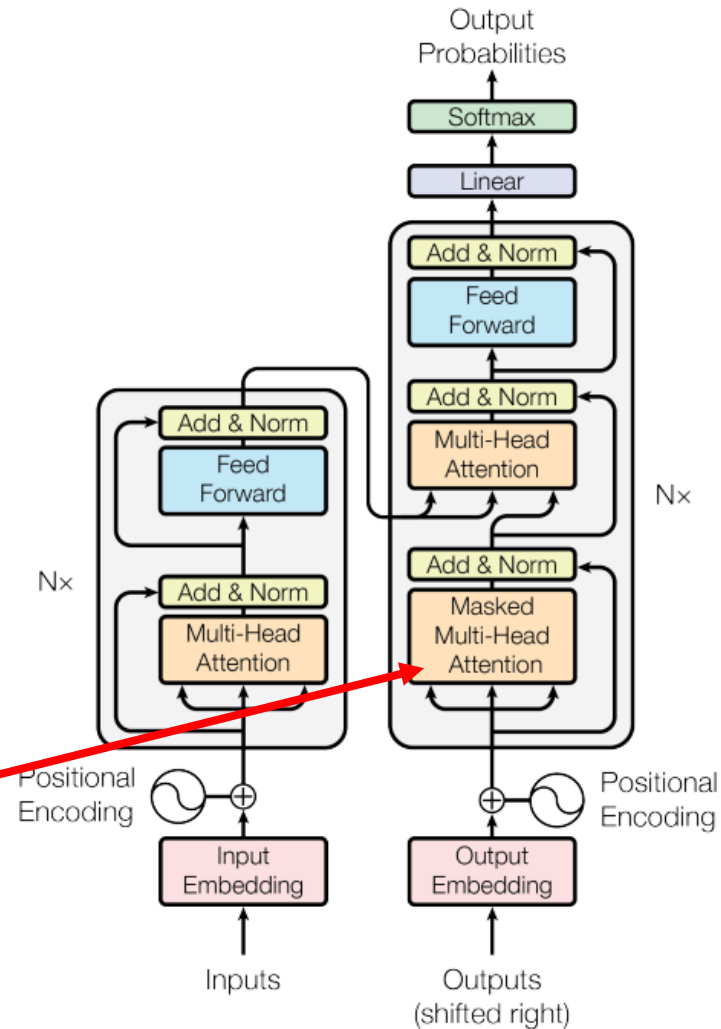


[Attention Is All You Need, Vaswani 2017]

Transformer – Full architecture

Decoder block components:

- ▶ Similar components as the encoder.
- ▶ Encoder output passed to multi-head attention.
- ▶ First multi-head attention sublayer is masked to prevent the model to “cheat” and look at what is coming next in the sentence.
 - ▶ Ensures the prediction at position i only depends on outputs at position less than i .



[Attention Is All You Need, Vaswani 2017]

Transformer – Full architecture

Several encoder and decoder blocks are stacked.

- ▶ In original Transformer paper: $N_x = 6$

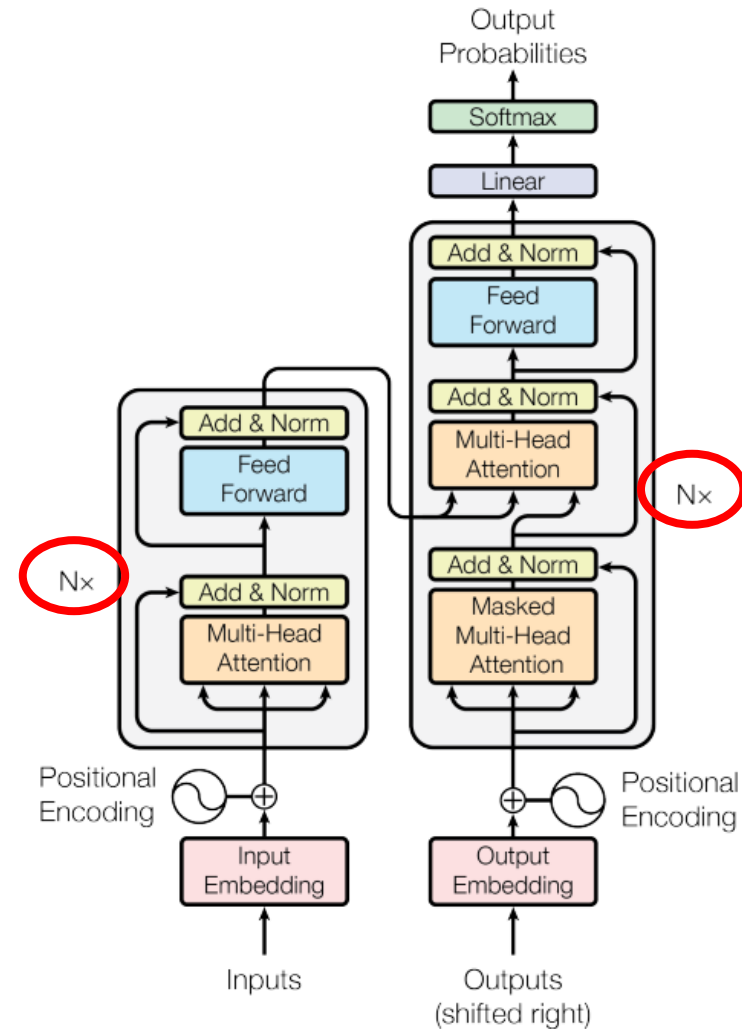
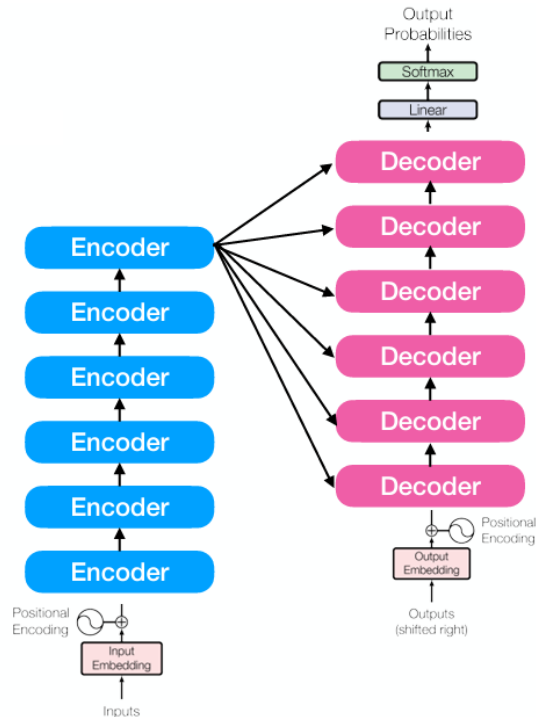


Image source: Sebastian Raschka, STAT 453: Intro to Deep Learning

[Attention Is All You Need, Vaswani 2017]

Transformer – Full architecture

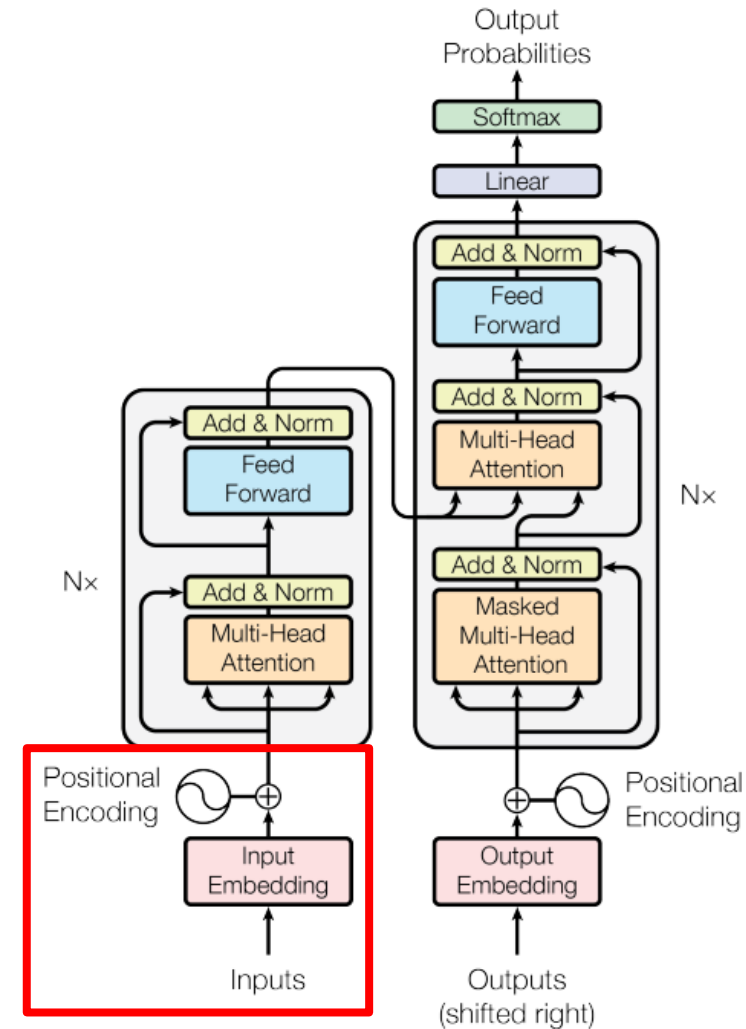
Input/output embedding:

- ▶ Original embeddings size: $d_{\text{model}} = 512$
- ▶ Embeddings are learned during the training process.

Positional encoding:

- ▶ RNNs were by design learning about the position of elements in the sequence.
- ▶ Transformers lose this information as they do not process the data sequentially!
- ▶ Positional encoding is used to retrieve the order information.

[Full explanation with example for positional embedding \(Hedu AI - Youtube\)](#)



[Attention Is All You Need, Vaswani 2017]

Transformer – Full architecture


Input/output embedding:

- ▶ Original embeddings size: $d_{\text{model}} = 512$
- ▶ Embeddings are learned during the training process.

Positional encoding:

- ▶ RNNs were by design learning about the position of elements in the sequence.
- ▶ Transformers lose this information as they do not process the data sequentially!
- ▶ Positional encoding is used to retrieve the order information.

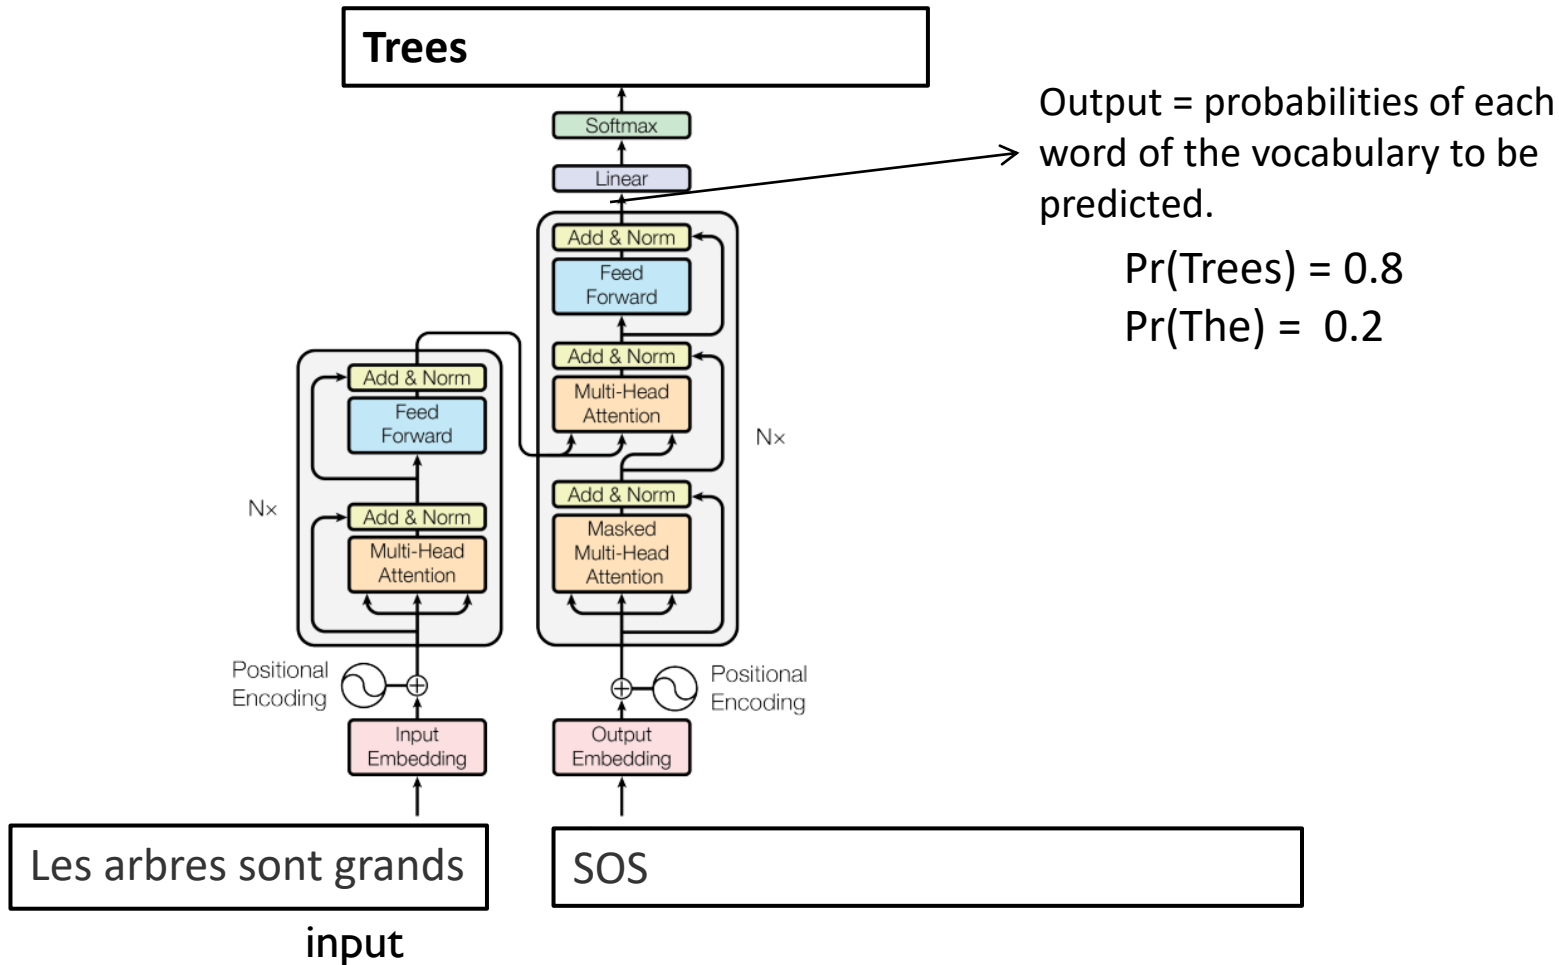
[Full explanation with example for positional embedding \(Hedu AI - Youtube\)](#)



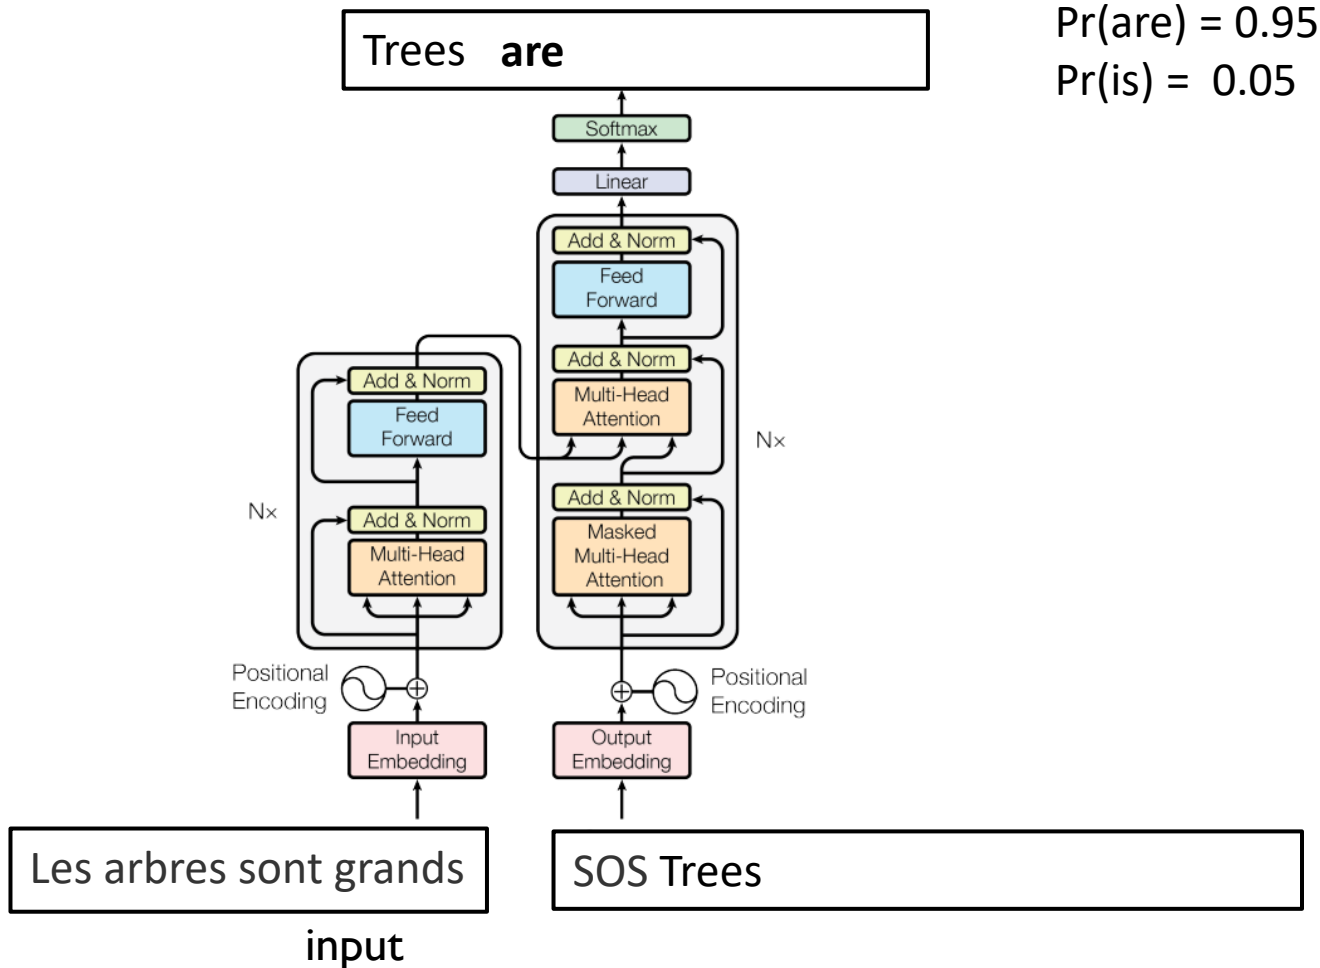
[Attention Is All You Need, Vaswani 2017]

Transformer – Translation (Inference)

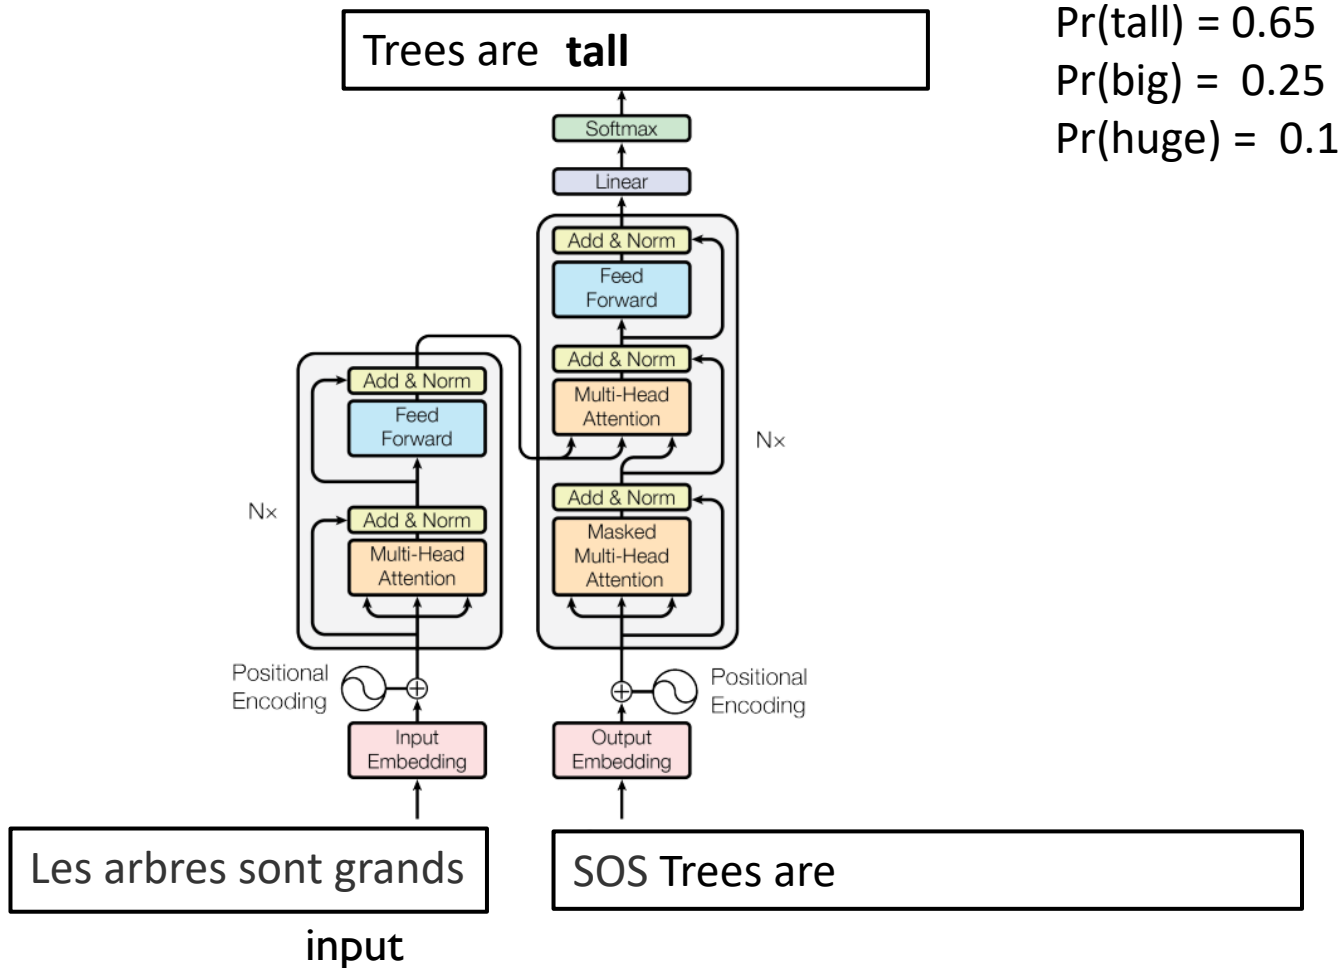
Predicted word = highest softmax probability.



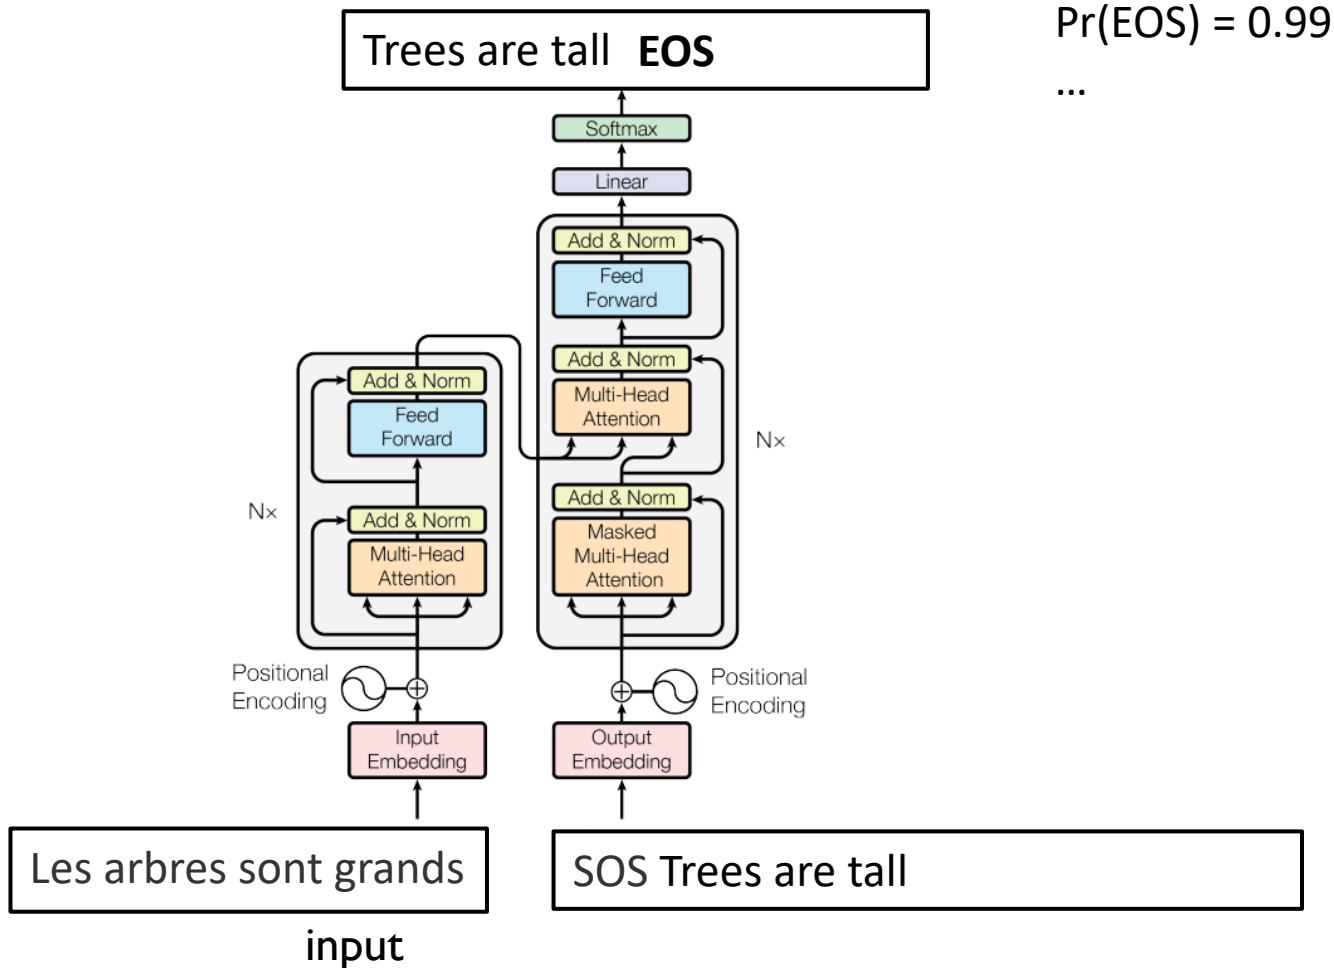
Transformer – Translation



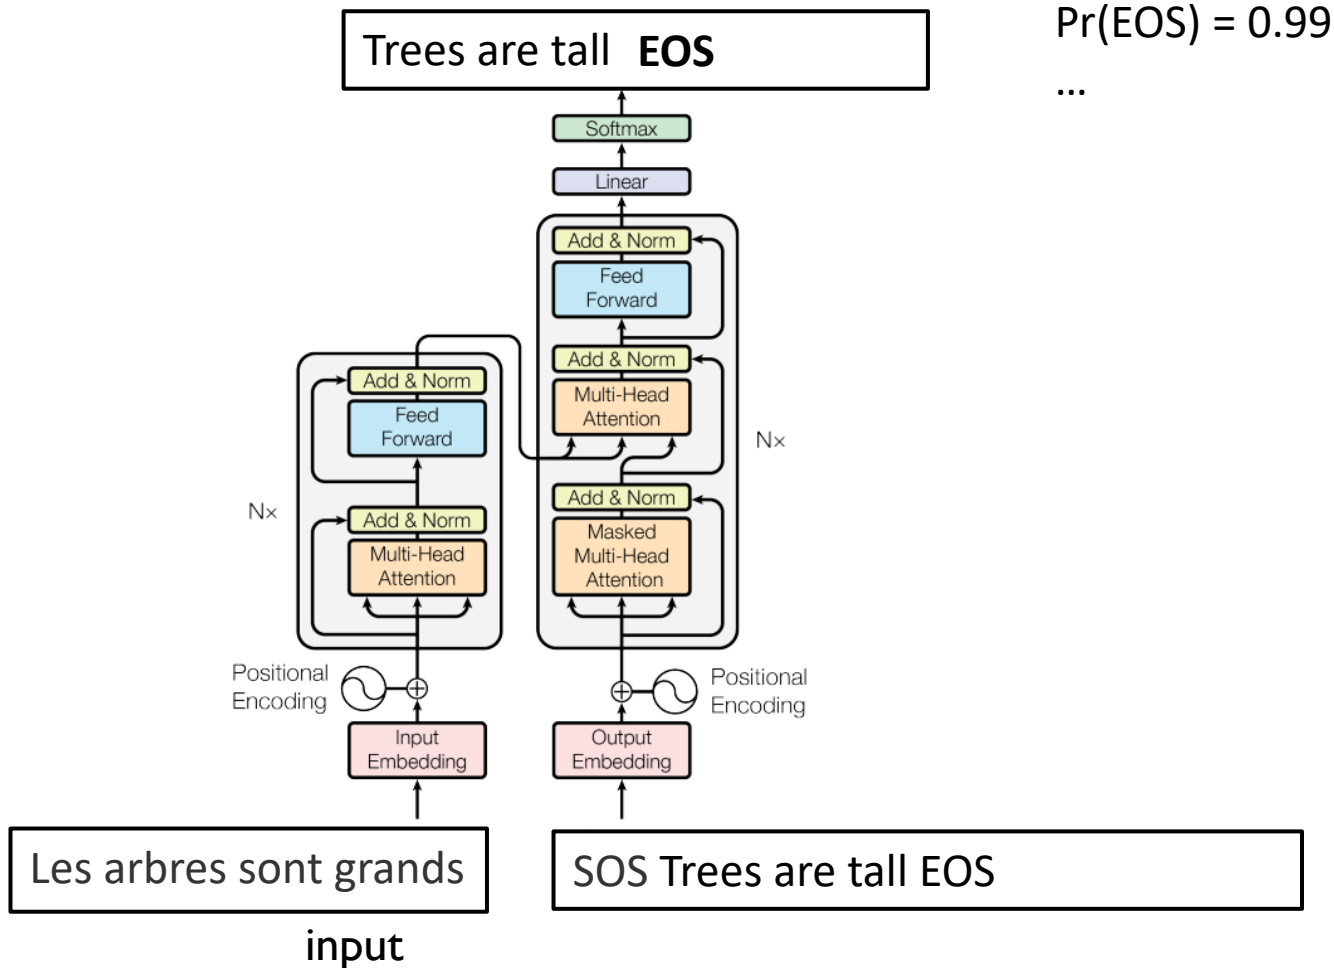
Transformer – Translation



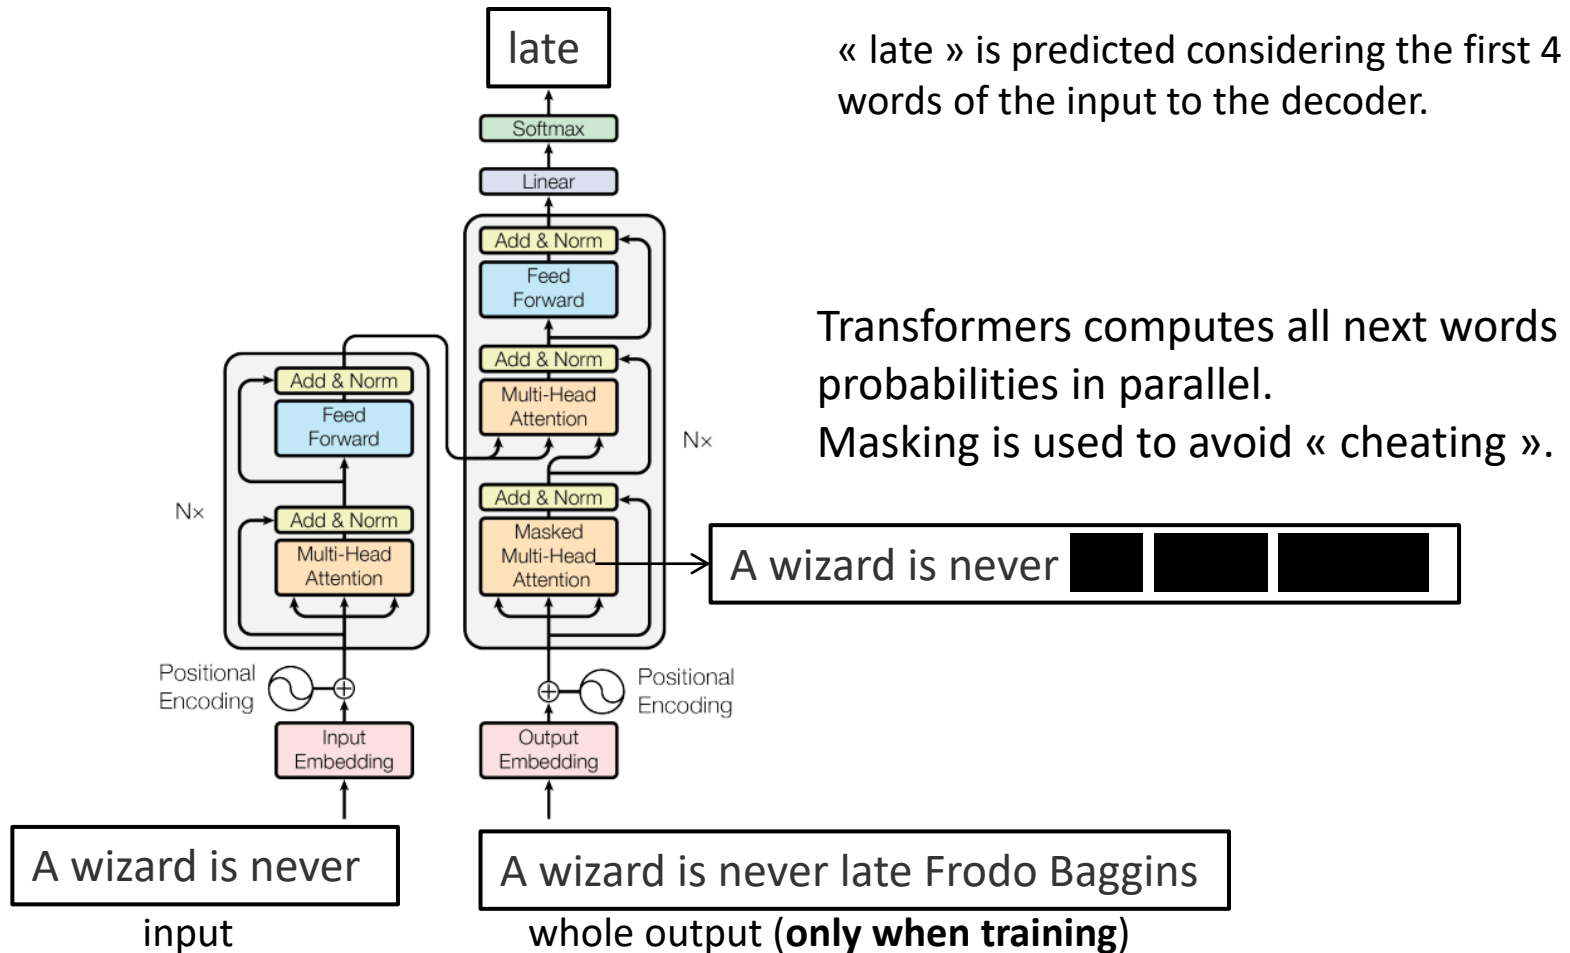
Transformer – Translation



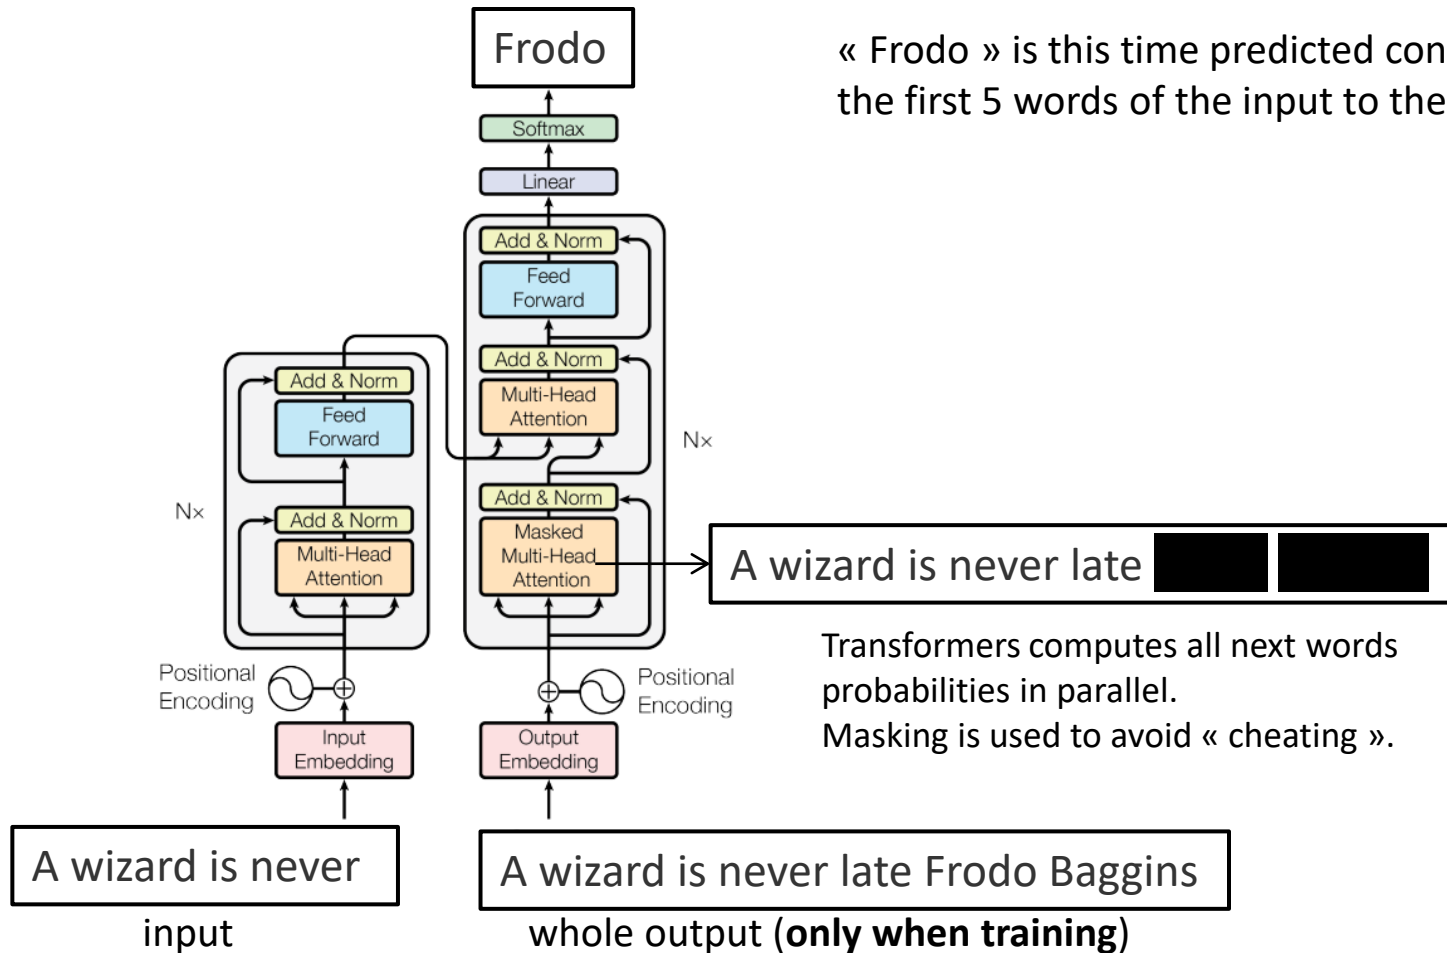
Transformer – Translation



Transformer – Next word prediction (Training)



Transformer – Next word prediction



« Frodo » is this time predicted considering the first 5 words of the input to the decoder.

Transformers computes all next words probabilities in parallel.
Masking is used to avoid « cheating ».

Vision Transformer (ViT)

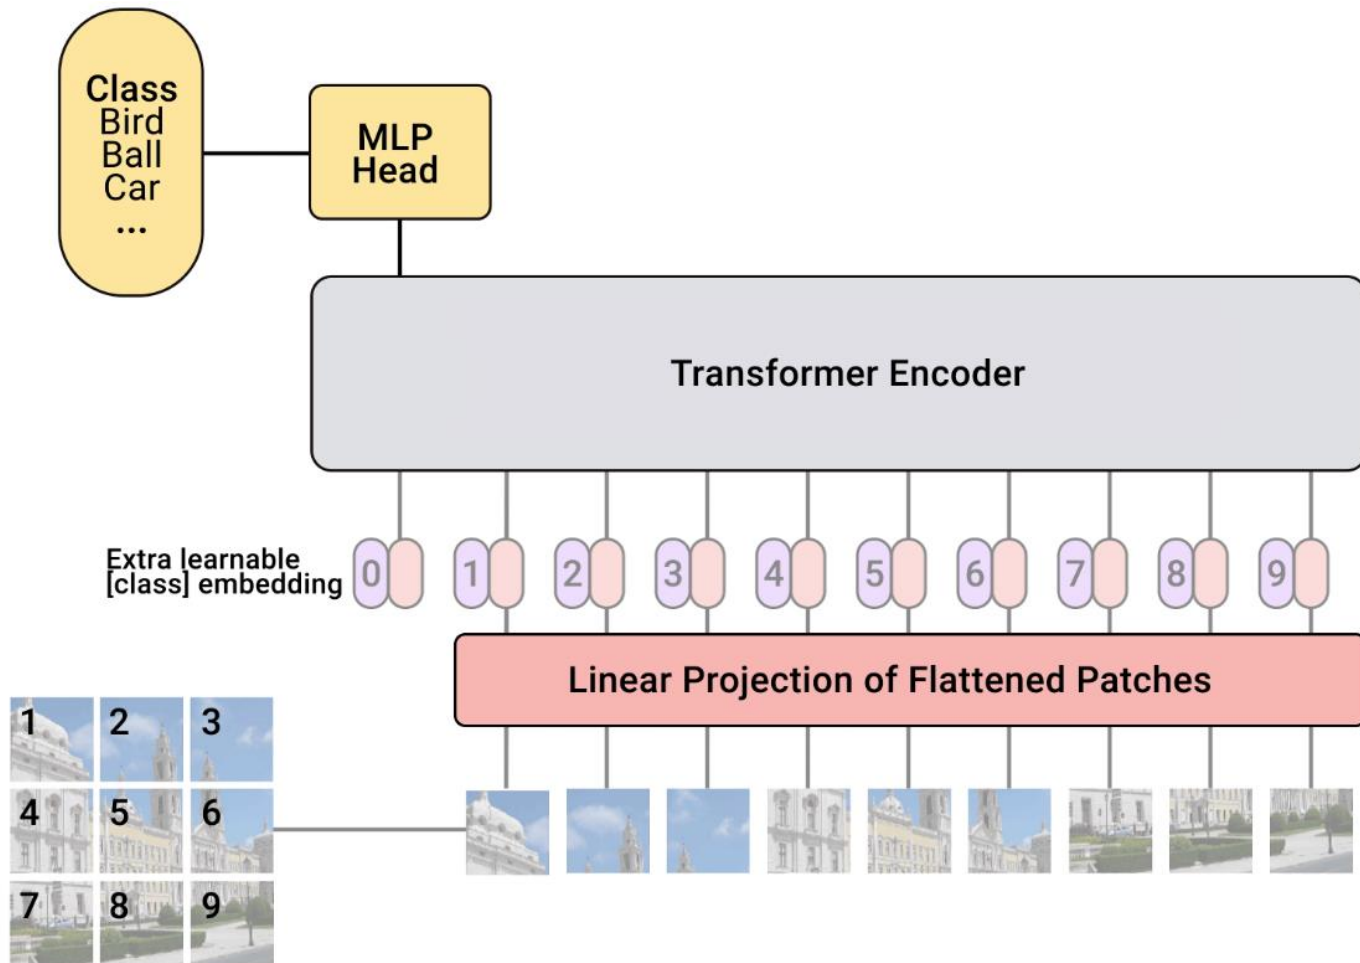


Image source: ai.googleblog.com/2020/12/transformers-for-image-recognition-at.html

Large language models

A **large language model** (LLM) is a general purpose language model consisting of a neural network with many parameters (typically billions of weights or more). LLMs trained on large quantities of unlabelled text perform well at a wide variety of tasks, a development which, since their emergence around 2018, has shifted the focus of natural language processing research away from the previous paradigm of training specialized supervised models for specific tasks.

Wikipedia

A **large language model** is an artificial neural network designed to analyze and generate natural language data. It is trained on vast amounts of text data and can perform various language tasks such as translation, summarization, and sentiment analysis. Large language models have revolutionized natural language processing, allowing machines to understand and generate human-like language with high accuracy.

ChatGPT

LLMs - A parameter story




Image source: huggingface.co/blog/large-language-models

LLMs are trained following 2 phases:

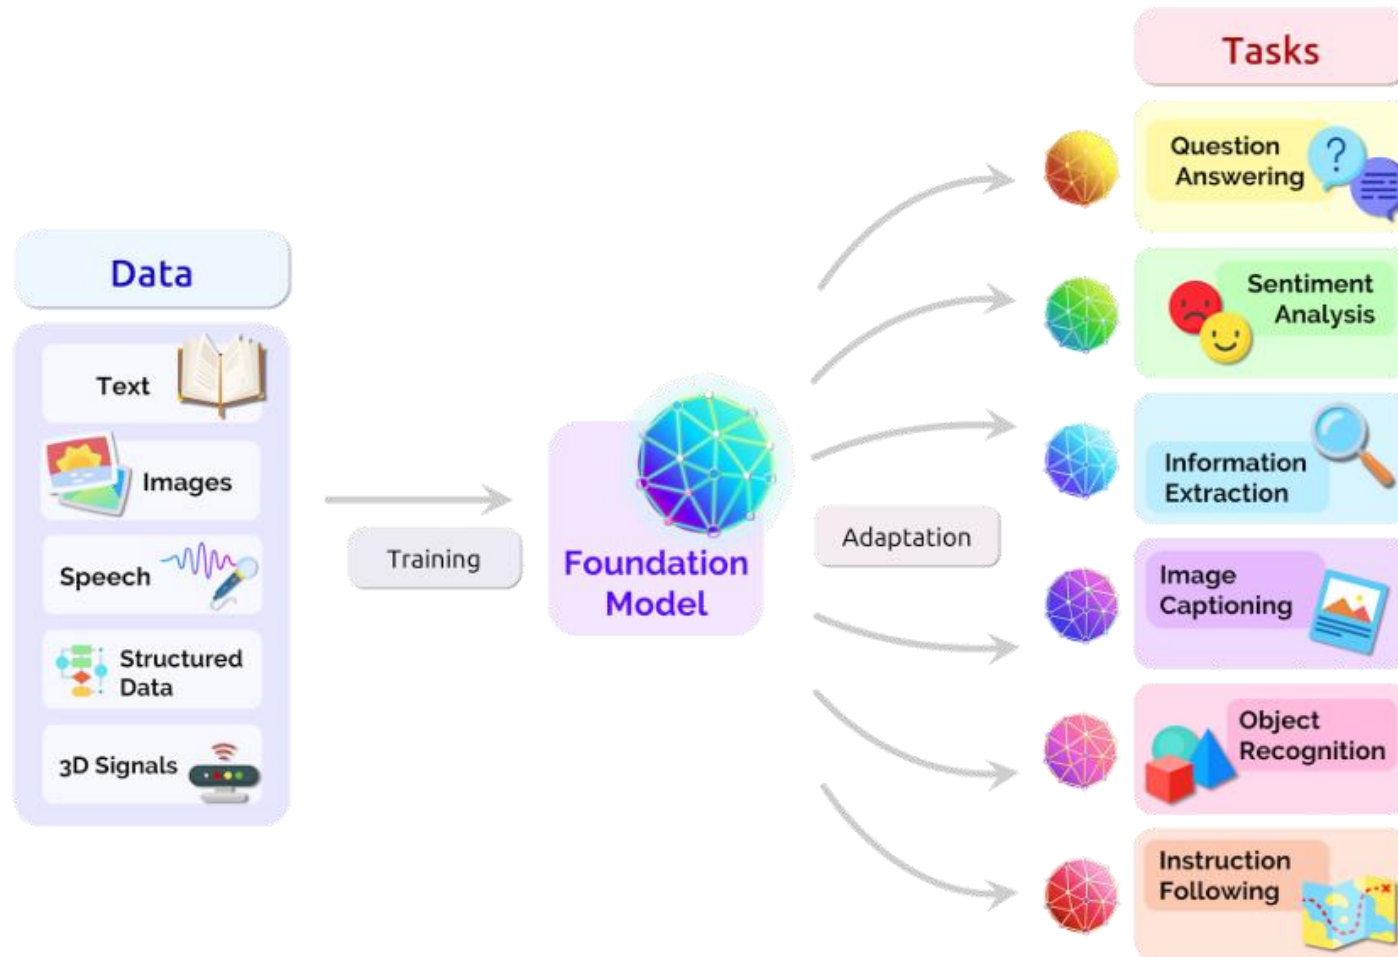
1. Pre-training

- ▶ Large amount of unlabelled data
- ▶ Self-supervised learning (learn one part of the input from another part of the input)
- ▶ General training (not task specific)
- ▶ Computationally expensive

2. Fine-tuning

- ▶ Labelled data
- ▶ Specific to down-stream task (e.g., translation, summarisation, Q&A, ...)
- ▶ Computationally cheaper

LLMs – Foundation models



[On the Opportunities and Risks of Foundation Models, Bommasani et al., 2022]

LLMs – Popular models

| LLM | DEVELOPER | POPULAR APPS THAT USE IT | # OF PARAMETERS | ACCESS |
|-----------------------------|---------------------------------|---|--|-------------|
| <u>GPT</u> | OpenAI | Microsoft, Duolingo, Stripe, Zapier, Dropbox, ChatGPT | 175 billion+ | API |
| <u>Gemini</u> | Google | Some queries on Bard | Nano: 1.8 & 3.25 billion; others unknown | API |
| <u>PaLM 2</u> | Google | Google Bard, Docs, Gmail, and other Google apps | 340 billion | API |
| <u>Llama 2</u> | Meta | Undisclosed | 7, 13, and 70 billion | Open source |
| <u>Vicuna</u> | LMSYS Org | Chatbot Arena | 7, 13, and 33 billion | Open source |
| <u>Claude 2</u> | Anthropic | Slack, Notion, Zoom | Unknown | API |
| <u>Stable Beluga</u> | Stability AI | Undisclosed | 7, 13, and 70 billion | Open source |
| <u>StableLM</u> | Stability AI | Undisclosed | 7, 13, and 70 billion | Open source |
| <u>Coral</u> | Cohere | HyperWrite, Jasper, Notion, LongShot | Unknown | API |
| <u>Falcon</u> | Technology Innovation Institute | Undisclosed | 1.3, 7.5, 40, and 180 billion | Open source |
| <u>MPT</u> | Mosaic | Undisclosed | 7 and 30 billion | Open source |
| <u>Mixtral 8x7B</u> | Mistral AI | Undisclosed | 46.7 billion | Open source |
| <u>XGen-7B</u> | Salesforce | Undisclosed | 7 billion | Open source |
| <u>Grok</u> | xAI | Grok Chatbot | Unknown | Chatbot |

LLMs – Features and tasks

| Model | Core differentiator | Pre-training objective | Parameters | Access | Information Extraction | Text Classification | Conversational AI | Summarization | Machine Translation | Content generation |
|-------------------|---|------------------------|------------|-------------|------------------------|---------------------|-------------------|---------------|---------------------|--------------------|
| BERT | First transformer-based LLM | AE | 370M | Source code | | | | | | |
| RoBERTa | More robust training procedure | AE | 354M | Source code | | | | | | |
| GPT-3 | Parameter size | AR | 175B | API | | | | | | |
| BART | Novel combination of pre-training objectives | AR and AE | 147M | Source code | | | | | | |
| GPT-2 | Parameter size | AR | 1.5B | Source code | | | | | | |
| T5 | Multi-task transfer learning | AR | 11B | Source code | | | | | | |
| LaMDA | Dialogue; safety and factual grounding | AR | 137B | No access | | | | | | |
| XLNet | Joint AE and AR | AE and AR | 110M | Source code | | | | | | |
| DistilBERT | Reduced model size via knowledge distillation | AE | 82M | Source code | | | | | | |
| ELECTRA | Computational efficiency | AE | 335M | Source code | | | | | | |
| PaLM | Training infrastructure | AR | 540B | No access | | | | | | |
| MT-NLG | Training infrastructure | AR and AE | 530B | API | | | | | | |
| UniLM | Optimised both for NLU and NLG | Seq2seq, AE and AR | 340M | Source code | | | | | | |
| BLOOM | Multilingual (46 languages) | AR | 176B | Source code | | | | | | |

AR = Autoregression

AE = Autoencoding

Seq2seq = Sequence-to-sequence



Highly appropriate

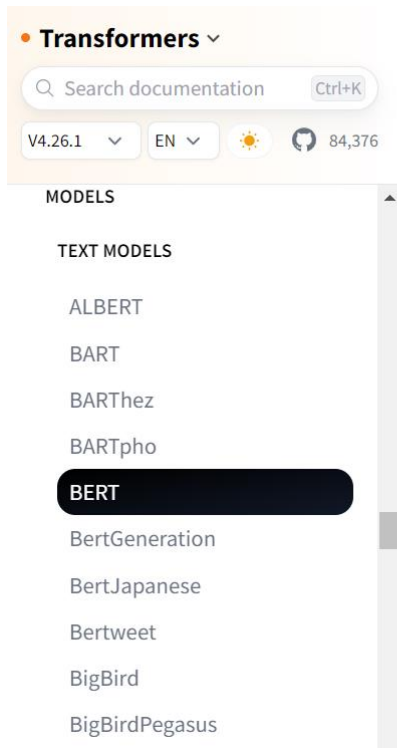
Appropriate

Somewhat appropriate

Image source: Janna Lipenkova - Choosing the right language model for your NLP use case

HuggingFace Transformers library

- ▶ Large collection of documentations and resources about models and datasets.



BERT

Overview

The BERT model was proposed in [BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding](#) by Jacob Devlin, Ming-Wei Chang, Kenton Lee and Kristina Toutanova. It's a bidirectional transformer pretrained using a combination of masked language modeling objective and next sentence prediction on a large corpus comprising the Toronto Book Corpus and Wikipedia.

The abstract from the paper is the following:

We introduce a new language representation model called BERT, which stands for Bidirectional Encoder Representations from Transformers. Unlike recent language representation models, BERT is designed to pre-train deep bidirectional representations from unlabeled text by jointly conditioning on both left and right context in all layers. As a result, the pre-trained BERT model can be fine-tuned with just one additional output layer to create state-of-the-art models for a wide range of tasks, such as question answering and language inference, without substantial task-specific architecture modifications.

https://huggingface.co/docs/transformers/v4.26.1/en/model_doc/index

- ▶ Chat-GPT is a chatbot based on the GPT-3.5 LLMs series.
- ▶ Fine-tuned on a variety of NLP tasks including translation, summarisation, Q&A and dialogue generation.


The model used for Chat-GPT is specially fine-tuned for chatbot applications, where the goal is to generate human-like responses to user inputs in a conversational manner.

► Also fine-tuned using reinforcement learning.

Step 1

Collect demonstration data and train a supervised policy.


A prompt is sampled from our prompt dataset.



A labeler demonstrates the desired output behavior.




This data is used to fine-tune GPT-3.5 with supervised learning.



Step 2

Collect comparison data and train a reward model.


A prompt and several model outputs are sampled.



A labeler ranks the outputs from best to worst.




This data is used to train our reward model.



Step 3

Optimize a policy against the reward model using the PPO reinforcement learning algorithm.

A new prompt is sampled from the dataset.



The PPO model is initialized from the supervised policy.



The policy generates an output.



The reward model calculates a reward for the output.



The reward is used to update the policy using PPO.



Image source: <https://openai.com/blog/chatgpt/>

Cost of training large language models

► Cost of training vs model size:

- \$2.5k- \$50k (110 million parameter model)
- \$10k- \$200k (340 million parameter model)
- \$80k- \$1.6m (1.5 billion parameter model)

[The cost of training nlp models: A concise overview, Sharir et al., 2020]

ChatGPT:

- Newest version (gpt-3.5-turbo): \$0.002 per 1000 tokens (10x less than a few months ago)
- Training GPT-3 consumed an estimated 1,287 MWh (~65k inhabitant city consumption per day in NZ) and produced 552 CO₂e (~80 Auckland-London return flights in economy class).

[The Carbon Footprint of Machine Learning Training Will Plateau, Then Shrink, Patterson et al., 2022]

- ▶ Transformers/Attention: [Attention Is All You Need, Vaswani et al., 2017]
<https://papers.nips.cc/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf>

- ▶ Popular models based on Transformers:
 - Google AI BERT [BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding”, Devlin et al., 2018]
<https://arxiv.org/pdf/1810.04805.pdf>
 - OpenAI GPT-2 [Language Models are Unsupervised Multitask Learners, Radford et al., 2018]
https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf
 - OpenAI GPT-3 [Language Models are Few-Shot Learners, Brown et al., 2020]
<https://arxiv.org/pdf/2005.14165.pdf>
 - Facebook AI BART [BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension, Lewis et al., 2019]
<https://arxiv.org/pdf/1910.13461>
 - Google AI T5 [Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer, Raffel et al., 2020]
<https://arxiv.org/pdf/1910.10683.pdf>

- NVIDIA Megatron-LM [Megatron-LM: Training Multi-Billion Parameter Language Models Using Model Parallelism, Shoeybi et al., 2020]
<https://arxiv.org/abs/1909.08053>
- Microsoft Turing & NVIDIA [Using DeepSpeed and Megatron to Train Megatron-Turing NLG 530B, A Large-Scale Generative Language Model, Smith et al., 2022]
<https://arxiv.org/pdf/2201.11990>

► Vision Transformers:

- [Image transformer, Parmar et al., 2018]
<http://proceedings.mlr.press/v80/parmar18a/parmar18a.pdf>
- [An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale, Dosovitskiy et al. , 2020]
<https://arxiv.org/pdf/2010.11929.pdf>
<https://ai.googleblog.com/2020/12/transformers-for-image-recognition-at.html>