



THE UNIVERSITY OF  
**AUCKLAND**  
Te Whare Wānanga o Tāmaki Makaurau  
NEW ZEALAND

# *Transfer Learning*

*COMPSCI 760*  
*2024 Semester 1*

*Olivier Graffeuille*

ogra439@aucklanduni.ac.nz



# Outline

- Tasks & Domains
- What is Transfer Learning?
- Transfer Learning Methods
  - Instance Transfer
  - Feature Transfer
  - Parameter Transfer
- Multi-Task Learning



# Outline

- ***Tasks & Domains***
- What is Transfer Learning?
- Transfer Learning Methods
  - Instance Transfer
  - Feature Transfer
  - Parameter Transfer
- Multi-Task Learning

## Prediction in Machine Learning


- Most ML techniques aim to predict a label from an input



*Classification*



*Object detection*



*Emotion Recognition*



*...cat age estimation?*

## Tasks (informally)

- A task is a function we are trying to learn
- i.e. what is the objective of the neural network?
  - Classify an image
  - Estimate the age of cat images
  - Detect cancer in 3D MRI scans
  - Speech to text

## Tasks (formally)

- **Task:**

A label space and a function.

$$T = \{\mathcal{Y}, f(\cdot)\}$$

Where  $\mathcal{Y}$  represents the label space

*e.g.  $\{0, 1\}$  for a binary classification problem, or any 19x19x16 dimensional matrix for YOLO*

And  $f(\cdot)$  is a function that maps from  $\mathcal{X} \rightarrow \mathcal{Y}$  (we're trying to learn this function)

*e.g. a model that predicts the most likely class/object position, given an input*

Note:

Discriminative models learn a function  $P(y | X)$

Generative models model  $P(X, y)$  instead

## Domains (informally)

- The possible inputs to our model
- i.e. what kind of data do we want our model to handle?
  - Any image
  - Images of leaves of trees
  - Hand-drawn digits
  - Audio of people speaking Arabic
  - Medical records

## Domains (formally)

- **Domain:**

A feature space and marginal probability density function.  $D = \{\mathcal{X}, p(X)\}$

Where  $\mathcal{X}$  represents the input space

*e.g.  $\mathbb{R}$  for a single variable model,  $\mathbb{R}^{28} \times \mathbb{R}^{28}$  for MNIST*

And  $P(X)$  represents the marginal distribution of  $X$

*e.g. a normal distribution for height, or a very complicated intractable distribution for MNIST*



## Classes, Domains and Tasks

- **Class:**

A distinct category or group that a data point is labelled as belonging to.

*E.g. “Cat”, “Dog”, “Happy”, “Popular”, “Correct”*

- **Domain:**

$$D = \{\mathcal{X}, p(X)\}$$

A feature space and marginal probability density function.

*E.g. Sunny photos of pets, painting of pets, audio recording of pet.*

- **Task:**

$$T = \{\mathcal{Y}, p(y|X)\}$$

A label space and conditional probability density function.

*E.g. Does the image contain a dog or a cat? How many dogs are in this image?*



# Outline

- Tasks & Domains
- ***What is Transfer Learning?***
- Transfer Learning Methods
  - Instance Transfer
  - Feature Transfer
  - Parameter Transfer
- Multi-Task Learning


## Transfer Learning (informally)

### Traditional ML


vs

### Transfer Learning

- Isolated, single task learning:
  - Knowledge is not retained or accumulated. Learning is performed w.o. considering past learned knowledge in other tasks



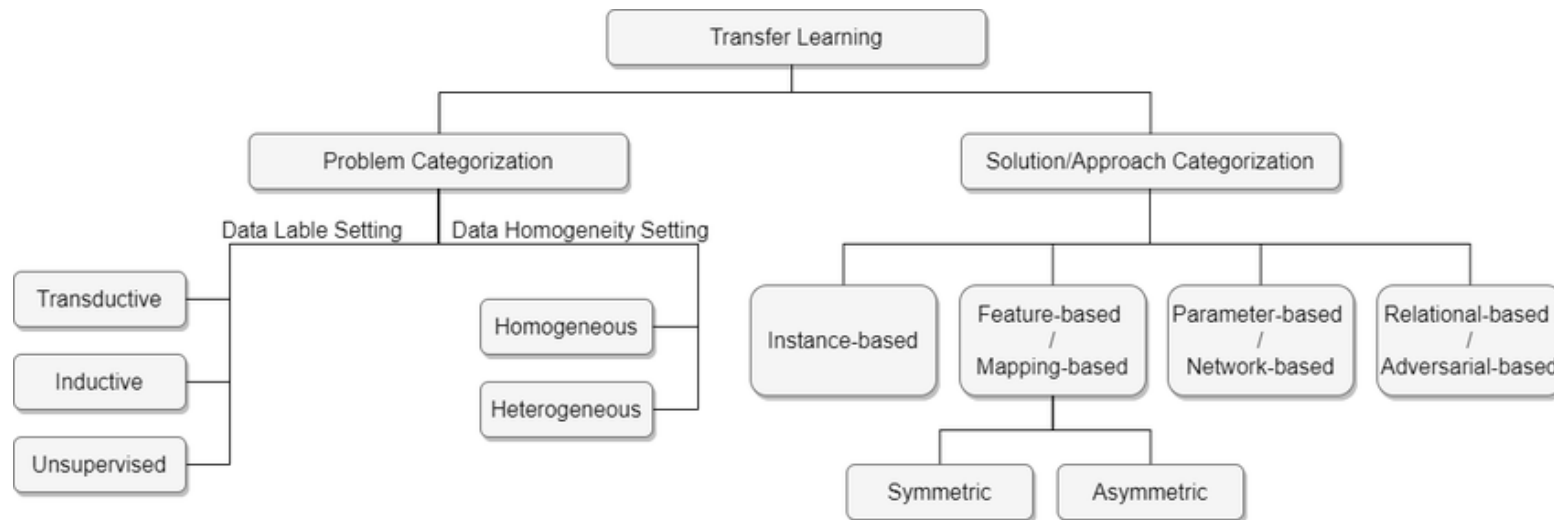
- Learning of a new tasks relies on the previous learned tasks:
  - Learning process can be faster, more accurate and/or need less training data



## Transfer Learning (formally)

- Given a source domain  $D_S$  with corresponding learning task  $T_S$  and a target domain  $D_T$  with corresponding learning task  $T_T$
- Transfer Learning (TL) is the process of improving the target predictive function  $f_T$  by using the related information from  $D_S$  and/or  $T_S$ , where  $D_S \neq D_T$  or  $T_S \neq T_T$ 
  - If  $D_S = D_T$  and  $T_S = T_T$  then the two tasks are just identical
- Source task generally has much more data than the target task

## Transfer Learning Taxonomy



## Change in domain

- $D_S \neq D_T$ 
  - $\mathcal{X}_S \neq \mathcal{X}_T$

### Heterogeneous TL

Different “kinds” of input


*e.g. Transferring from one kind of sensor to another*

- $\mathcal{X}_S = \mathcal{X}_T$  but,  $P(X_S) \neq P(X_T)$

### Domain adaptation

Same “kind” of input but different distribution

*e.g. Transferring from real images to cartoon images*



## Change in task

- $T_S \neq T_T$

- $y_S \neq y_T$

Different “kinds” of labels

*e.g. Transferring image classification to image object detection*

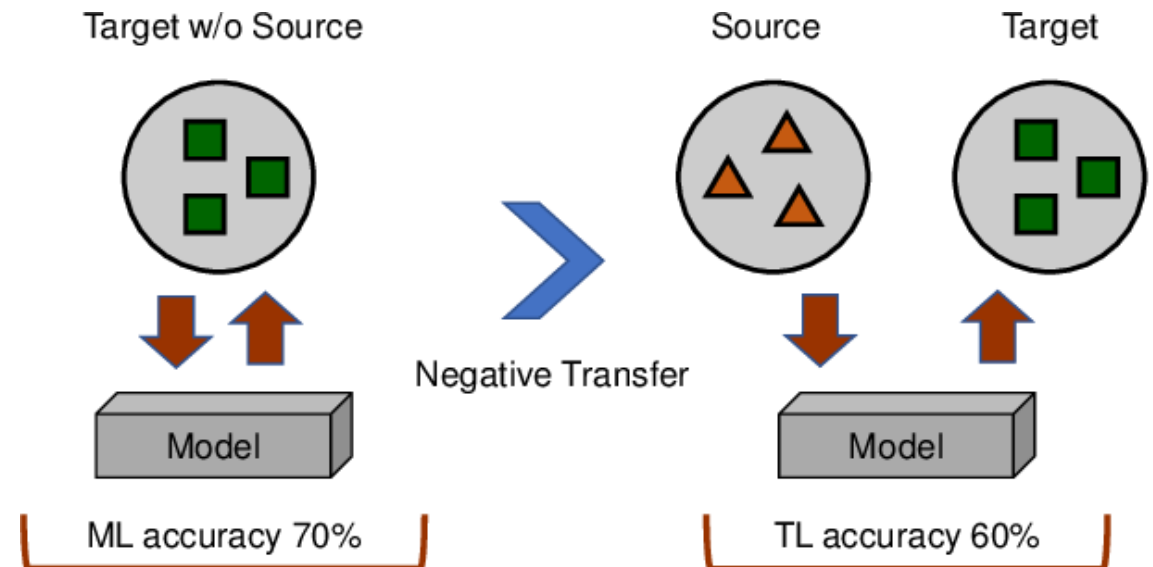
- $f_S(.) \neq f_T(.)$  i.e.  $P_S(y|x) \neq P_T(y|x)$

Relationship between input and label is different

*e.g. Transferring from one language to another (“sale” means dirty in French)*





## Negative Transfer

- Sometimes transfer learning can reduce performance on target task
  - This is called **Negative Transfer**
- Transfer learning is not guaranteed to improve performance!



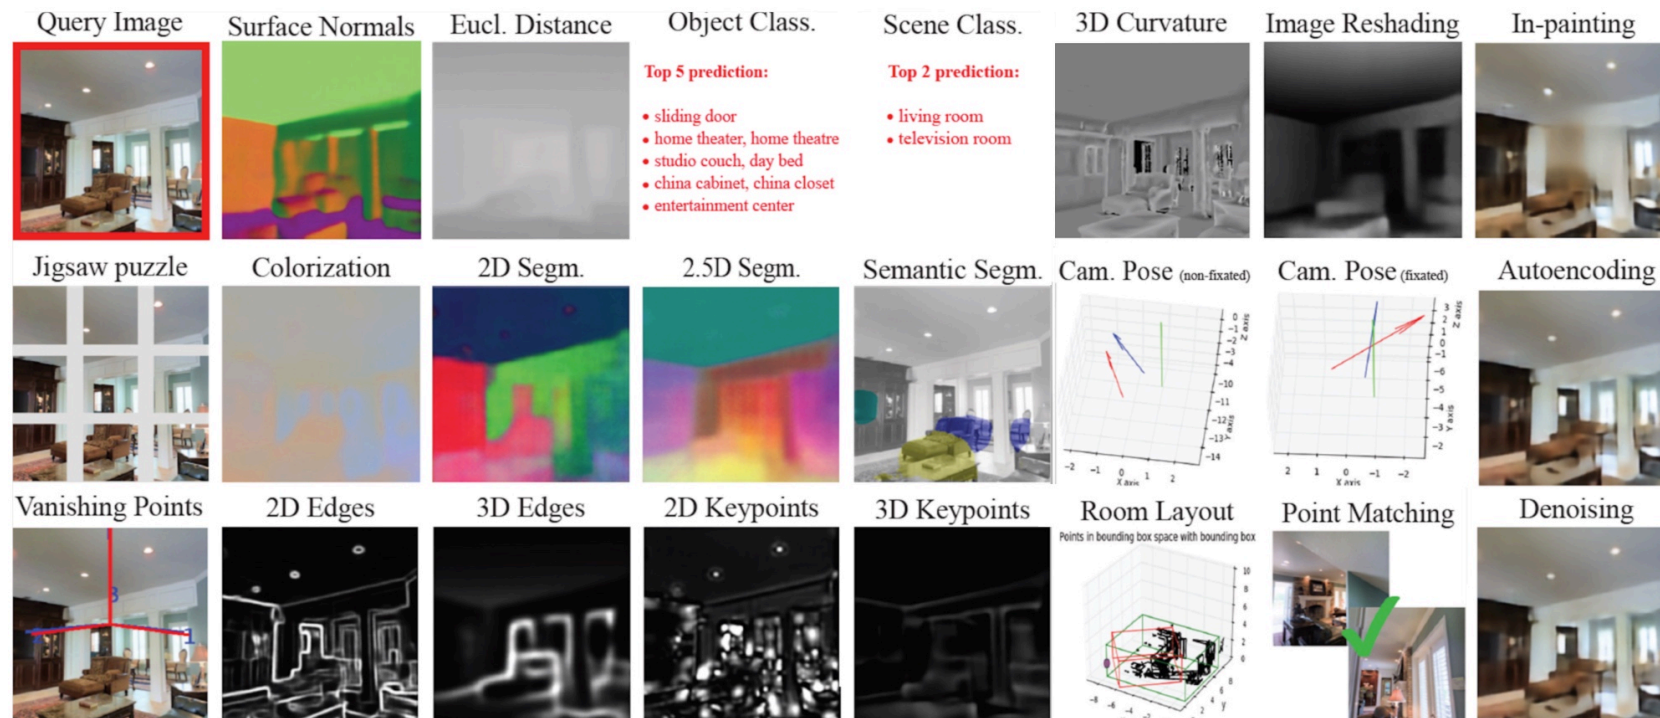


## Negative Transfer

- Assumption in transfer learning: *tasks are related*
  - Transferring from Japanese alzheimers detection model to NZ alzheimers detection model? 
  - Transferring from image classification (imagenet) to car object detection? 
  - Transferring from an NLP task (e.g. Q&A) to object detection of cars? 
  - Transferring from different language for an NLP task (English -> Chinese)? 
  - Transferring from global climate model to ecology sustainability model?

## Negative Transfer

- Demo of transfer of image tasks  
<http://taskonomy.stanford.edu/transfers/>






# Outline

- Tasks & Domains
- What is Transfer Learning?
- ***Transfer Learning Methods***
  - *Instance Transfer*
  - Feature Transfer
  - Parameter Transfer
- Multi-Task Learning

## Instance Transfer

- Simplest solution:  
can we re-use some data samples from the source domain on the target domain directly?
- Re-weight each source data instance
- Assumptions:  
Source and target domain have same feature and label spaces  
Source and target domain have same task function  
Different marginal probability distributions



## Instance Transfer

- Assumptions:

$$\mathcal{Y}_S = \mathcal{Y}_T$$

$$P_S(y|x) = P_T(y|x)$$

$$\mathcal{X}_S = \mathcal{X}_T$$

$$P(X_S) \approx P(X_T)$$

What transfer learning setting is this?

- Techniques: weight samples from source domain based on:
  - based on improvement in performance (TrAdaBoost)
  - based on relative probability of occurrence  $\frac{P(X_T)}{P(X_S)}$




# Outline

- Tasks & Domains
- What is Transfer Learning?
- Transfer Learning Methods
  - Instance Transfer
  - *Feature Transfer*
  - Parameter Transfer
- Multi-Task Learning

## Feature Transfer

- Find a common feature space for source and task domain  $\mathcal{X}_S \neq \mathcal{X}_T$ 
  - *e.g. MRI scans and blood tests of patient*  
*or: time series data + tabular data*  
*or: medical records from different hospitals which measure different variables*
- This allows us to use the same classifier for the source and target model





# Outline

- Tasks & Domains
- What is Transfer Learning?
- Transfer Learning Methods
  - Instance Transfer
  - Feature Transfer
  - *Parameter Transfer*
- Multi-Task Learning





## Parameter Transfer


- Transfer parameters from model trained on source dataset to target dataset
  - Intuition: models encode knowledge in their parameters  
So, we can transfer parameters to transfer knowledge
  - Transferring knowledge from source model to target model
- Decision trees can transfer decision nodes
- Neural networks can transfer model weights



## Fine Tuning

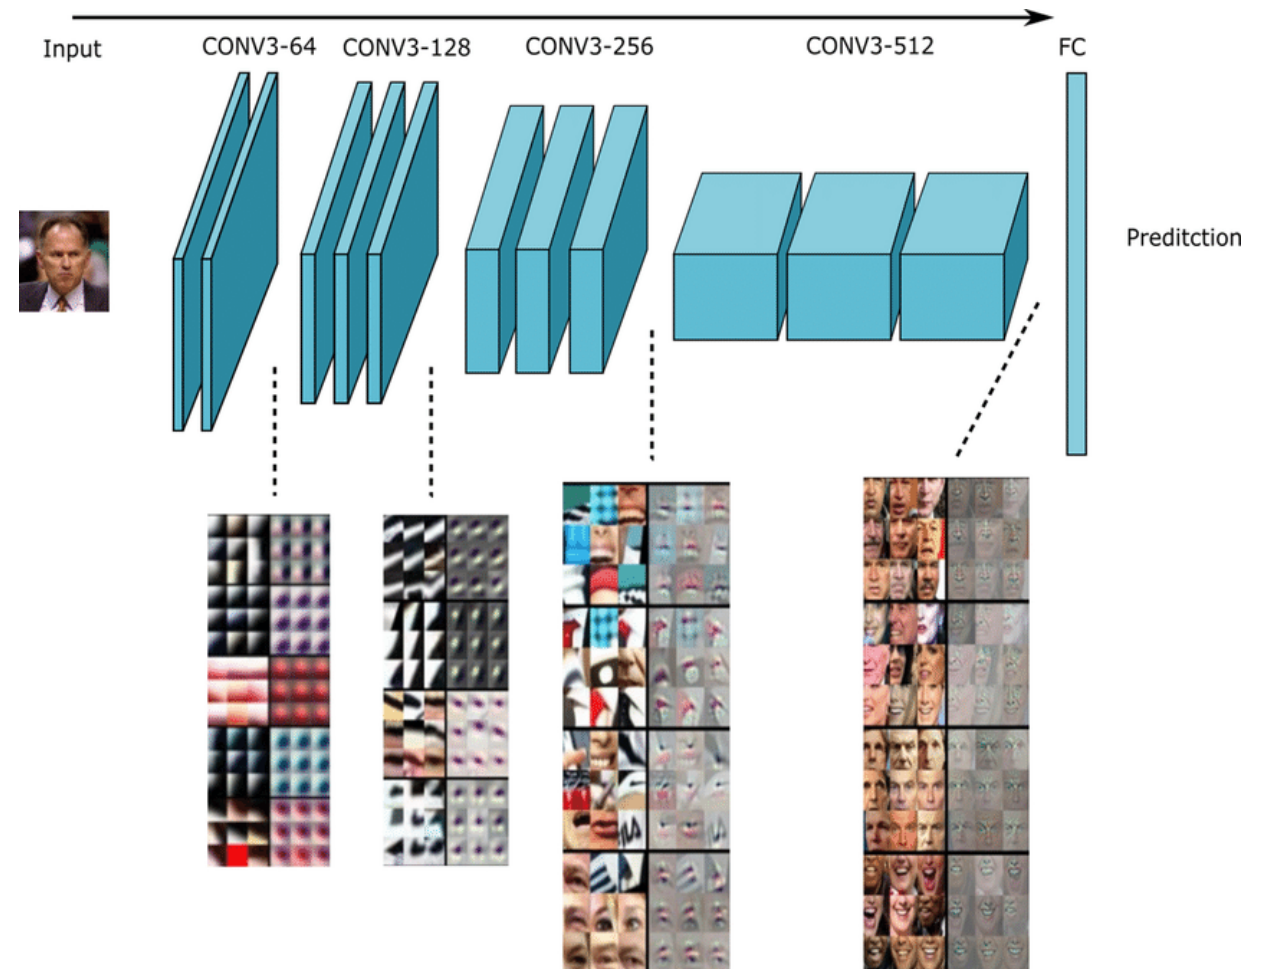
- Train model on source data, then fine-tune on target data
- Challenge: limited target data, careful to not overfit
- Techniques:
  - Limiting the number of training epochs
  - Use a smaller learning rate
  - Freeze some layers
  - Some combination (gradual unfreezing of select layers)

## Which layers to transfer?




## Which layers to transfer?


- It depends
- Computer vision:  
Usually transfer shallow  
(early) layers



## Which layers to transfer?




## Which layers to transfer? Trade-off





## Pre-Trained Models

- When training neural networks, we can either:
  - Select a random initialisation for the network parameters (noise)
  - Select an initialisation from a pre-trained model
- Popular pre-trained models:
  - Computer vision
    - VGG16
    - ResNet-50
    - ViT
  - NLP
    - Bert
    - Llama 3



## Adapters

- Small NN architecture which is inserted into pre-trained models
  - The pre-trained model is entirely frozen
  - The adapter is trained normally
- Example: LoRA
  - New weight matrix added (parallel to existing layers)
  - New weight matrix has many fewer parameters





- Many pre-trained models can serve as the starting point for training. These pre-trained models give a more reliable architecture and save time and resources.
- Transfer learning:
  - used in scenarios where there is not enough data for training or when we want better results in a short amount of time.
  - involves selecting a source model similar to the target domain, adapting the source model to the target model before transferring the knowledge, and training the source model to achieve the target model.
- It is common to fine-tune the higher-level layers of the model while freezing the lower levels as the basic knowledge is the same that is transferred from the source task to the target task of the same domain.
- In tasks with a small amount of data, if the source model is too similar to the target model, there might be an issue of overfitting. To prevent the transfer learning model from overfitting, it is essential to tune the learning rate, freeze some layers from the source model, or add linear classifiers while training the target model can help avoid this issue.




# Outline


- Tasks & Domains
- What is Transfer Learning?
- Transfer Learning Methods
  - Instance Transfer
  - Feature Transfer
  - Parameter Transfer
- ***Multi-Task Learning***

## Multi-Task Learning

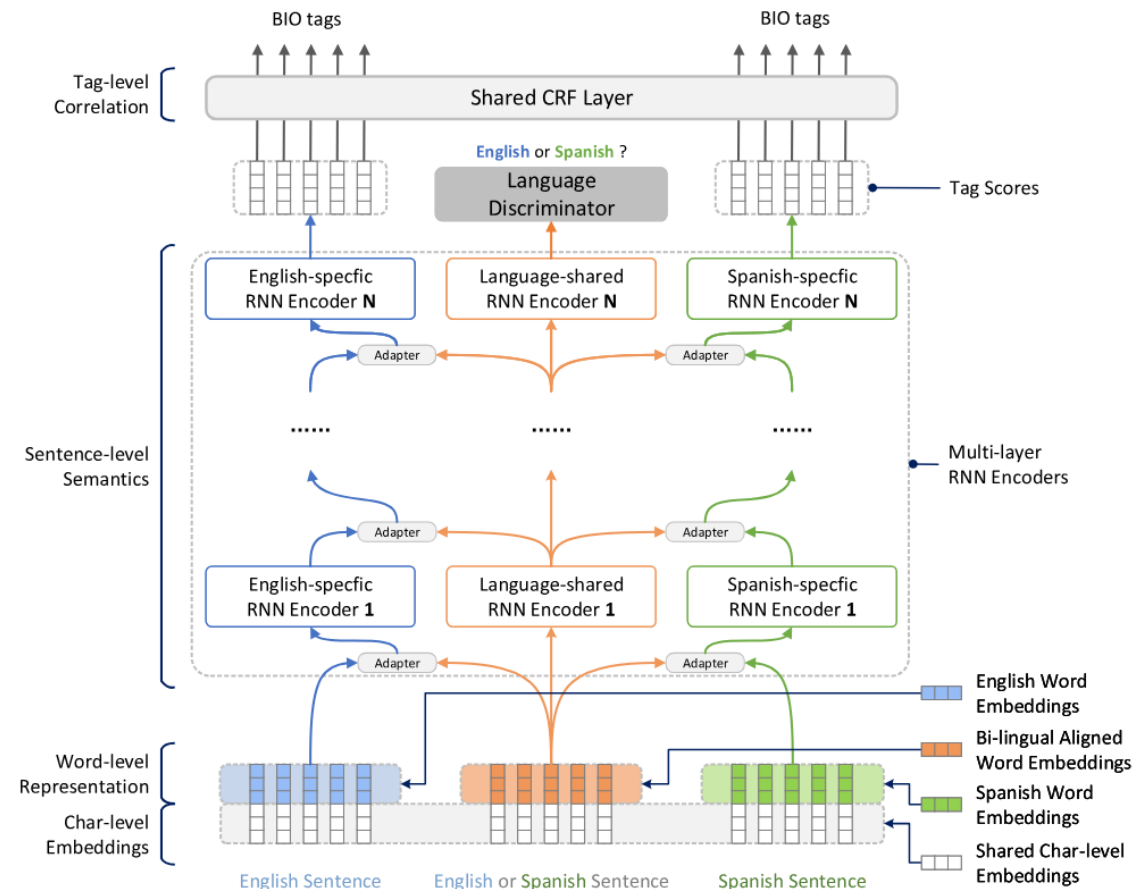
- Jointly train multiple related tasks
  - Improves performance of all tasks by knowledge transfer
  - Reduce inference cost by sharing network weights
- Multi-task Learning (MTL) is a subset of transfer learning
  - We care about performance of all tasks
  - Tasks are trained jointly



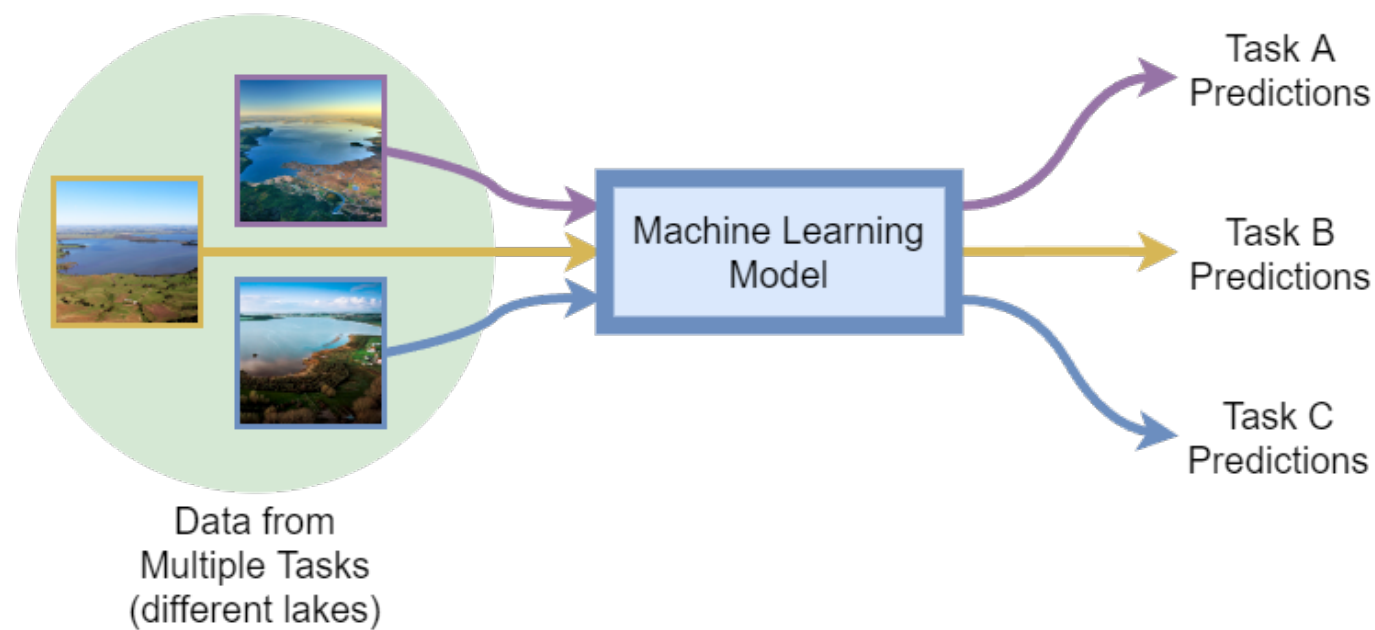
## Examples



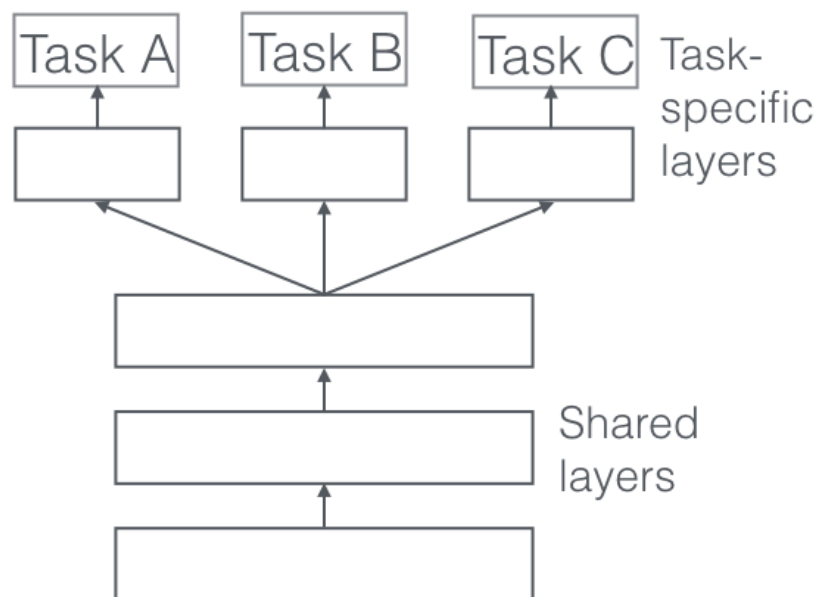
## Examples



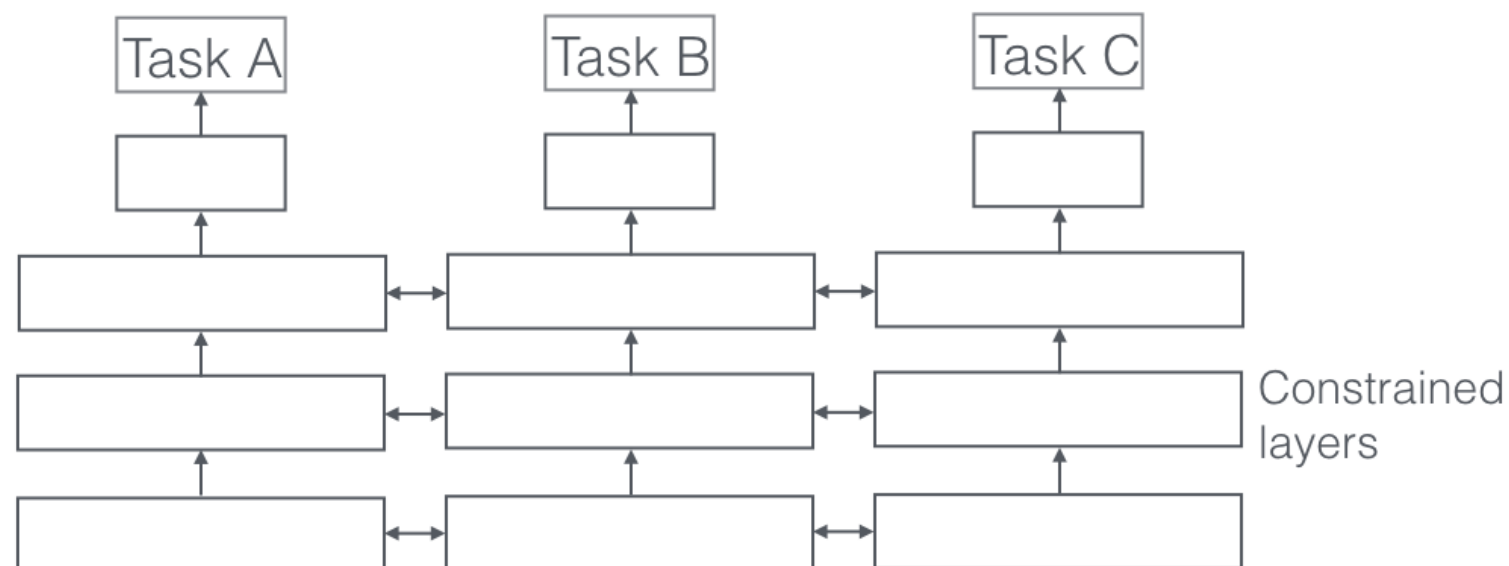
## Examples



## Techniques




Hard parameter sharing



Soft parameter sharing

## Hard-Parameter Sharing


Single-task learning architecture:





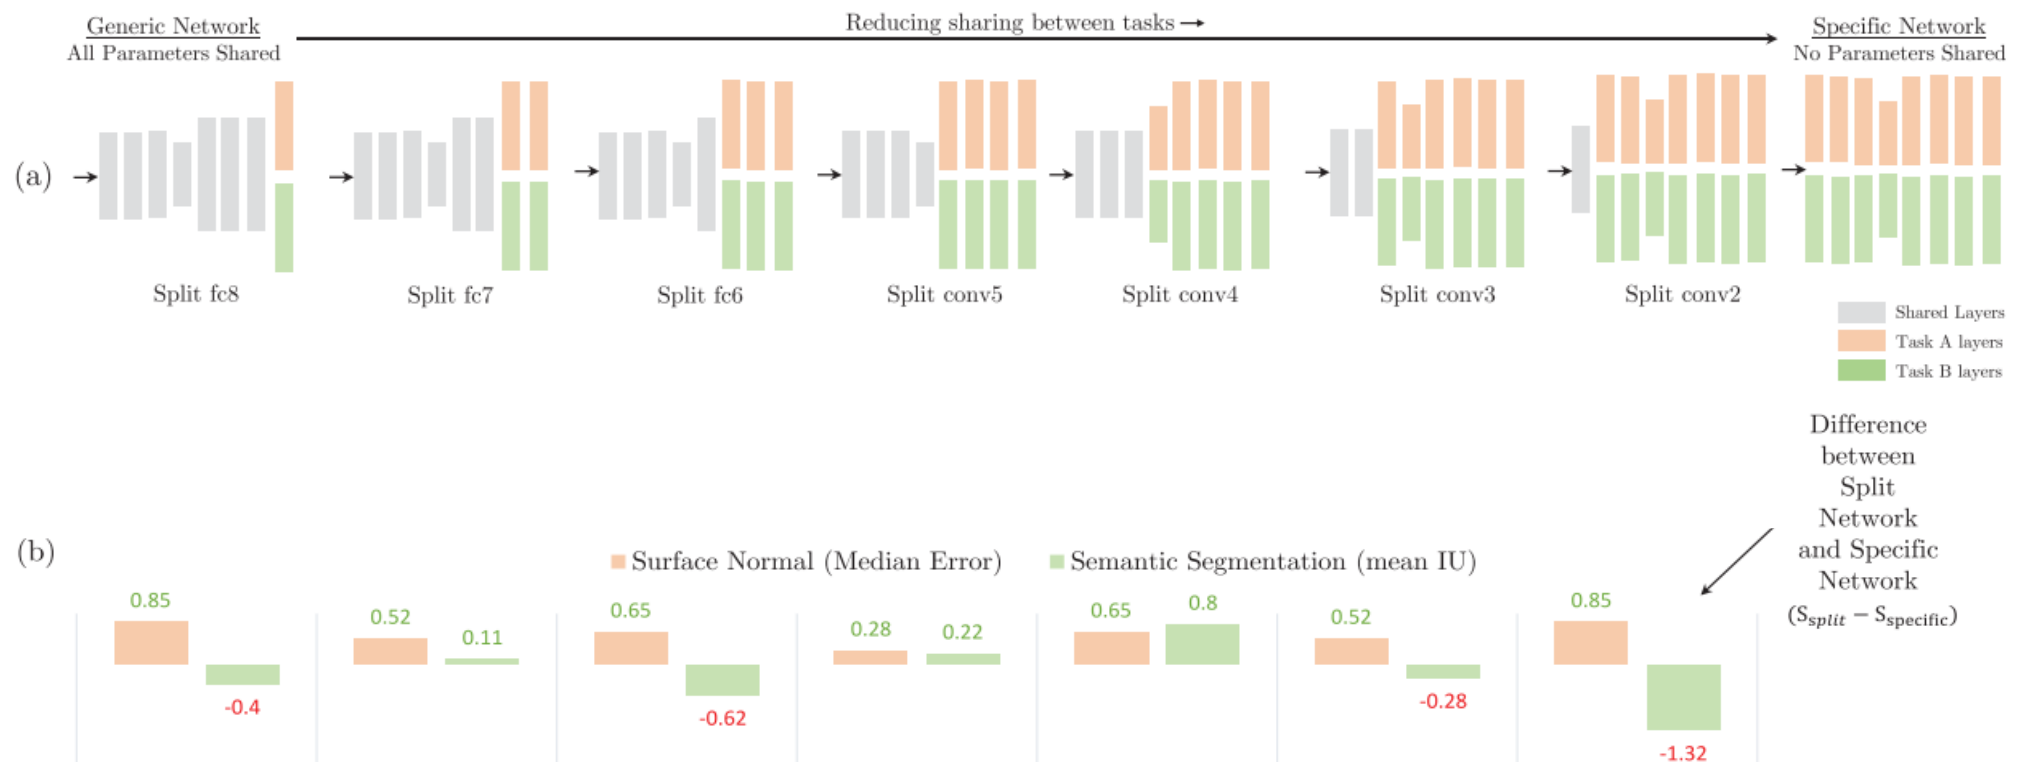
## Hard-Parameter Sharing

Multi-task learning architecture: (hard-sharing)



## Hard-Parameter Sharing

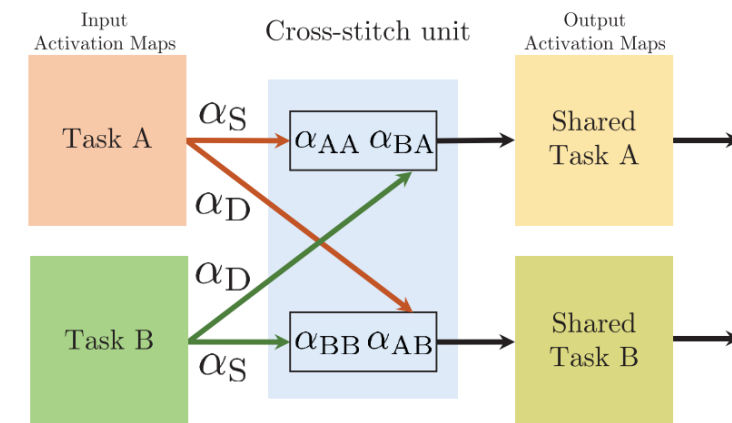

What layers should we share?



# Multi-Task Learning

## Cross-Stitch Networks

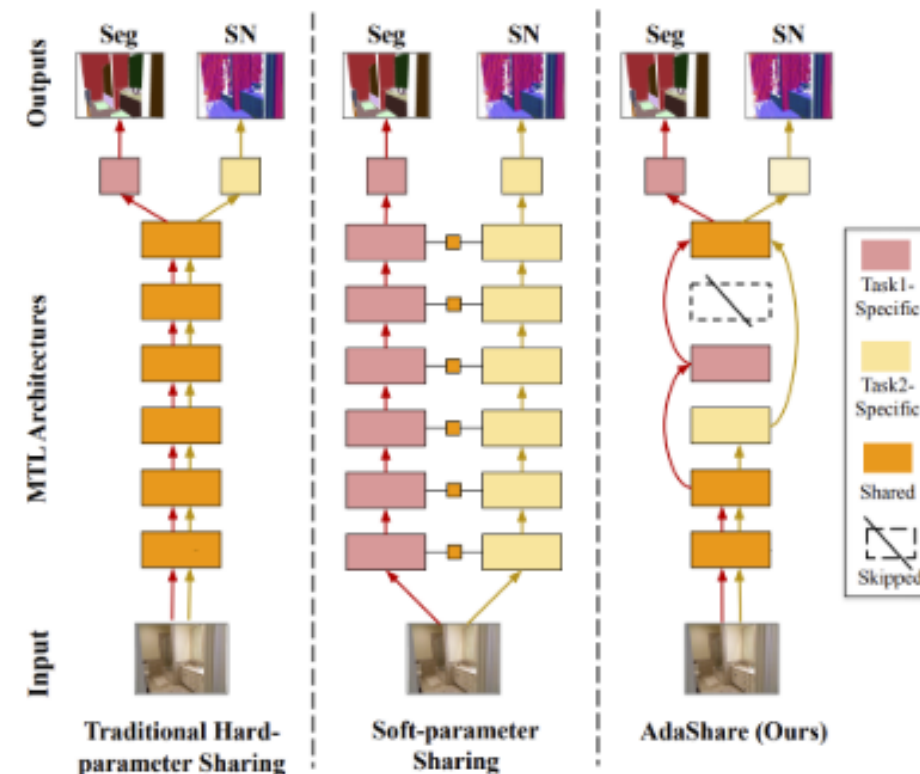
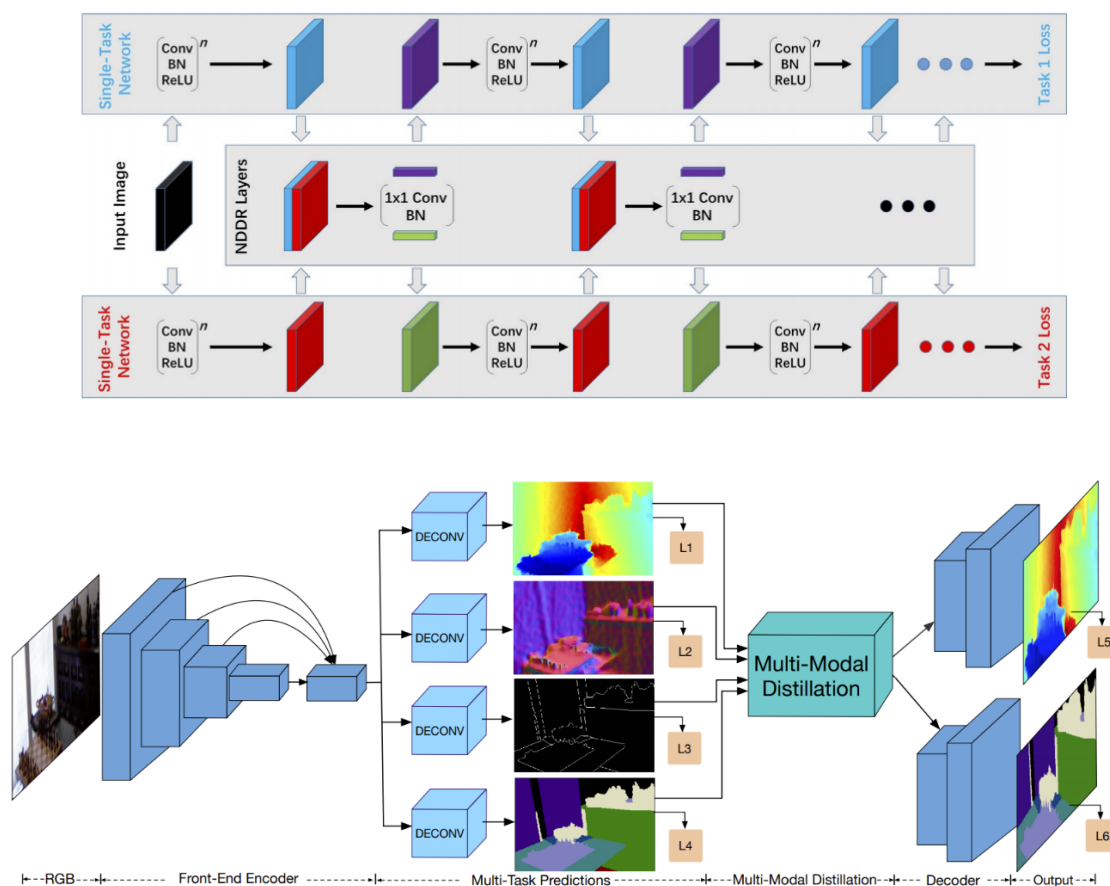
What layers should we share?  
**Let the model decide!**



“Cross-stitch” layers compute linear combinations of intermediate representations

# Multi-Task Learning

## Other MTL architectures



## Multi-Task Learning Optimisation

Q: How do we train a model with multiple tasks?

A: We have multiple loss functions!

Q: How do we weigh these loss functions?

A: [1000 different proposed techniques]

$$\min_{\theta} L_i(f(x_i; \theta_{sh}, \theta_i), y_i),$$

*STL optimisation*

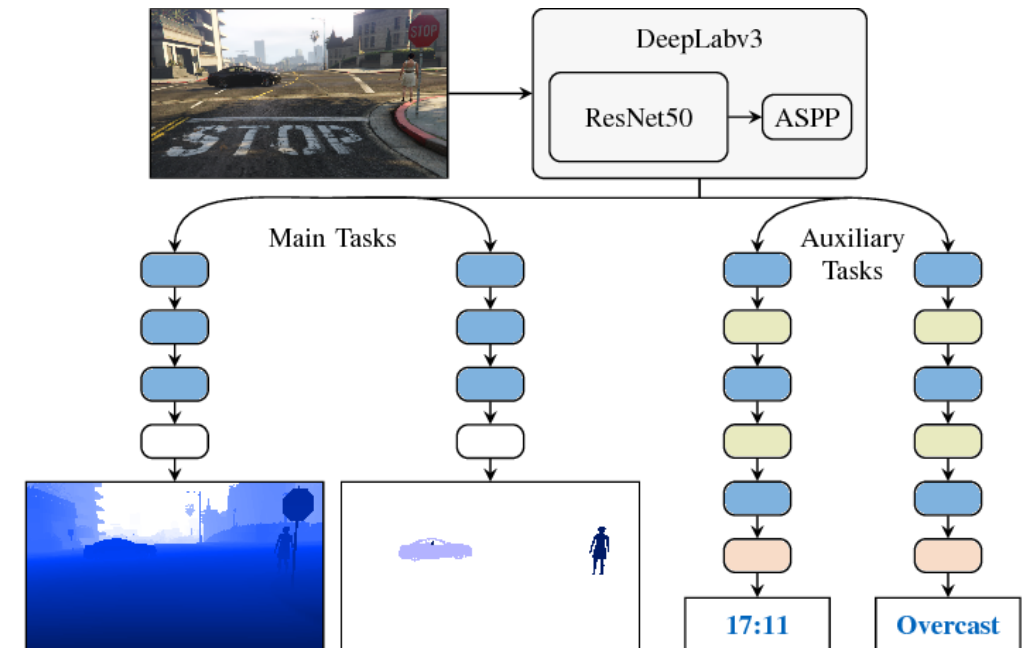
$$\min_{\theta} \sum_{i=1}^K \lambda_i \cdot L_i(f(x_i; \theta_{sh}, \theta_i), y_i),$$

*MTL optimisation*

*Task-specific loss weight coefficient*

## Auxiliary Learning

- If learning multiple related tasks together improves learning, why not learn extra (auxiliary) tasks along with the task(s) we care about?



**Legend:**

■ Convolutional   ■ Max pooling   ■ Fully connected   □ Bilinear interpolation

## MTL in the context of LLMs

- Traditionally, NLP models learnt multiple tasks simultaneously
  - *e.g. QA + text classification*
- Next-token prediction (used to train LLMs) is a single task but enables these models to excel in many tasks
  - This reduces the importance of MTL for NLP