



# **COMPSCI 761: ADVANCED TOPICS IN ARTIFICIAL INTELLIGENCE**

## **LOCAL SEARCH**

**Anna Trofimova, August 2022**

# TODAY

- CSP Definition recap
- CSP Backtracking
- CSP Inference



# SYSTEMATIC VS LOCAL SEARCH



- Systematic search strategies
  - Frontier maintains all unexpanded successors of expanded nodes.
  - Traverse the search space of a problem instance in a systematic manner.
  - Guarantee **completeness**, i.e., that eventually either a solution is found, or determine that the solution does not exist.
- Local search strategies (also called Iterative Improvement):
  - Frontier maintains some unexpanded successors of expanded nodes.
  - Start at some state and “move” from present location(s) to neighbouring location(s). The moves are determined by the present location(s).
  - Do not guaranteed completeness.

# WHAT DOES SEARCH RETURNS?

- In Systematic Search a solution can be a path or a state!
- In Local Search a state is a solution!
  - Since it doesn't maintain a path ("only a current state") it can't return a path
  - Can use a goal test or partial goal state
    - (but not a complete goal state – why?)

# LOCAL SEARCH ALGORITHMS

- In many optimisation problems, ***path*** is irrelevant; the goal state ***is*** the solution
- Then state space = set of “complete” configurations;
  - find ***configuration satisfying constraints***, e.g., n-queens problem; or, find ***optimal configuration***, e.g., travelling salesperson problem



- In such cases, can use ***iterative improvement*** algorithms: keep a single “current” state, try to improve it
- Constant space, suitable for online as well as offline search

# A GENERIC LOCAL SEARCH ALGORITHM

## The Generic LocalSearch Procedure

**INPUT:** A search problem with features  $X_1, \dots, X_k$

**OUTPUT:** A solution  $(x_1, \dots, x_k)$

Choose  $(x_1, \dots, x_k)$  using initialisation function

**while** stopping criterion is not met **do**

    Select a successor  $(y_1, \dots, y_k)$


    Update  $(x_1, \dots, x_k)$  to  $(y_1, \dots, y_k)$

**end while**

**return**  $(x_1, \dots, x_k)$


# EXAMPLE: N-QUEENS

- Iterative Improvement
  - assign all variables randomly in the beginning (thus violating several constraints),
  - change one variable at a time, trying to reduce the number of violations at each step.
  - Greedy Search with  $h$  = number of constraints violated



# TRAVELING SALESPERSON PROBLEM

Start with any complete tour, perform pairwise exchanges



Variants of this approach get within 1% of optimal very quickly with thousands of cities




# STOPPING CRITERIA

- **Heuristic value** – You can use the heuristic value as a termination test in some problems. (e.g., n-queens)
- **Goal test** – You can use a goal test (like in Systematic Search). This will work in Sudoku or in N-queens.
  - You cannot have a “complete goal state” because then the problem will already be solved
- **Number of runs** – Sometimes you just need to specify when to stop. For instance, in the Traveling Salesperson Problem.

# LOCAL SEARCH HEURISTICS

- A local search heuristic,  $h: S \rightarrow R$ .
- This allows the algorithm to make moves maximising or minimising the value.
- It does not have to be admissible or consistent.
- It can be involved in the goal test or termination condition, or not.



# ITERATIVE BEST IMPROVEMENT

**Iterative best improvement (IBI)** is a local search strategy that always selects a successor that minimises (or maximises) a heuristic function which is typically the loss (or gain), i.e., it performs **greedy choice**.

Successor A

$$A' = \operatorname{argmin} h(A') \quad \text{or} \quad A' = \operatorname{argmax} h(A')$$

where  $h:S \rightarrow \mathbb{R}$  is a heuristic function

Characteristic:

Tie break is typically random.



# HILL CLIMBING OR GREEDY DESCENT

For minimisation the method is called greedy descent (similar to gradient descent – we will see later)




For maximisation the method is called hill-climbing



# HILL-CLIMBING

*"Like climbing Everest in thick fog with amnesia"*

- Simple, general idea:
  - Start wherever
  - Repeat: move to the **best** neighbouring state
  - If no neighbours better than current, quit






# HILL-CLIMBING SEARCH

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
```


# HILL-CLIMBING WITH 8-QUEENS

- Randomly generated 8-queens starting states...
  - Gets stuck 86% of the time
  - Solves only 14% of problem instances
- However
  - 4 steps on average when it succeeds
  - 3 steps on average when it gets stuck
  - (for a state space with  $8^8 \approx 17$  million states)



# HILL-CLIMBING DIFFICULTIES

Problem: depending on initial state, can get stuck in local maxima





Note: these difficulties apply to all local search algorithms, and usually become much worse as the search space becomes higher dimensional




# HILL-CLIMBING PROBLEMS

- Local maxima
- Plateaus
- Diagonal ridges




# INVERTED VIEW FOR GREEDY DESCENT



- Sometimes, have to go sideways or even backwards in order to make progress towards the actual solution.


# HILL-CLIMBING WITH SIDE-WAY MOVES

- **Hill-Climbing with sideways moves** uses the same procedures as hill-climbing with the following differences:
  - When no uphill successors (ones that increase loss strictly) are present, move “sideways” (successors that have the same value).
  - Introduce a parameter  $m$  such that only  $\leq m$  such sideways moves can be performed consecutively.
- Loop on at maxima
  - Limit number of consecutive sideways steps
- Escape from shoulders



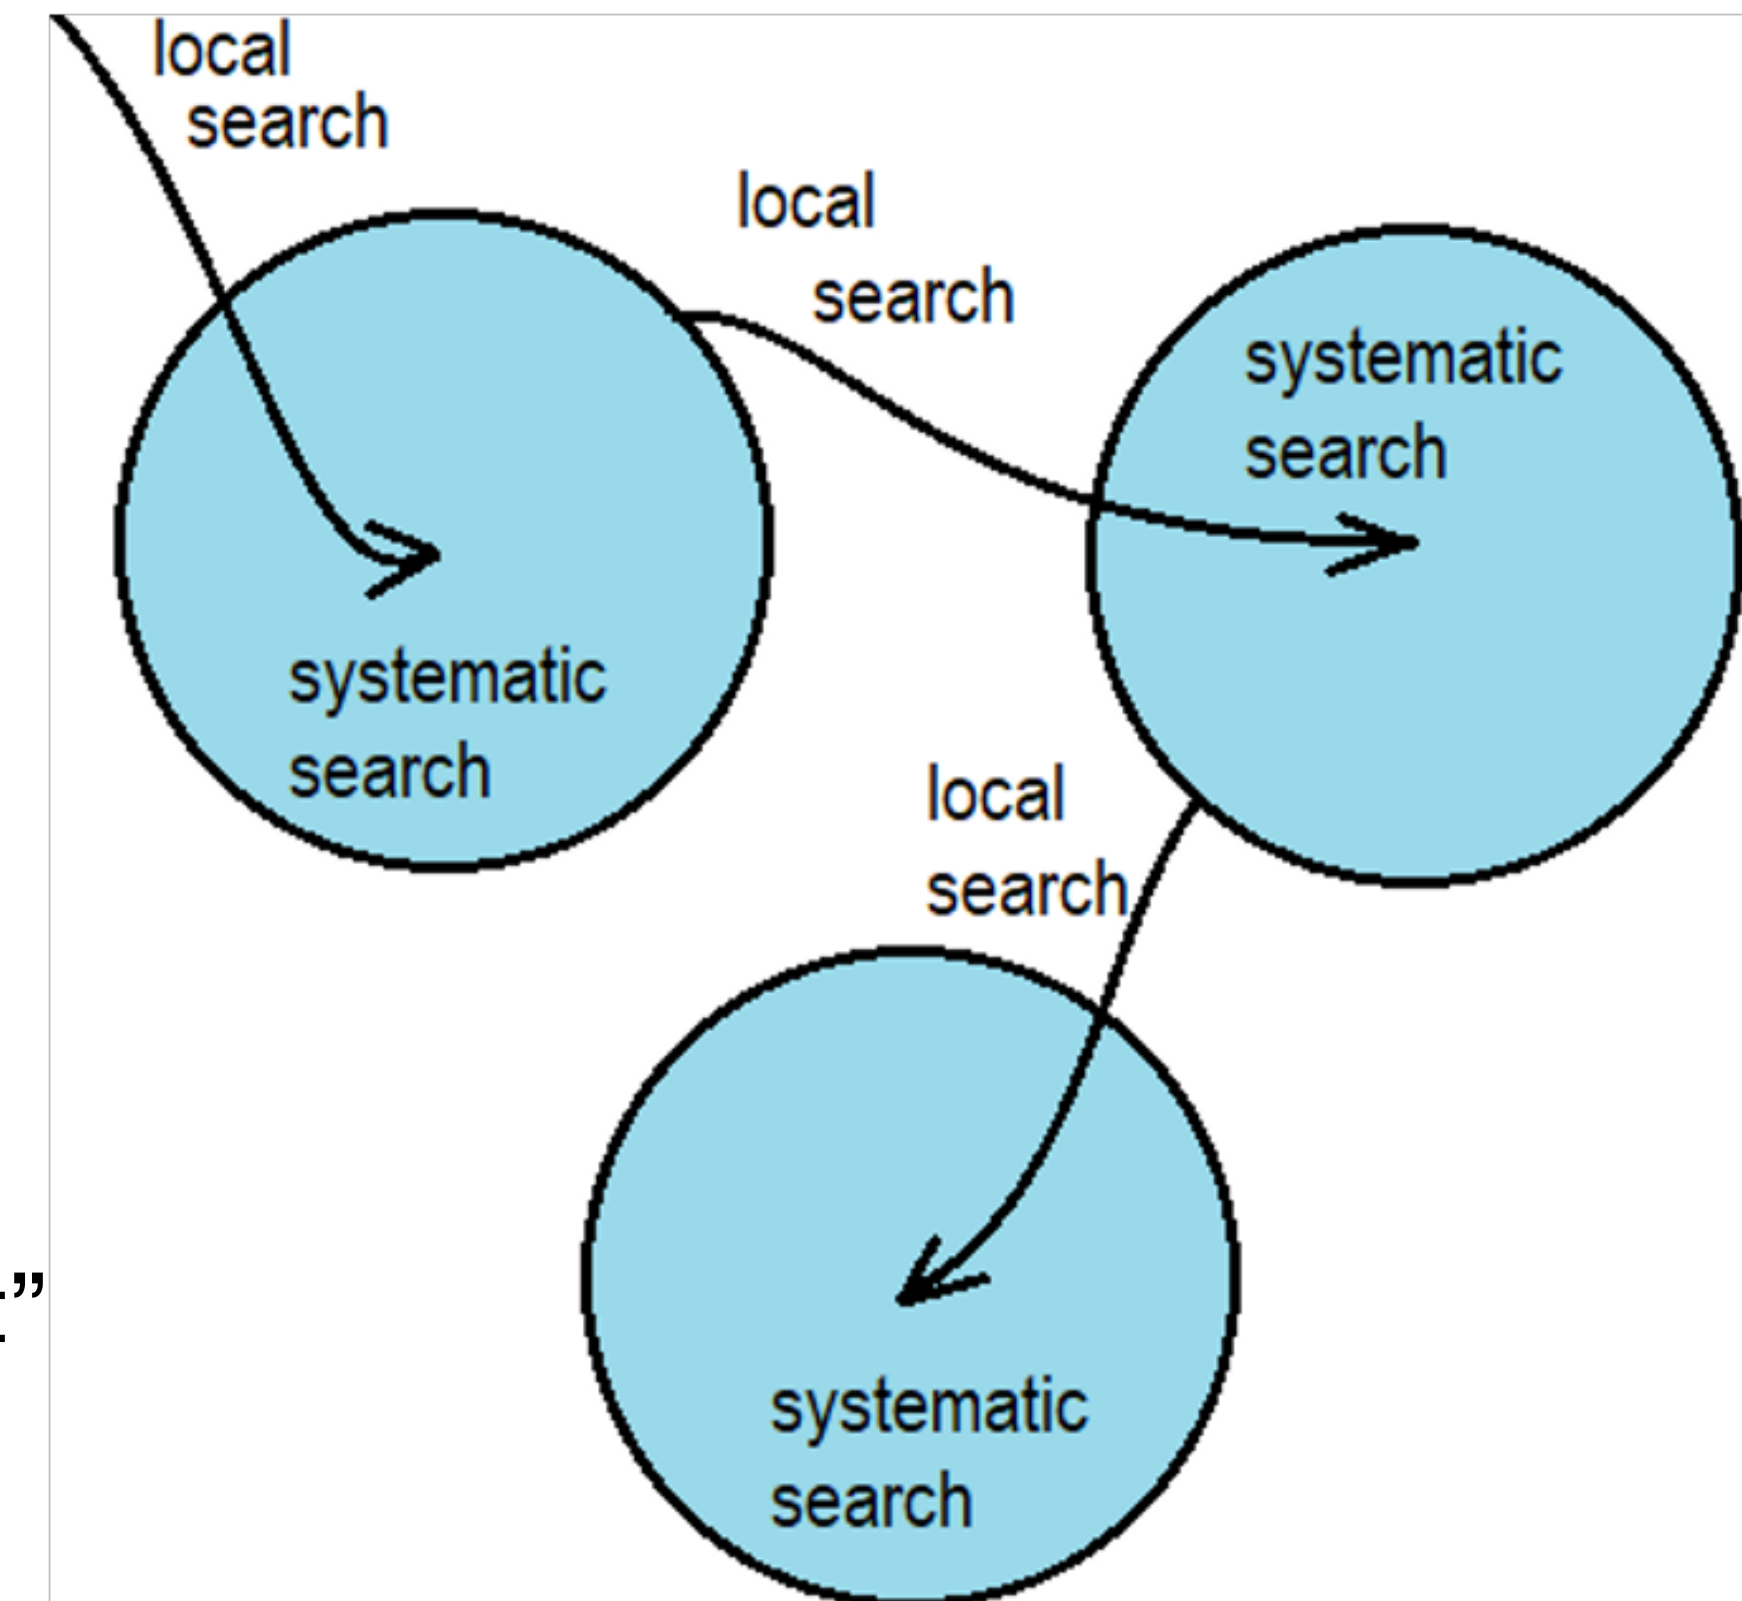
# HILL-CLIMBING WITH SIDE WAY MOVES ON THE 8-QUEENS PROBLEM

- No sideways moves:
  - Succeeds w/ prob.  $p = 0.14$
  - Average number of moves per trial:
    - 4 when succeeding, 3 when getting stuck
  - Expected total number of moves needed:
    - $3(1-p)/p + 4 \approx 22$  moves
- Allowing 100 sideways moves:
  - Succeeds w/ prob.  $p = 0.94$
  - Average number of moves per trial:
    - 21 when succeeding, 65 when getting stuck
  - Expected total number of moves needed:
    - $65(1-p)/p + 21 \approx 25$  moves



# ENFORCED HILL-CLIMBING

- Perform breadth first search from a local optima
  - to find the next state with better h function
  - good At Escaping Shoulders/local Optima
- Typically,
  - prolonged periods of exhaustive search
  - bridged by relatively quick periods of hill-climbing
  - Sometimes called “Variable Neighbourhood Descent”
- Middle ground b/w local and systematic search



# TABU SEARCH

- Prevent returning quickly to the same state
  - Keep fixed length queue (“tabu list”)
  - Add most recent state to queue; drop oldest
  - Never make the step that is currently tabu’ed
- Properties:
  - As the size of the tabu list grows, hill-climbing will asymptotically become “non-redundant” (won’t look at the same state twice)
  - In practice, a reasonable sized tabu list (say 100 or so) improves the performance of hill climbing in many problems
  - The list can allow the algorithm to escape some of the local minima



# SUMMARY

Hill Climbing – chose best child

Hill Climbing with Sideways Moves – chose best child or equal child

Tabu Search – keep list of nodes you have been to and don't go back

Enforced Hill Climbing – use breadth first search until you find a “better” node