



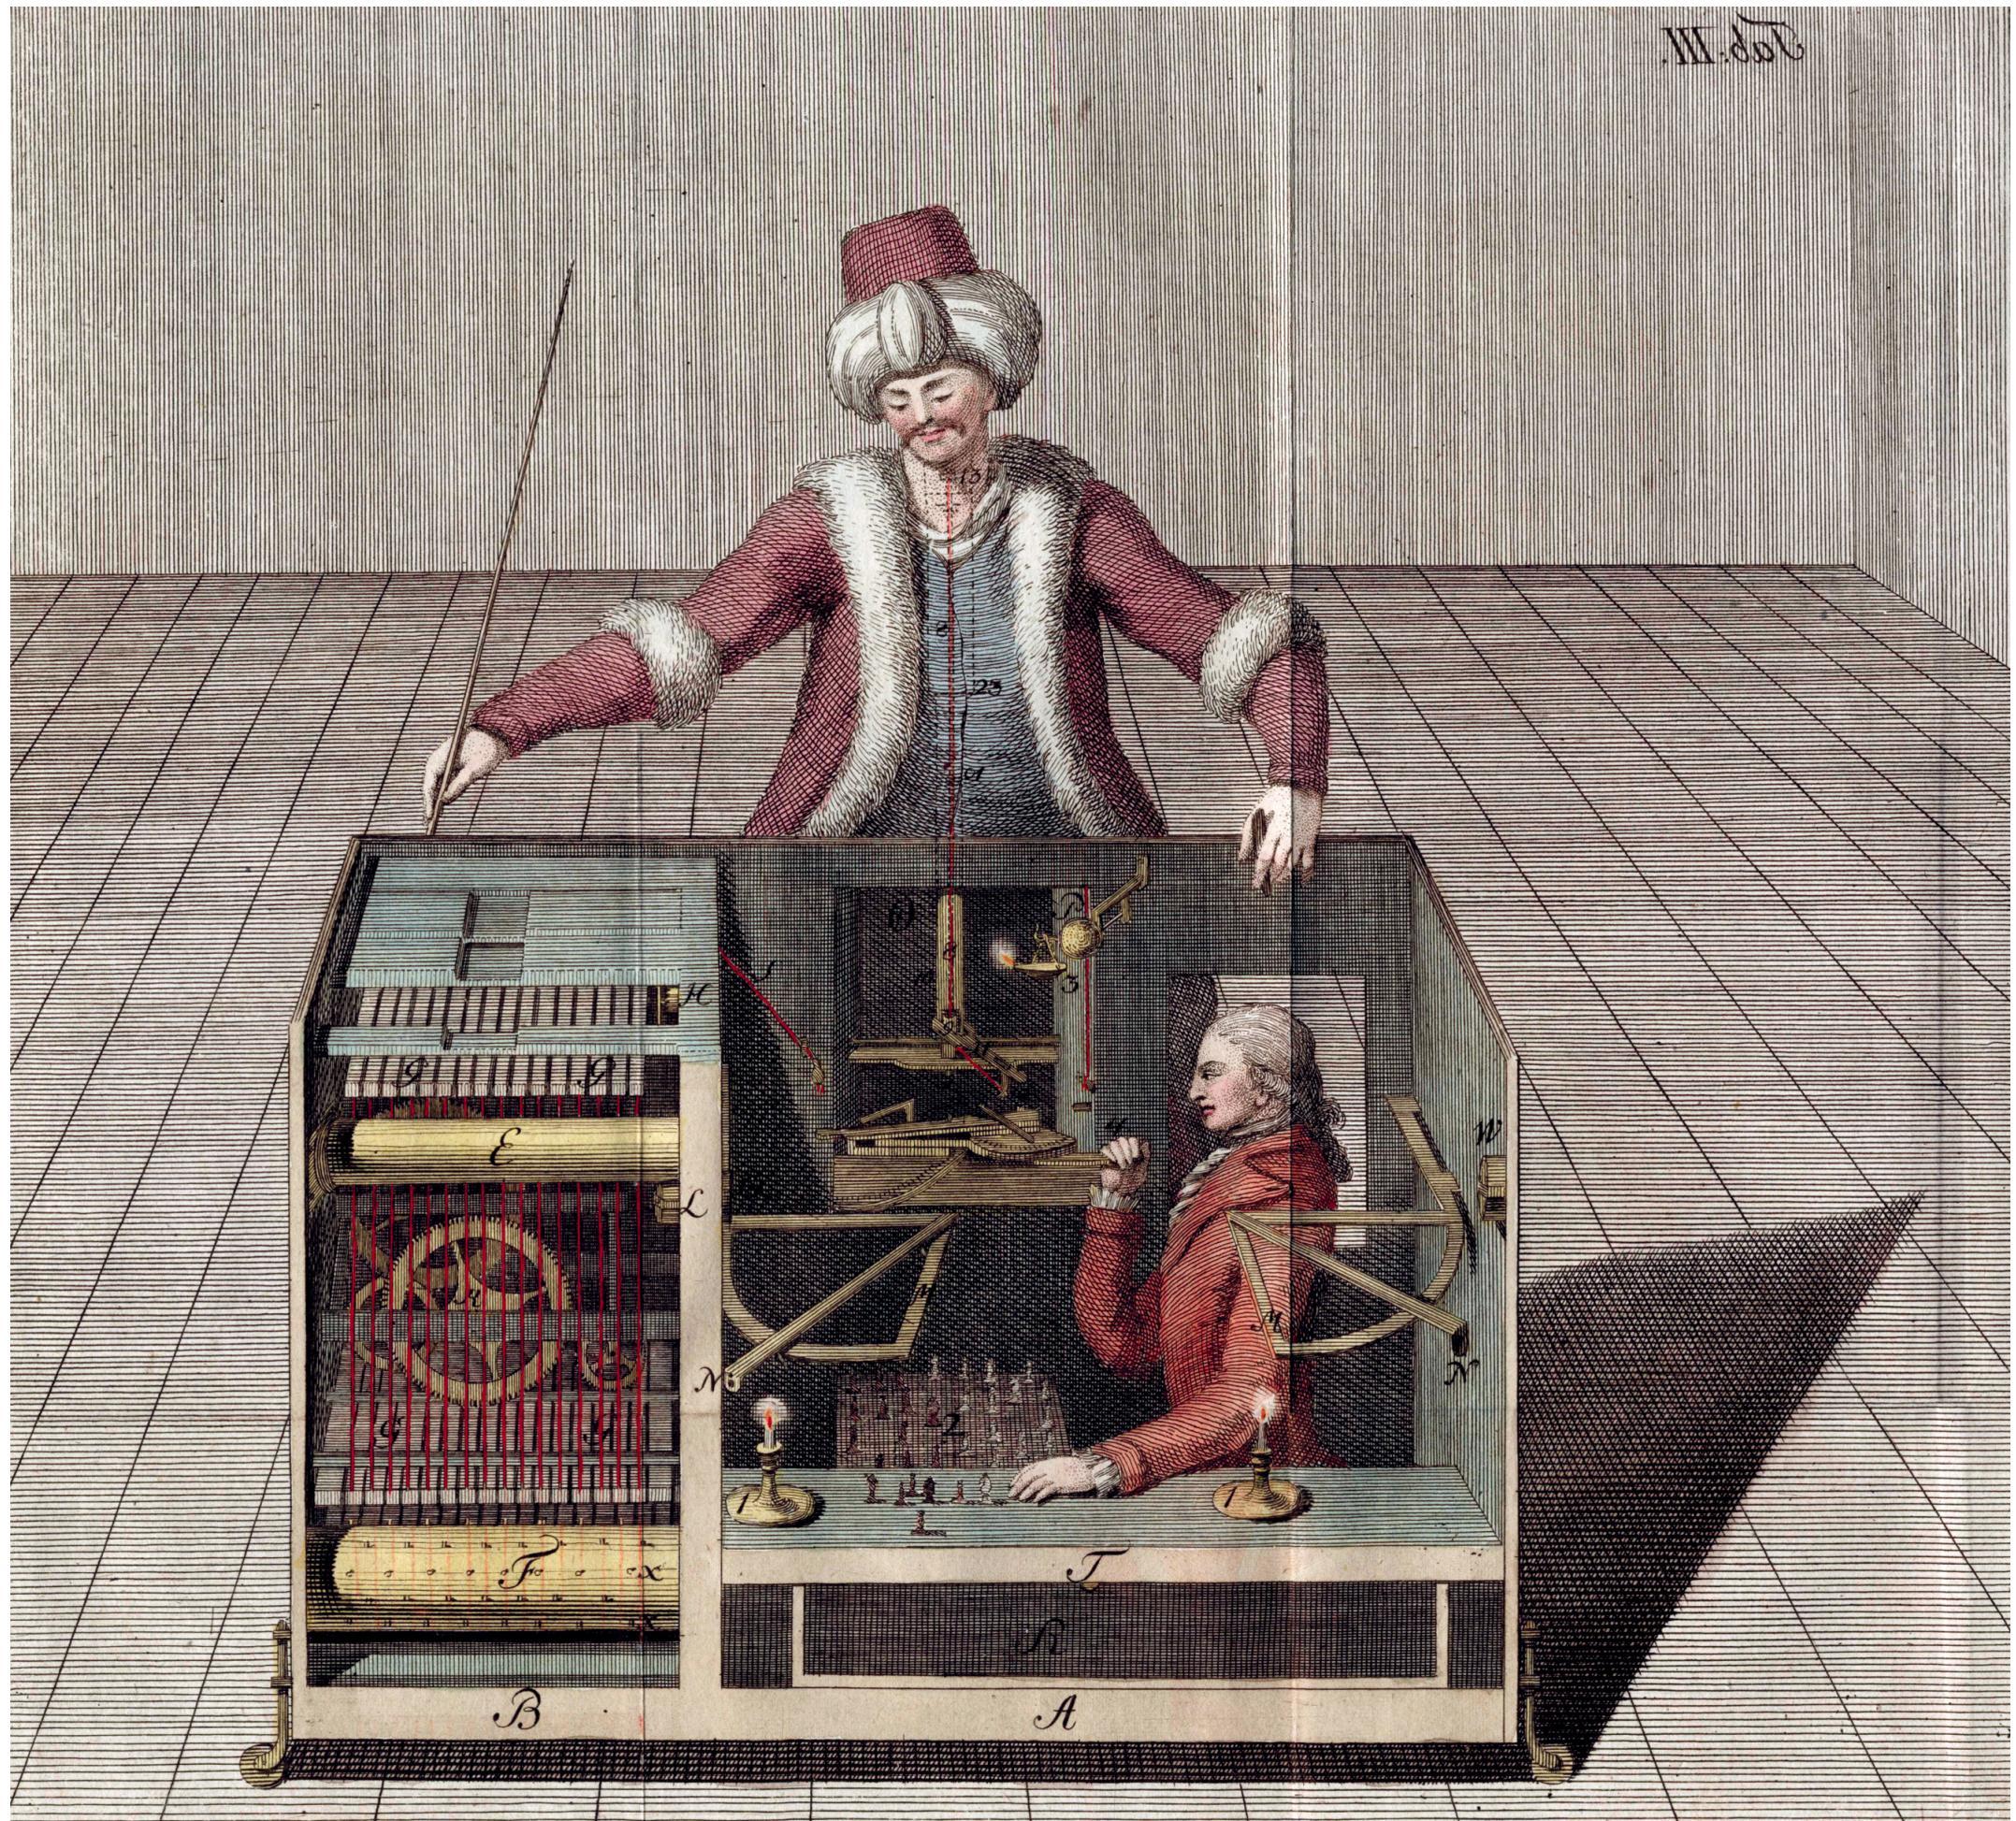
# **COMPSCI 761: ADVANCED TOPICS IN ARTIFICIAL INTELLIGENCE GAMES**

**Anna Trofimova, August 2022**

# TODAY

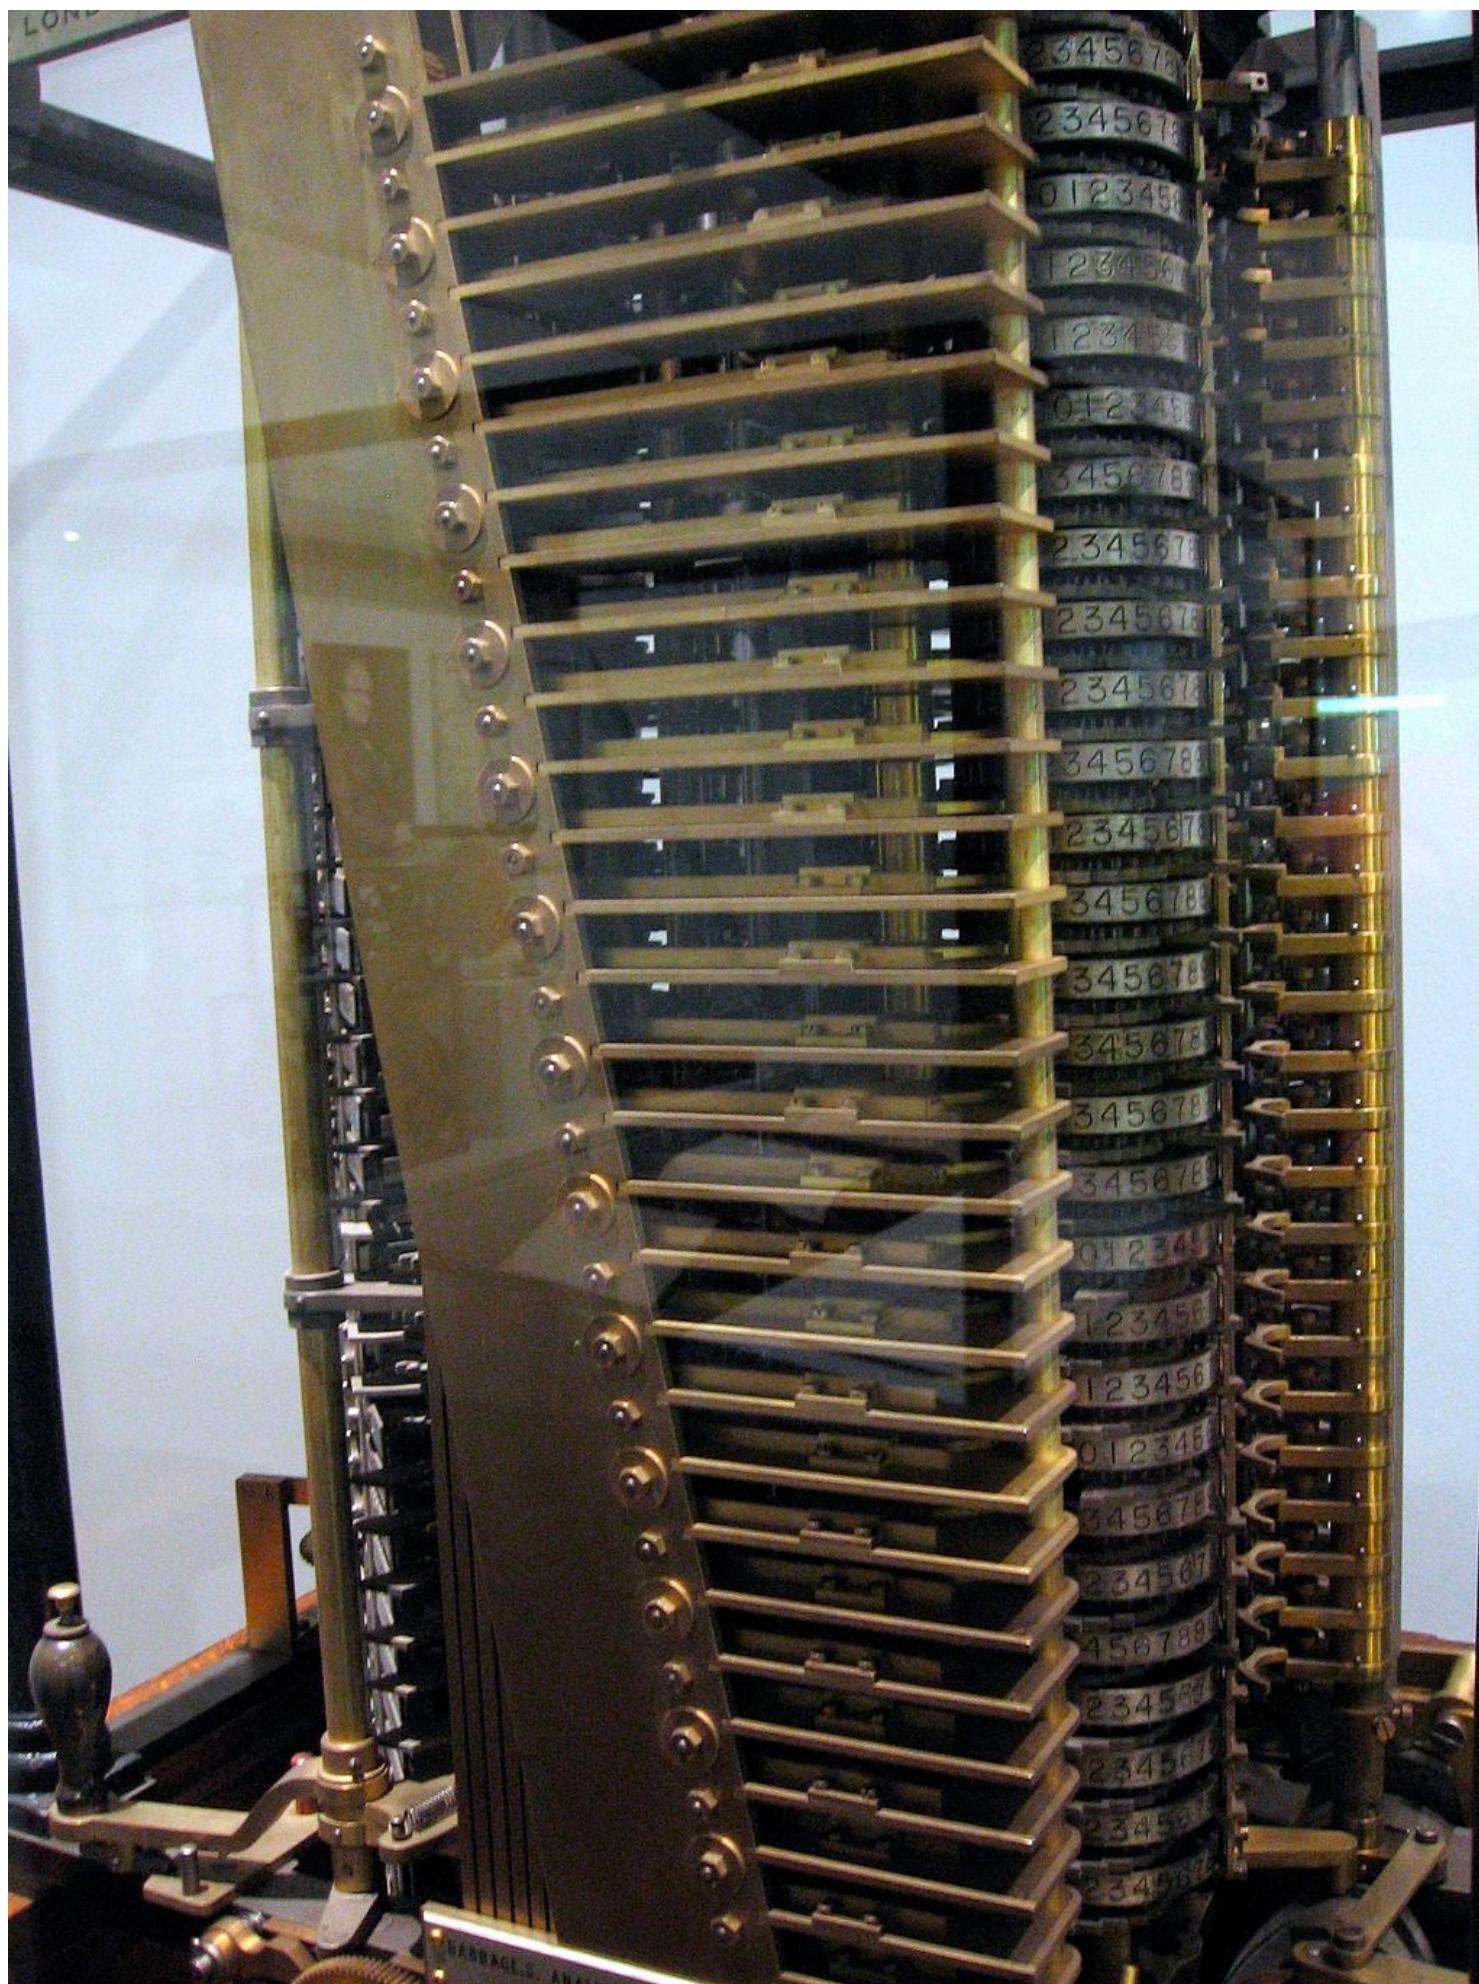
- Games history and motivation
- Minimax
- $\alpha$ - $\beta$  pruning

# MECHANICAL TURK

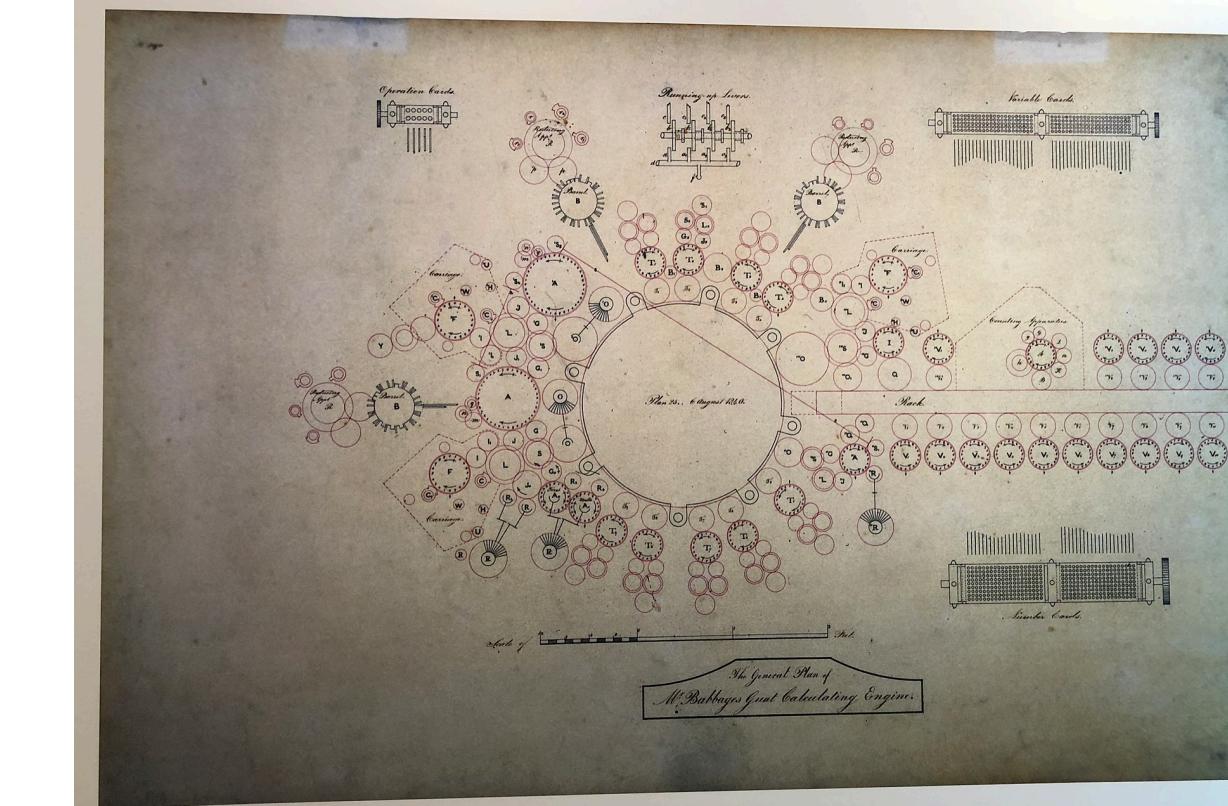


- Constructed and unveiled in 1770 by Wolfgang von Kempelen for Maria Theresa of Austria
- A mechanical illusion - a machine with a human player inside
- Was destroyed by fire in 1854 in Chinese Museum of Charles Willson Peale, US.
- How did the human operator knew the move of the opponent?

# THE ANALYTICAL ENGINE

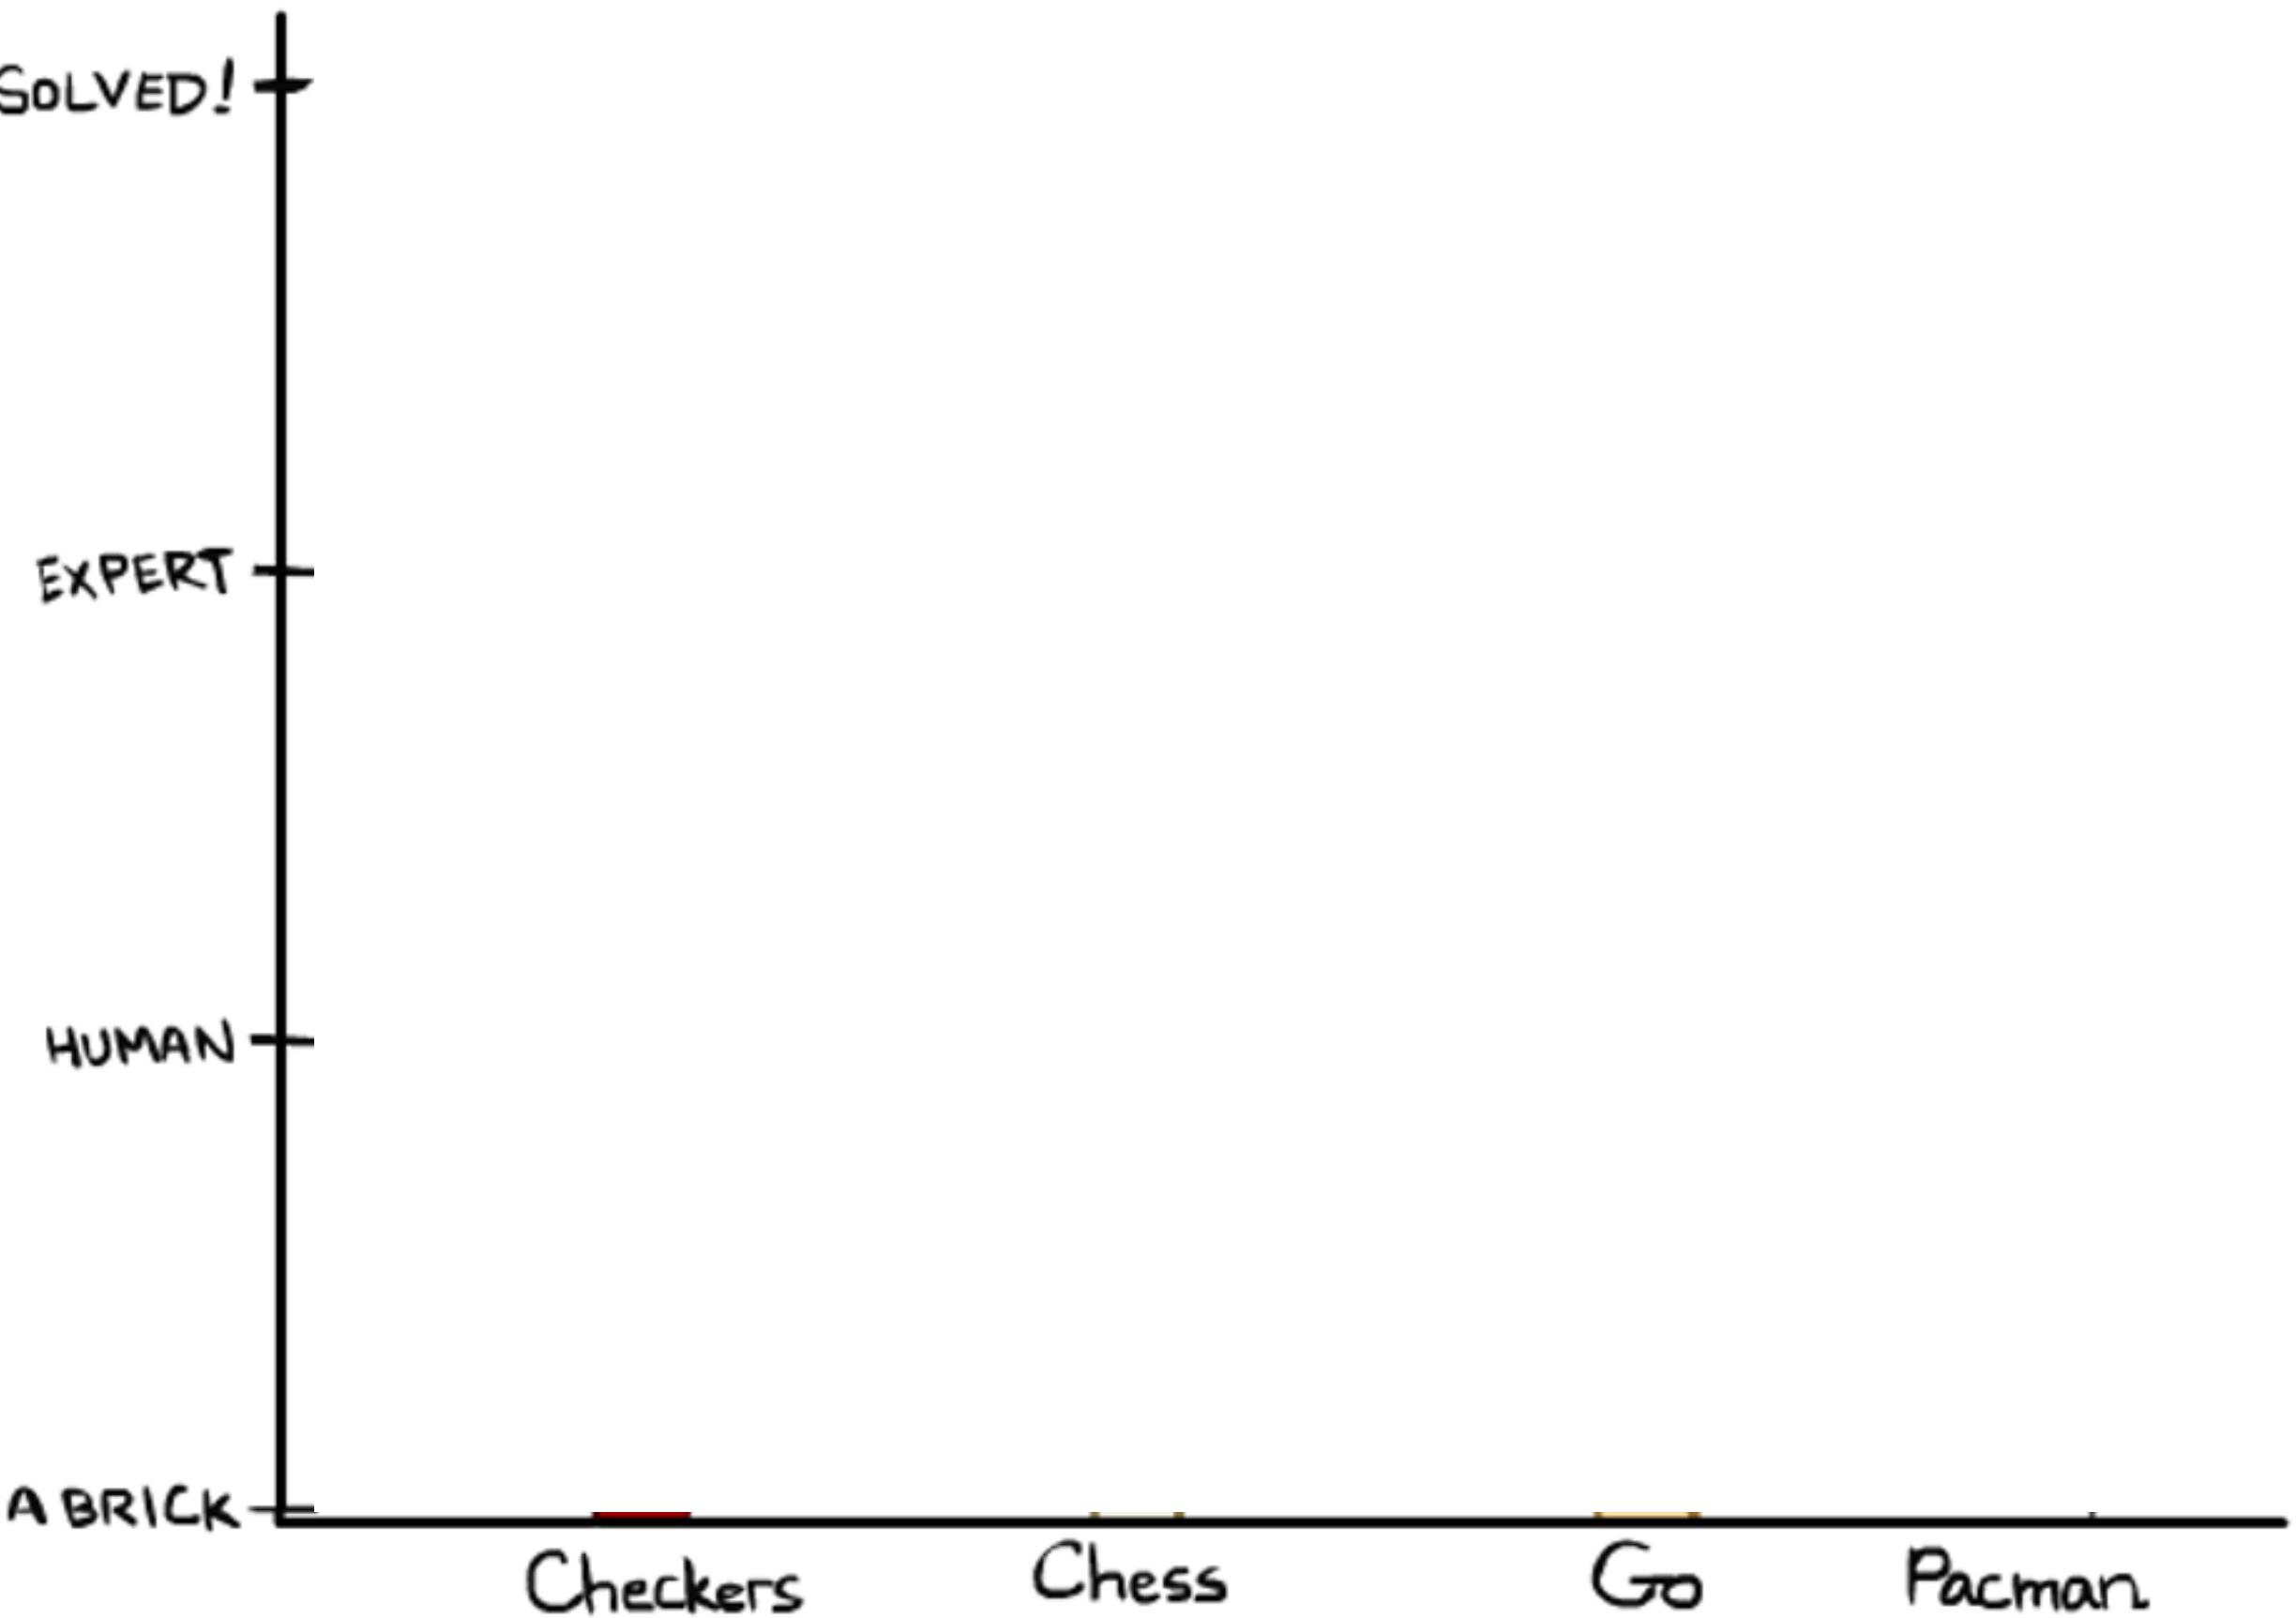


- Babbage (1791 - 1871) is a computer pioneer, mathematician, engineer, inventor, philosopher
  - was working on The Analytical Engine, a mechanical general-purpose computer with arithmetic logic units and control flow
  - documented the tic-tac-toe automaton in his 1864 autobiography, *Passages from the Life of a Philosopher*
    - But he never build the machine as it wasn't profitable



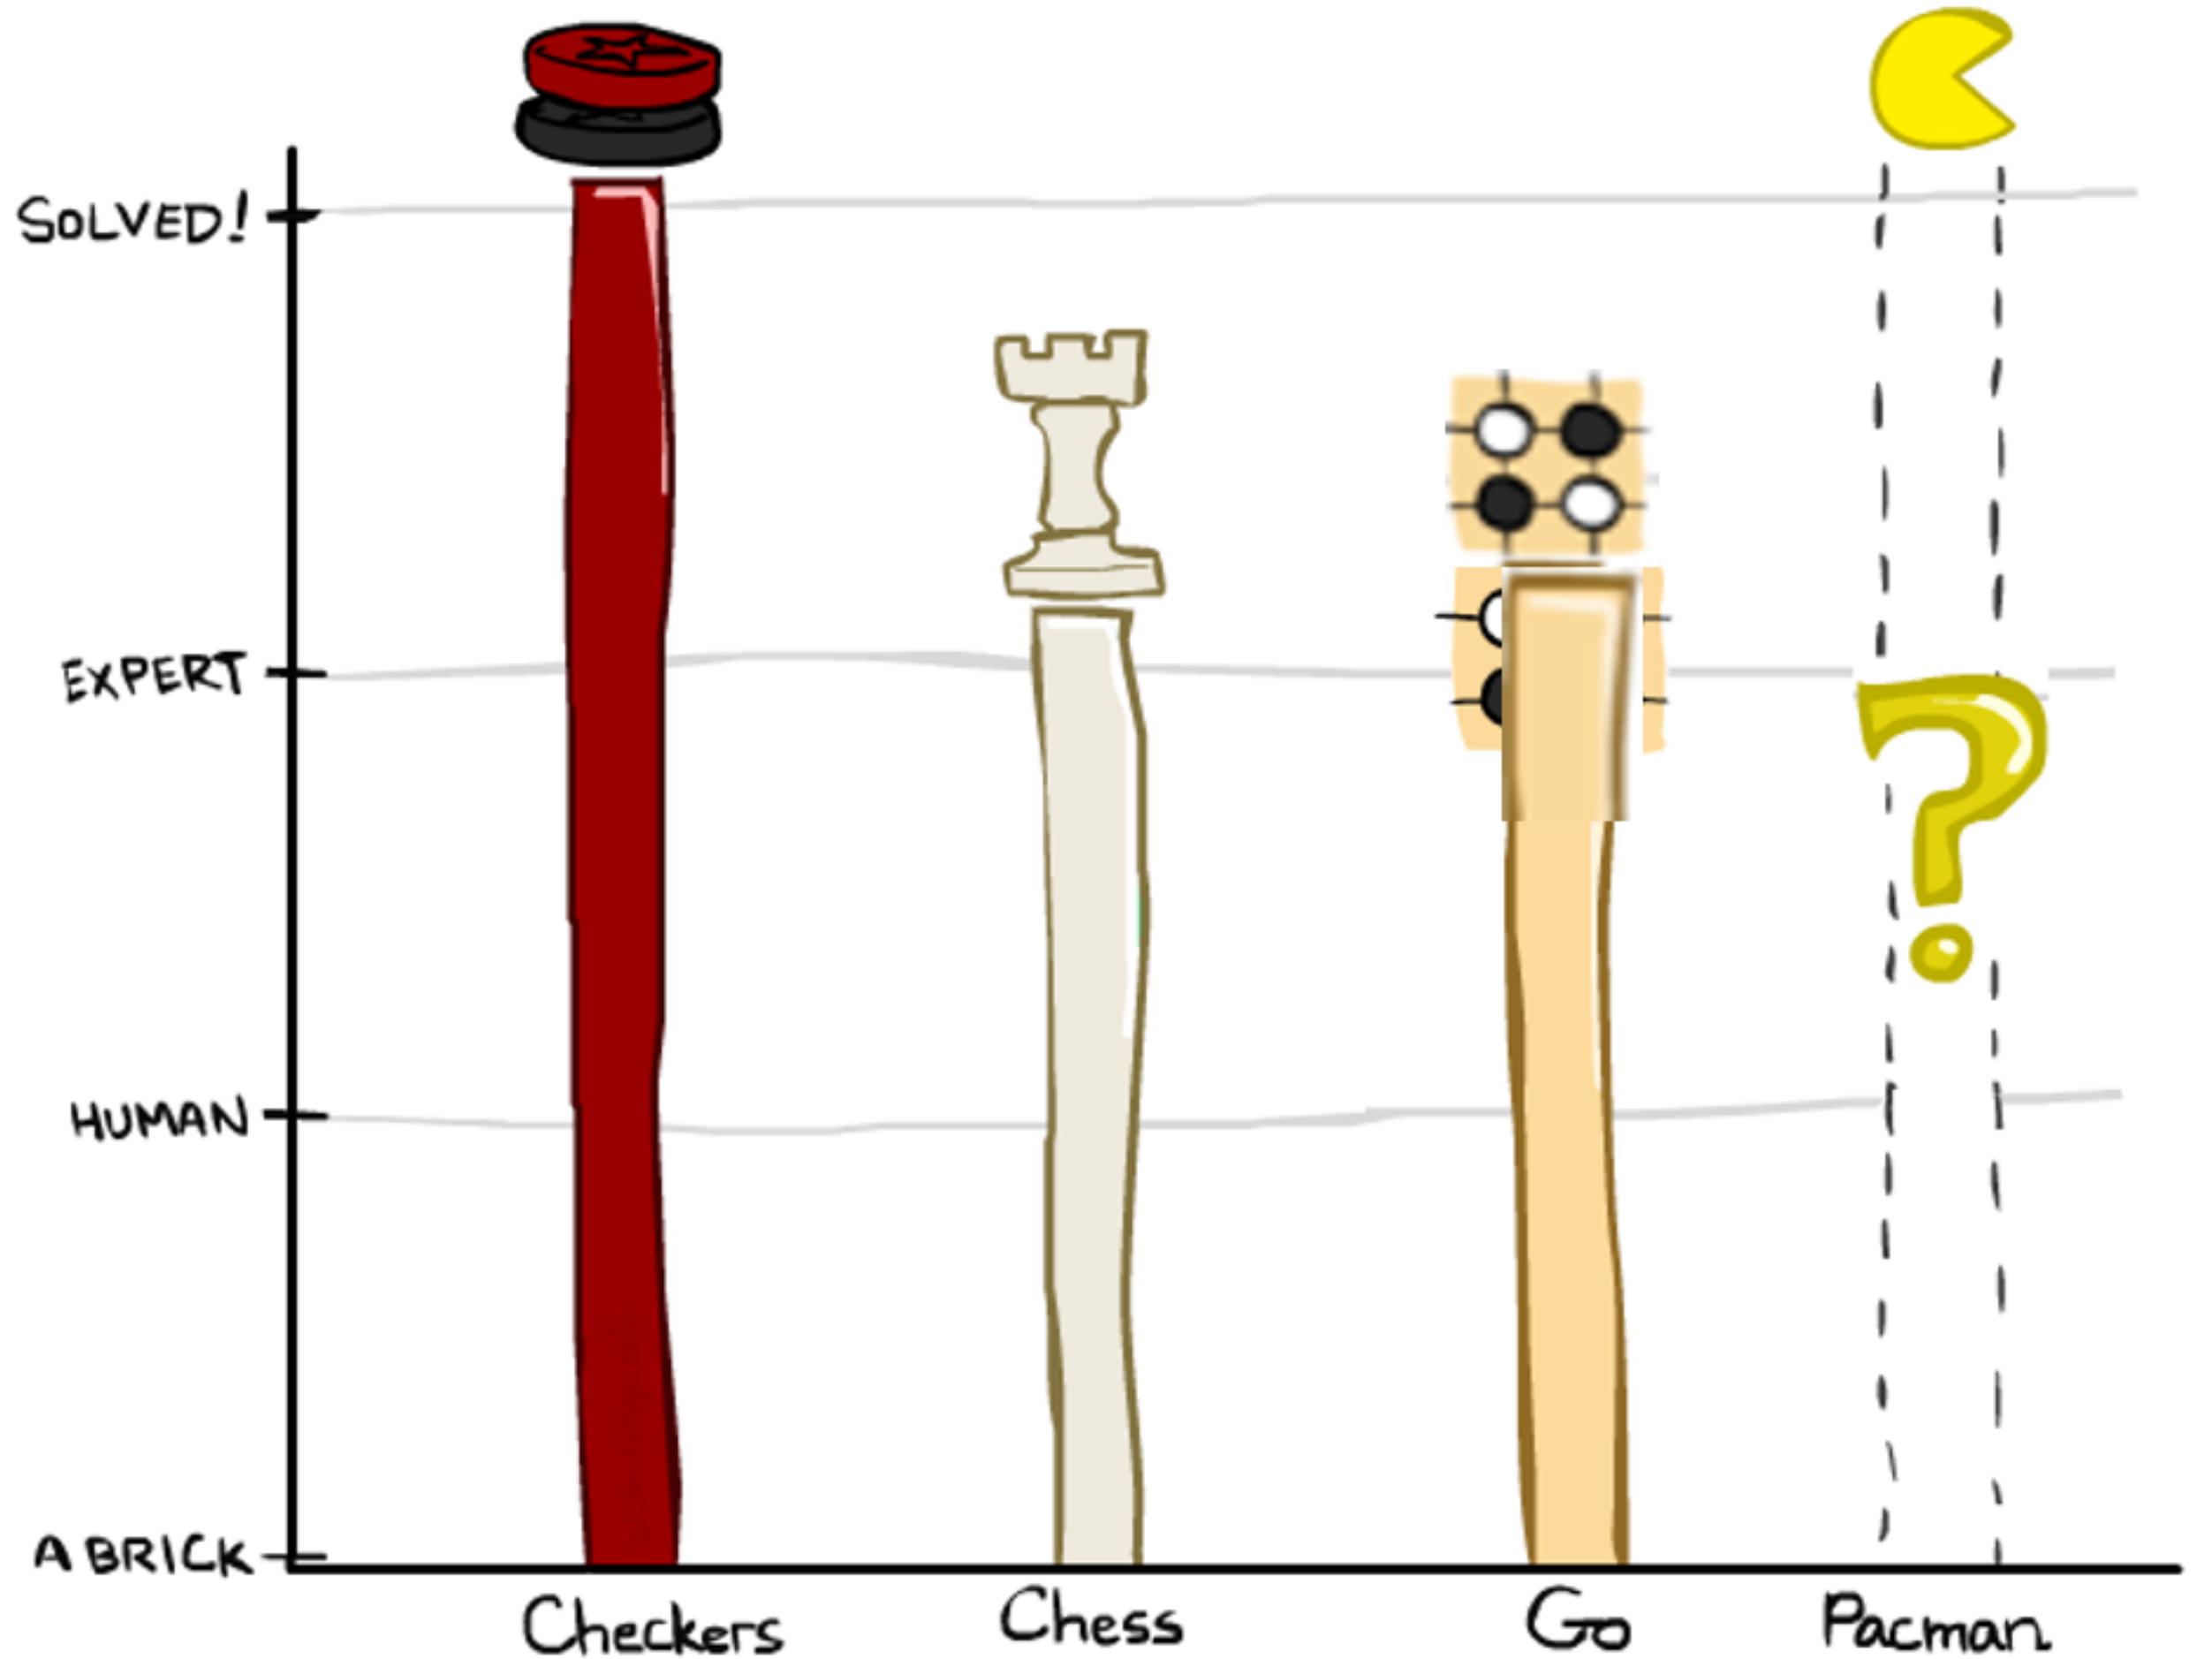
# A BRIEF HISTORY

- Checkers:
  - 1950: First computer player.
  - 1959: Samuel's self-taught program.
  - 1994: First computer world champion: Chinook defeats Tinsley
  - 2007: Checkers solved! Endgame database of 39 trillion states
- Chess:
  - 1945-1960: Zuse, Wiener, Shannon, Turing, Newell & Simon, McCart
  - 1960s onward: gradual improvement under “standard model”
  - 1997: Deep Blue defeats human champion Gary Kasparov
  - 2021: Stockfish rating 3551 (vs 2870 for Magnus Carlsen).
- Go:
  - 1968: Zobrist's program plays legal Go, barely ( $b>300!$ )
  - 1968-2005: various ad hoc approaches tried, novice level
  - 2005-2014: Monte Carlo tree search -> strong amateur
  - 2016-2017: AlphaGo defeats human world champions



# A BRIEF HISTORY

- Checkers:
  - 1950: First computer player.
  - 1959: Samuel's self-taught program.
  - 1994: First computer world champion: Chinook defeats Tinsley
  - 2007: Checkers solved! Endgame database of 39 trillion states
- Chess:
  - 1945-1960: Zuse, Wiener, Shannon, Turing, Newell & Simon, McCarty
  - 1960s onward: gradual improvement under "standard model"
  - 1997: Deep Blue defeats human champion Gary Kasparov
  - 2021: Stockfish rating 3551 (vs 2870 for Magnus Carlsen).
- Go:
  - 1968: Zobrist's program plays legal Go, barely ( $b>300!$ )
  - 1968-2005: various ad hoc approaches tried, novice level
  - 2005-2014: Monte Carlo tree search -> strong amateur
  - 2016-2017: AlphaGo defeats human world champions



# PACMAN - BEHAVIOUR FROM COMPUTATION



# GAMES: MOTIVATION

- Games are a form of multi-agent environment
  - What do other agents do and how do they affect our success?
  - Cooperative vs. competitive multi-agent environments.
  - Competitive multi-agent environments give rise to adversarial search a.k.a. *games*
- Why study games?
  - Games are fun!
  - Historical role in AI
  - Studying games teaches us how to deal with other agents trying to foil our plans
  - *Huge* state spaces – Games are *hard*!
  - Nice, clean environment with clear criteria for success

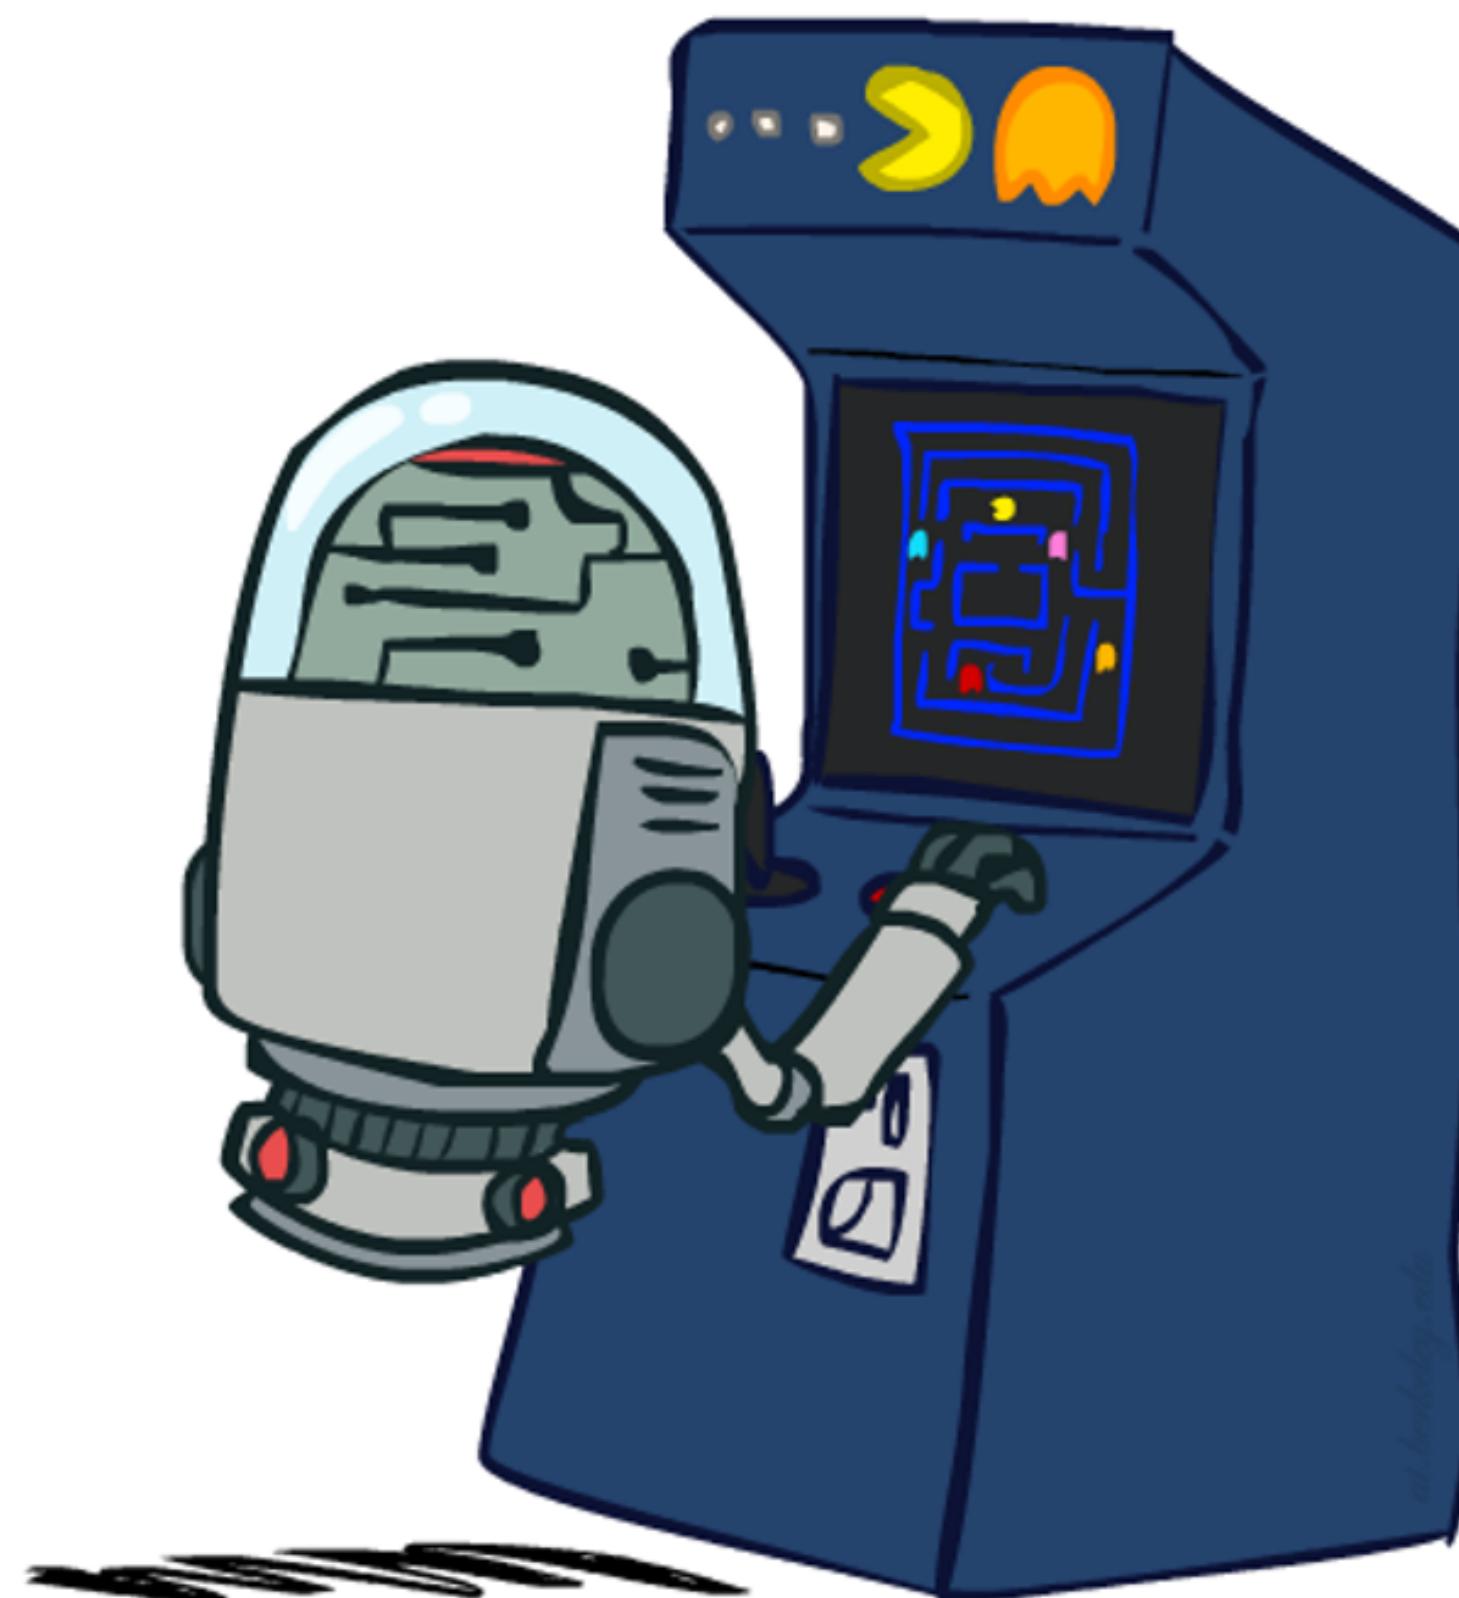
# TYPES OF GAMES

- Game = task environment with > 1 agent
- Axes:
  - Deterministic or stochastic?
  - Perfect information (fully observable)?
  - One, two, or more players?
  - Turn-taking or simultaneous?
  - Zero sum?
- Want algorithms for calculating a ***contingent plan*** (a.k.a. **strategy or policy**) which recommends a move for every possible eventuality

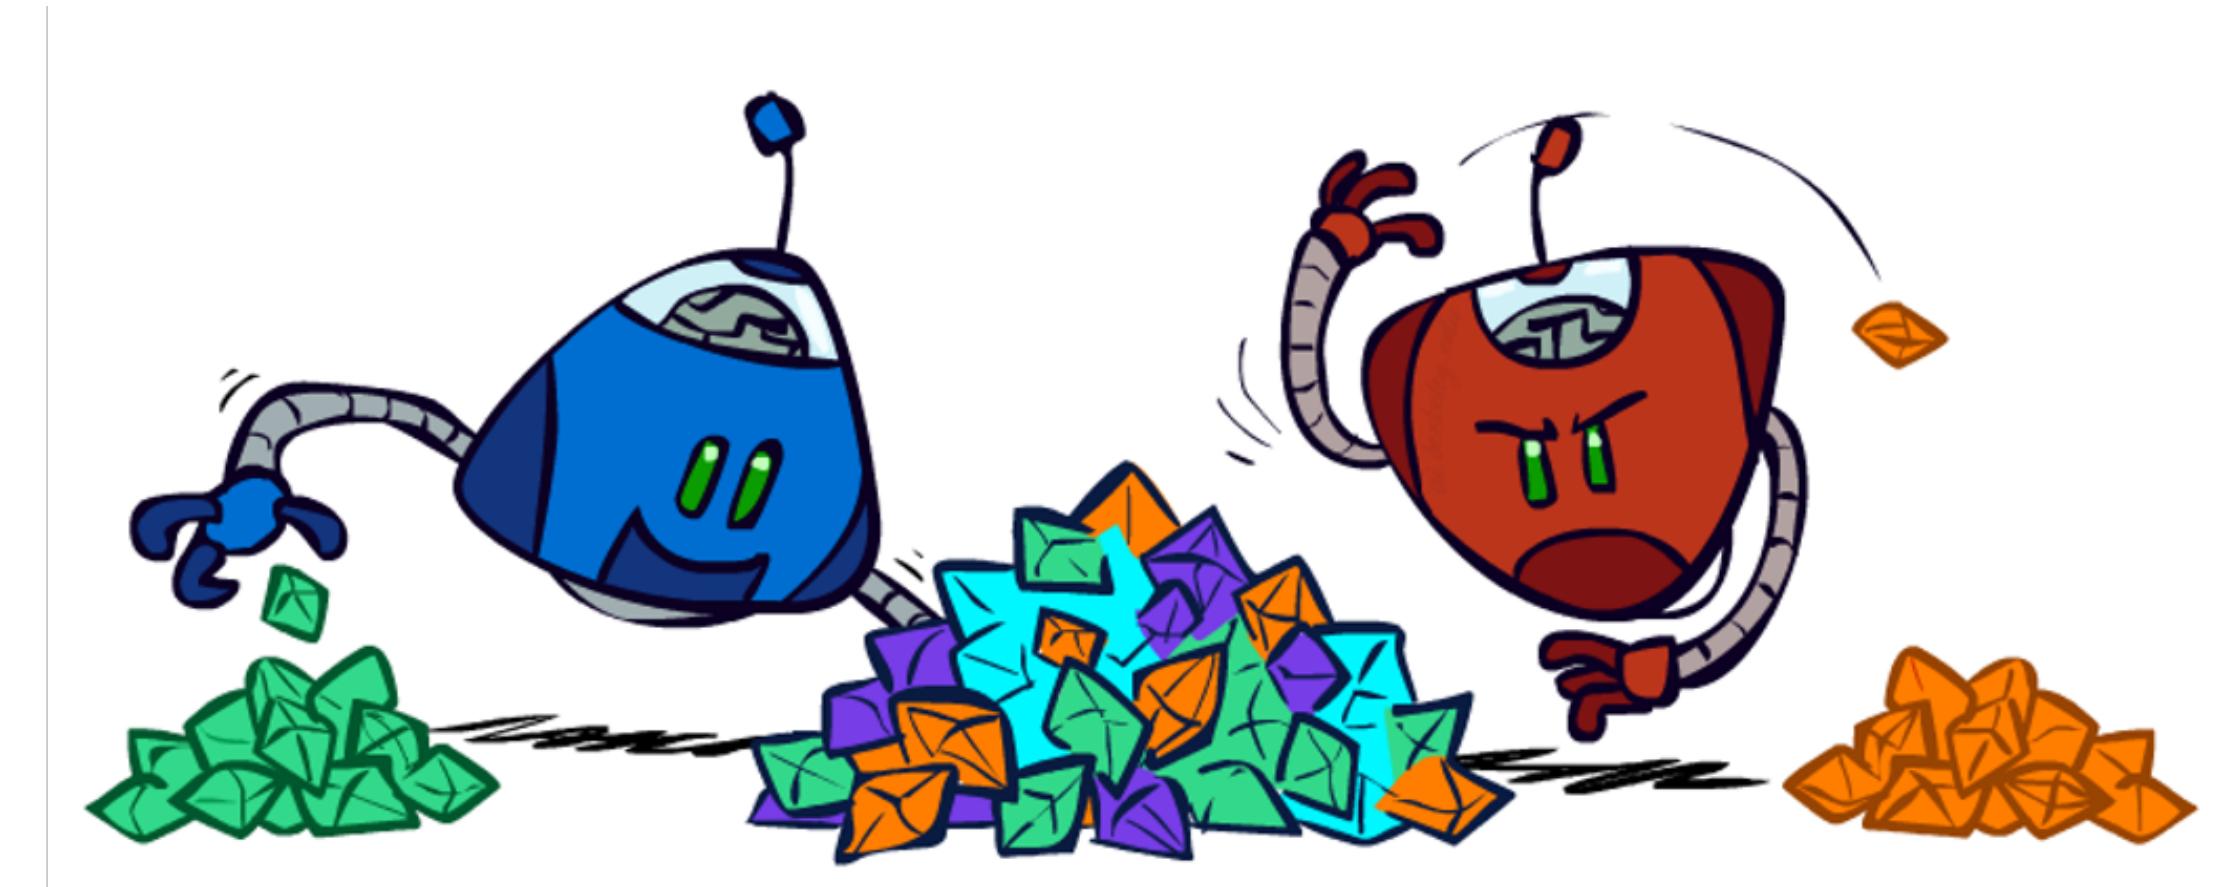
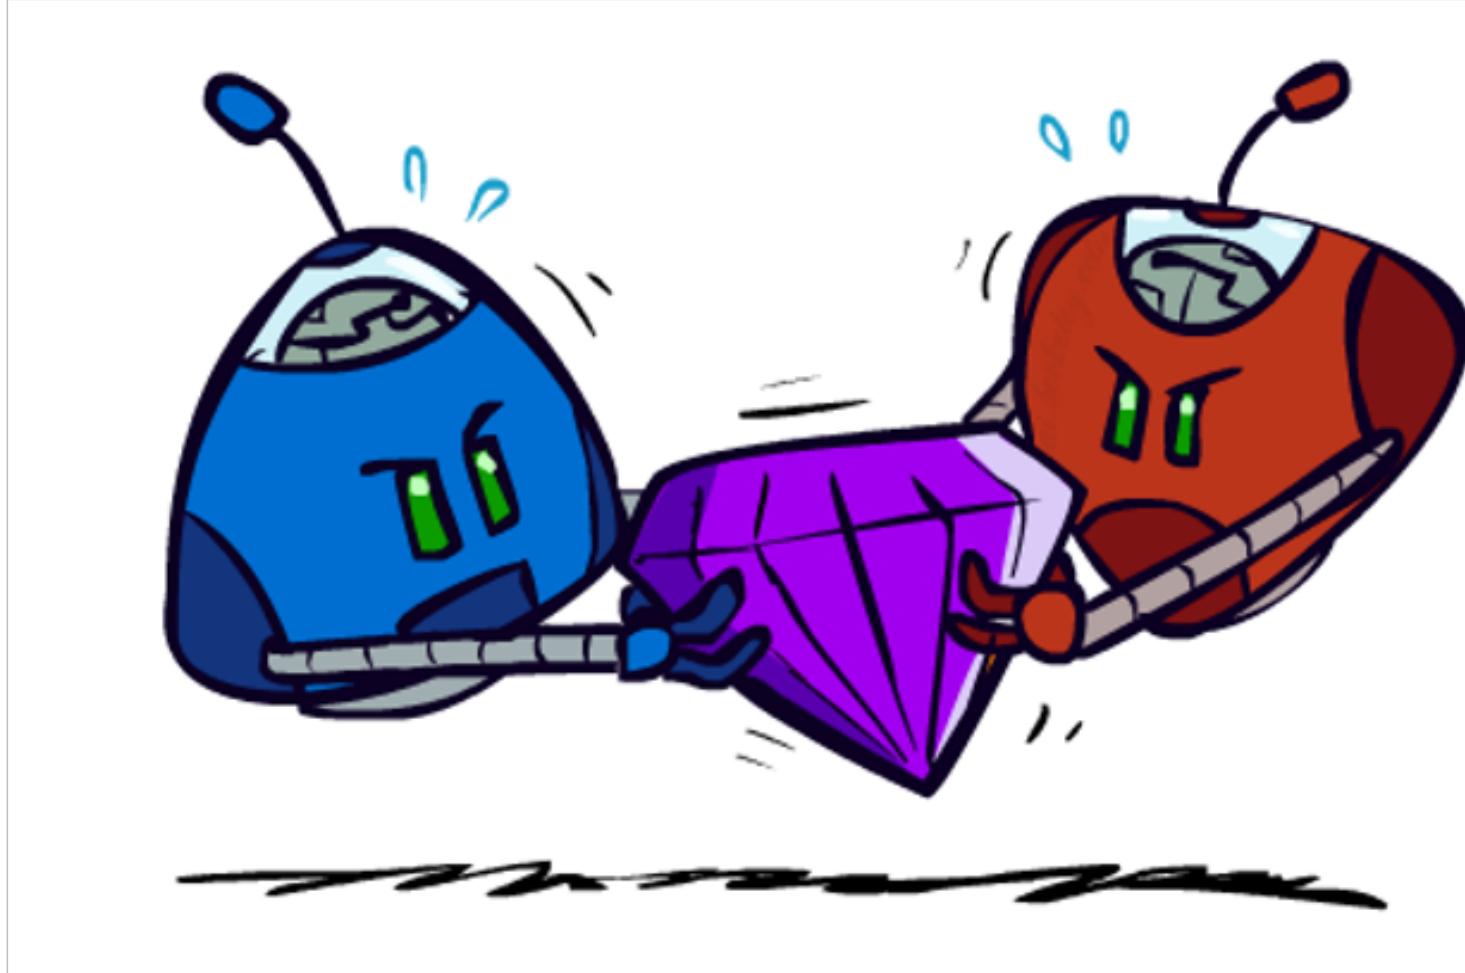


# DETERMINISTIC GAMES

- Many possible formalisations, one is:
  - States:  $S$  (start at  $s_0$ )
  - Players:  $P=\{1\dots N\}$  (usually take turns)
  - Actions:  $A$  (may depend on player / state)
  - Transition Function:  $S \times A \rightarrow S$
  - Terminal Test:  $S \rightarrow \{t, f\}$
  - Terminal Utilities:  $S \times P \rightarrow R$
- Solution for a player is a **policy**:  $S \rightarrow A$



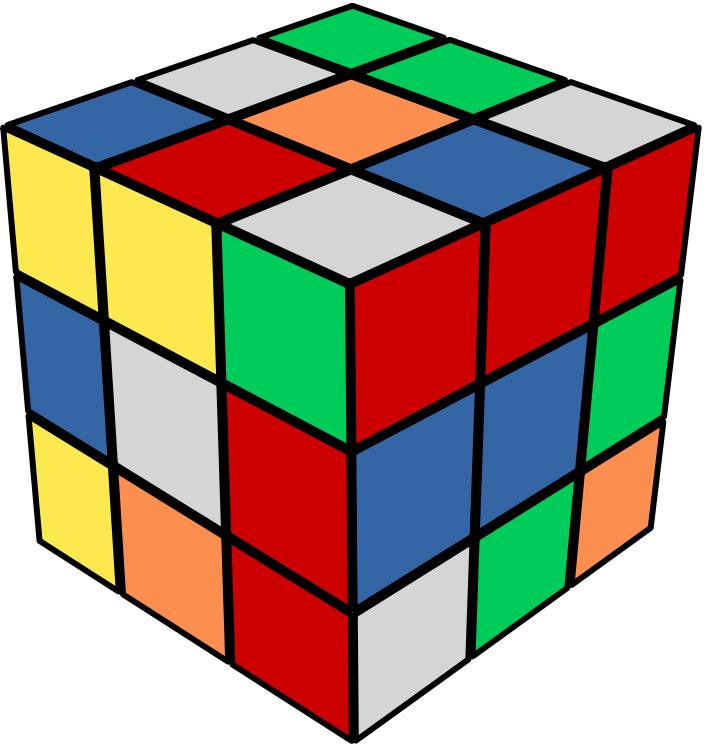
# ZERO-SUM GAMES



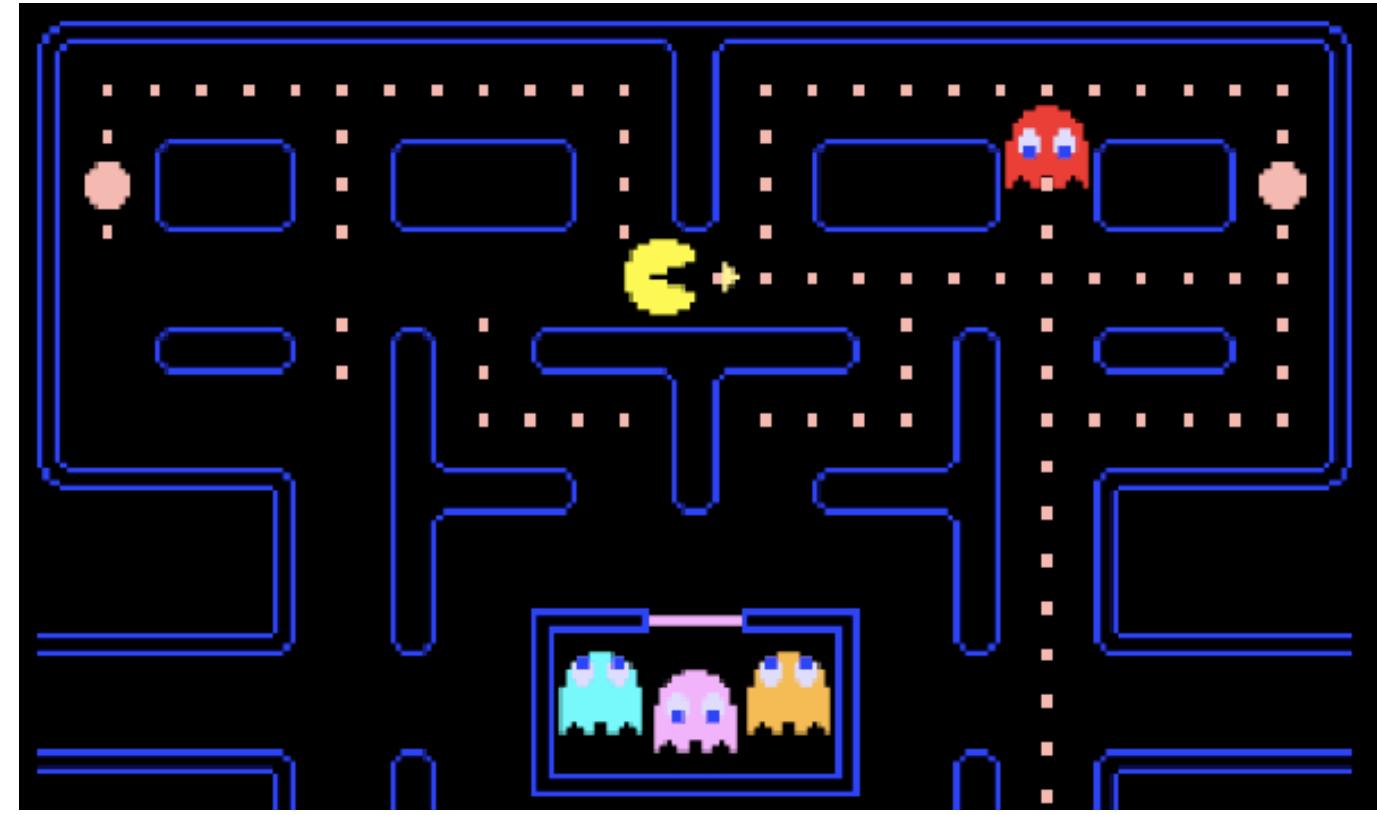
- Zero-Sum Games
  - Agents have **opposite** utilities
  - Pure competition:
  - One **maximizes**, the other **minimizes**

- General Games
  - Agents have **independent** utilities
  - Cooperation, indifference, competition, shifting alliances, and more are all possible

# 8-QUEENS WITH RANDOM RESTARTS

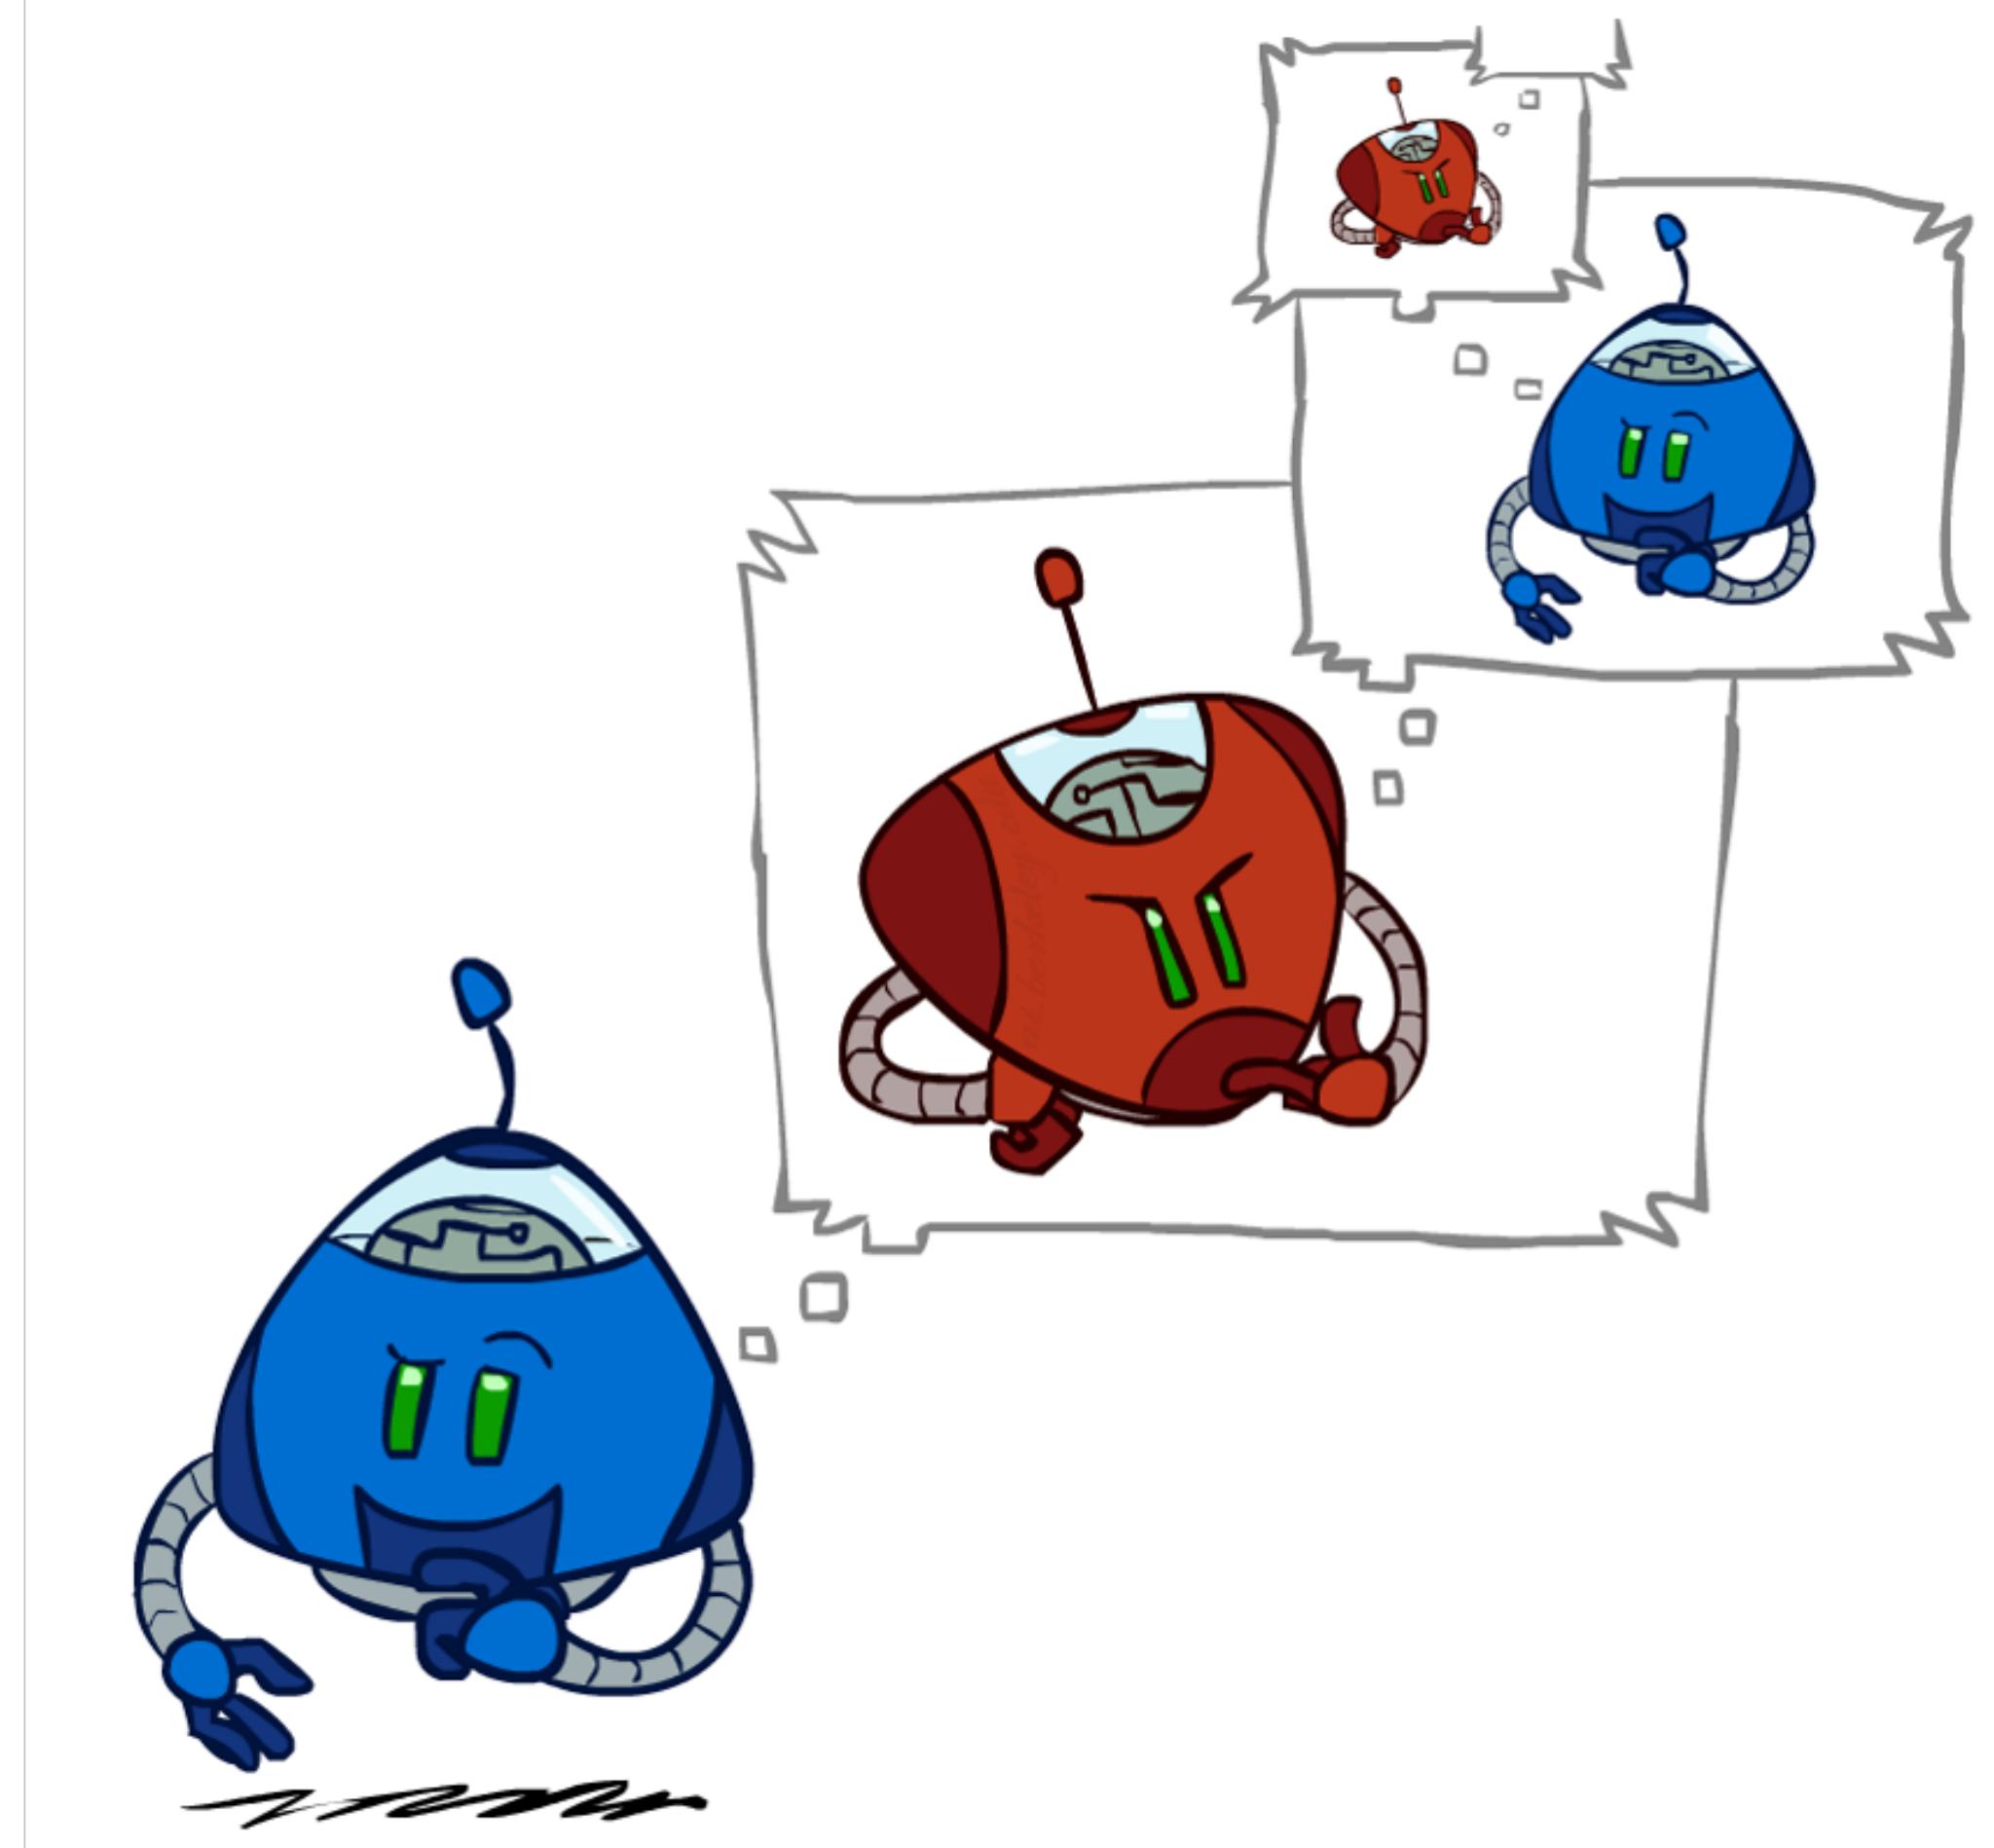


5	3		7					
6			1	9	5			
	9	8				6		
8			6				3	
4		8	3				1	
7			2			6		
	6			2	8			
		4	1	9			5	
			8		7	9		



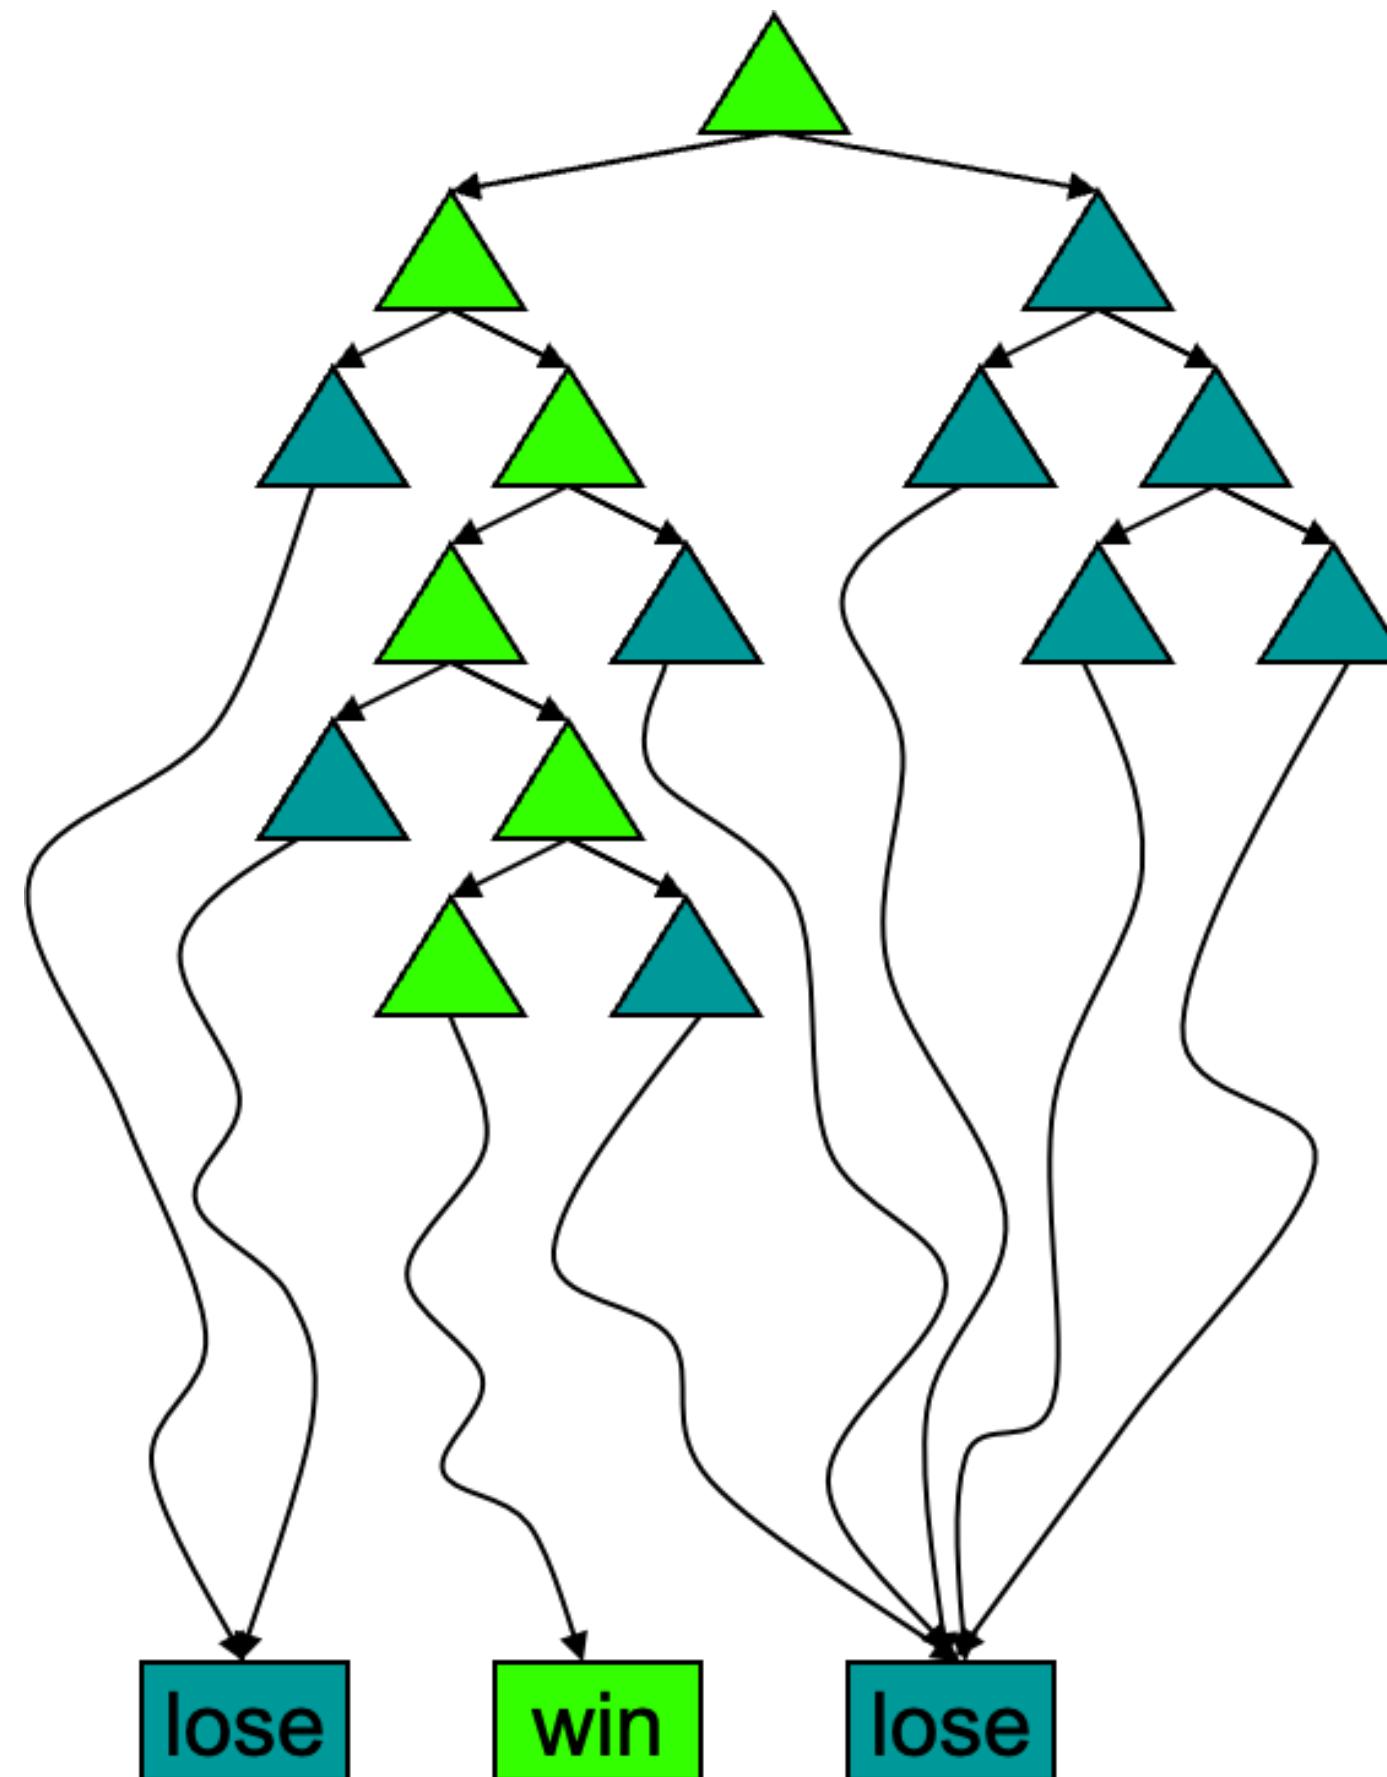
- "Unpredictable" opponent → specifying a move for every possible opponent reply
- Time limits → unlikely to find optimal solution, must approximate

# ADVERSARIAL SEARCH

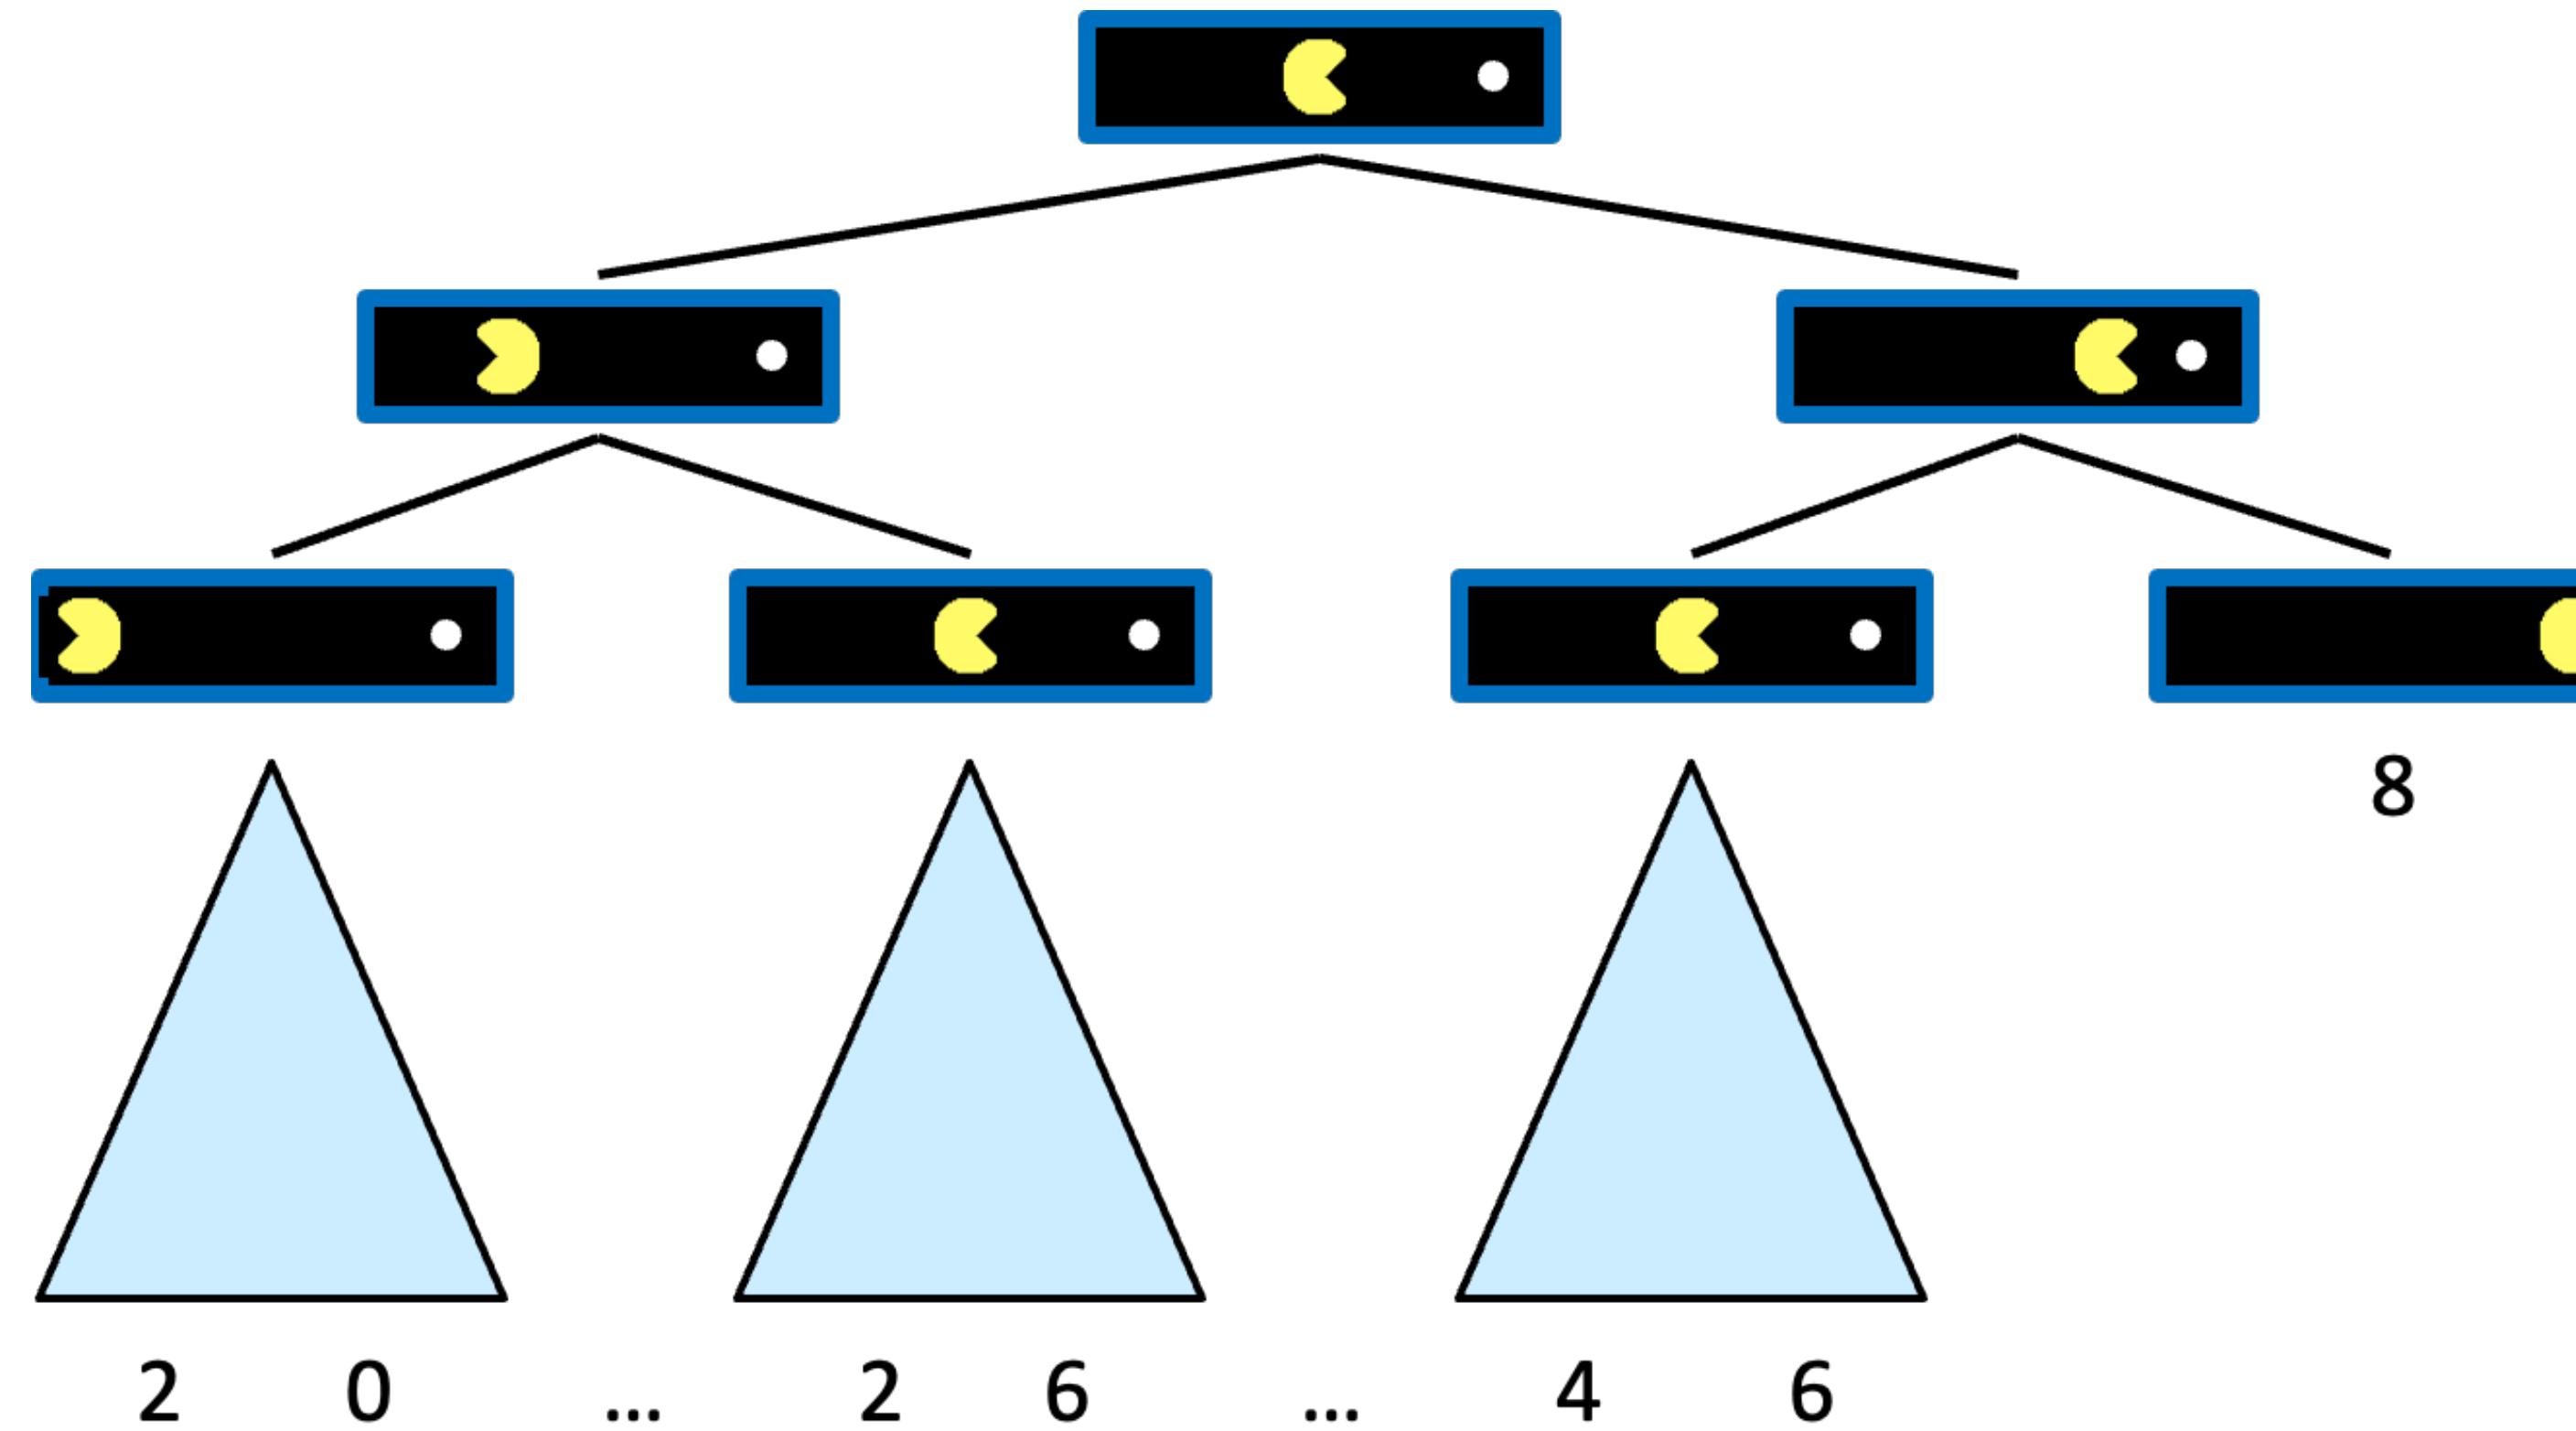


# DETERMINISTIC SINGLE-PLAYER?

- Deterministic, single player, perfect information:
    - Know the rules
    - Know what actions do
    - Know when you win
    - E.g. Freecell, 8-Puzzle, Rubik's cube
  - ... it's just search!
  - Slight reinterpretation:
    - Each node stores the best outcome it can reach
    - This is the maximal outcome of its children
    - Note that we don't store path sums as before
  - After search, can pick move that leads to best node

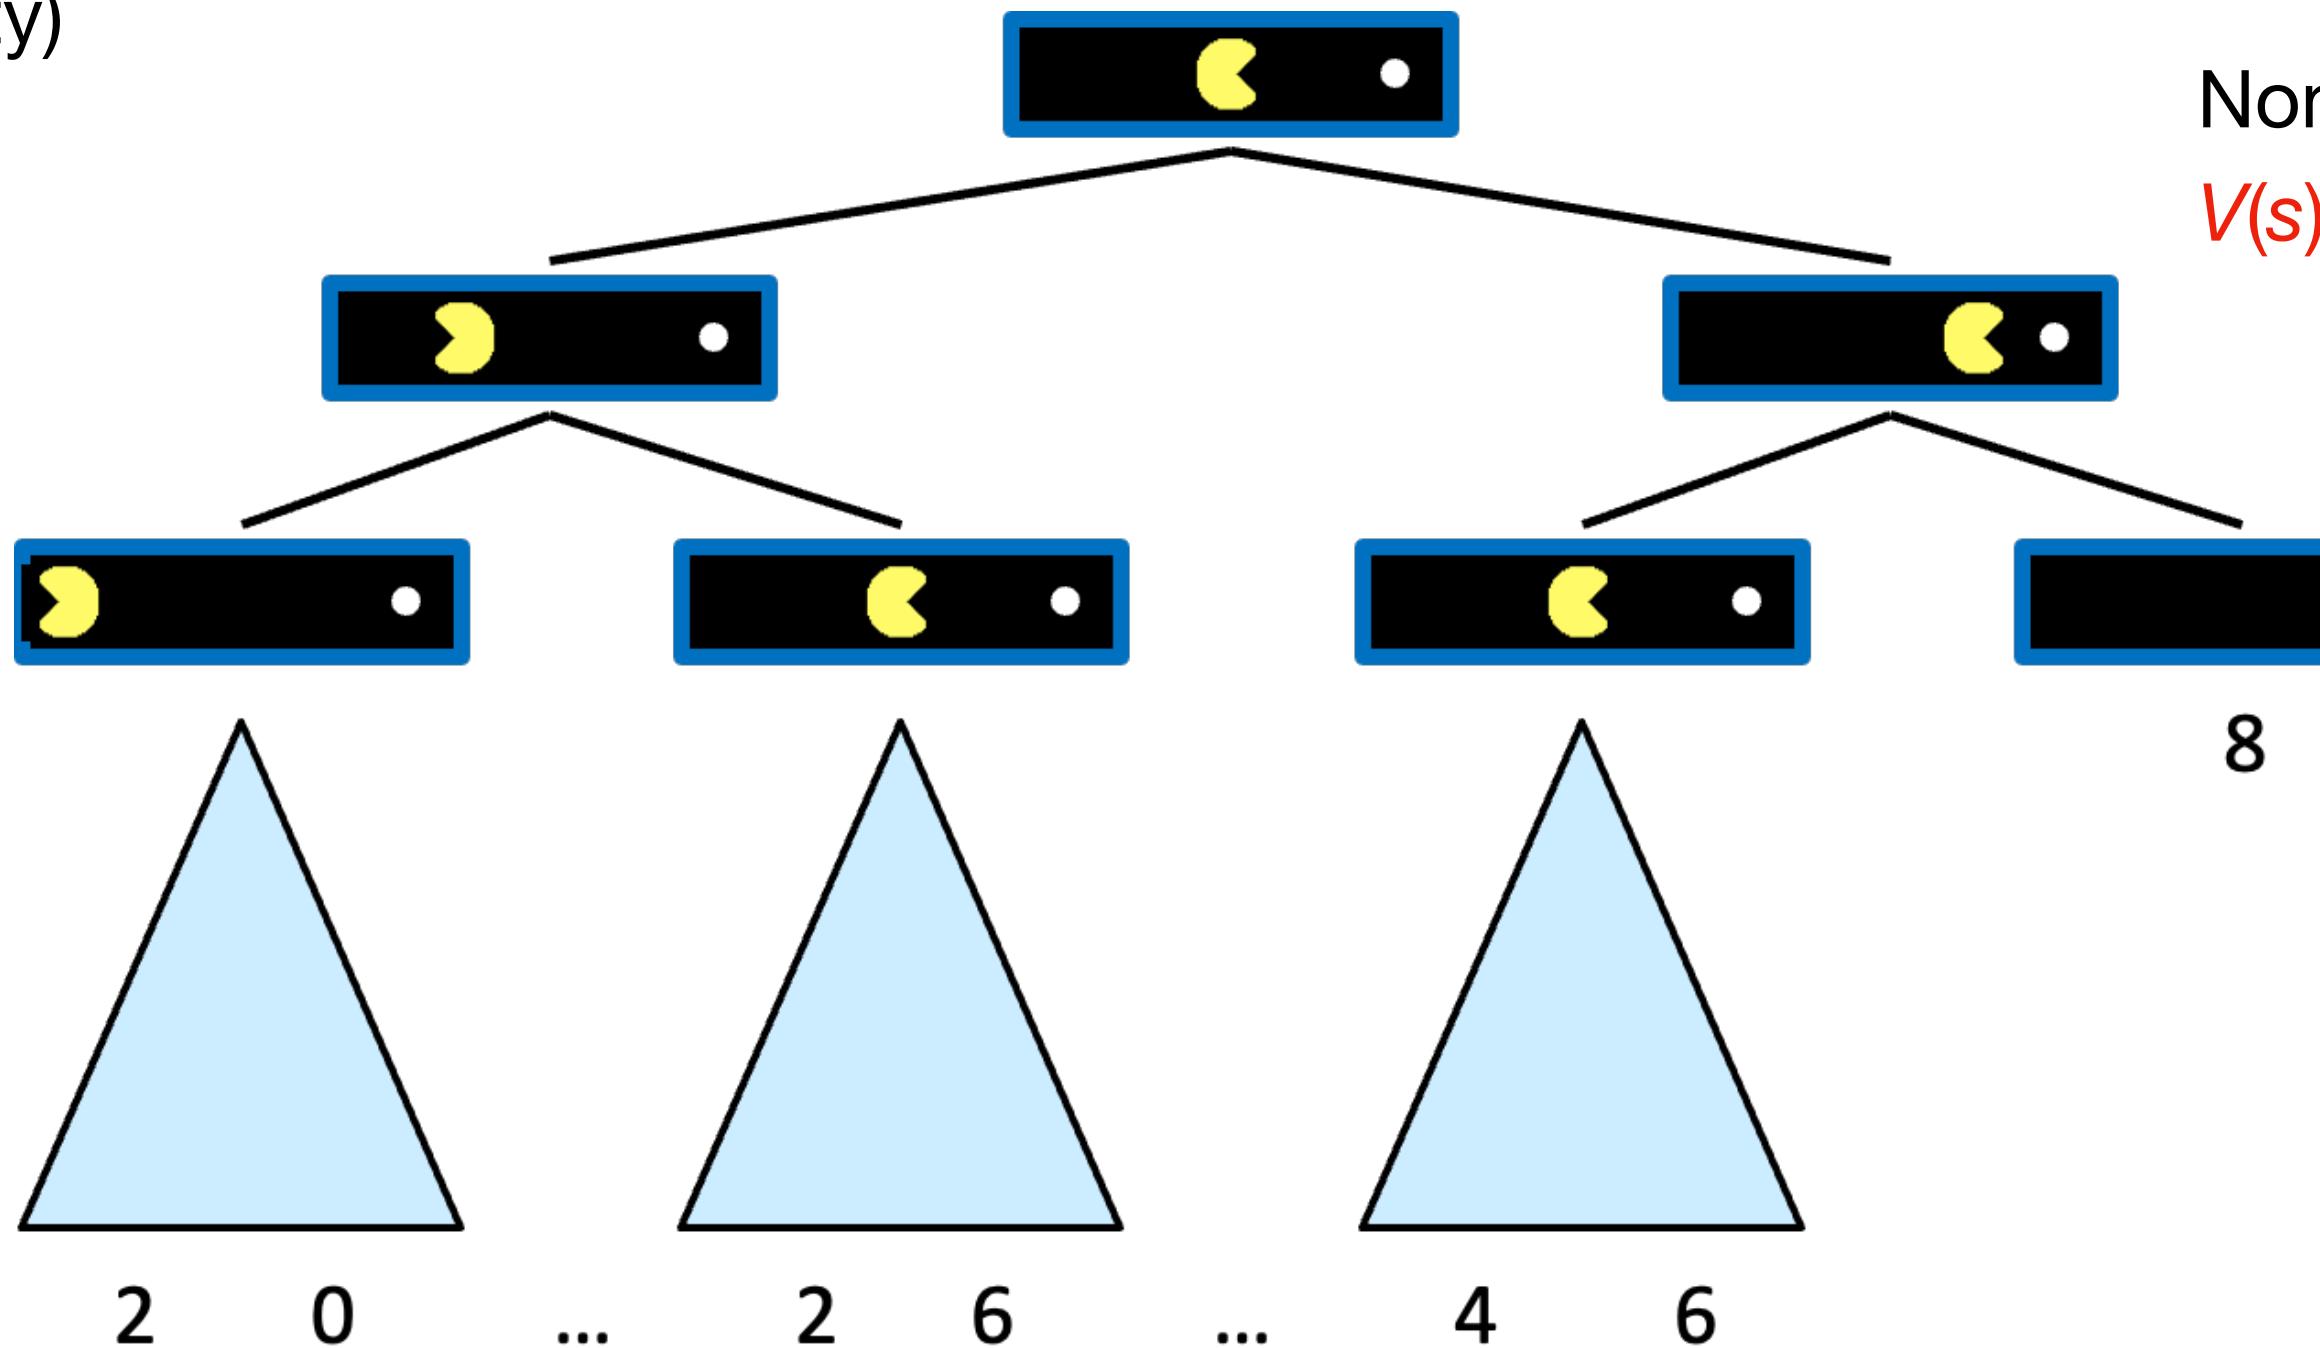


# PACEMAN SINGLE -PLAYER SEARCH TREE



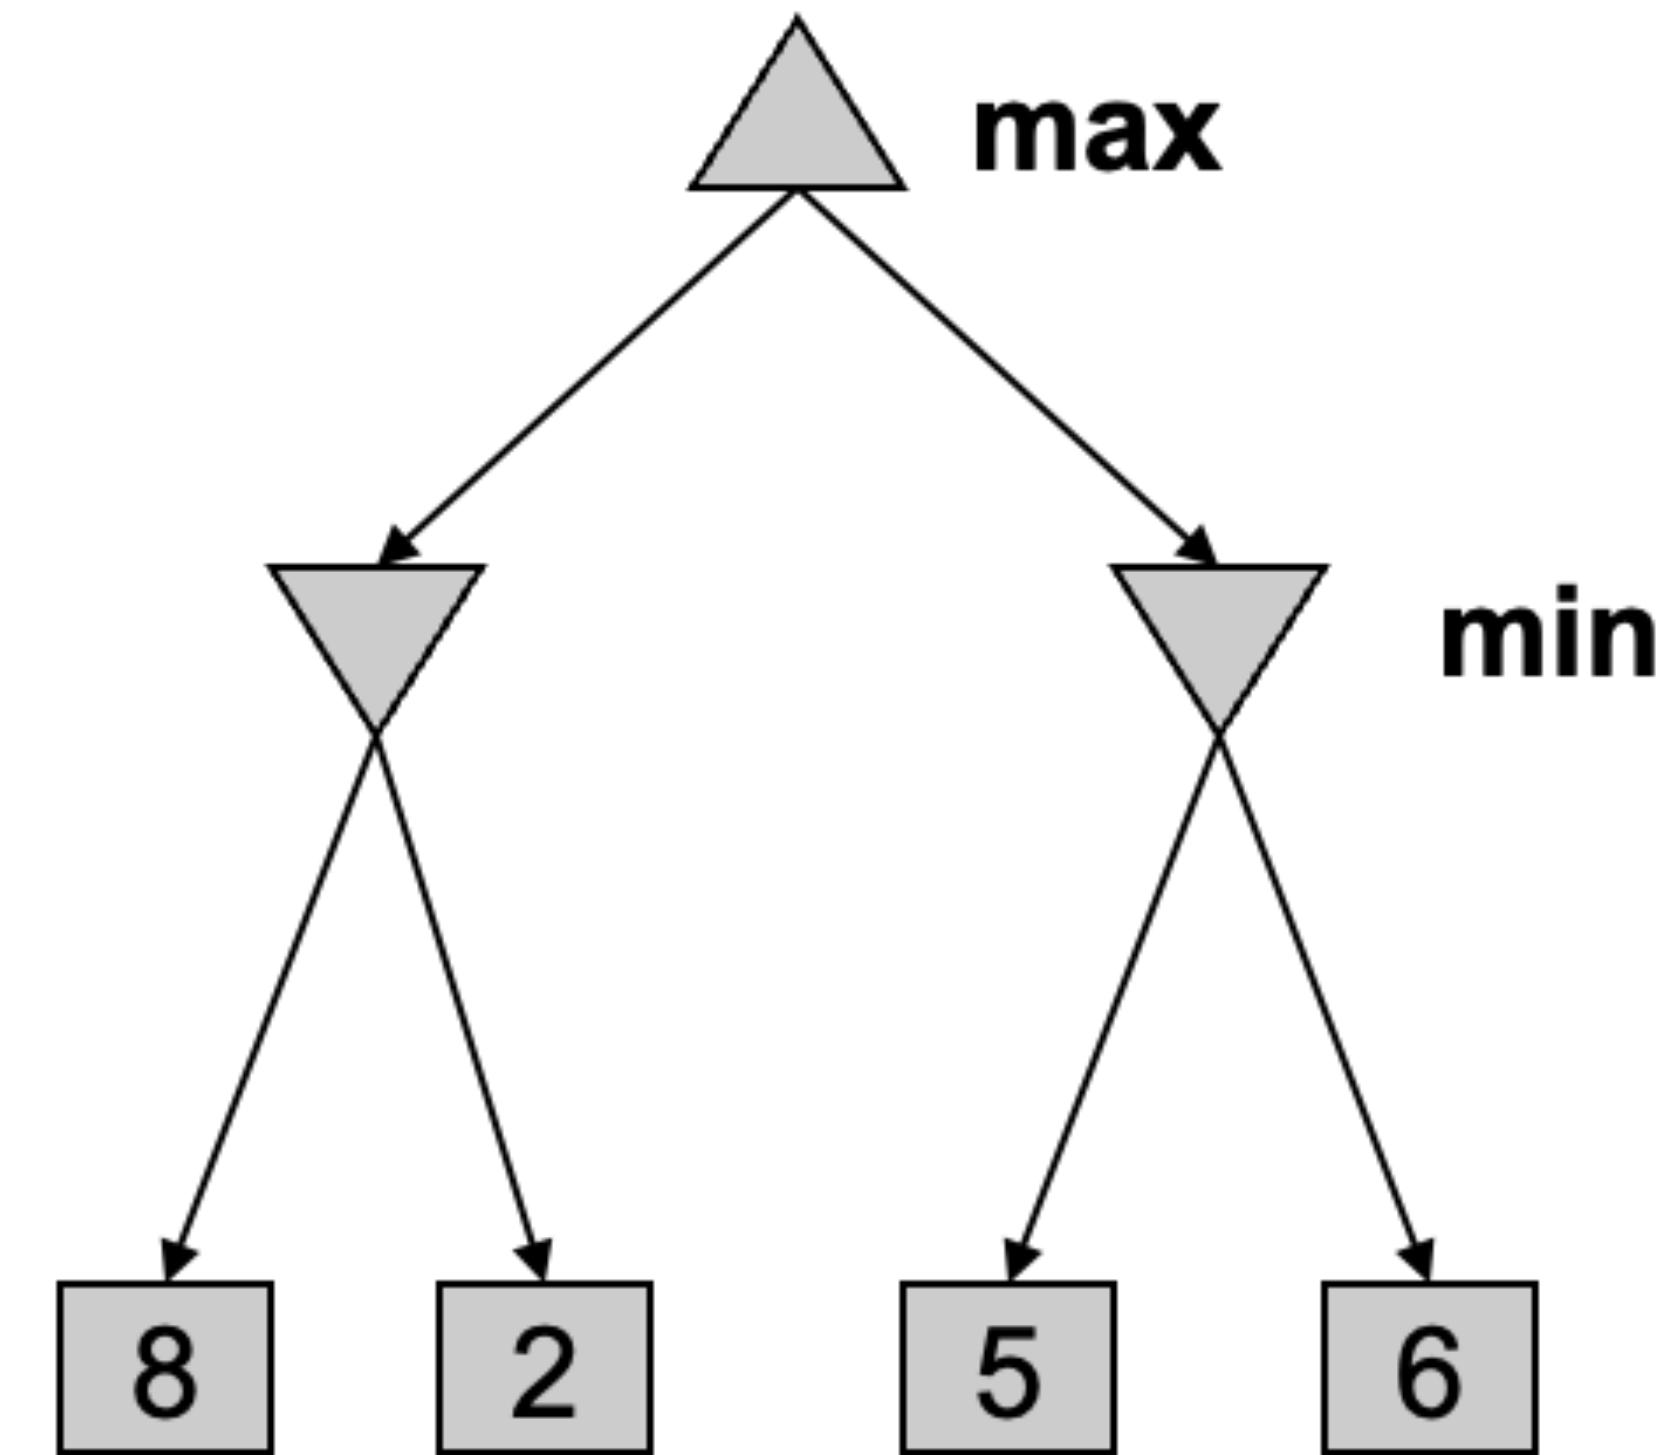
# VALUE OF A STATE

Value of a state: The best achievable outcome (utility) from that state

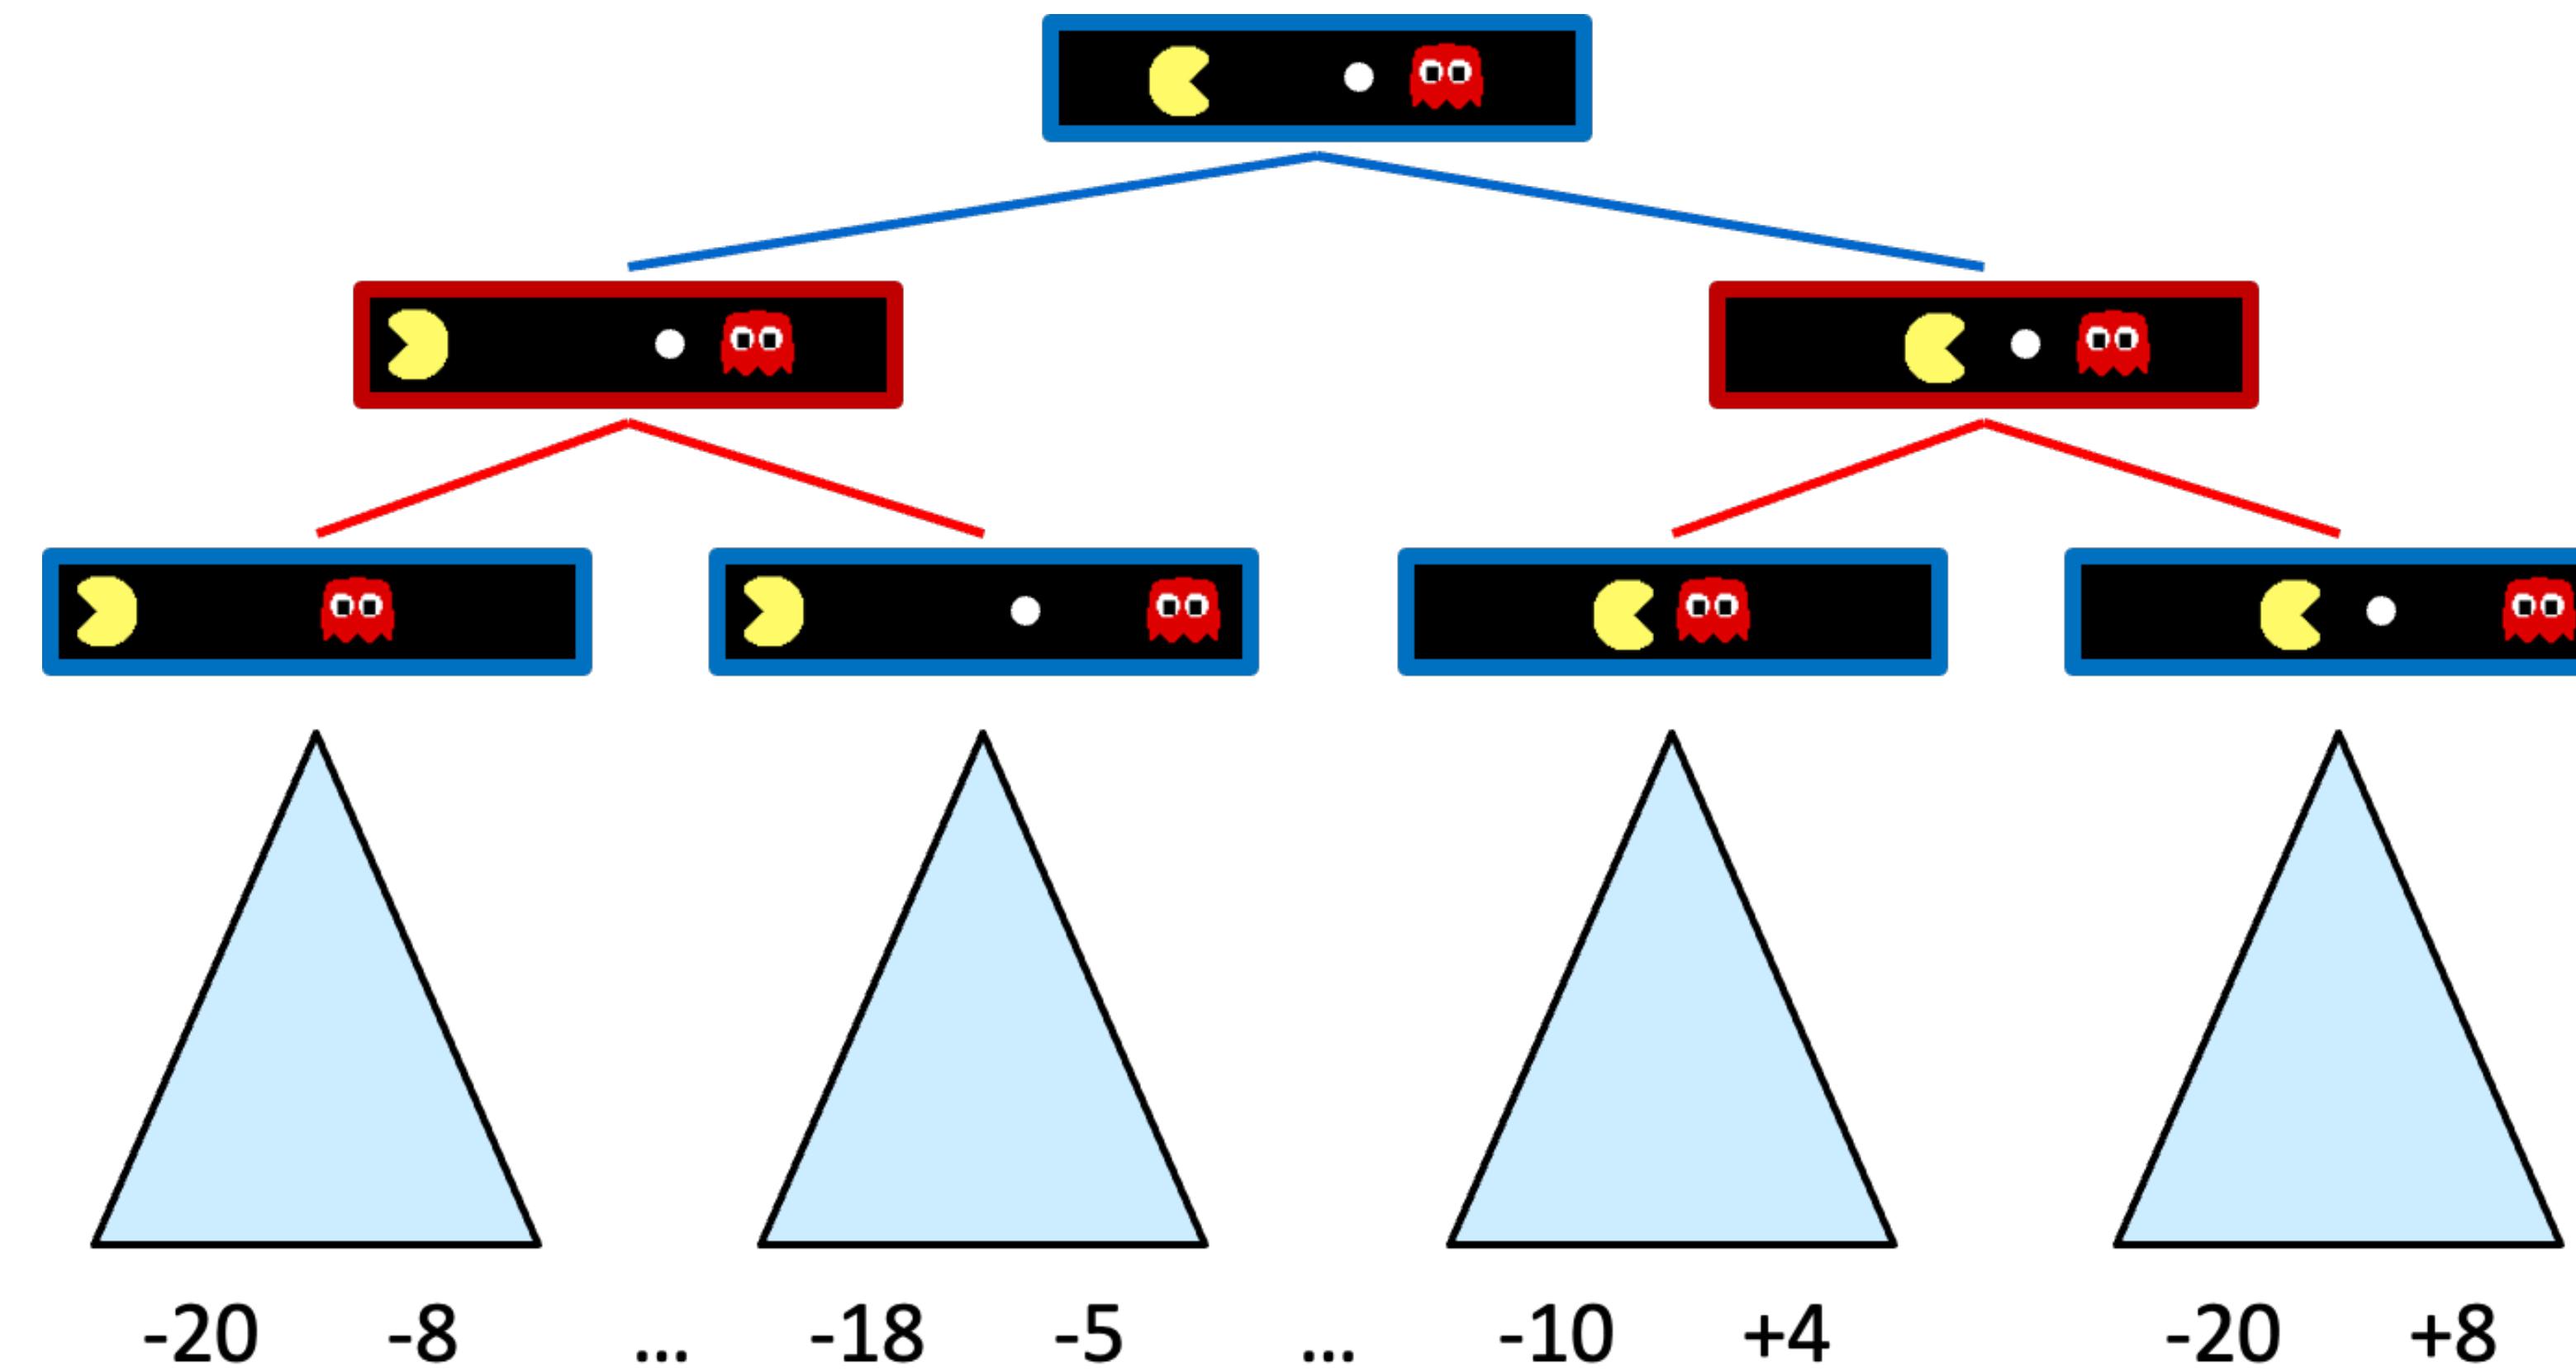


# DETERMINISTIC TWO-PLAYER

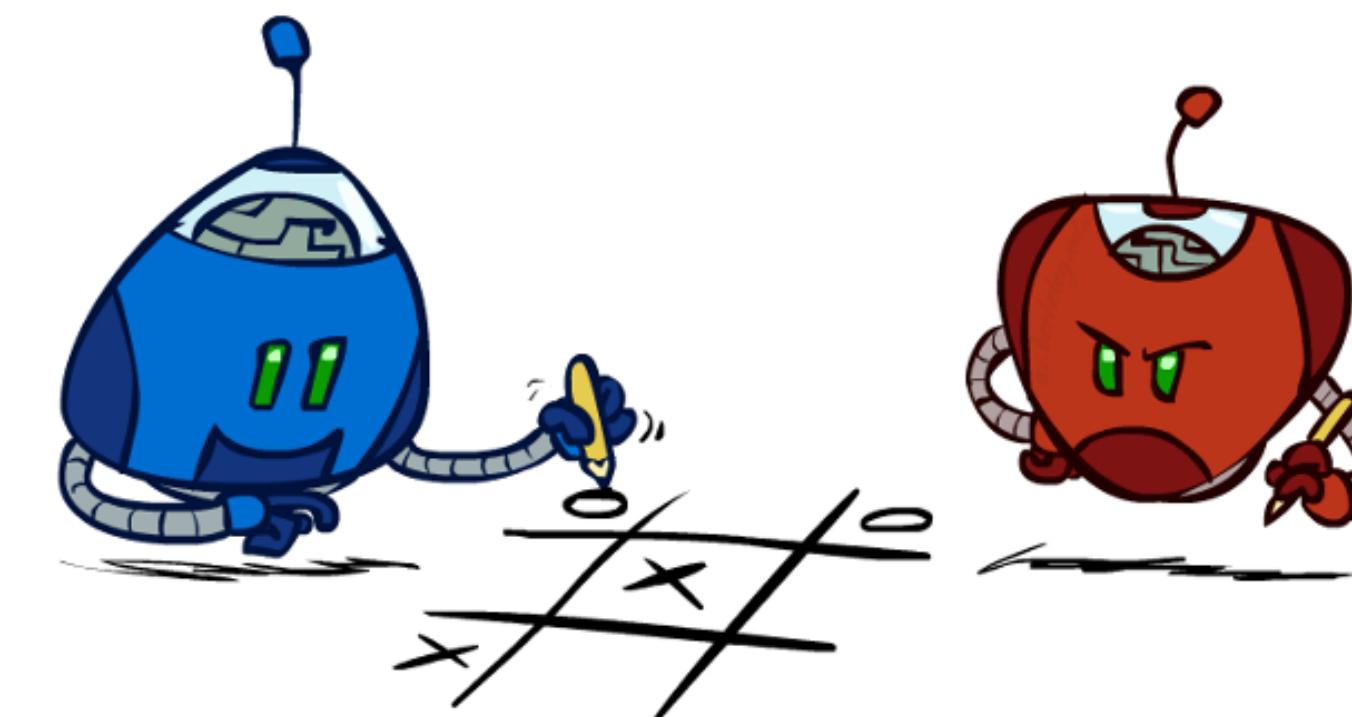
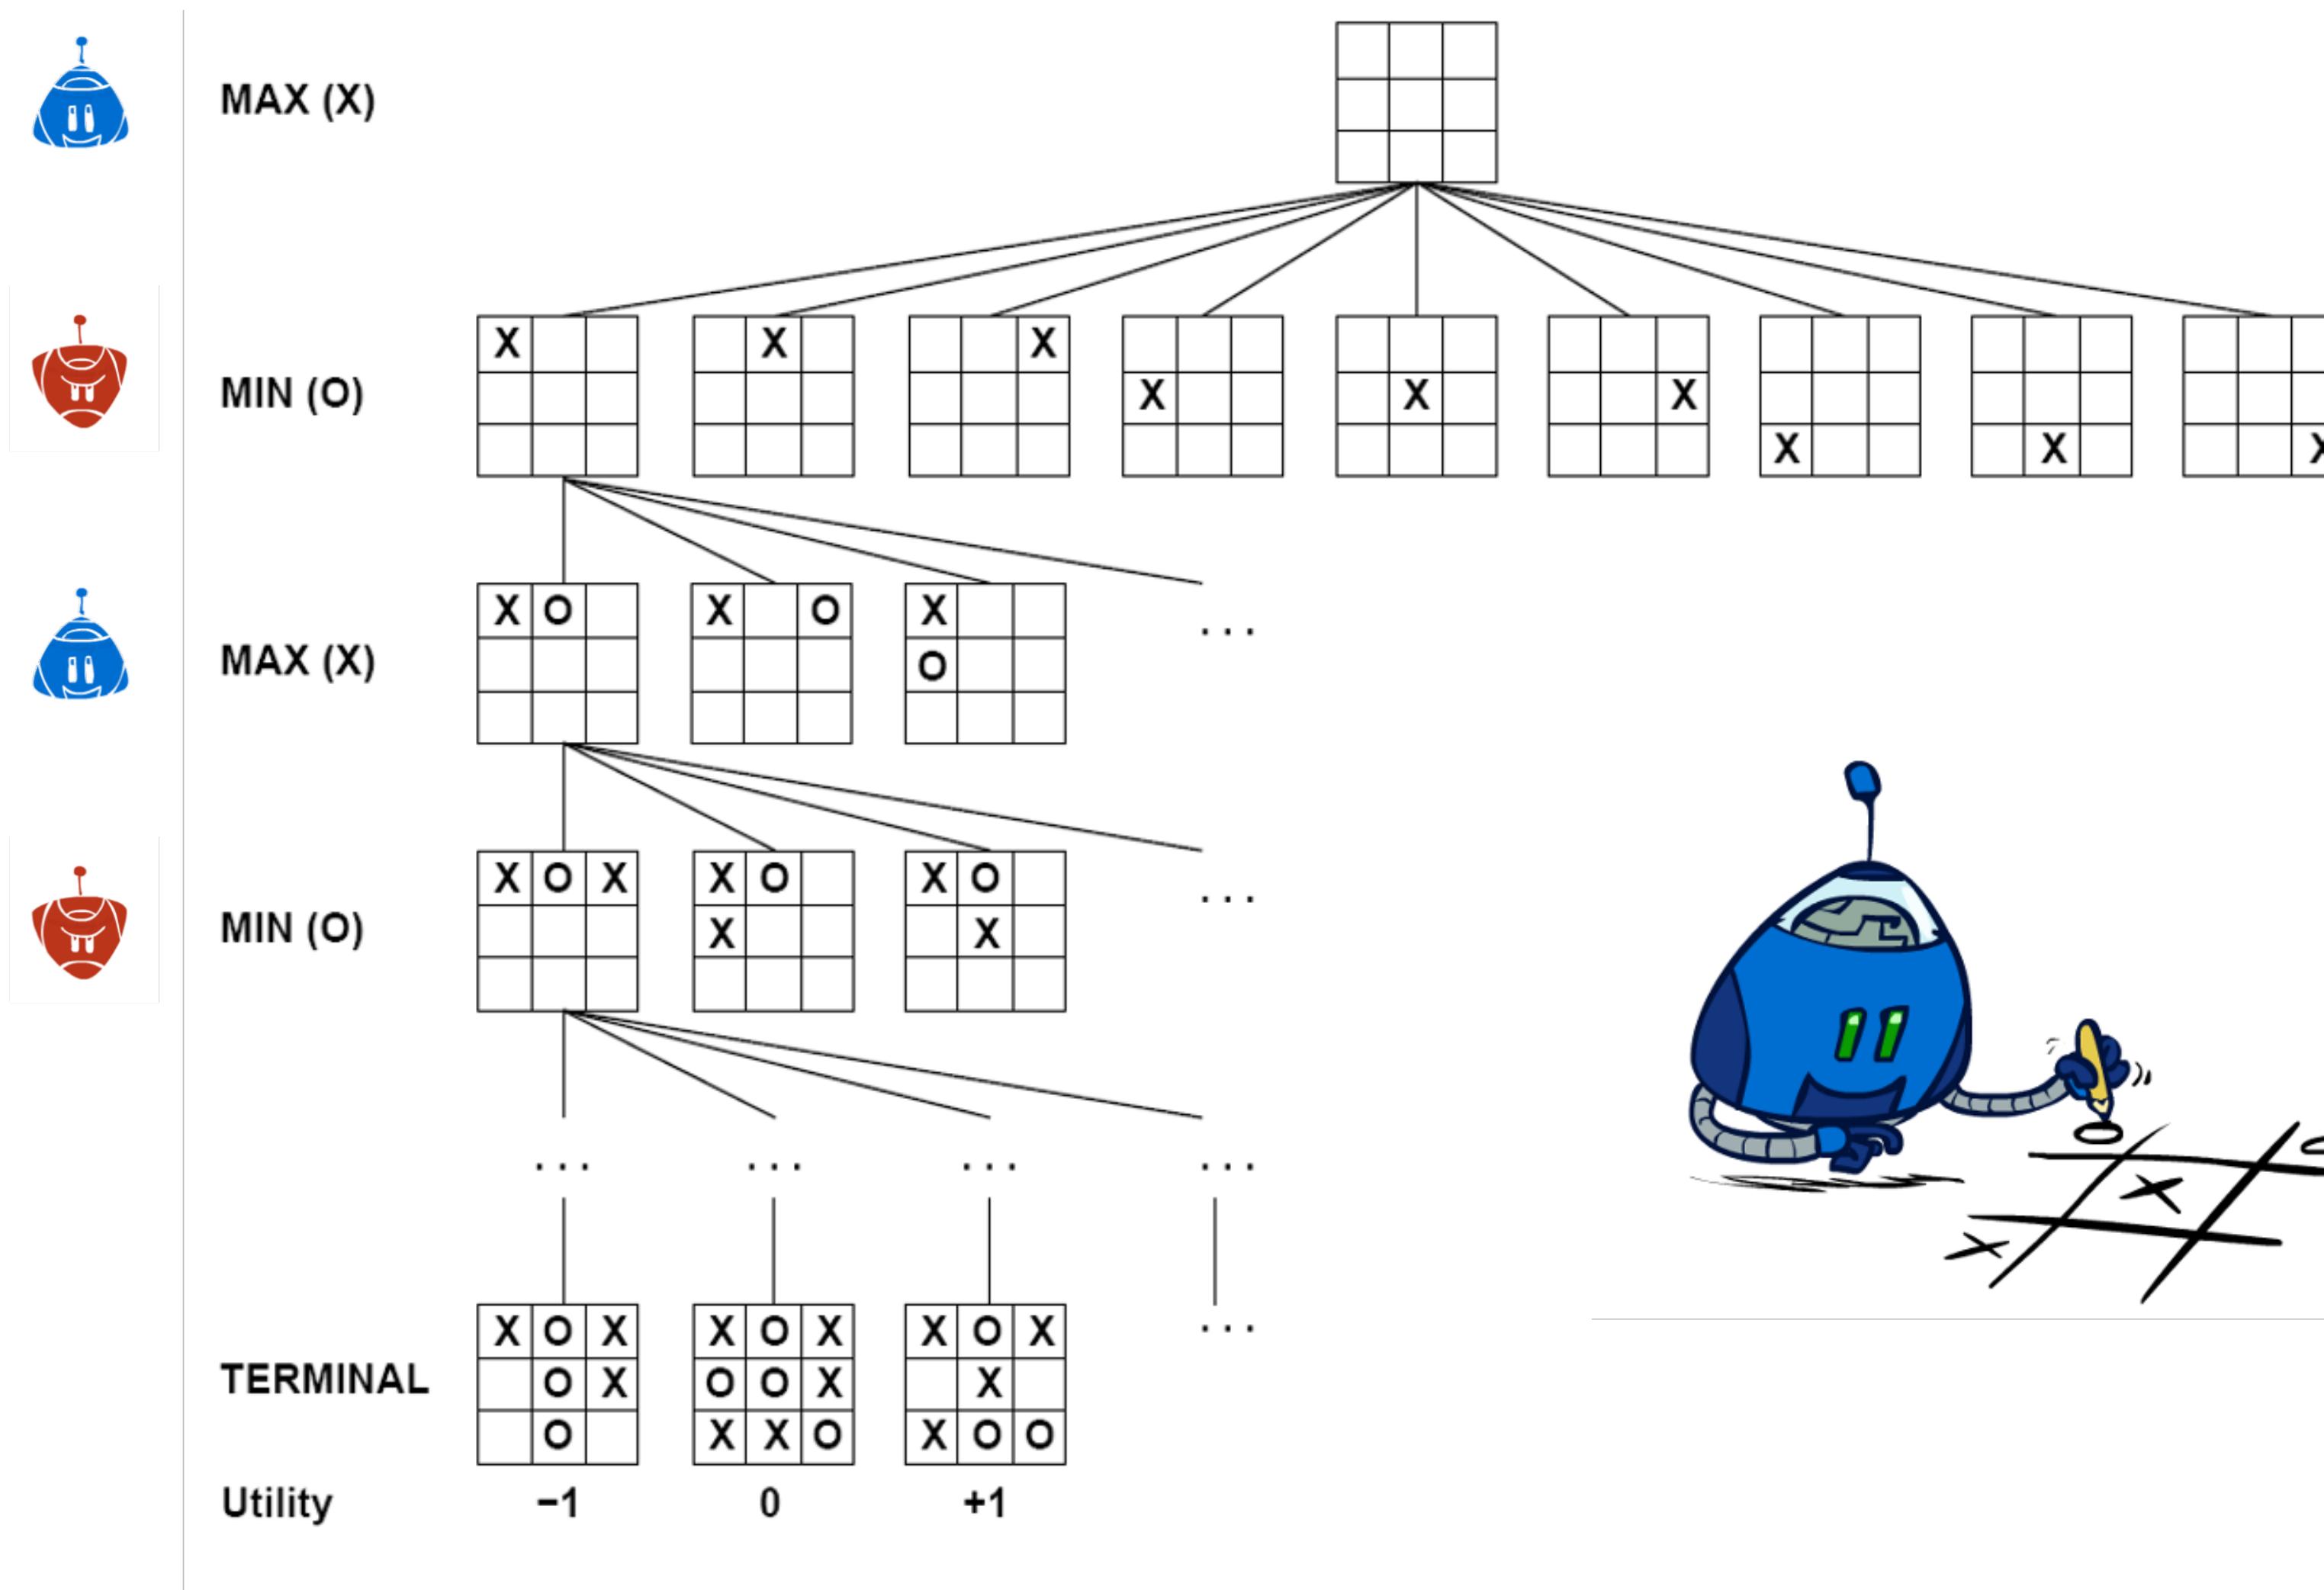
- E.g. tic-tac-toe, chess, checkers
- Minimax search
  - A state-space search tree
  - Players alternate
  - Each layer, or ply, consists of a round of moves
  - Choose move to position with highest **minimax value** = best achievable utility against best play
- Zero-sum games
  - One player maximizes result
  - The other minimizes result



# ADVERSARIAL GAME TREE



# TIC-TAC-TOE GAME TREE



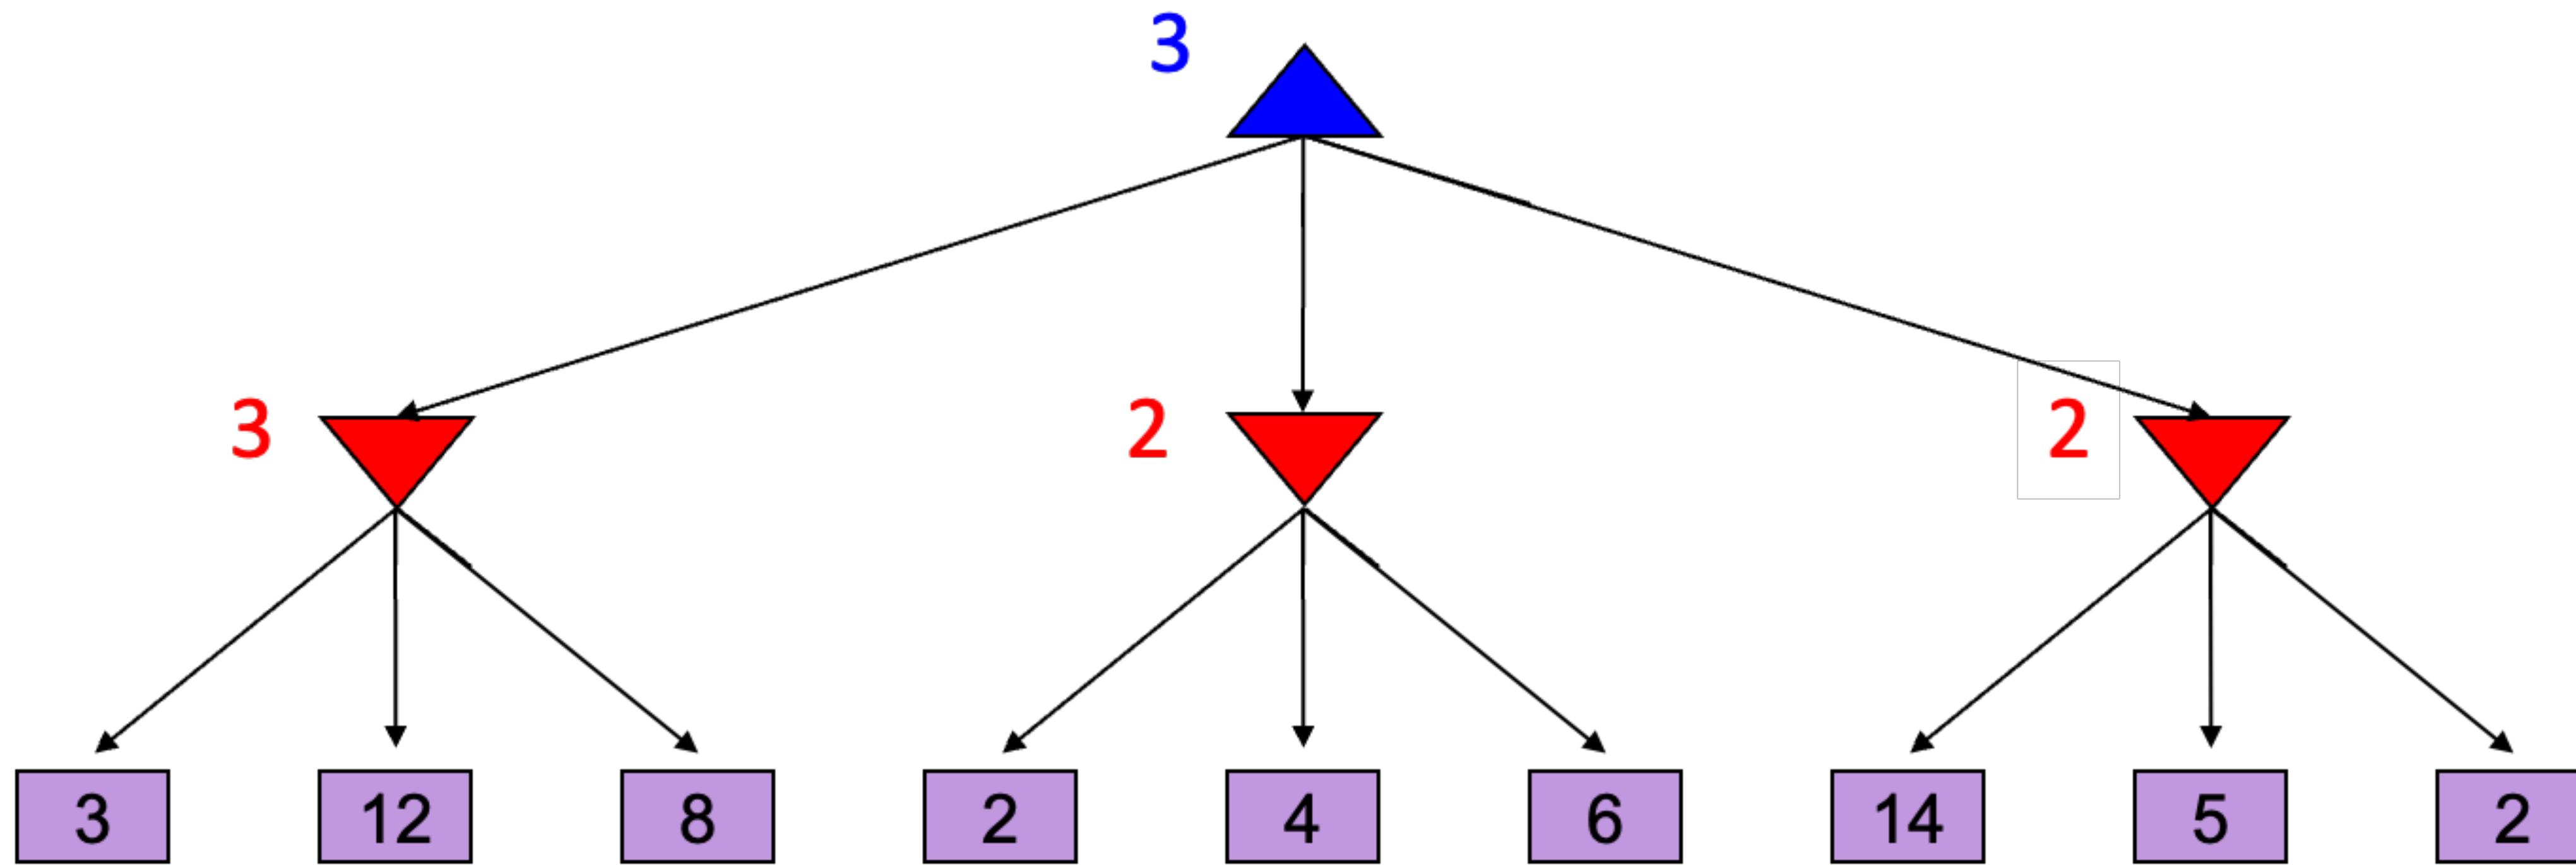
# MINIMAX SEARCH

```
function minimax-decision(s) returns an action  
    return the action a in Actions(s) with the highest  
        minimax_value(Result(s,a))
```



```
function minimax_value(s) returns a value  
    if Terminal-Test(s) then return Utility(s)  
    if Player(s) = MAX then return maxa in Actions(s) minimax_value(Result(s,a))  
    if Player(s) = MIN then return mina in Actions(s) minimax_value(Result(s,a))
```

# MINIMAX EXAMPLE

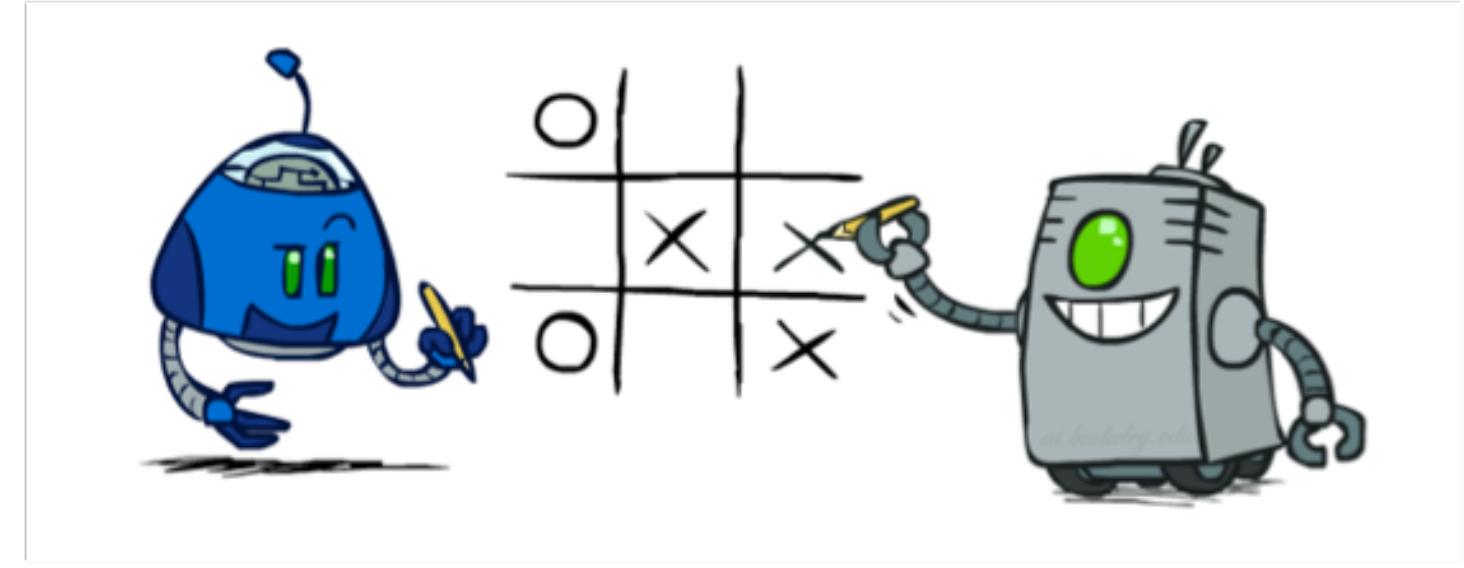
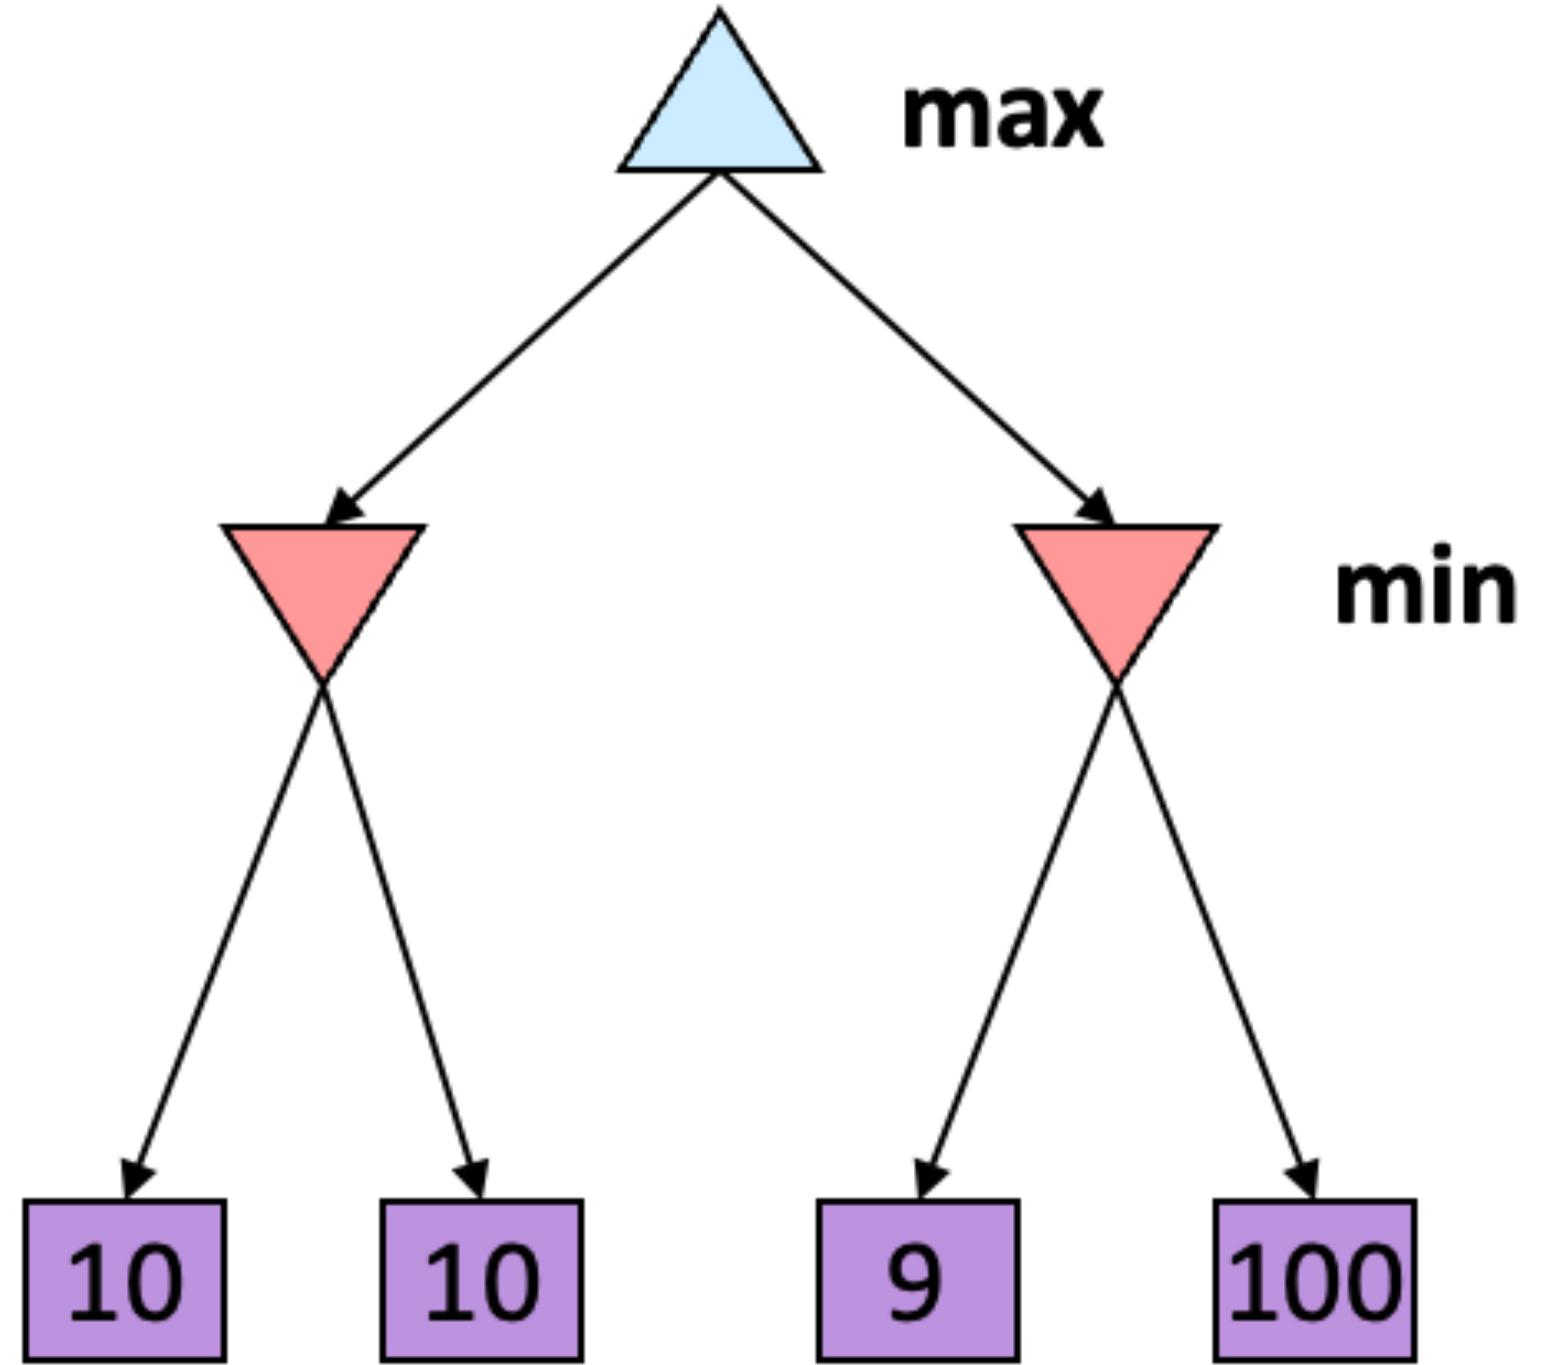
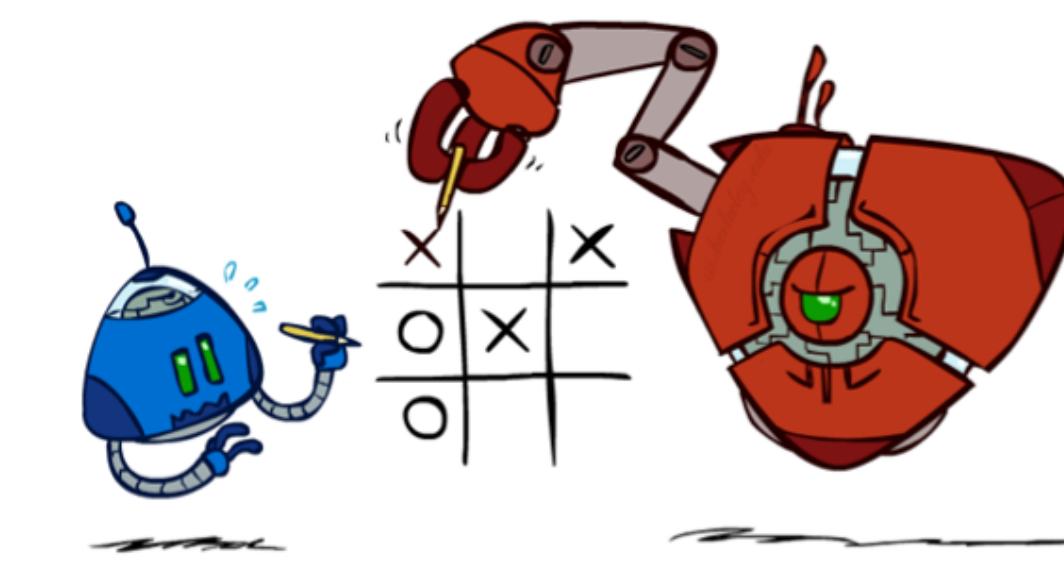


# MINIMAX PROPERTIES

- Optimal against a perfect player. Otherwise?

# MINIMAX PROPERTIES

- Optimal against a perfect player. Otherwise?



# MINIMAX PROPERTIES

- Optimal against a perfect player. Otherwise?
- Time complexity?
  - $O(b^m)$
  - $m$  = maximum depth of search tree,  $b$  = branching factor
- Space complexity?
  - $O(bm)$
- For chess,  $b \sim 35$ ,  $m \sim 100$ 
  - Exact solution is completely infeasible
  - But, do we need to explore the whole tree?

“

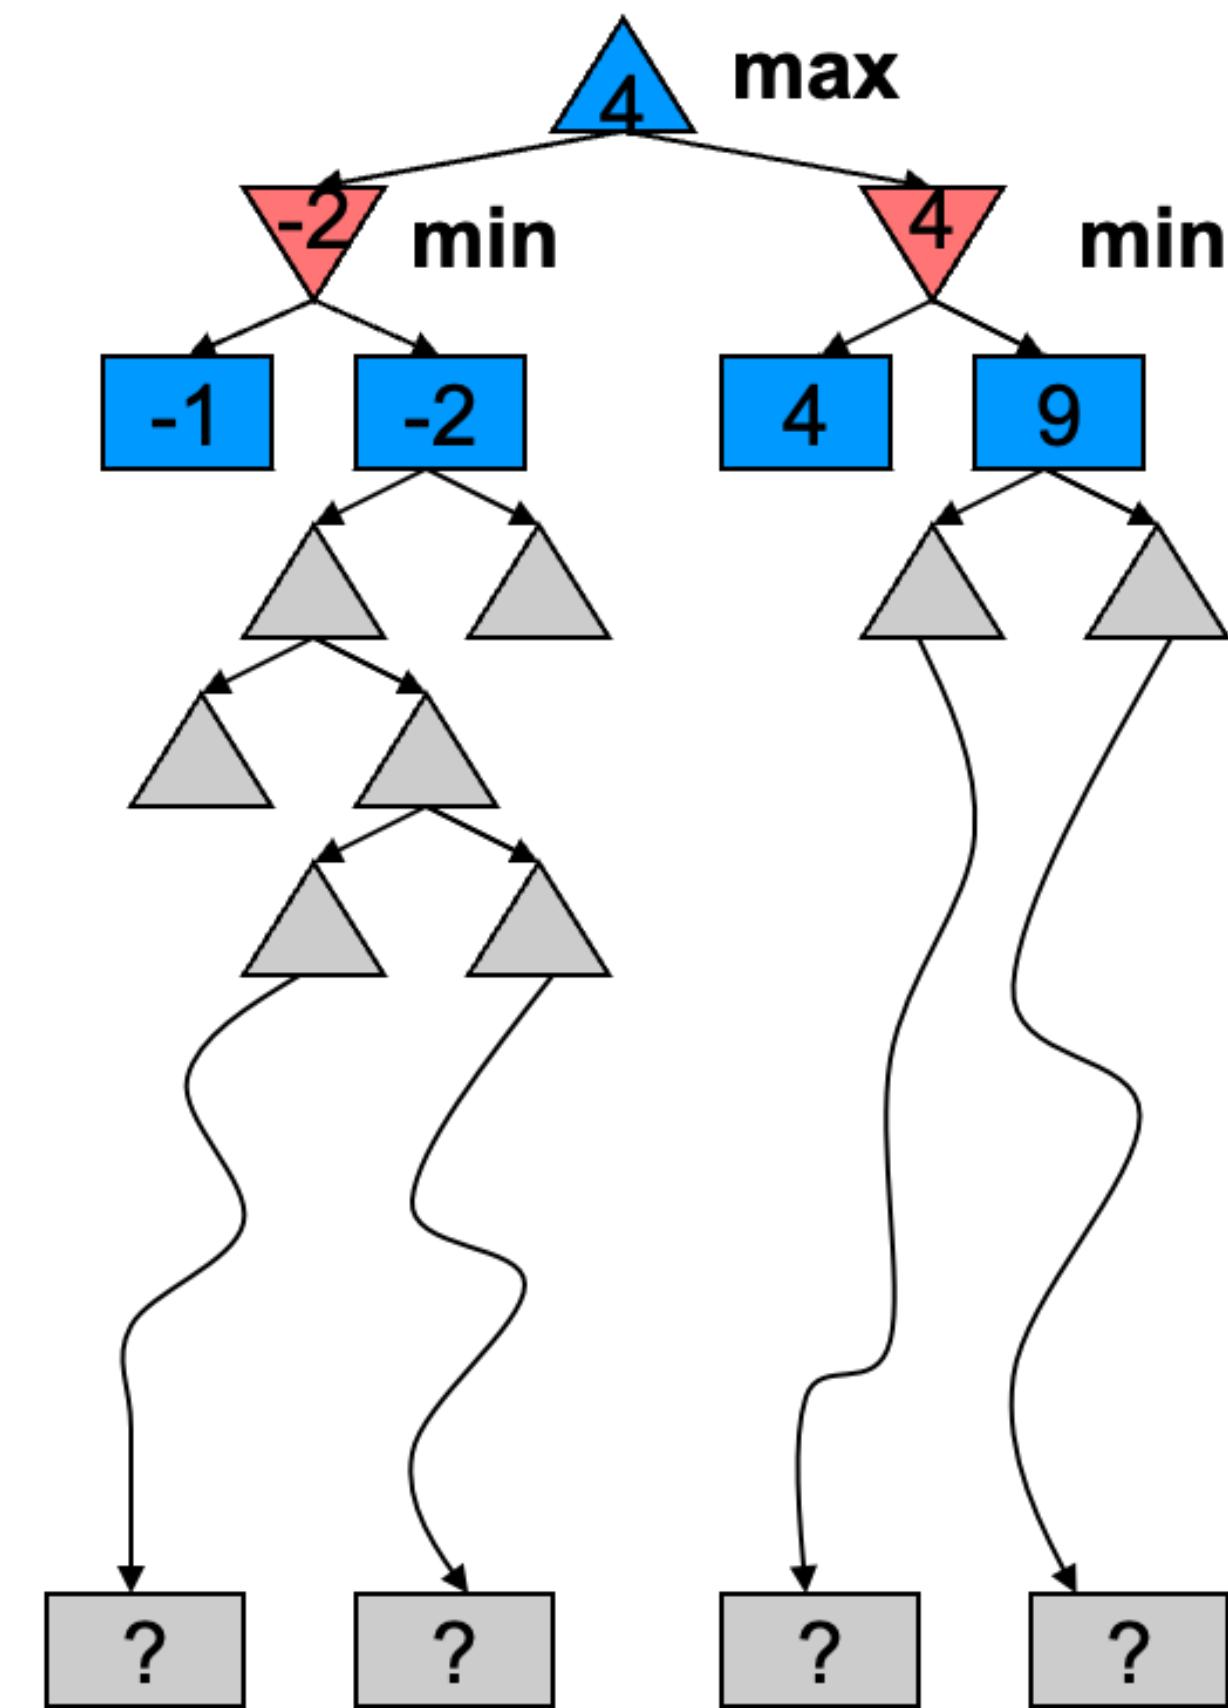
The whole question of making an automaton play any game depended upon the possibility of the machine being able to *represent all the myriads of combinations* relating to it.

”

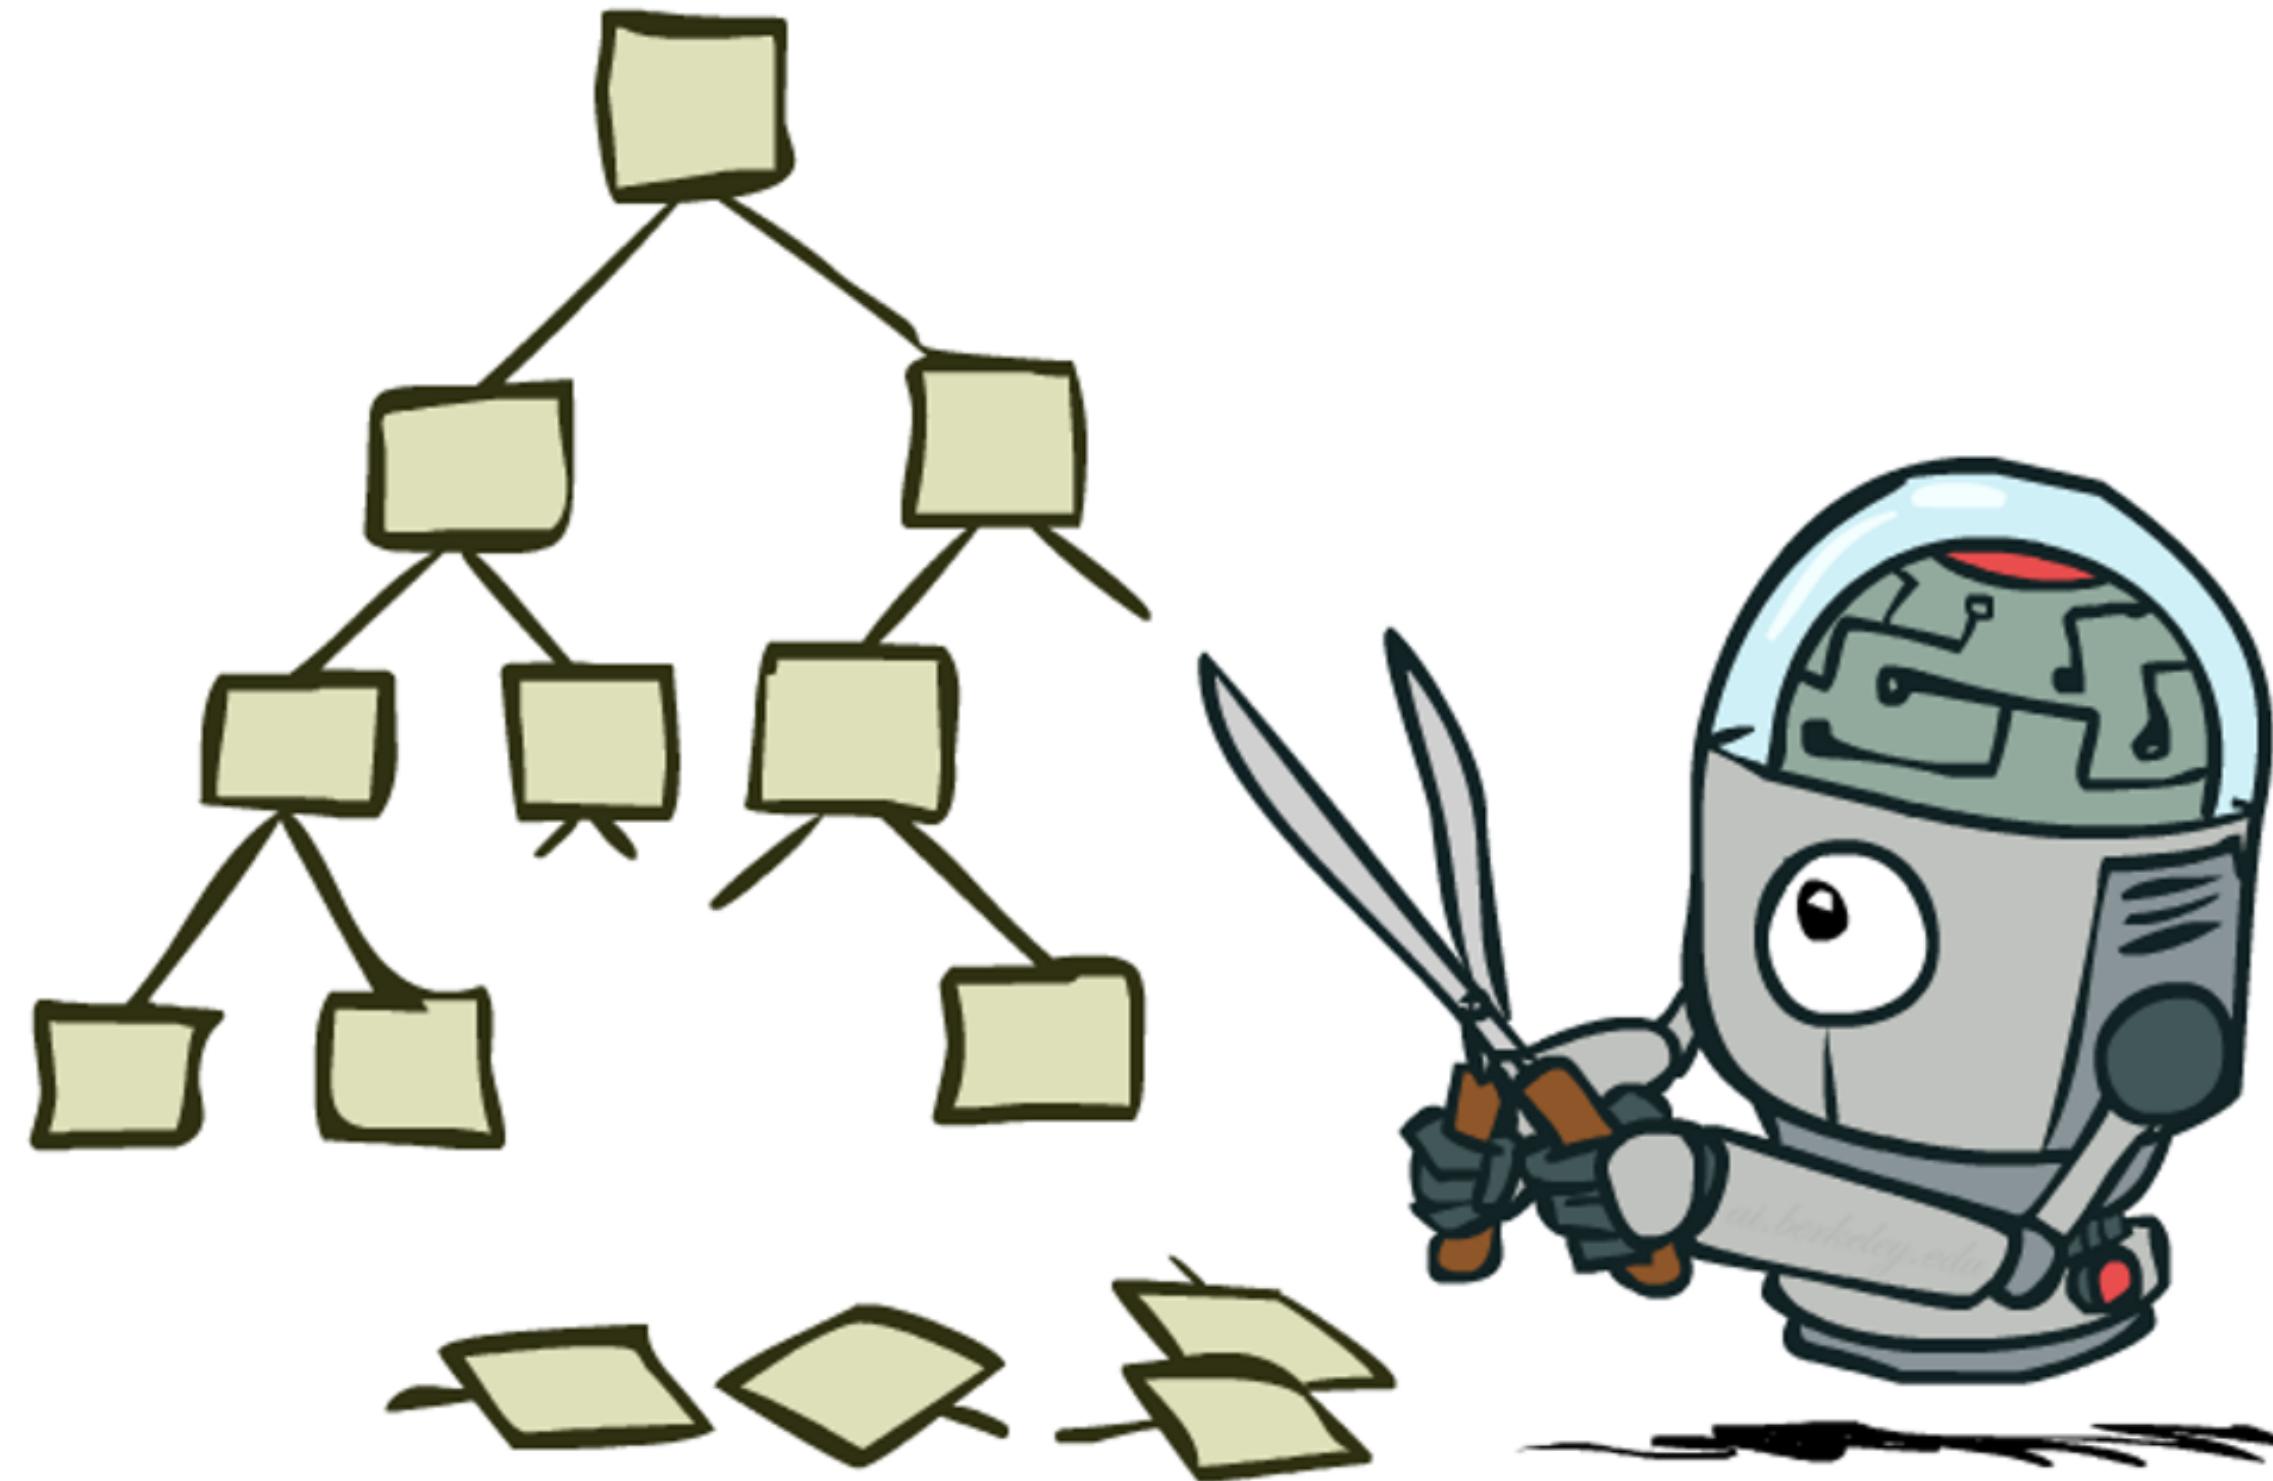
- Charles Babbage

# RESOURCE LIMITS

- Cannot search to leaves
- Limited search
  - Instead, search a limited depth of the tree
  - Replace terminal utilities with an eval function for non-terminal positions
- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Example:
  - Suppose we have 100 seconds, can explore 10K nodes / sec
  - So can check 1M nodes per move
  - $\alpha\beta$  reaches about depth 8 – decent chess program



# GAME TREE PRUNING



# DEPTH MATTERS

- Evaluation functions are always imperfect
- Deeper search => better play (usually)
- Or, deeper search gives same quality of play with a less accurate evaluation function

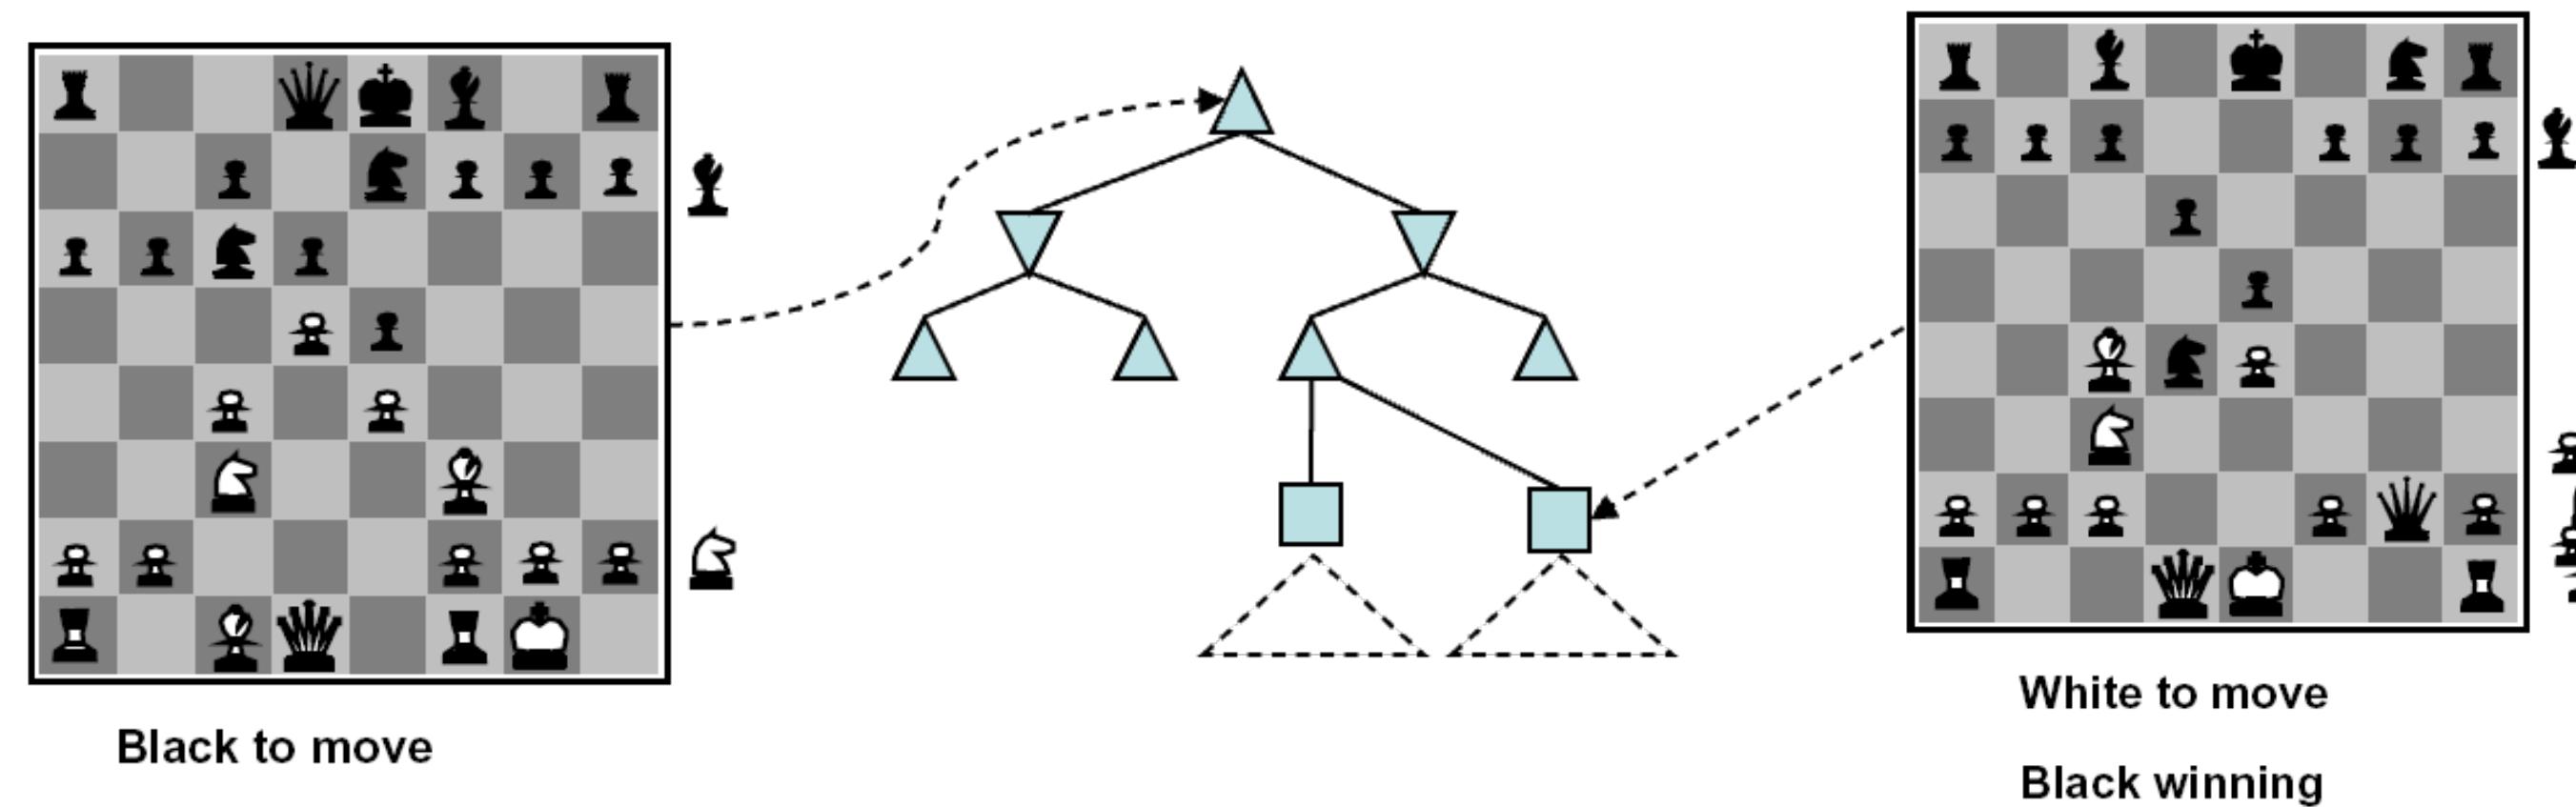


# ITERATIVE DEEPENING

- Iterative deepening uses DFS as a subroutine:
  1. Do a DFS which only searches for paths of length 1 or less. (DFS gives up on any path of length 2)
  2. If “1” failed, do a DFS which only searches paths of length 2 or less.
  3. If “2” failed, do a DFS which only searches paths of length 3 or less.
- ....and so on.
- This works for single-agent search as well!
- Why do we want to do this for multiplayer games?

# EVALUATION FUNCTION

- Function which scores non-terminals



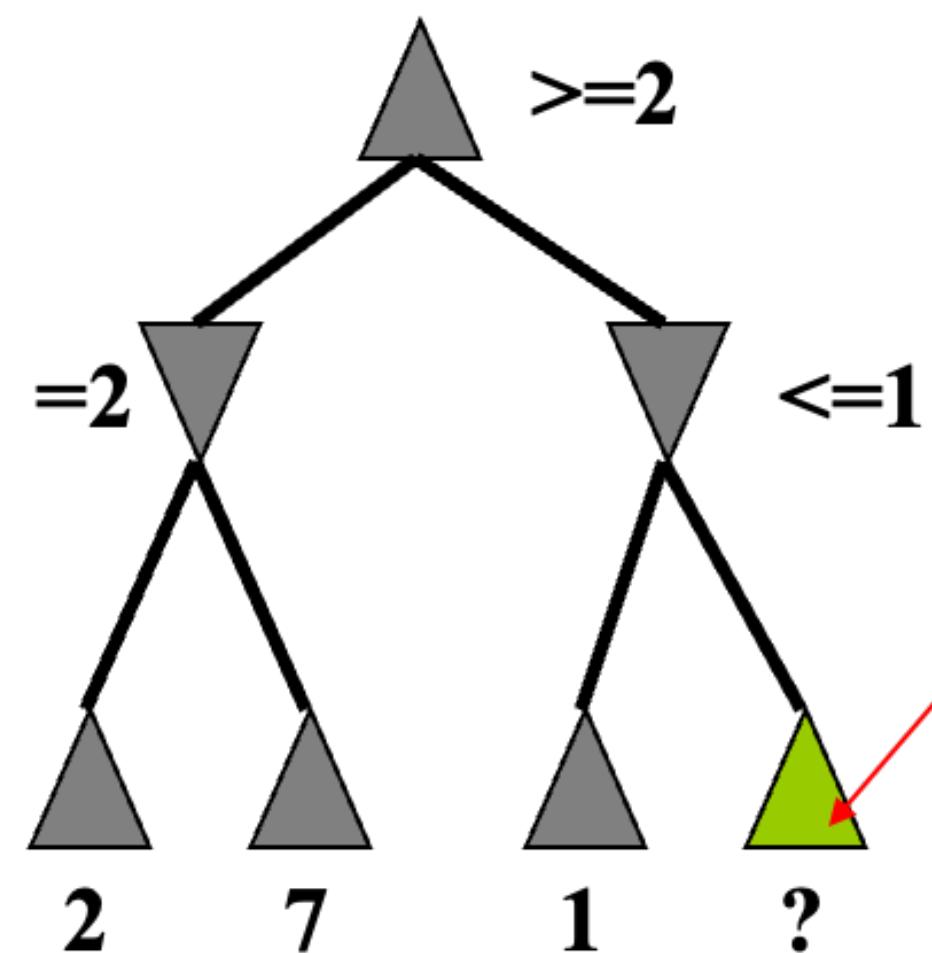
- Ideal function: returns the utility of the position
- In practice: typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g.  $f_1(s) = (\text{num white queens} - \text{num black queens})$ , etc.
- Or a more complex nonlinear function (e.g., NN) trained by self-play RL

# $\alpha$ - $\beta$ PRUNING

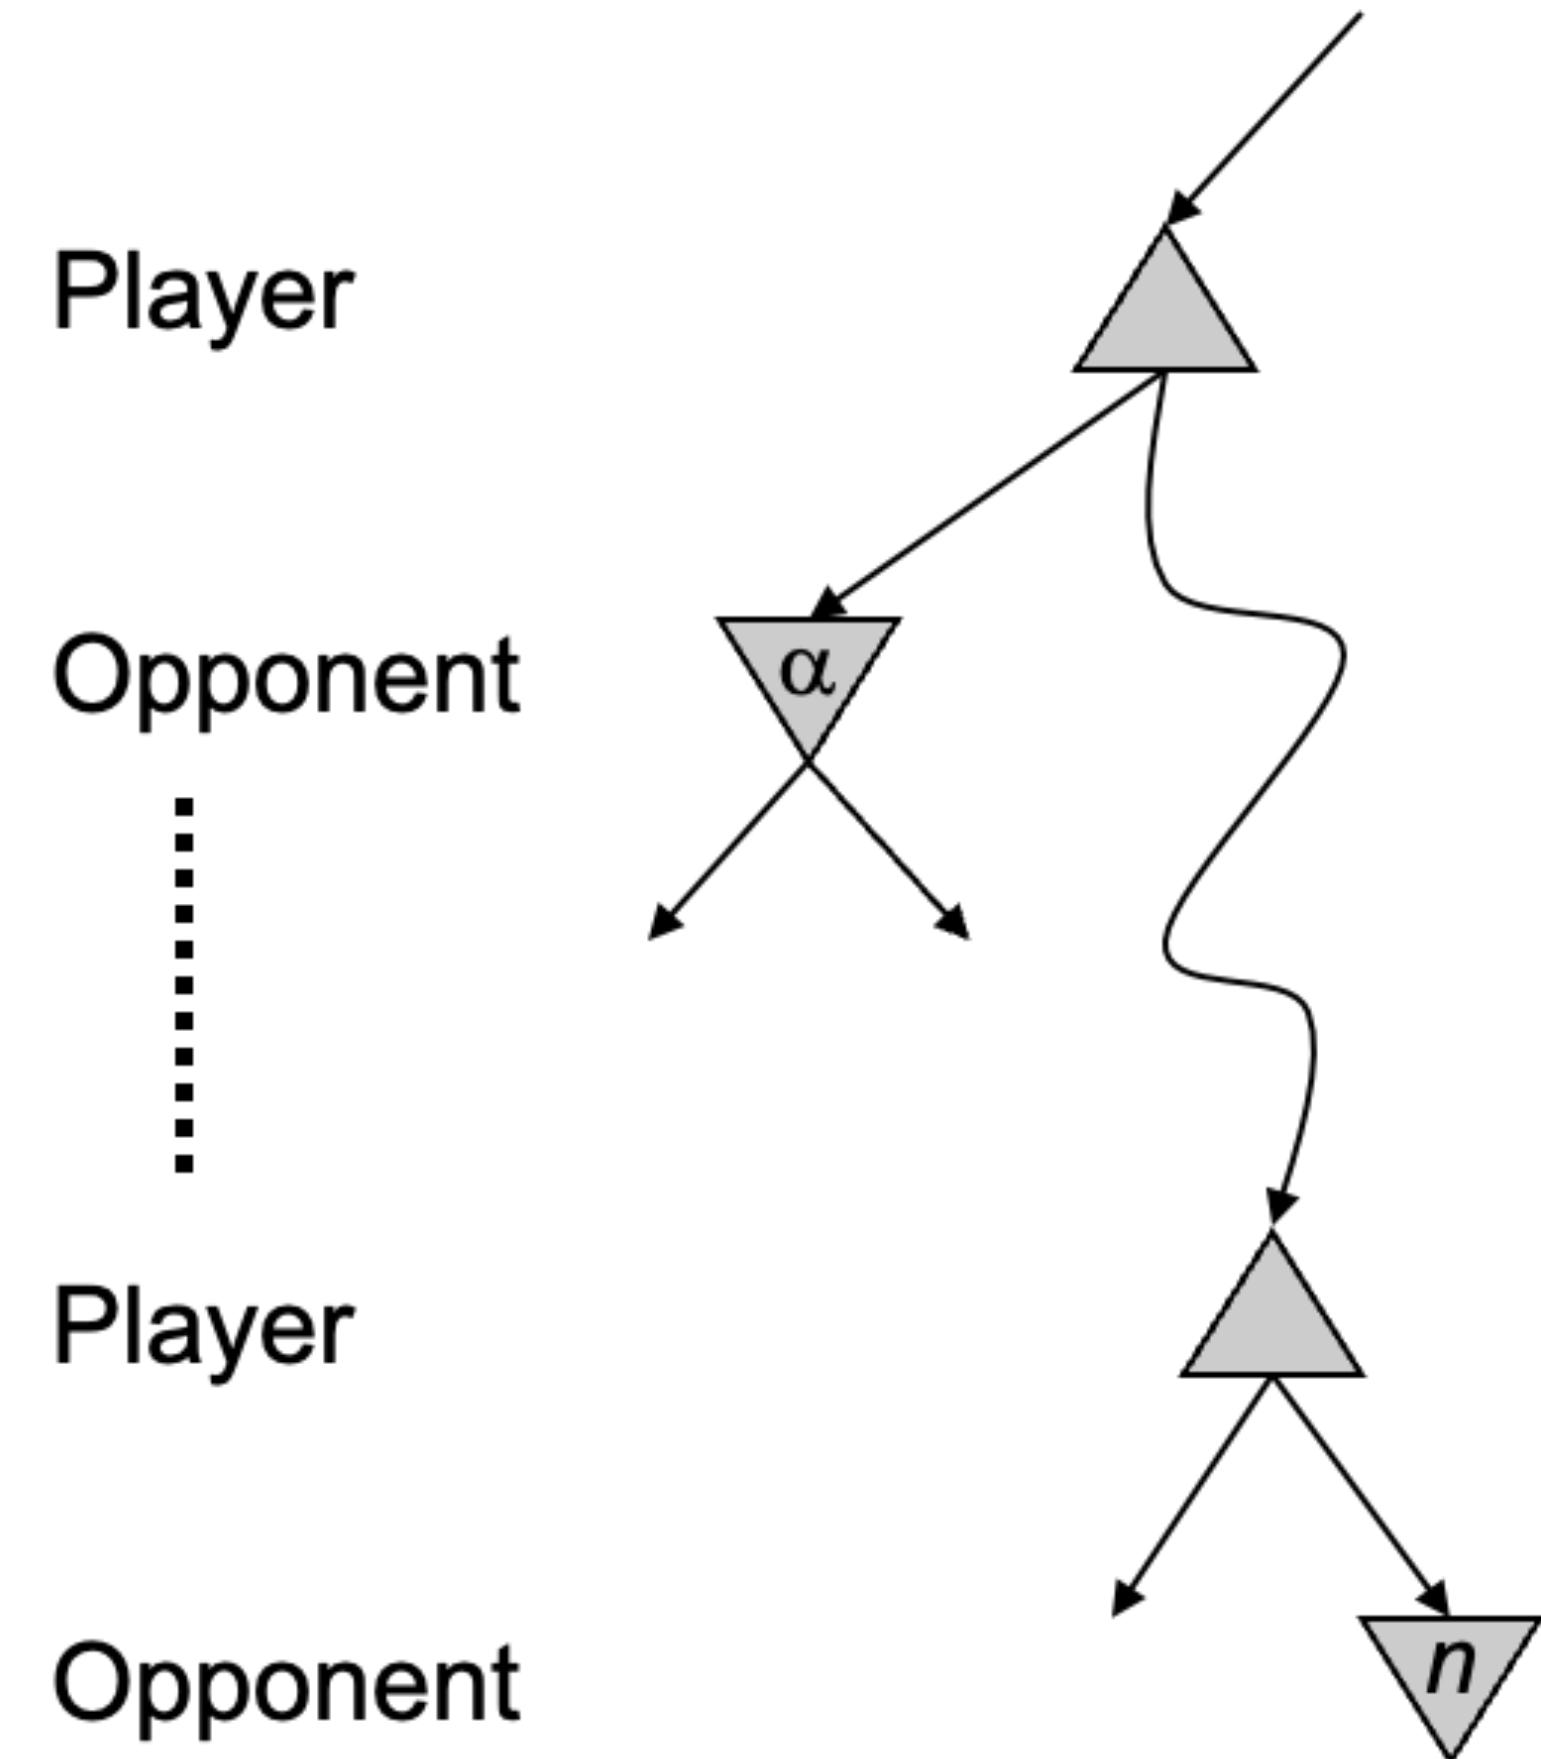
- A way to improve the performance of the Minimax Procedure
- Basic idea: “If you have an idea which is surely bad, don’t take the time to see how truly awful it is” ~ Pat Winston



- We don't need to compute the value at this node.
- No matter what it is it can't effect the value of the root node.

# $\alpha$ - $\beta$ PRUNING

- General case (pruning children of MIN node)
  - We're computing the MIN-VALUE at some node  $n$
  - We're looping over  $n$ 's children
  - $n$ 's estimate of the children's min is dropping
  - Who cares about  $n$ 's value? MAX
  - Let  $a$  be the best value that MAX can get so far at any choice point along the current path from the root
  - If  $n$  becomes worse than  $a$ , MAX will avoid it, so we can prune  $n$ 's other children (it's already bad enough that it won't be played)
- Pruning children of MAX node is symmetric
  - Let  $\beta$  be the best value that MIN can get so far at any choice point along the current path from the root



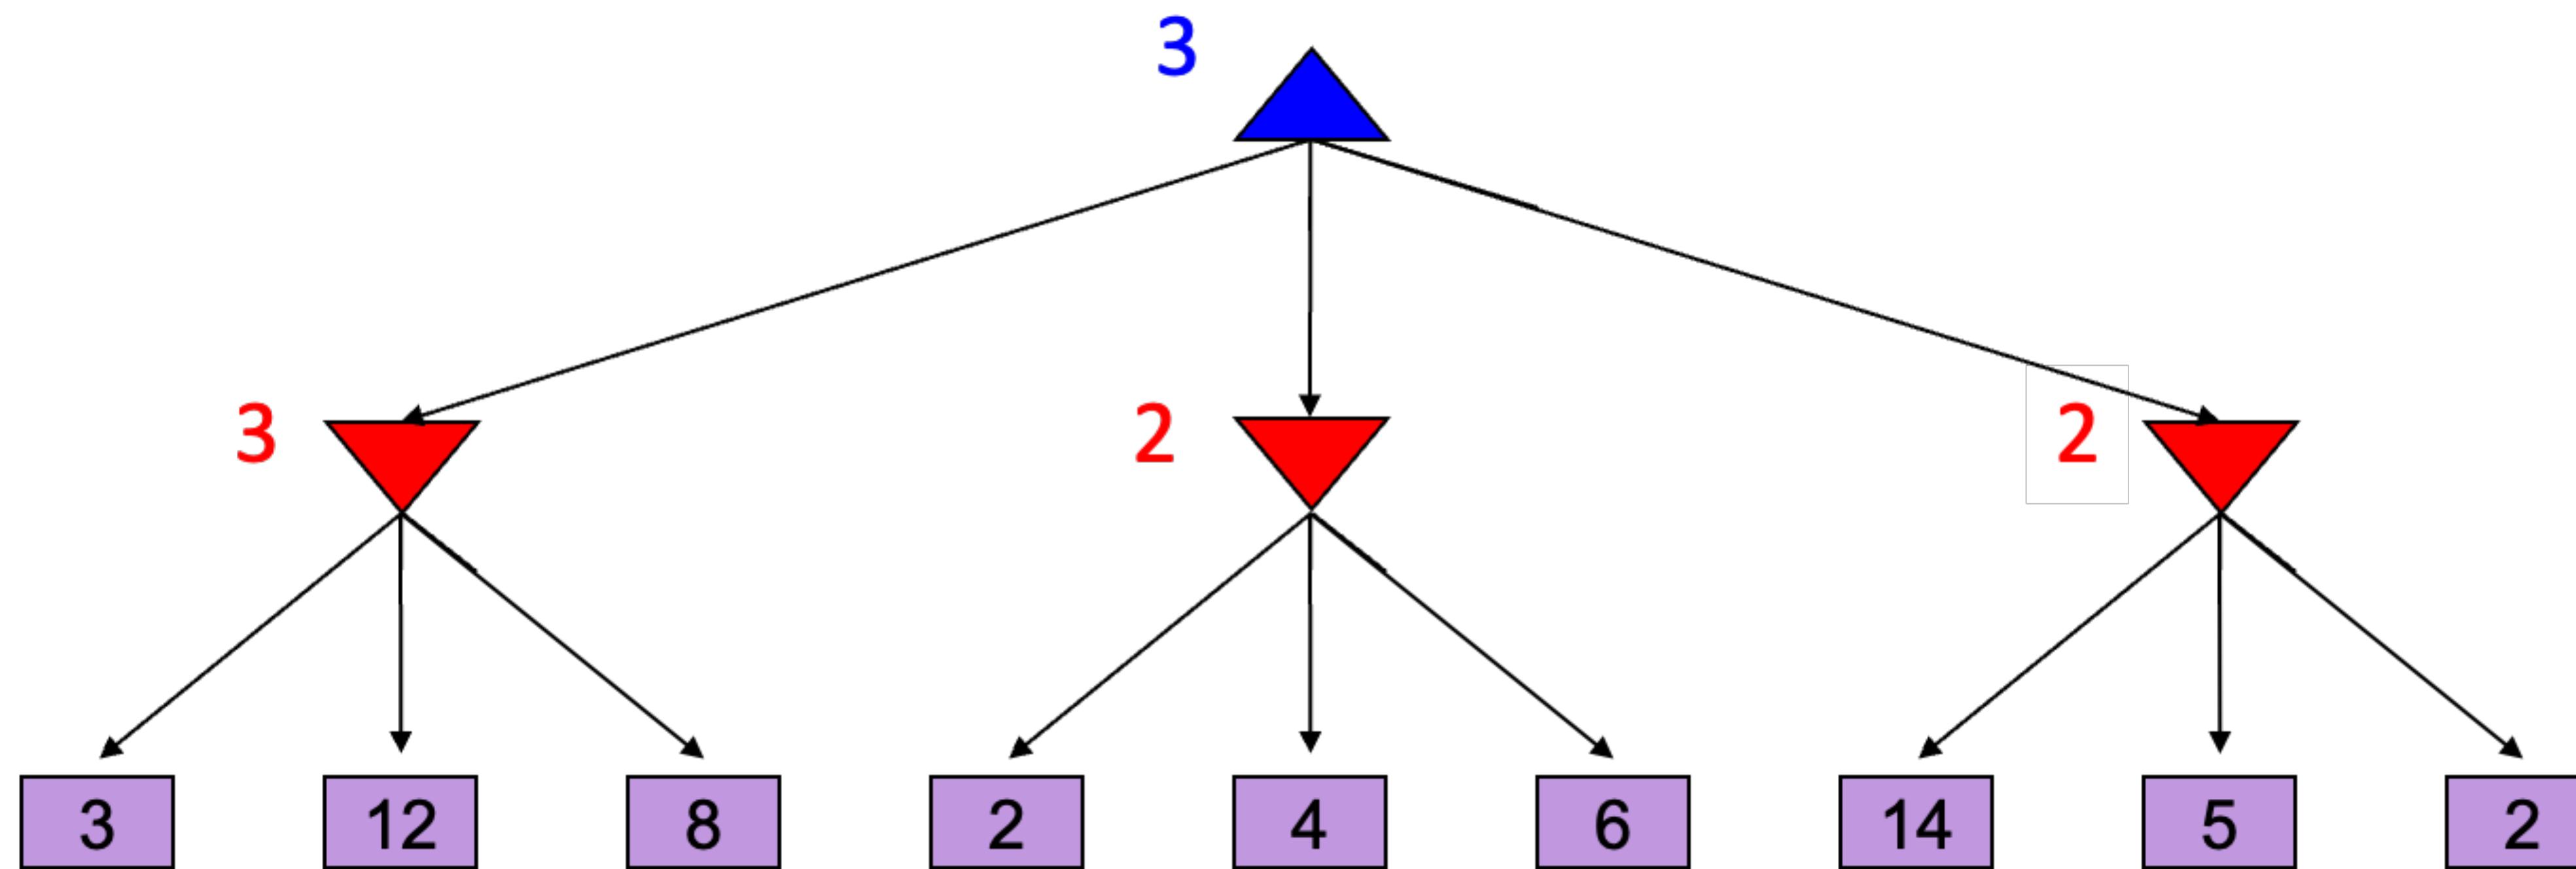
# $\alpha$ - $\beta$ PRUNING ALGORITHM

$\alpha$ : MAX's best option on path to root  
 $\beta$ : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, value(successor,  $\alpha$ ,  $\beta$ ))  
        if v ≥  $\beta$   
            return v  
     $\alpha$  = max( $\alpha$ , v)  
    return v
```

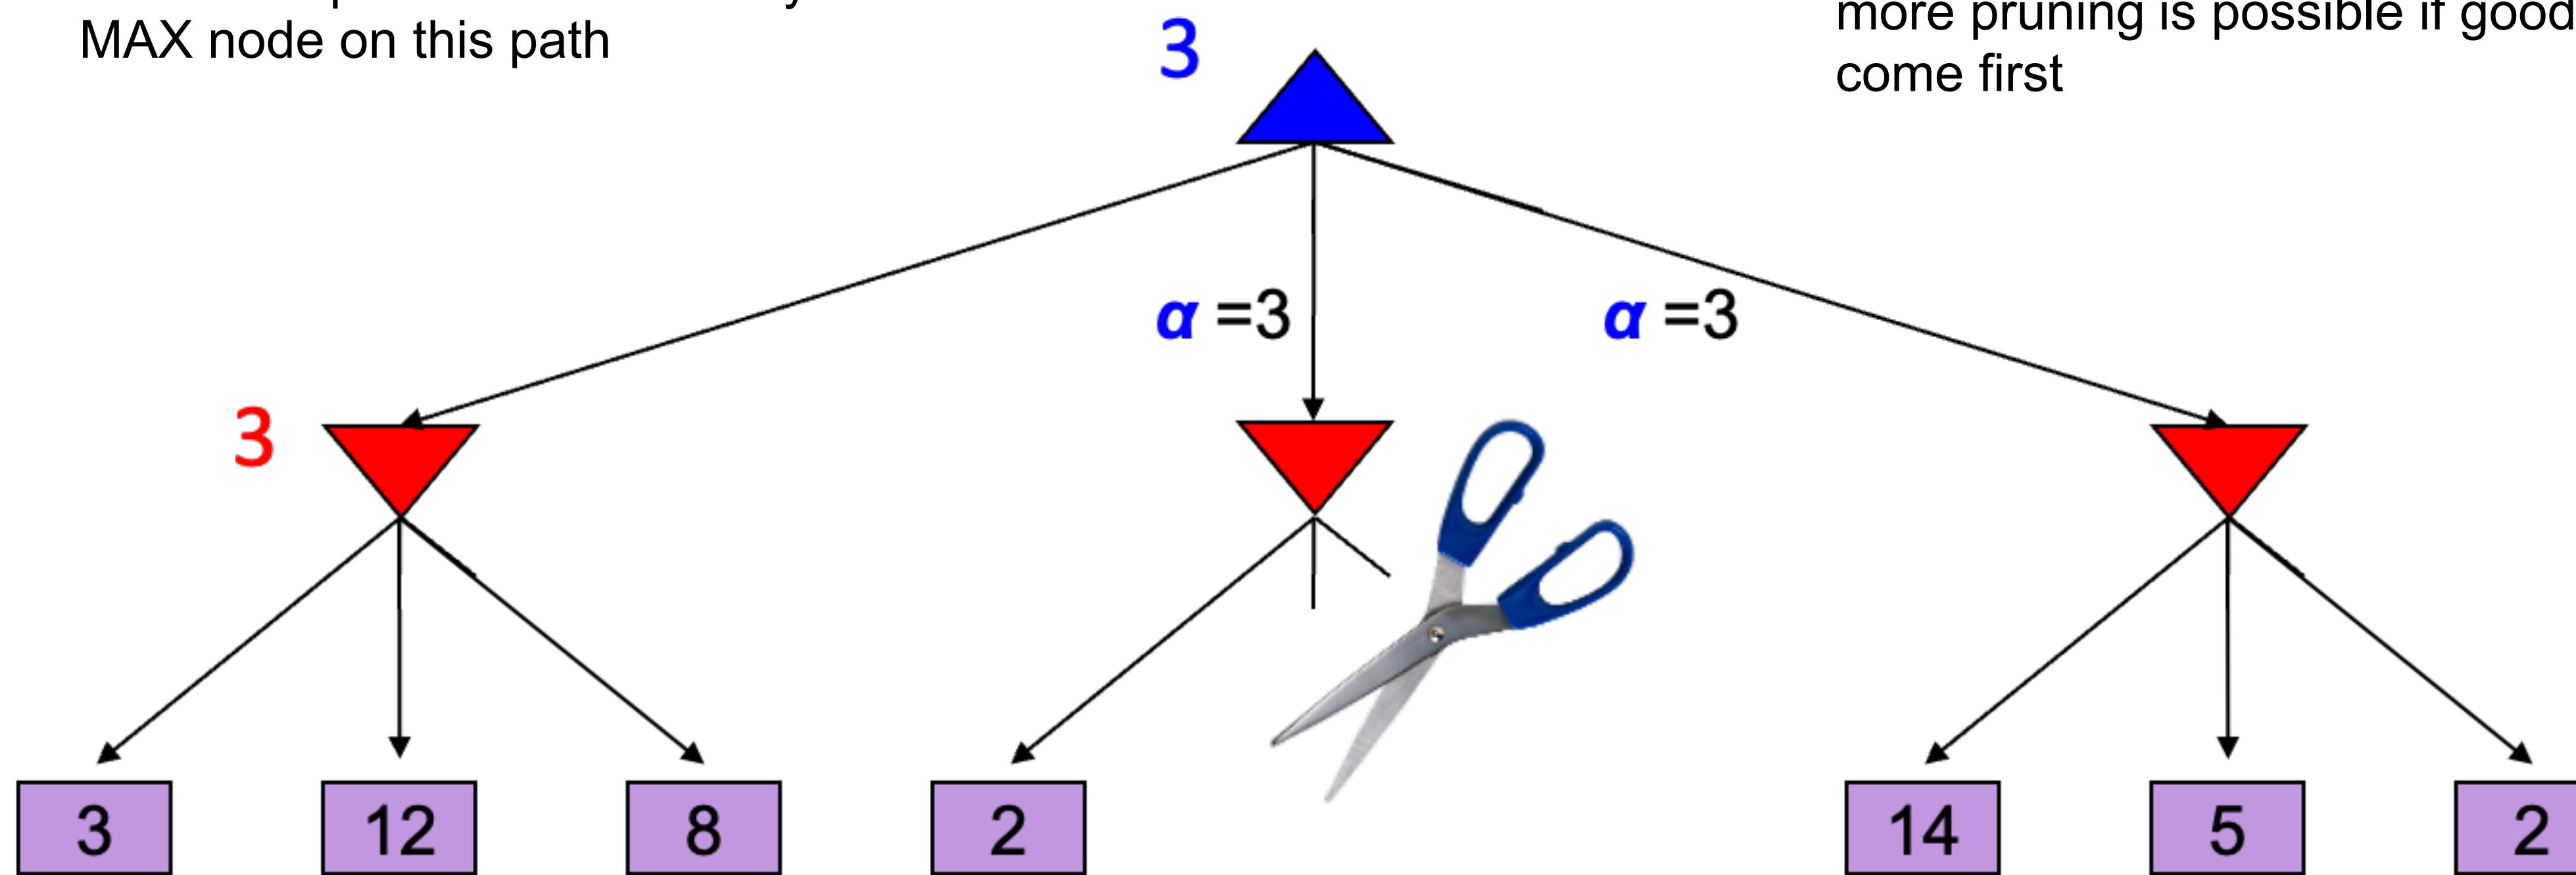
```
def min-value(state ,  $\alpha$ ,  $\beta$ ):  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, value(successor,  $\alpha$ ,  $\beta$ ))  
        if v ≤  $\alpha$   
            return v  
     $\beta$  = min( $\beta$ , v)  
    return v
```

# MINIMAX EXAMPLE



# $\alpha$ - $\beta$ PRUNING EXAMPLE

$\alpha$  = best option so far from any MAX node on this path

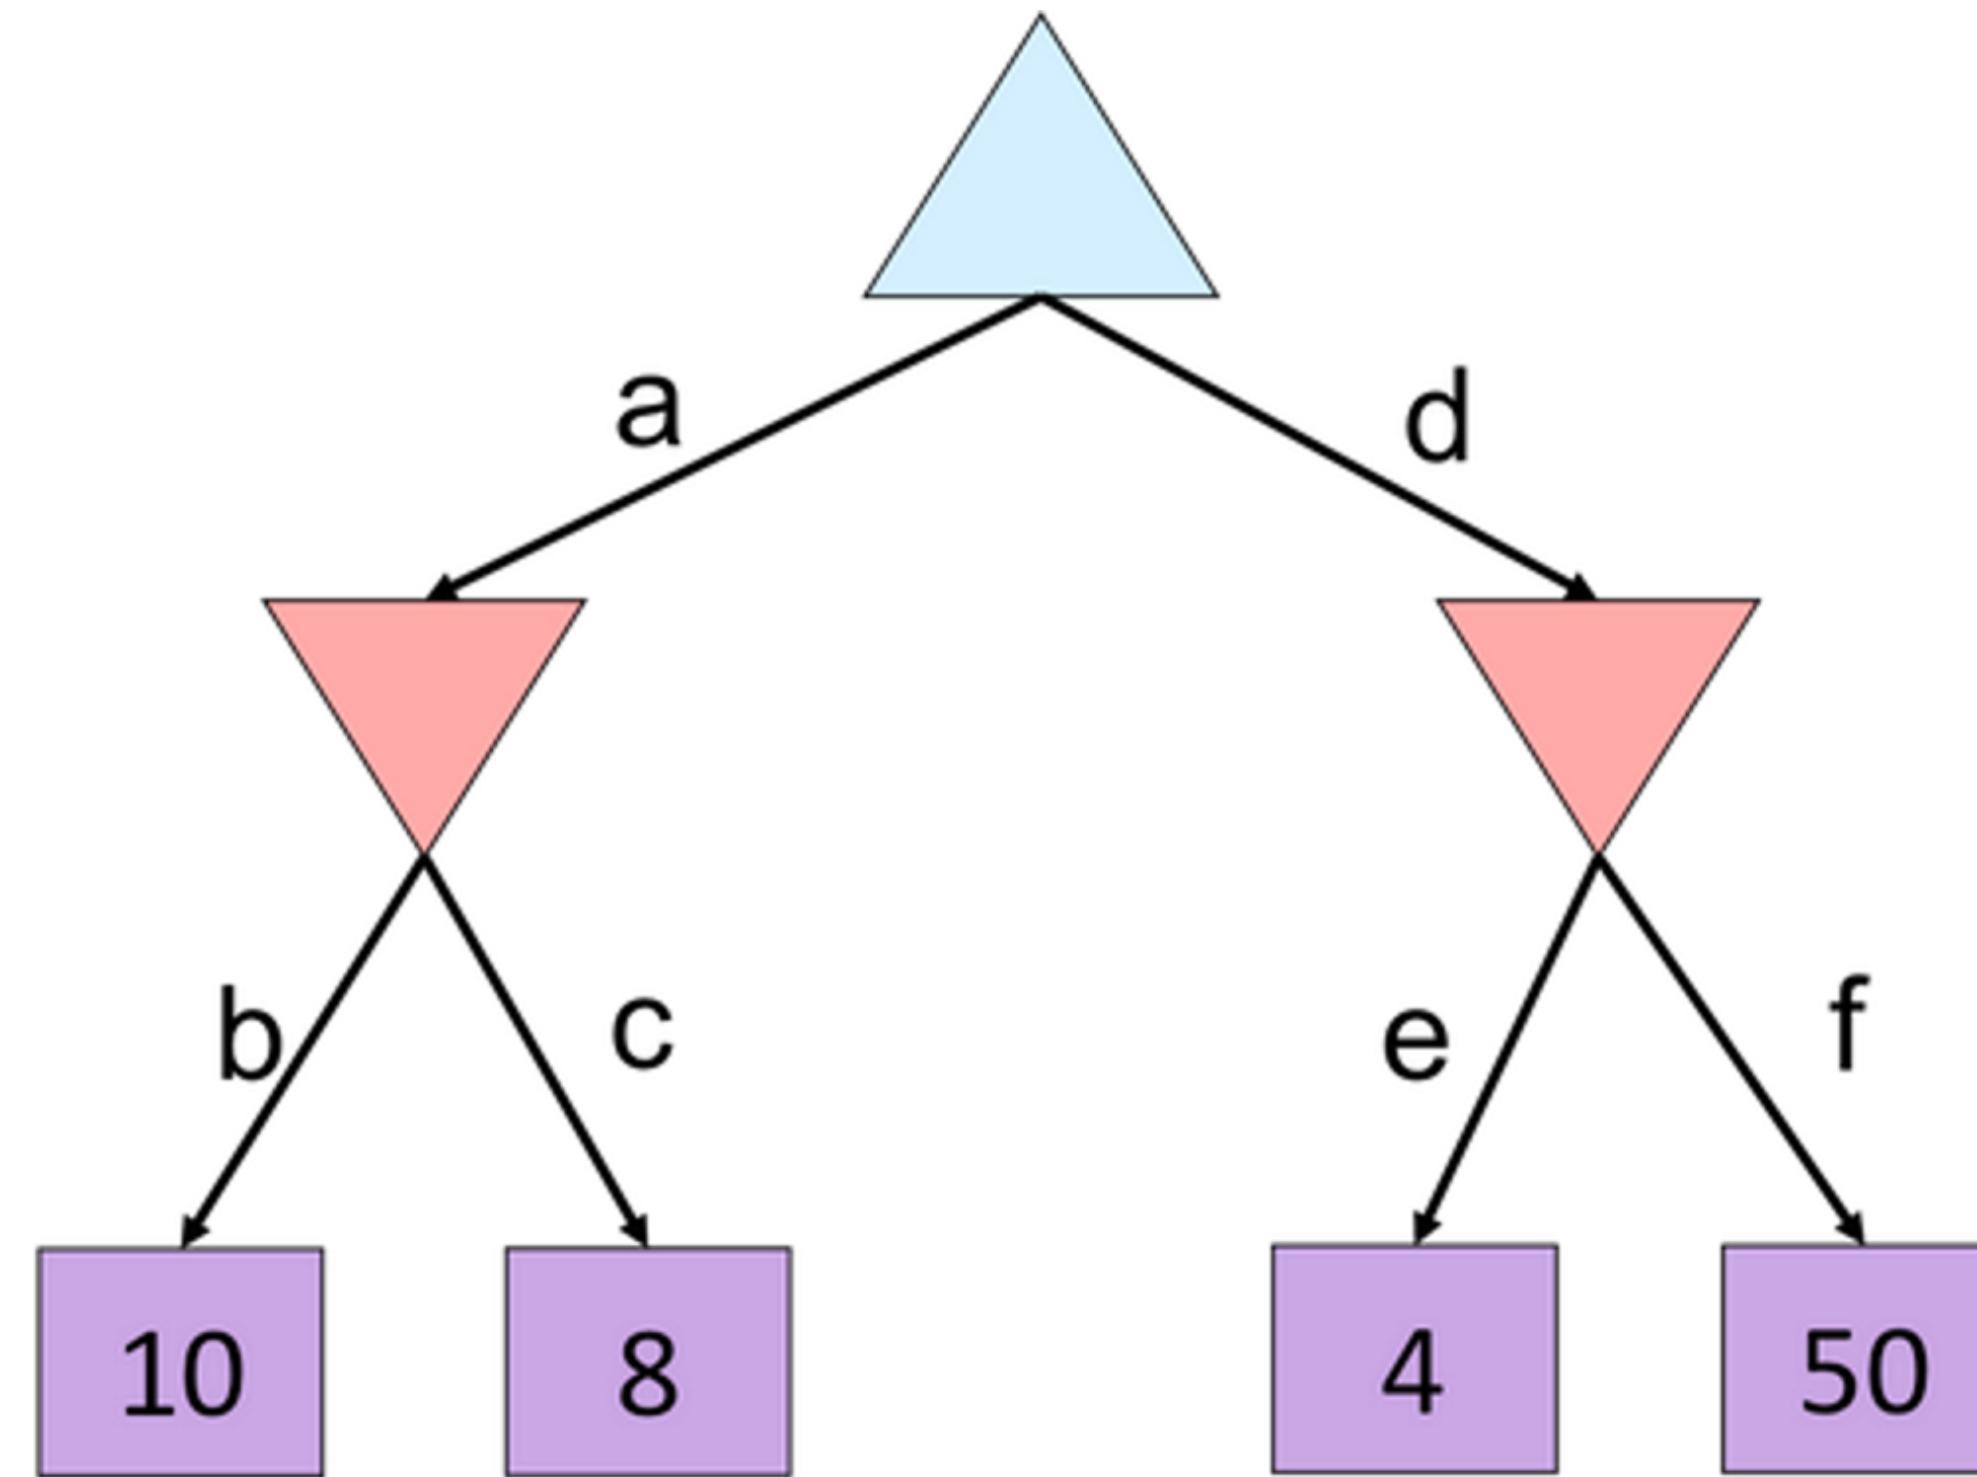


***The order of generation matters:***  
more pruning is possible if good moves come first

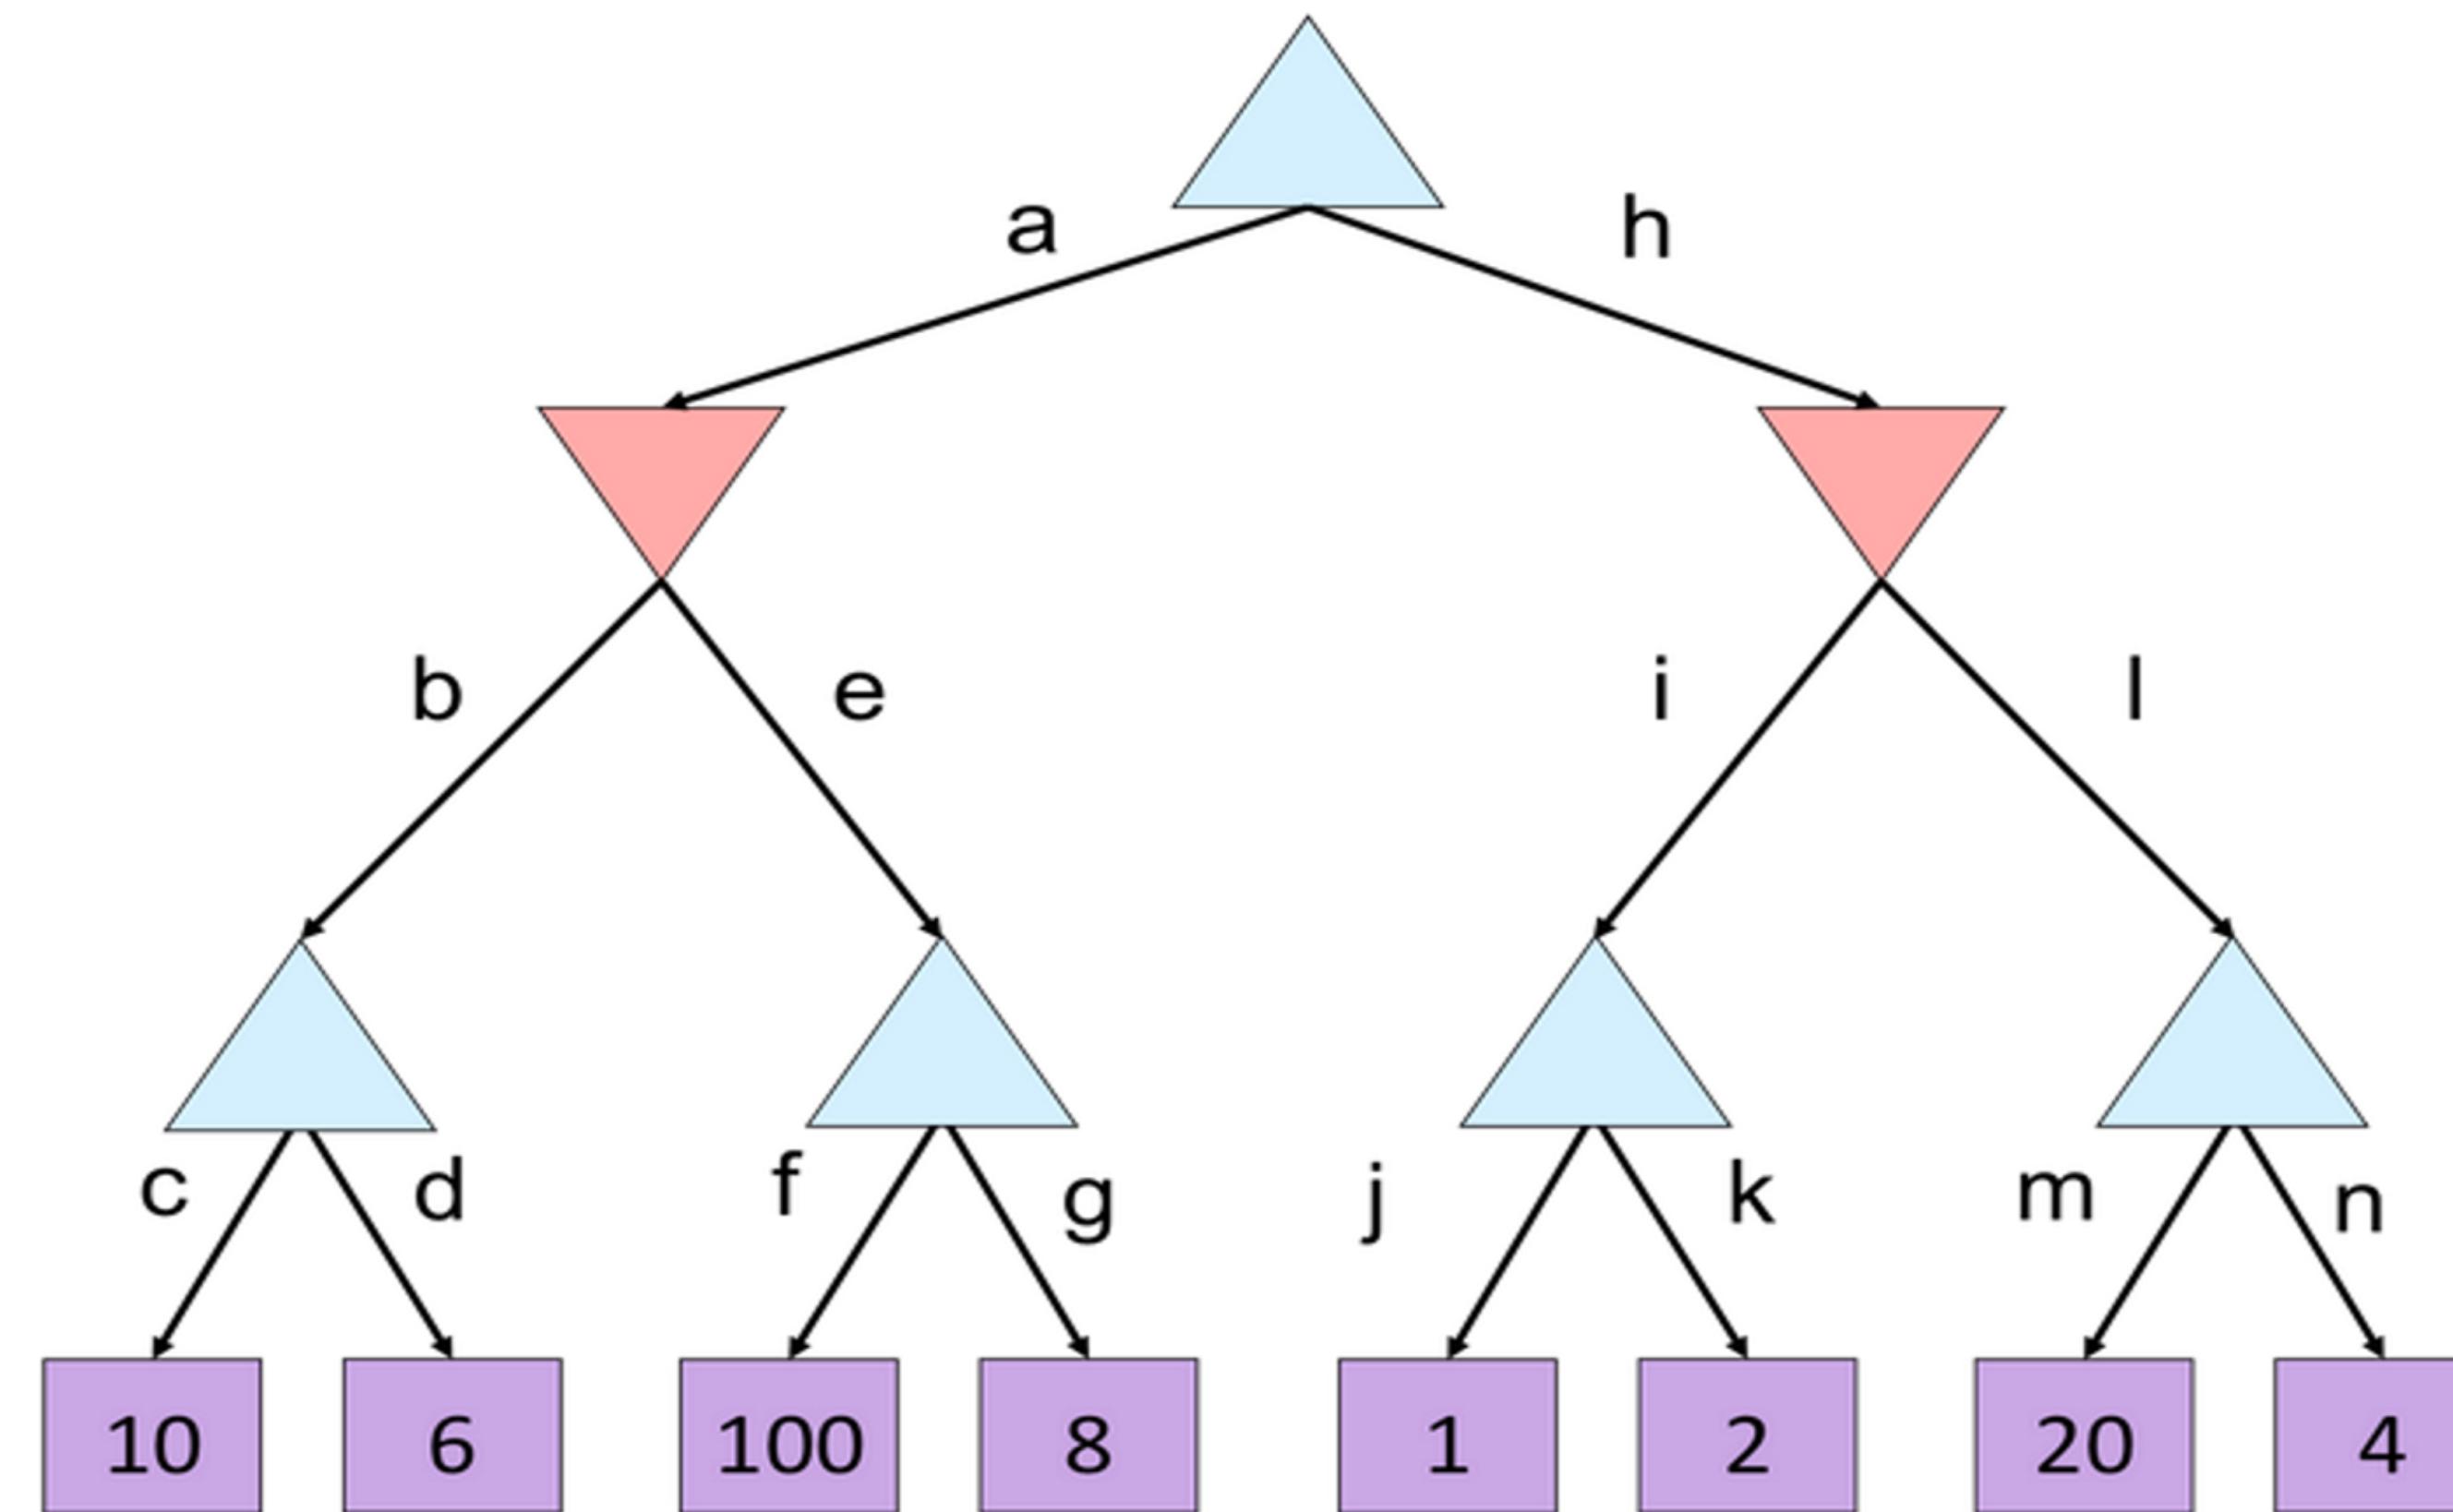
# $\alpha$ - $\beta$ PRUNING PROPERTIES

- Pruning has **no effect** on final result
- Good move ordering improves effectiveness of pruning
- With “perfect ordering”:
  - Time complexity drops to  $O(b^{m/2})$
  - Doubles solvable depth
  - Full search of, e.g. chess, is still hopeless!
- A simple example of **metareasoning**, here reasoning about which computations are relevant
- For chess: only  $35^{50}$  instead of  $35^{100}!!$  Yaaay!!!!

# $a$ - $\beta$ PRUNING QUIZ



# $\alpha$ - $\beta$ PRUNING QUIZ 2



# SUMMARY

- Games require decisions when optimality is impossible
  - Bounded-depth search and approximate evaluation functions
- Games force efficient use of computation
  - Alpha-beta pruning
- Game playing has produced important research ideas
  - Reinforcement learning
  - Iterative deepening
  - Rational metareasoning