




COMPSCI 761: ADVANCED TOPICS IN ARTIFICIAL INTELLIGENCE


ADVERSARIAL SEARCH I

Anna Trofimova, August 2022

RECAP: SEARCH PROBLEM VS CSP VS GAME




5	3			7				
6			1	9	5			
	9	8					6	
8				6				3
4			8		3			1
7				2				6
	6					2	8	
			4	1	9			5
				8			7	9




- "Unpredictable" opponent → specifying a move for every possible opponent reply
- Time limits → unlikely to find optimal solution, must approximate

RECAP: ZERO-SUM GAMES




- Zero-Sum Games
 - Agents have **opposite** utilities
 - Pure competition:
 - One **maximizes**, the other **minimizes**



- General Games
 - Agents have **independent** utilities
 - Cooperation, indifference, competition, shifting alliances, and more are all possible

RECAP: DETERMINISTIC TWO-PLAYER

- E.g. tic-tac-toe, chess, checkers
- Minimax search
 - A state-space search tree
 - Players alternate
 - Each layer, or ply, consists of a round of moves
 - Choose move to position with highest **minimax value** = best achievable utility against best play
- Zero-sum games
 - One player maximizes result
 - The other minimizes result




RECAP: MINIMAX PROPERTIES

- Optimal against a perfect player. Otherwise?
- Time complexity?
 - $O(b^m)$
 - m = maximum depth of search tree, b = branching factor
- Space complexity?
 - $O(bm)$
- For chess, $b \sim 35$, $m \sim 100$
 - Exact solution is completely infeasible
 - But, do we need to explore the whole tree?

TODAY


- α - β pruning
- Expectimax

GAME TREE PRUNING



α - β PRUNING


- A way to improve the performance of the Minimax Procedure
- Basic idea: “If you have an idea which is surely bad, don’t take the time to see how truly awful it is” ~ Pat Winston



- We don't need to compute the value at this node.
- No matter what it is it can't effect the value of the root node.

α - β PRUNING

- General case (pruning children of **MIN** node)
 - We're computing the **MIN-VALUE** at some node n
 - We're looping over n 's children
 - n 's estimate of the children's min is dropping
 - Who cares about n 's value? **MAX**
 - Let α be the best value that **MAX** can get so far at any choice point along the current path from the root
 - If n becomes worse than α , MAX will avoid it, so we can prune n 's other children (it's already bad enough that it won't be played)
- Pruning children of **MAX** node is symmetric
 - Let β be the best value that **MIN** can get so far at any choice point along the current path from the root



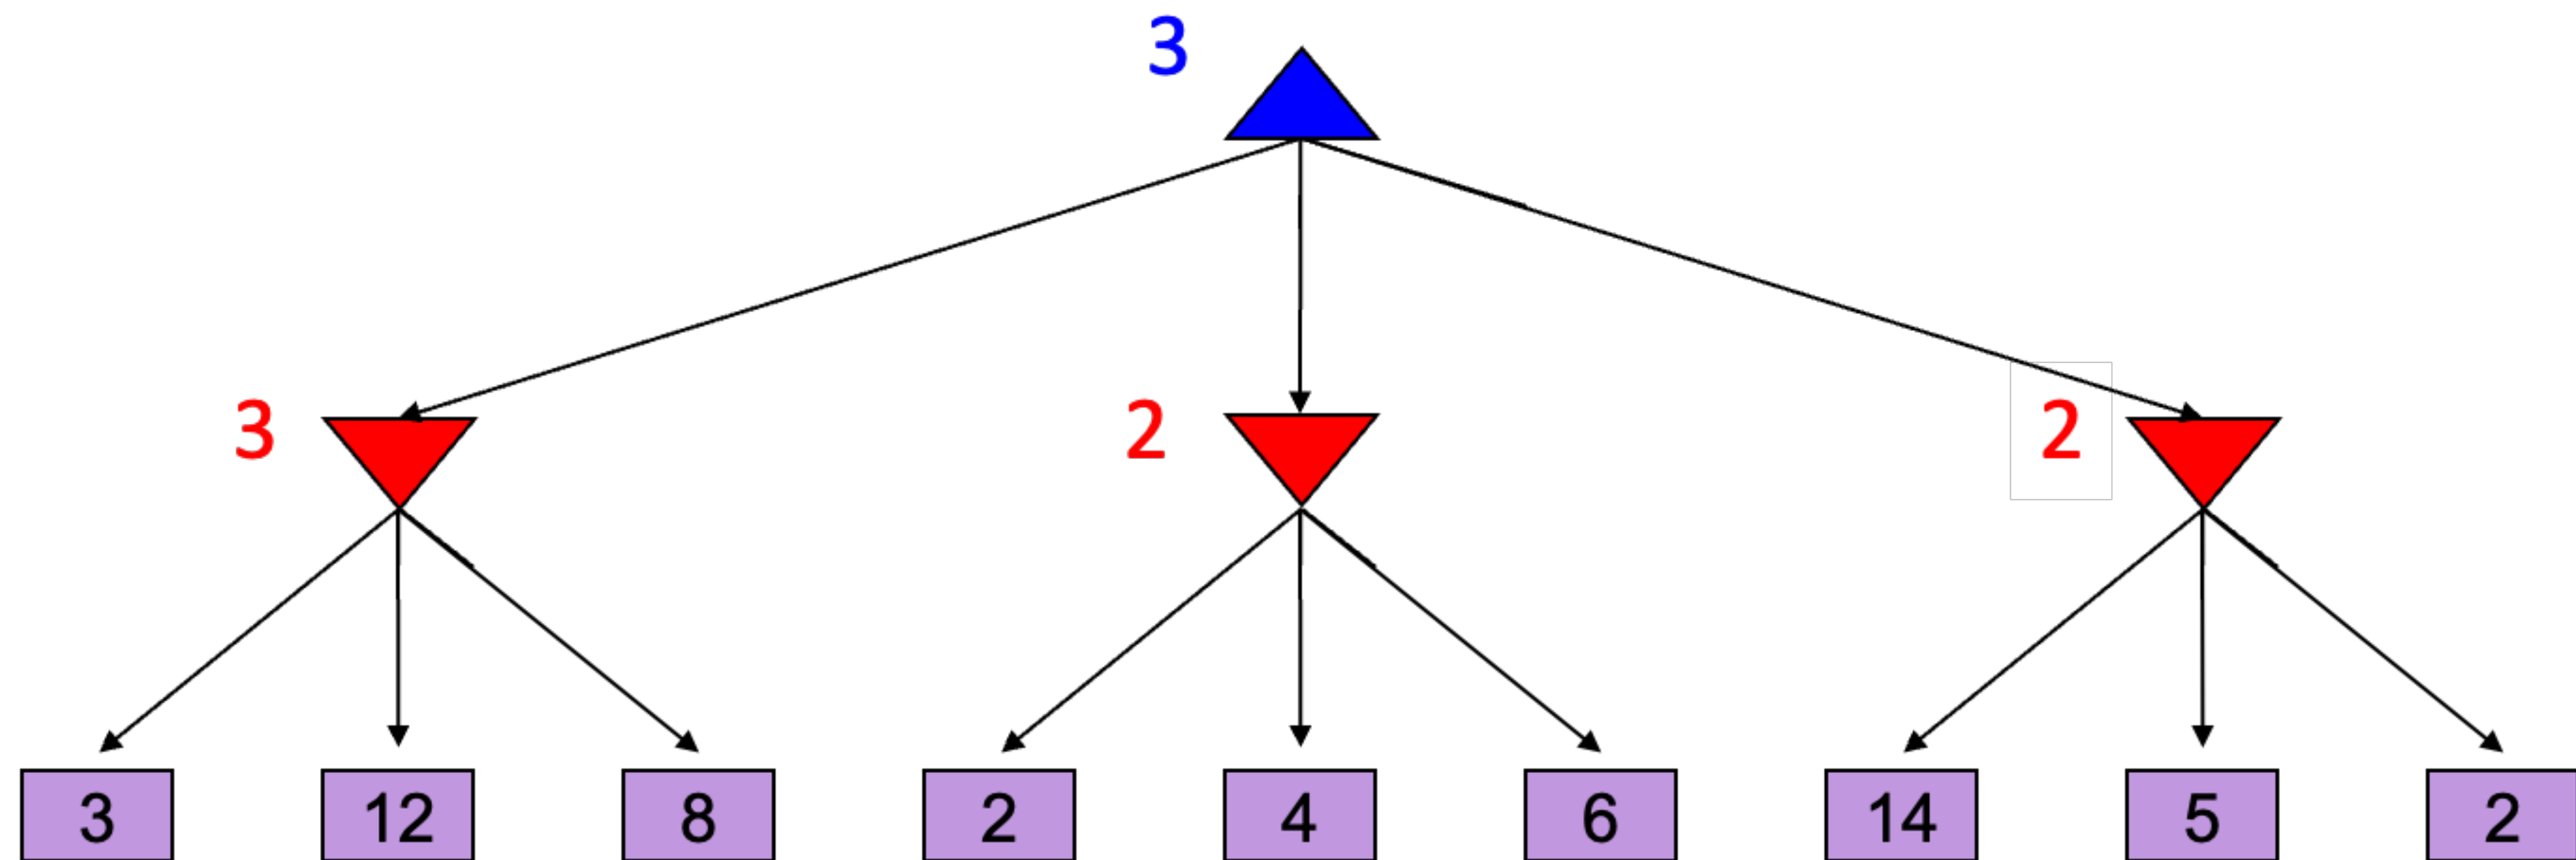
α - β PRUNING ALGORITHM

α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$   
            return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$   
            return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```


MINIMAX EXAMPLE




α - β PRUNING EXAMPLE

α = best option so far from any MAX node on this path

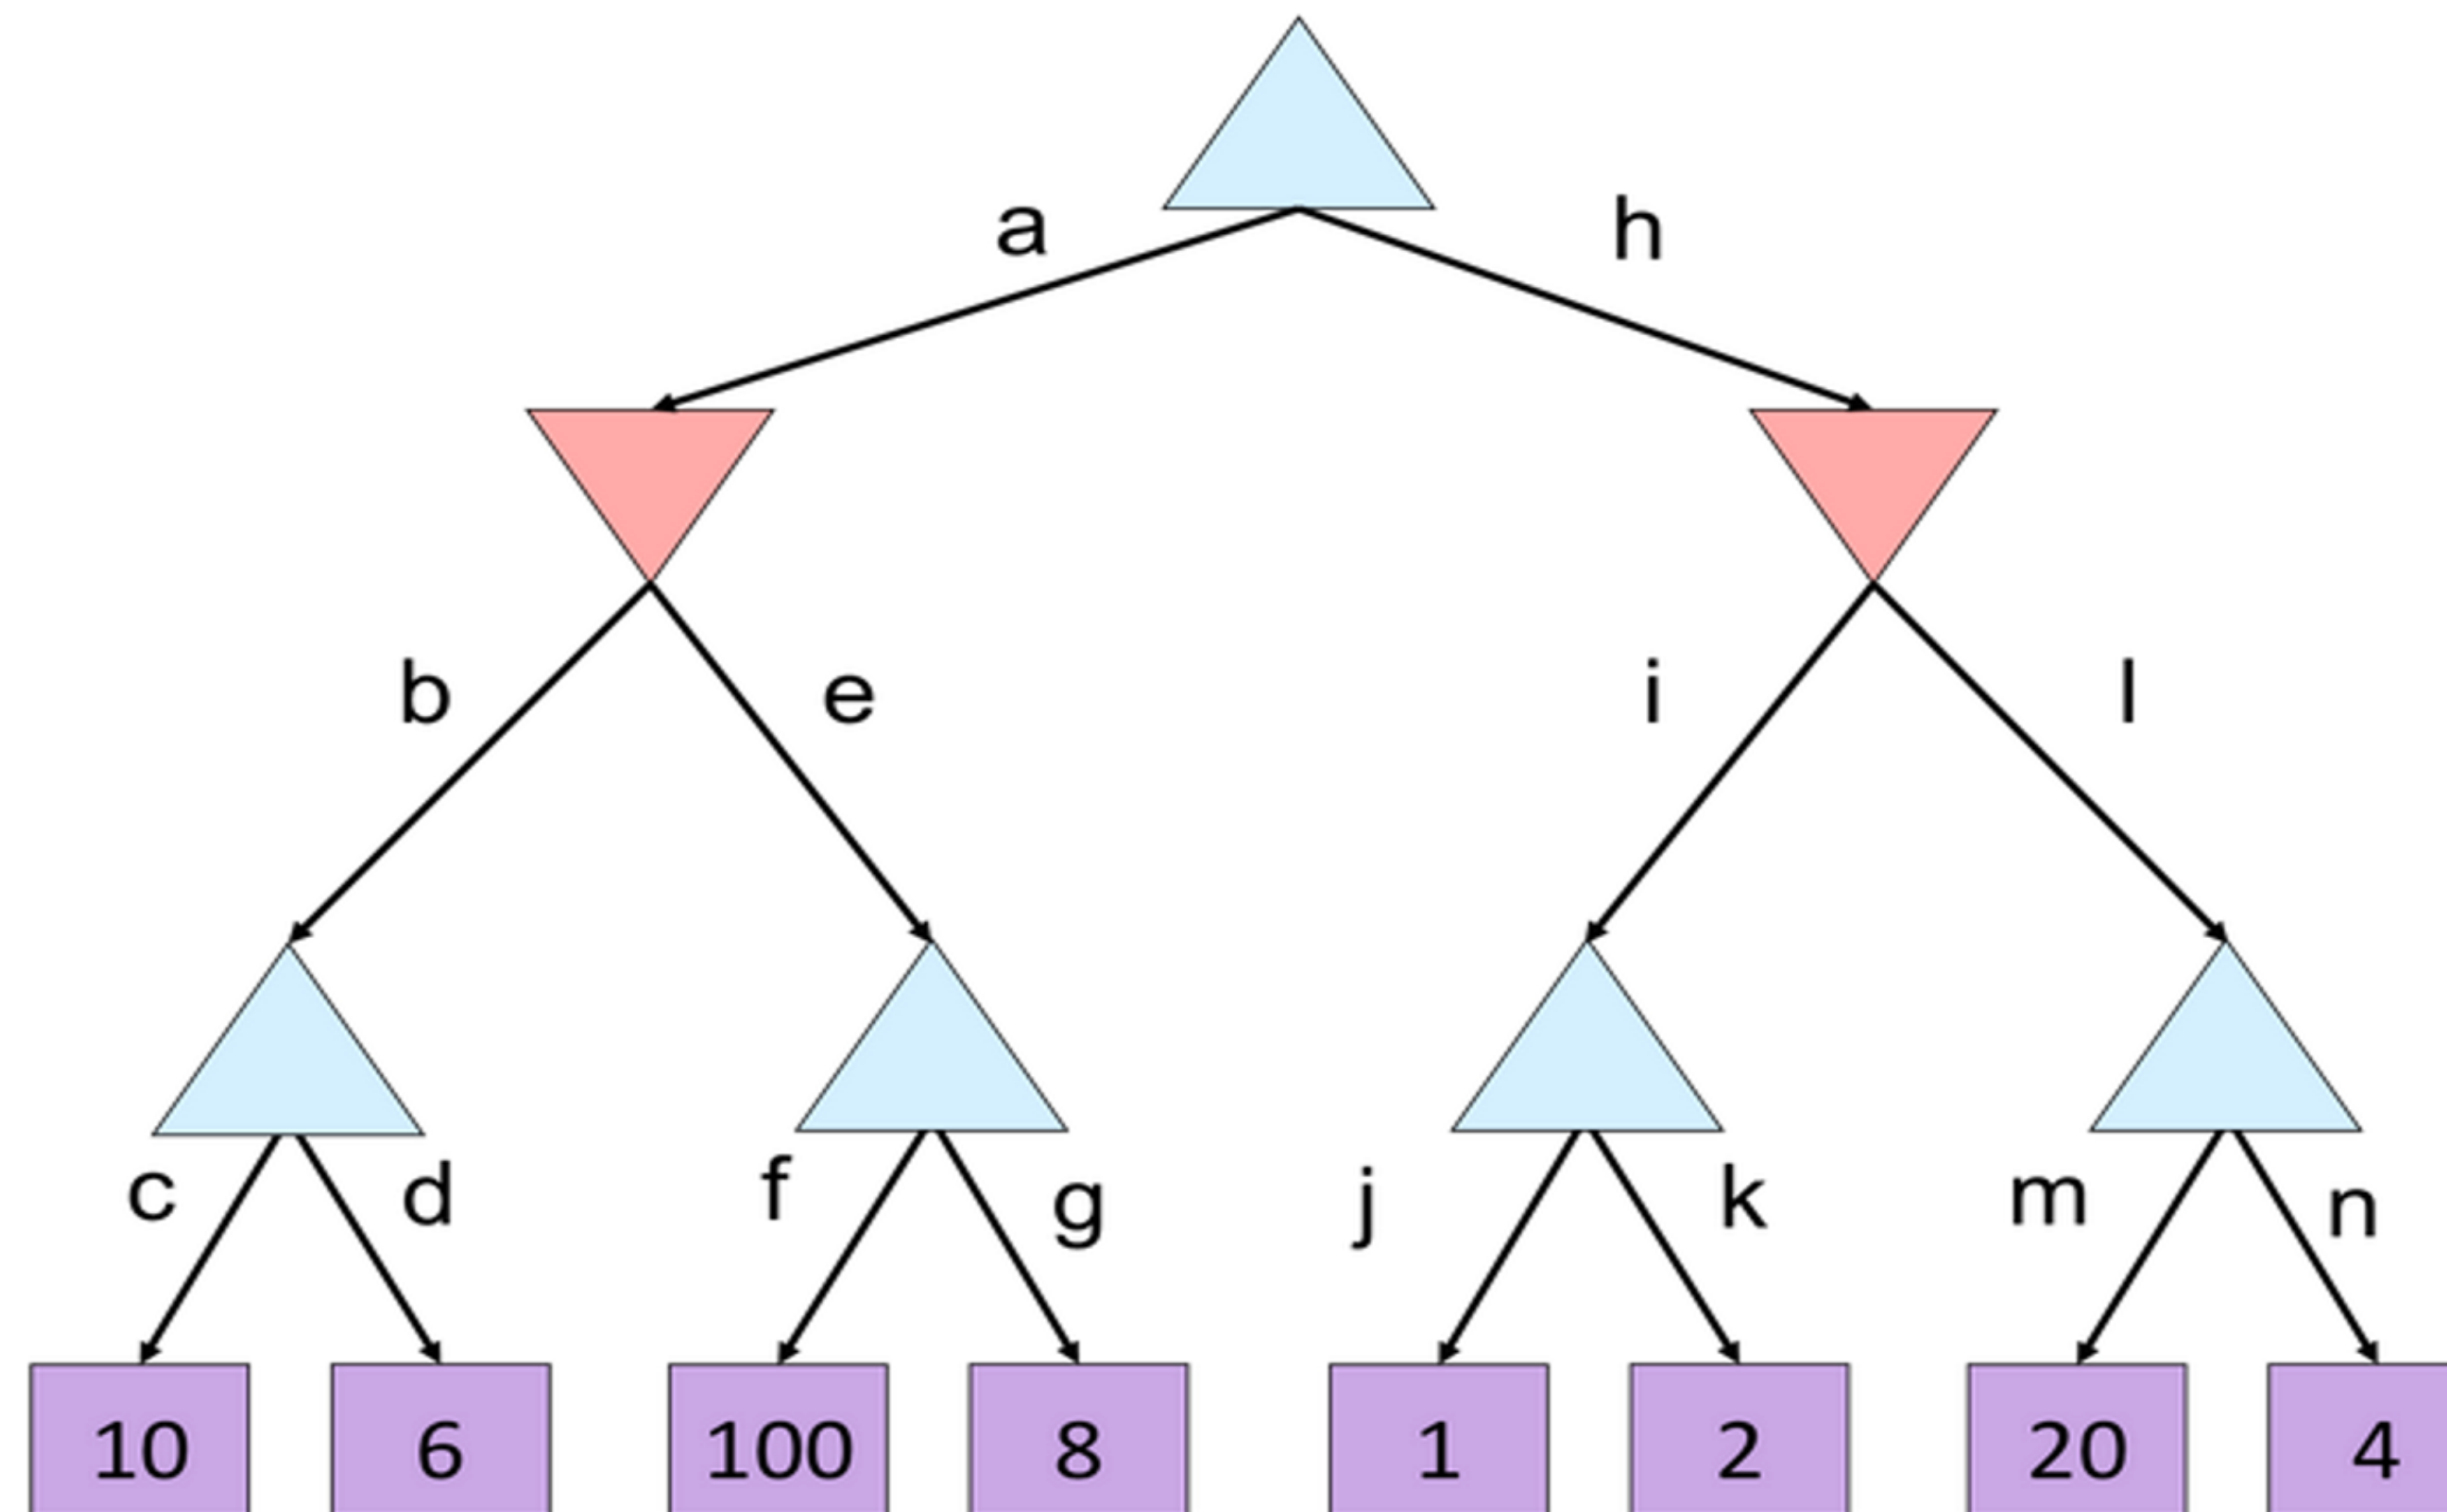
The order of generation matters:
more pruning is possible if good moves come first



α - β PRUNING QUIZ



α - β PRUNING QUIZ 2



α - β PRUNING PROPERTIES

- Pruning has **no effect** on final result
- Good move ordering improves effectiveness of pruning
- With “perfect ordering”:
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth
 - Full search of, e.g. chess, is still hopeless!
- A simple example of **metareasoning**, here reasoning about which computations are relevant
- For chess: only 35^{50} instead of 35^{100} !! Yaaay!!!! Still not feasible...



“


The whole question of making an automaton play any game depended upon the possibility of the machine being able to *represent all the myriads of combinations* relating to it.

”

- Charles Babbage


RESOURCE LIMITS

- Cannot search to leaves
- Limited search
 - Instead, search a limited depth of the tree
 - Replace terminal utilities with an eval function for non-terminal positions
- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Example:
 - Suppose we have 100 seconds, can explore 10K nodes / sec
 - So can check 1M nodes per move
 - α - β reaches about depth 8 – decent chess program



EVALUATION FUNCTION

- Function which scores non-terminals




- Ideal function: returns the utility of the position
- In practice: typically weighted linear sum of features:

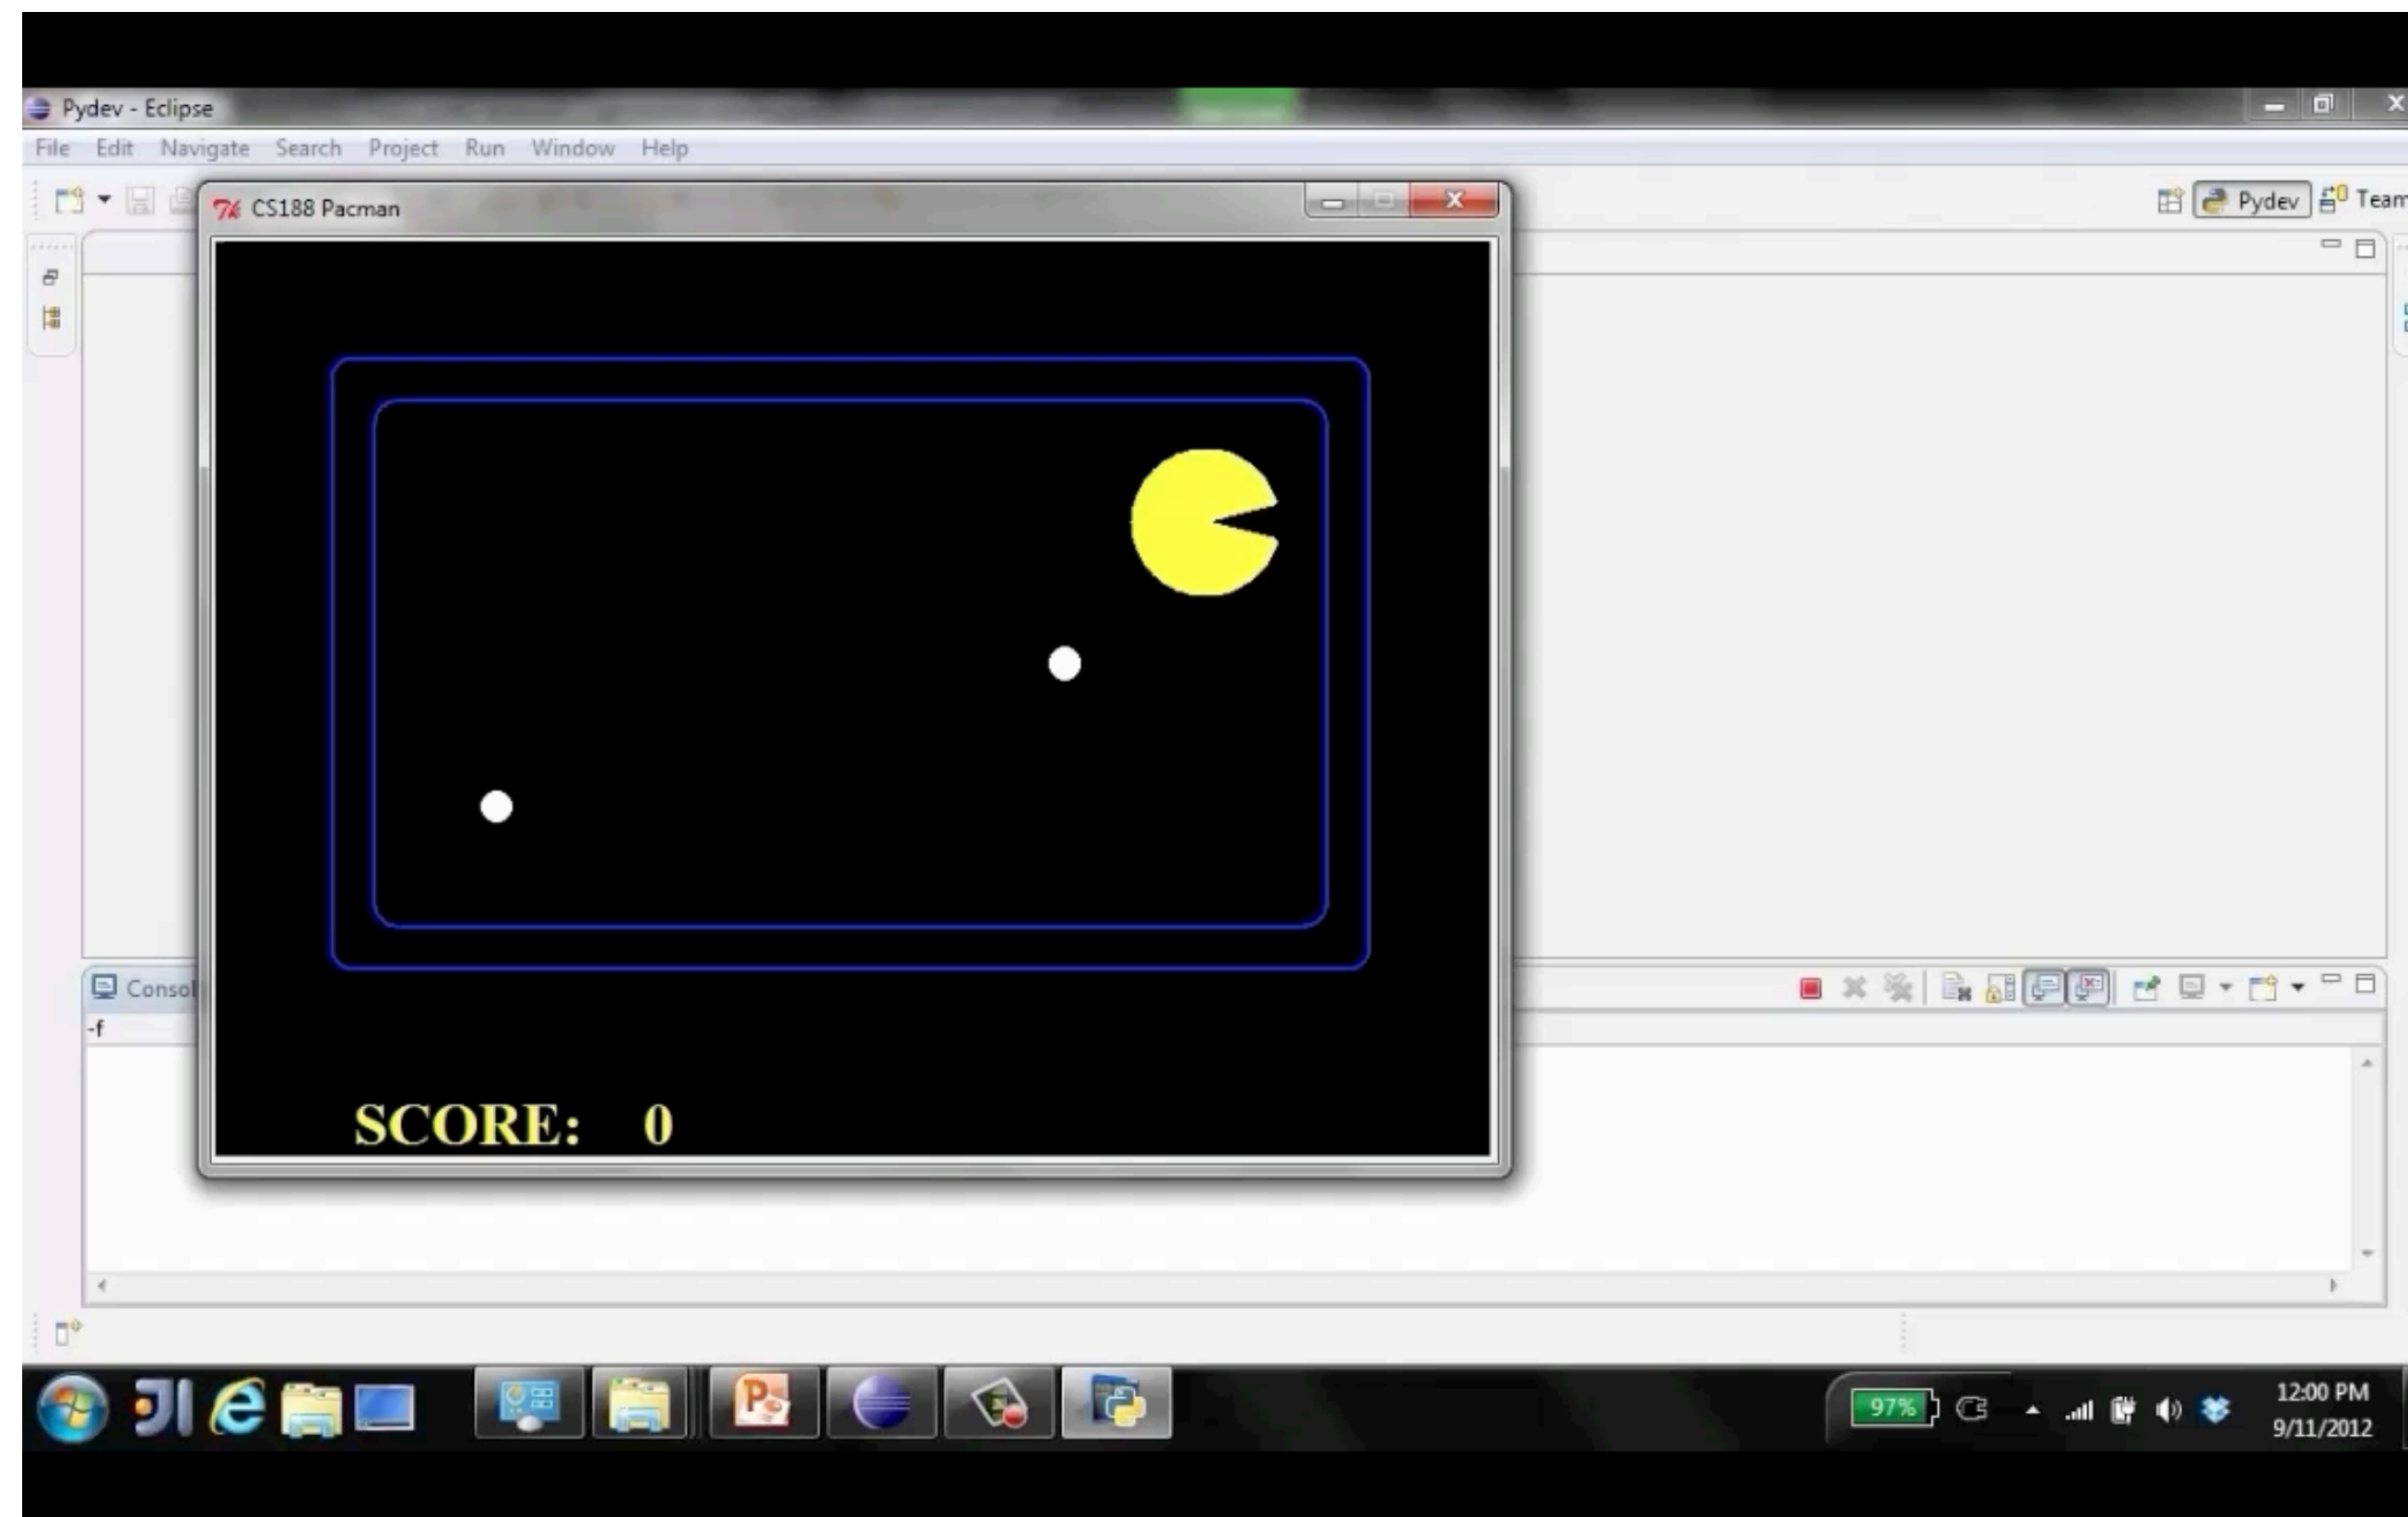
$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g. $f_1(s) = (\text{num white queens} - \text{num black queens})$, etc.
- Or a more complex nonlinear function (e.g., NN) trained by self-play RL

EVALUATION FUNCTION FOR PACMAN?




VIDEO OF DEMO THRASHING ($D=2$)



VIDEO OF DEMO THRASHING (D=2) FIXED



Why Pacman Starves



- A danger of replanning agents!
 - He knows his score will go up by eating the dot now (west, east)
 - He knows his score will go up just as much by eating the dot later (east, west)
 - There are no point-scoring opportunities after eating the dot (within the horizon, two here)
 - Therefore, waiting seems just as good as eating: he may go east, then back west in the next round of replanning!

DEPTH MATTERS

- Evaluation functions are always imperfect
- Deeper search => better play (usually)
- Or, deeper search gives same quality of play with a less accurate evaluation function




ITERATIVE DEEPENING

- Iterative deepening uses DFS as a subroutine:
 1. Do a DFS which only searches for paths of length 1 or less. (DFS gives up on any path of length 2)
 2. If “1” failed, do a DFS which only searches paths of length 2 or less.
 3. If “2” failed, do a DFS which only searches paths of length 3 or less.
-and so on.
- This works for single-agent search as well!
- Why do we want to do this for multiplayer games?




THE STORY SO FAR

- Focus on two-player, zero-sum, deterministic, observable, turn-taking games
- Minimax defines rational behavior
- Recursive DFS implementation: space complexity $O(bm)$, time complexity $O(b^m)$
- Alpha-beta pruning with good node ordering reduces time complexity to $O(b^{m/2})$
- Still nowhere close to solving chess, let alone Go or StarCraft

GAMES WITH UNCERTAIN OUTCOMES






WORST-CASE VS. AVERAGE CASE





WORST-CASE VS. AVERAGE CASE


- Idea: uncertain outcomes controlled by chance, not an adversary!




CHANCE OUTCOMES IN TREES



Tictactoe, chess
Minimax



Tetris, investing
Expectimax



Backgammon, Monopoly
Expectiminimax

EXPECTIMAX SEARCH

- Why wouldn't we know what the result of an action will be?
 - Explicit randomness: rolling dice
 - Unpredictable opponents: the ghosts respond randomly
 - Actions can fail: when moving a robot, wheels might slip
- Values should now reflect average-case (expectimax) outcomes, not worst-case (minimax) outcomes
- **Expectimax search:** compute the average score under optimal play
 - Max nodes as in minimax search
 - Chance nodes are like min nodes but the outcome is uncertain
 - Calculate their expected utilities
 - I.e. take weighted average (expectation) of children
- Later, we'll learn how to formalize the underlying uncertain-result problems as **Markov Decision Processes**

MINIMAX SEARCH

```
function minimax-decision(s) returns an action
  return the action a in Actions(s) with the highest
    minimax_value(Result(s,a))
```



```
function minimax_value(s) returns a value
  if Terminal-Test(s) then return Utility(s)
  if Player(s) = MAX then return maxa in Actions(s) minimax_value(Result(s,a))
  if Player(s) = MIN then return mina in Actions(s) minimax_value(Result(s,a))
```

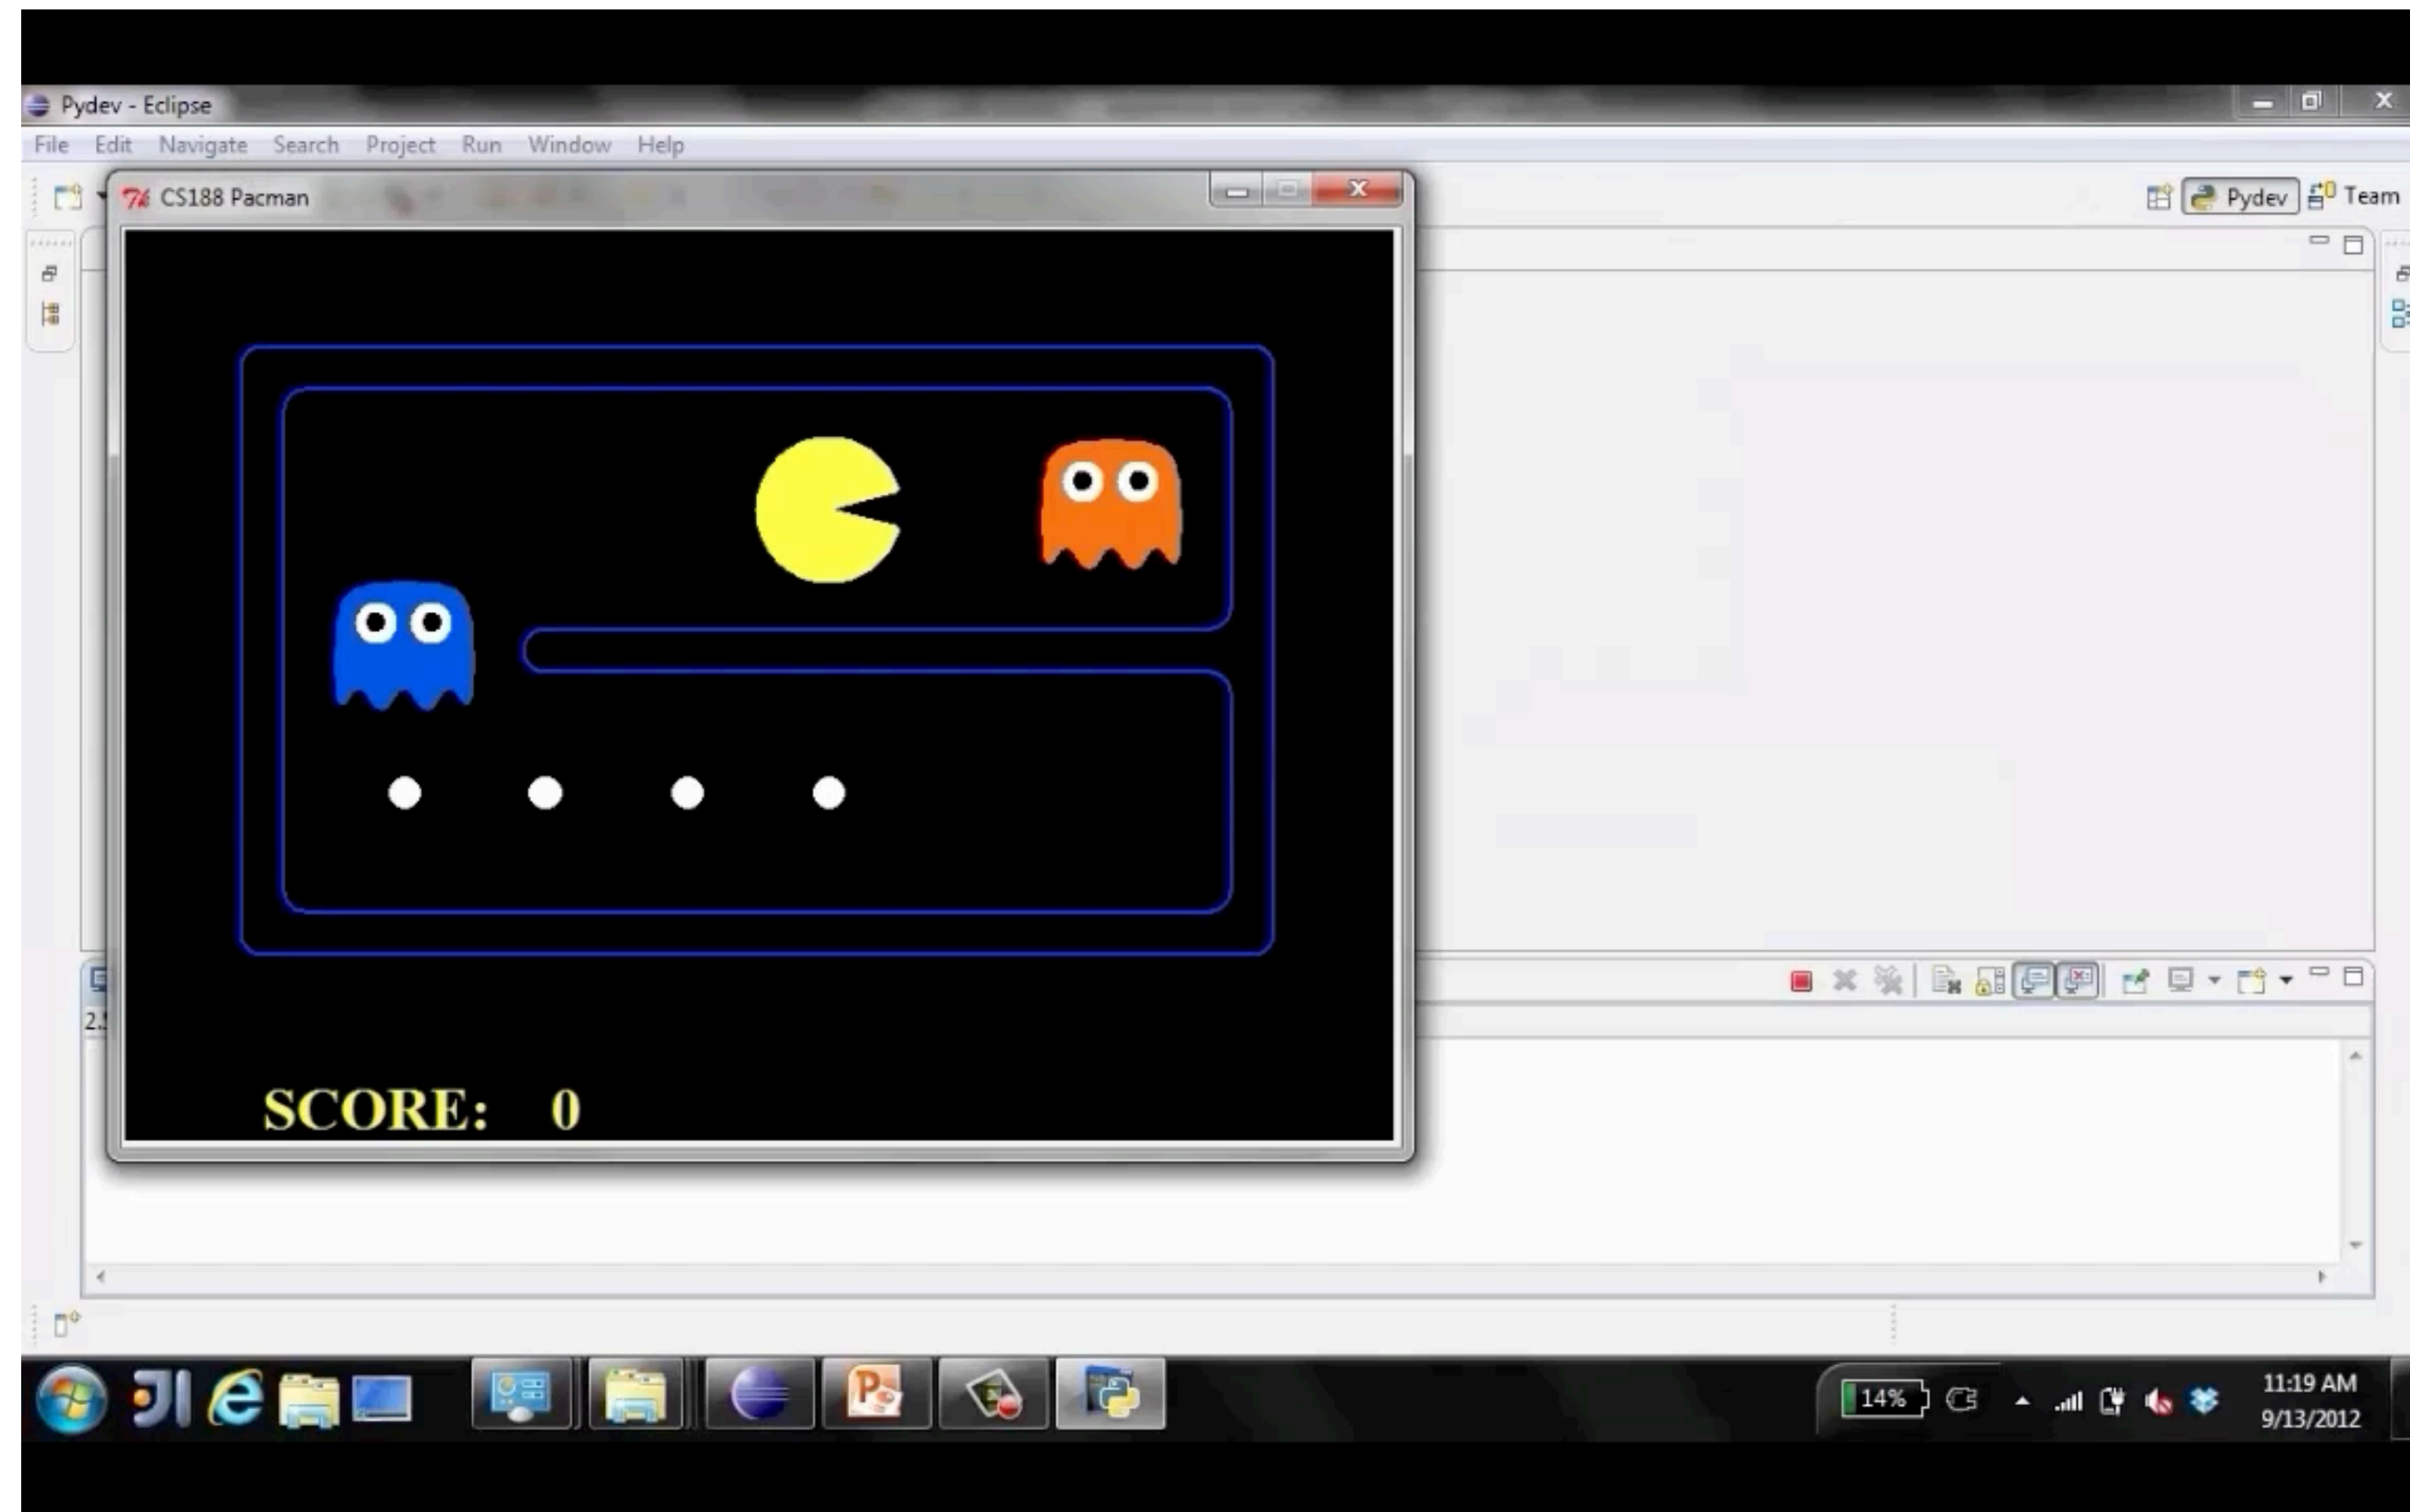

EXPECTIMAX SEARCH

function **decision**(s) returns an action
return the action **a** in **Actions**(s) with the highest
value(**Result**(s,a))

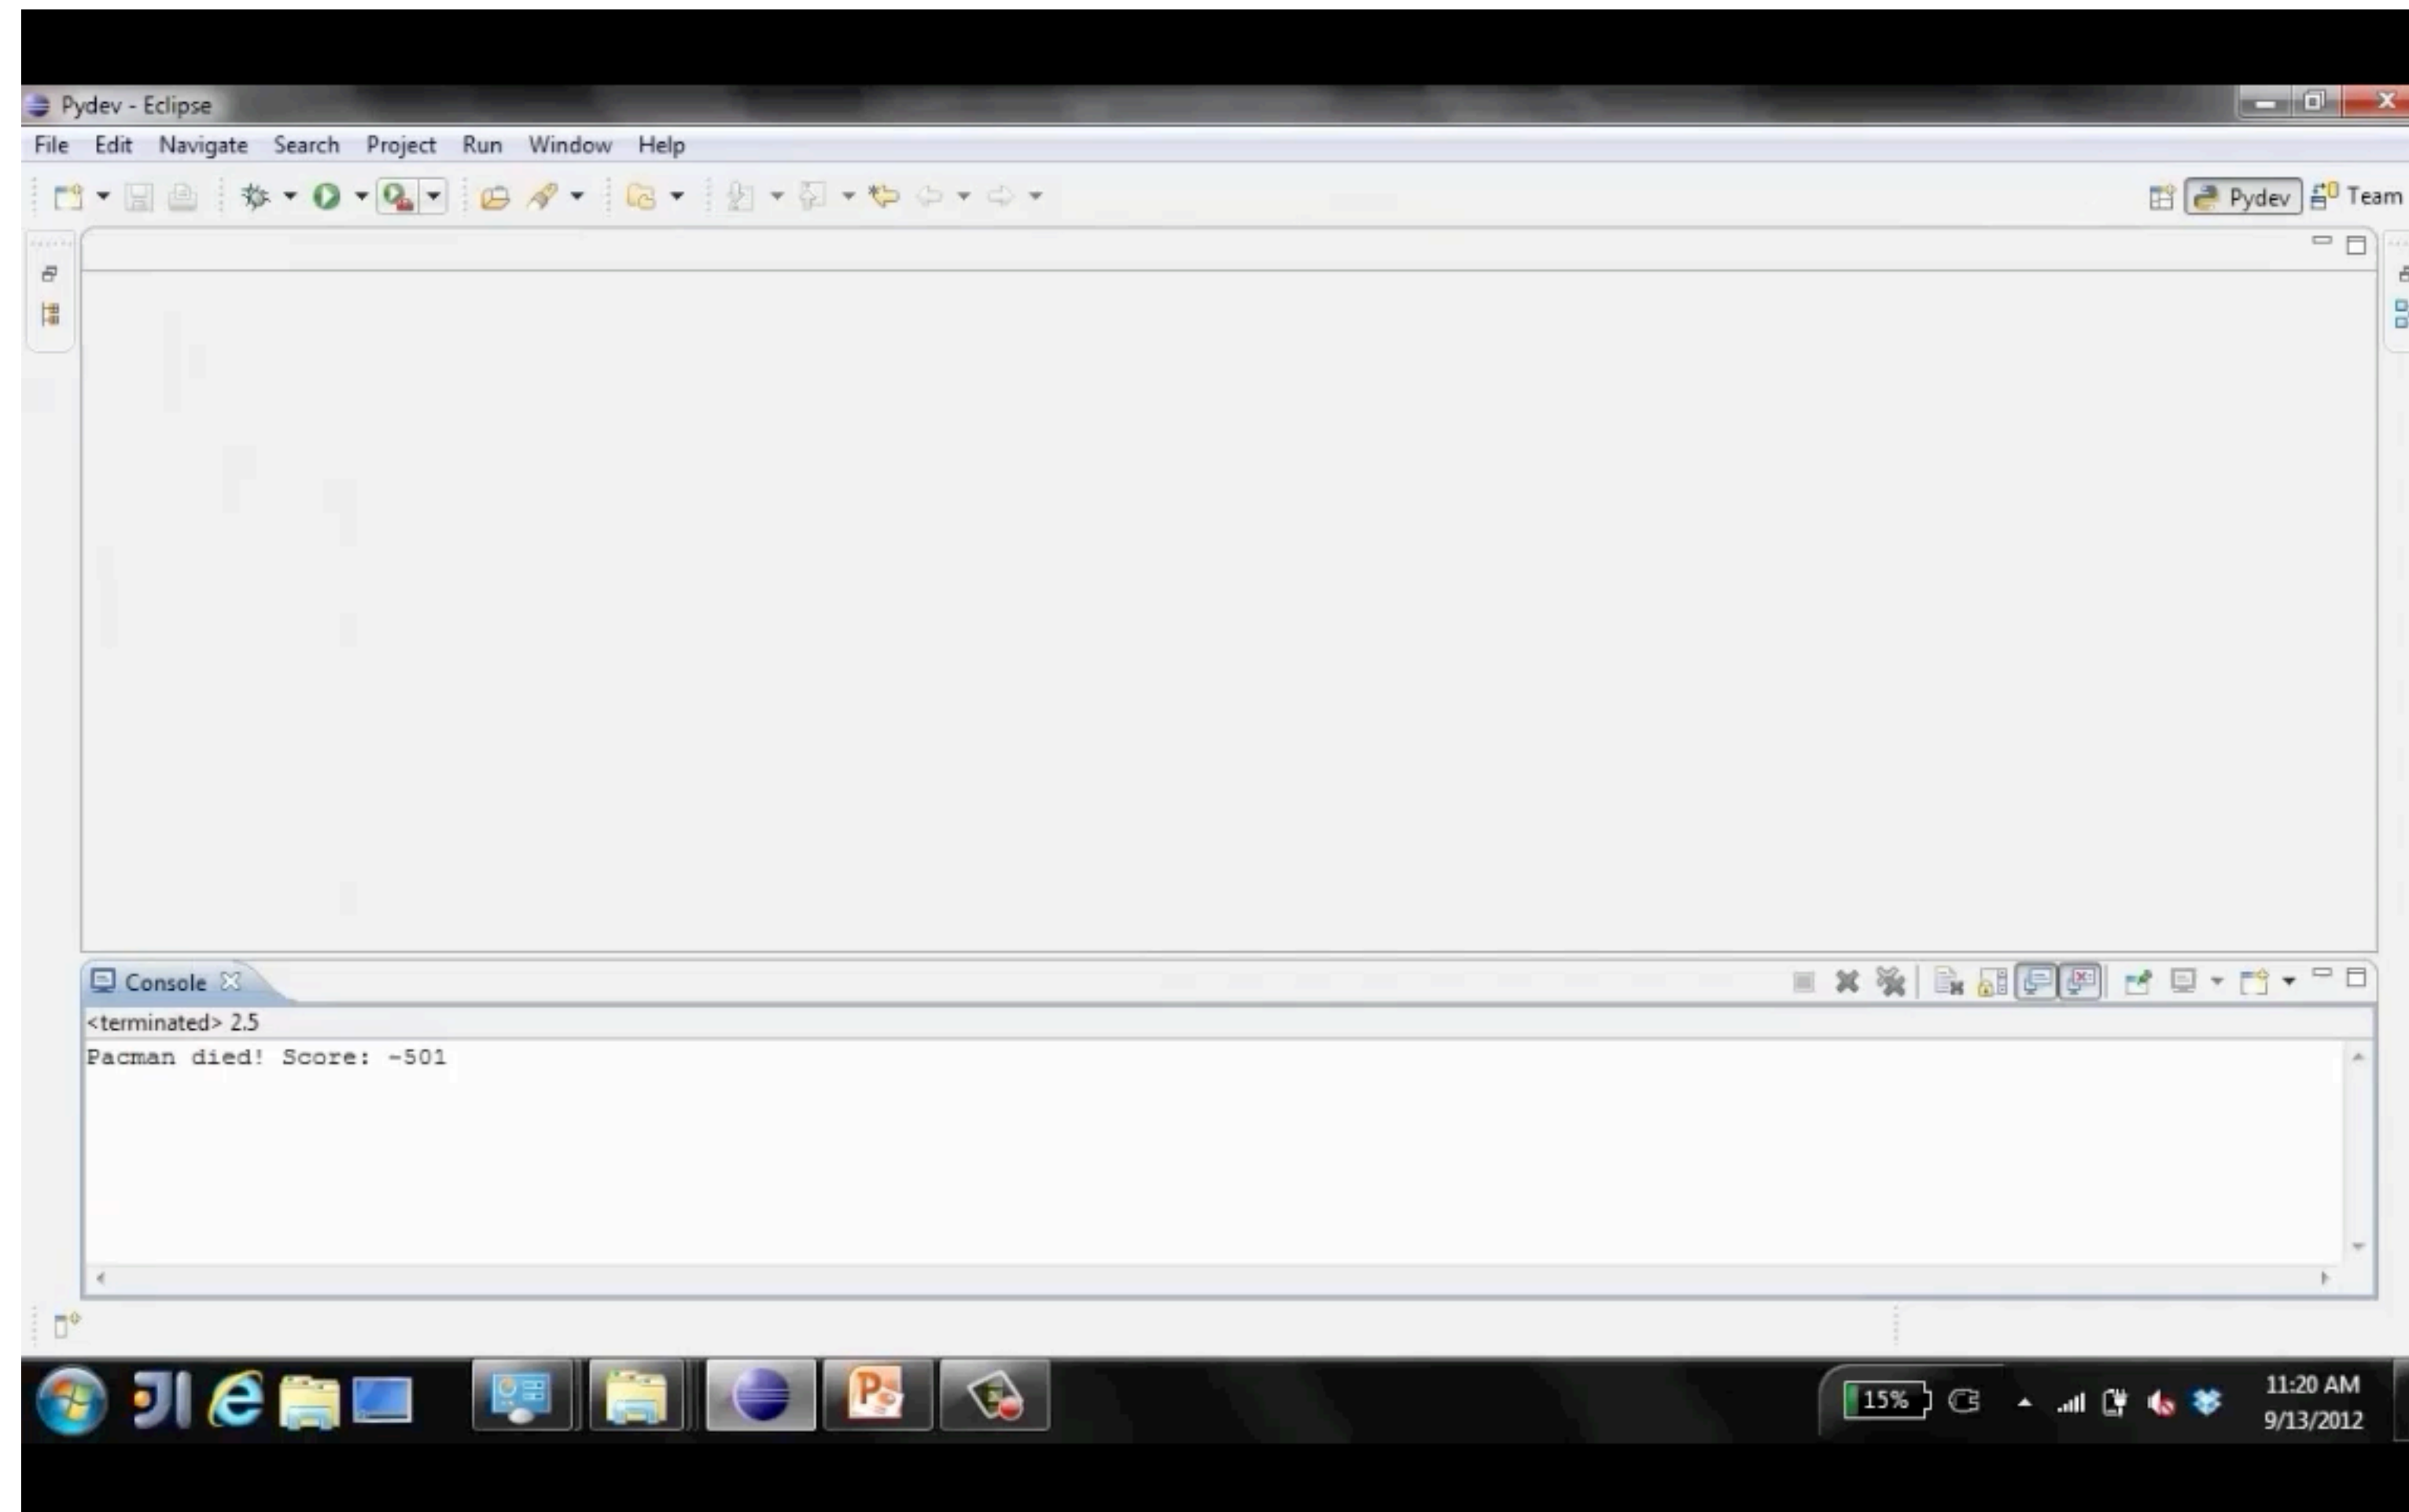


function **value**(s) returns a value
if **Terminal-Test**(s) then return **Utility**(s)
if **Player**(s) = **MAX** then return **max**_{a in **Actions**(s)} **value**(**Result**(s,a))
if **Player**(s) = **MIN** then return **min**_{a in **Actions**(s)} **value**(**Result**(s,a))
if **Player**(s) = **CHANCE** then return **sum**_{r in **chanceEvent**(s)} **Pr**(r) * **value**(**Result**(s,r))

DEMO MINIMAX VS EXPECTIMAX




DEMO MINIMAX VS EXPECTIMAX




EXPECTIMAX PSEUDOCODE

```
def exp-value(state):  
    initialize v = 0  
    for each successor of state:  
        p = probability(successor)  
        v += p * value(successor)  
    return v
```



$$v = (1/2) (8) + (1/3) (24) + (1/6) (-12) = 10$$

EXPECTIMAX EXAMPLE



SUMMARY

- Games require decisions when optimality is impossible
 - Bounded-depth search and approximate evaluation functions
- Games force efficient use of computation
 - Alpha-beta pruning
- Game playing has produced important research ideas
 - Reinforcement learning
 - Iterative deepening
 - Rational metareasoning