

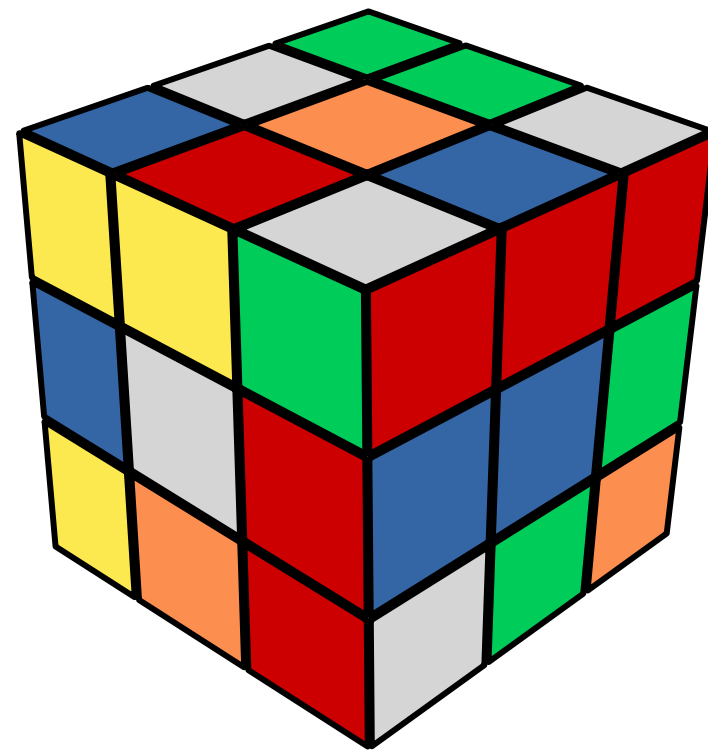


COMPSCI 761: ADVANCED TOPICS IN ARTIFICIAL INTELLIGENCE


SUMMARY II

Anna Trofimova, July 2022

SEARCH PROBLEM VS CSP VS GAME



5	3			7				
6			1	9	5			
	9	8					6	
8				6				3
4			8		3			1
7				2				6
	6					2	8	
			4	1	9			5
				8			7	9



CONSTRAINT SATISFACTION PROBLEMS (CSPS)

- Constraint Satisfaction Problems are defined by a set of variables X_i , each with a domain D_i of possible values, and a set of constraints C that specify allowable combinations of values.
- The aim is to find an assignment of the variables X_i from the domains D_i in such a way that none of the constraints C are violated.
 - i.e. all of the constraints C are satisfied

TYPES OF CONSTRAINTS

- Unary constraints involve a single variable
 - $M \neq 0$
- Binary constraints involve pairs of variables
 - $SA \neq WA$
- Higher-order constraints involve 3 or more variables
 - $Y = D + E$ or $Y = D + E - 10$
- Inequality constraints on Continuous variables
 - $\text{EndJob1} + 5 \leq \text{StartJob3}$
- Soft constraints (Preferences)
 - 11am lecture is better than 8am lecture!

EXAMPLE: MAP-COLOURING ★

Variables: WA, NT, Q, NSW, V, SA, T

Domains: $D_i = \{\text{red, green, blue}\}$

Constraints: adjacent regions must have different colours **e.g. $WA \neq NT$, etc.**



BACKTRACKING SEARCH SPACE PROPERTIES

The search space for this Depth First Search has certain very specific properties:

- If there are n variables, every solution will occur at exactly depth n .
- Variable assignments are commutative
[WA = red then NT = green] same as [NT = green then WA = red]

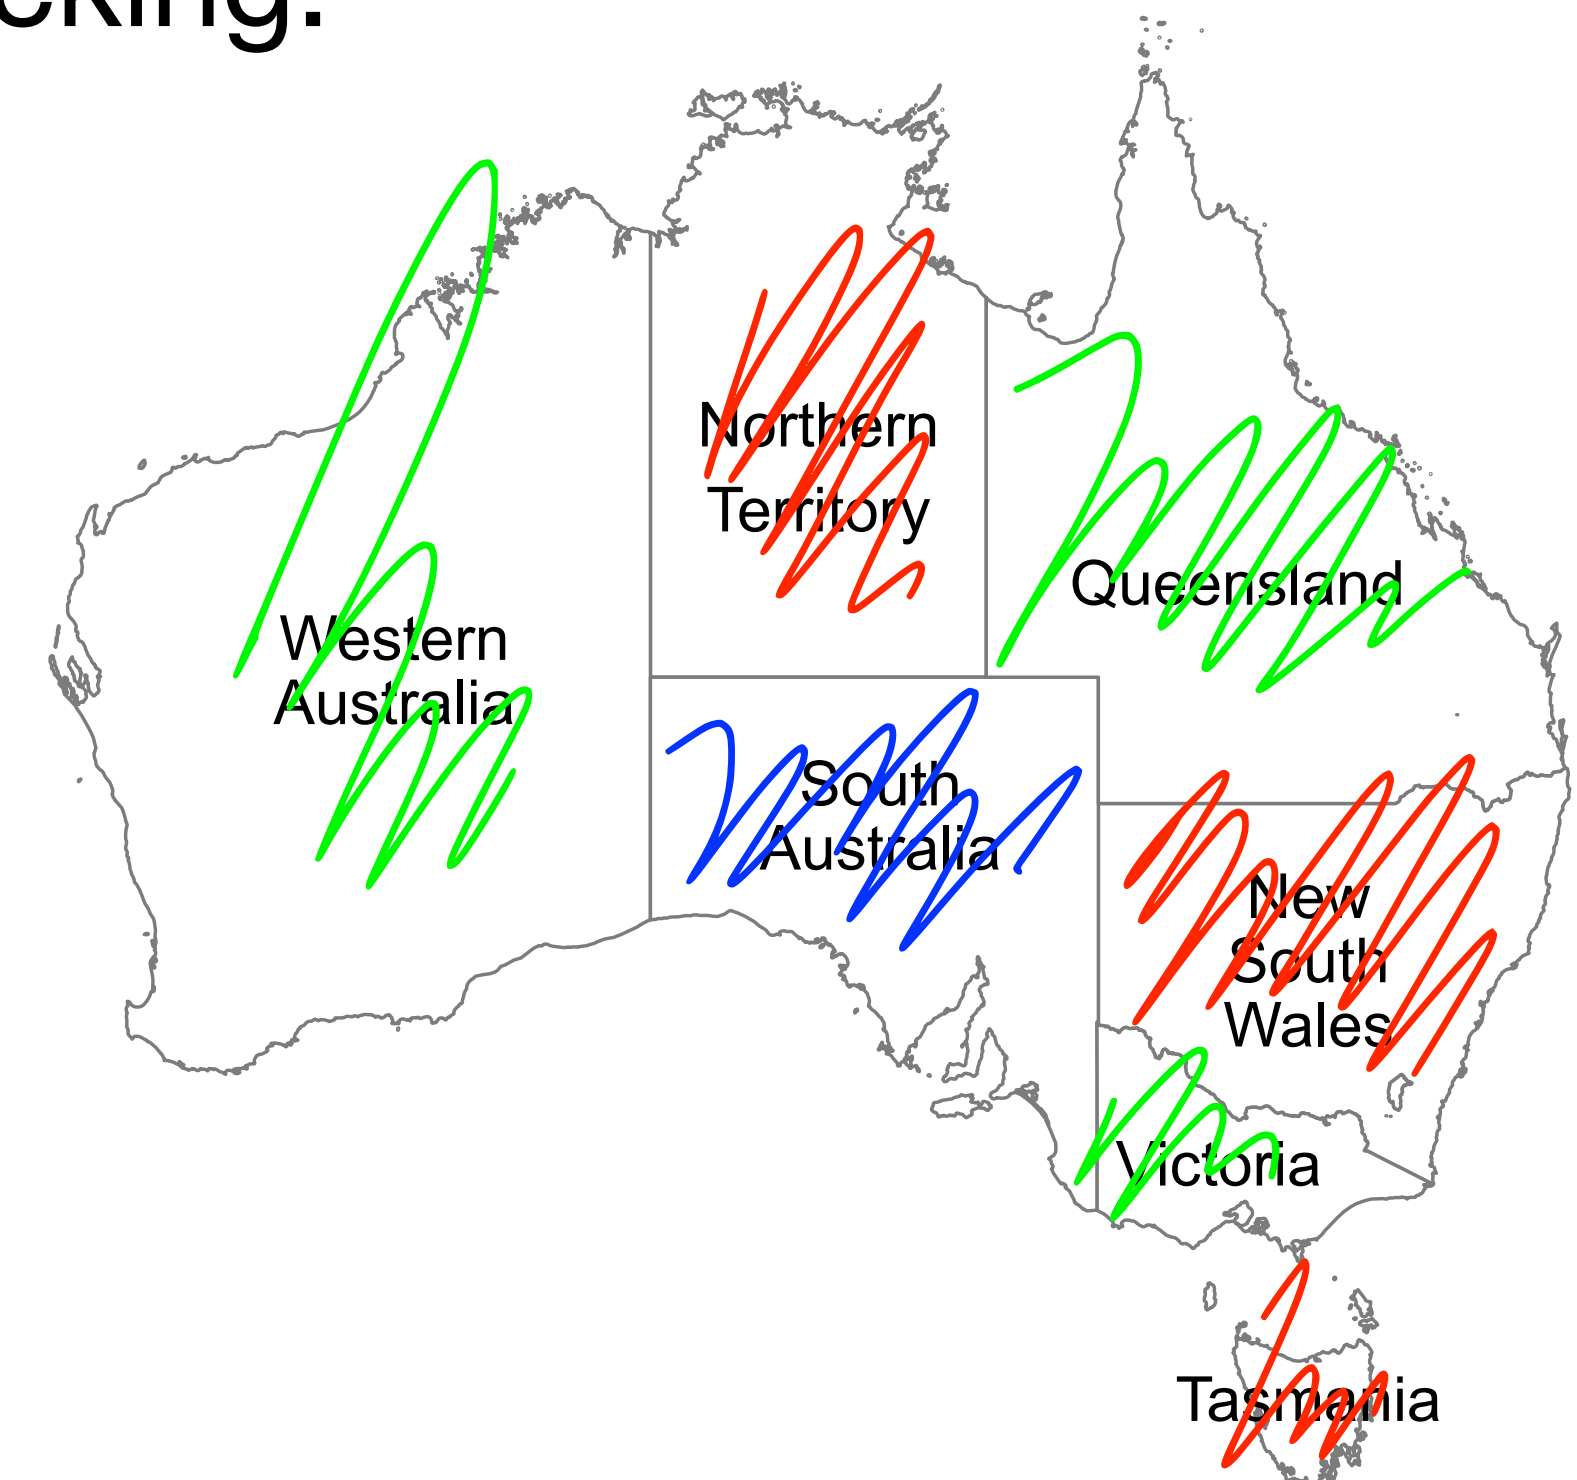
EXAMPLE: BACKTRACKING ★

- Given that WA is green, NT is red, apply backtracking choosing the remaining states in the alphabetical order and assigning colours in the following order {r, g, b}. Select the correct outcome of the backtracking:
 - a. No solution
 - b. NSW = r, Q = g, SA = b, WA = g, T = r
 - c. NSW = r, Q = b, SA = r, WA = g, T = r
 - c. NSW = r, Q = b, SA = r, WA = b, T = r

EXAMPLE: BACKTRACKING ★

- Given that WA is green, NT is red, apply backpacking choosing the remaining states in the alphabetical order and assigning colours in the following order {r, g, b}. Select the correct outcome of the backtracking:

NSW, NT, Q, SA, T, V, WA
rgb, R, rgb, rgb, rgb, rgb, G
R, R, rgb, rgb, rgb, rgb, G
R, R, G, rgb, rgb, rgb, G
R, R, G, B, rgb, rgb, G
R, R, G, B, R, rgb, G
R, R, G, B, R, G, G
=> R, G, B, R, G



EXAMPLE: BACKTRACKING ★

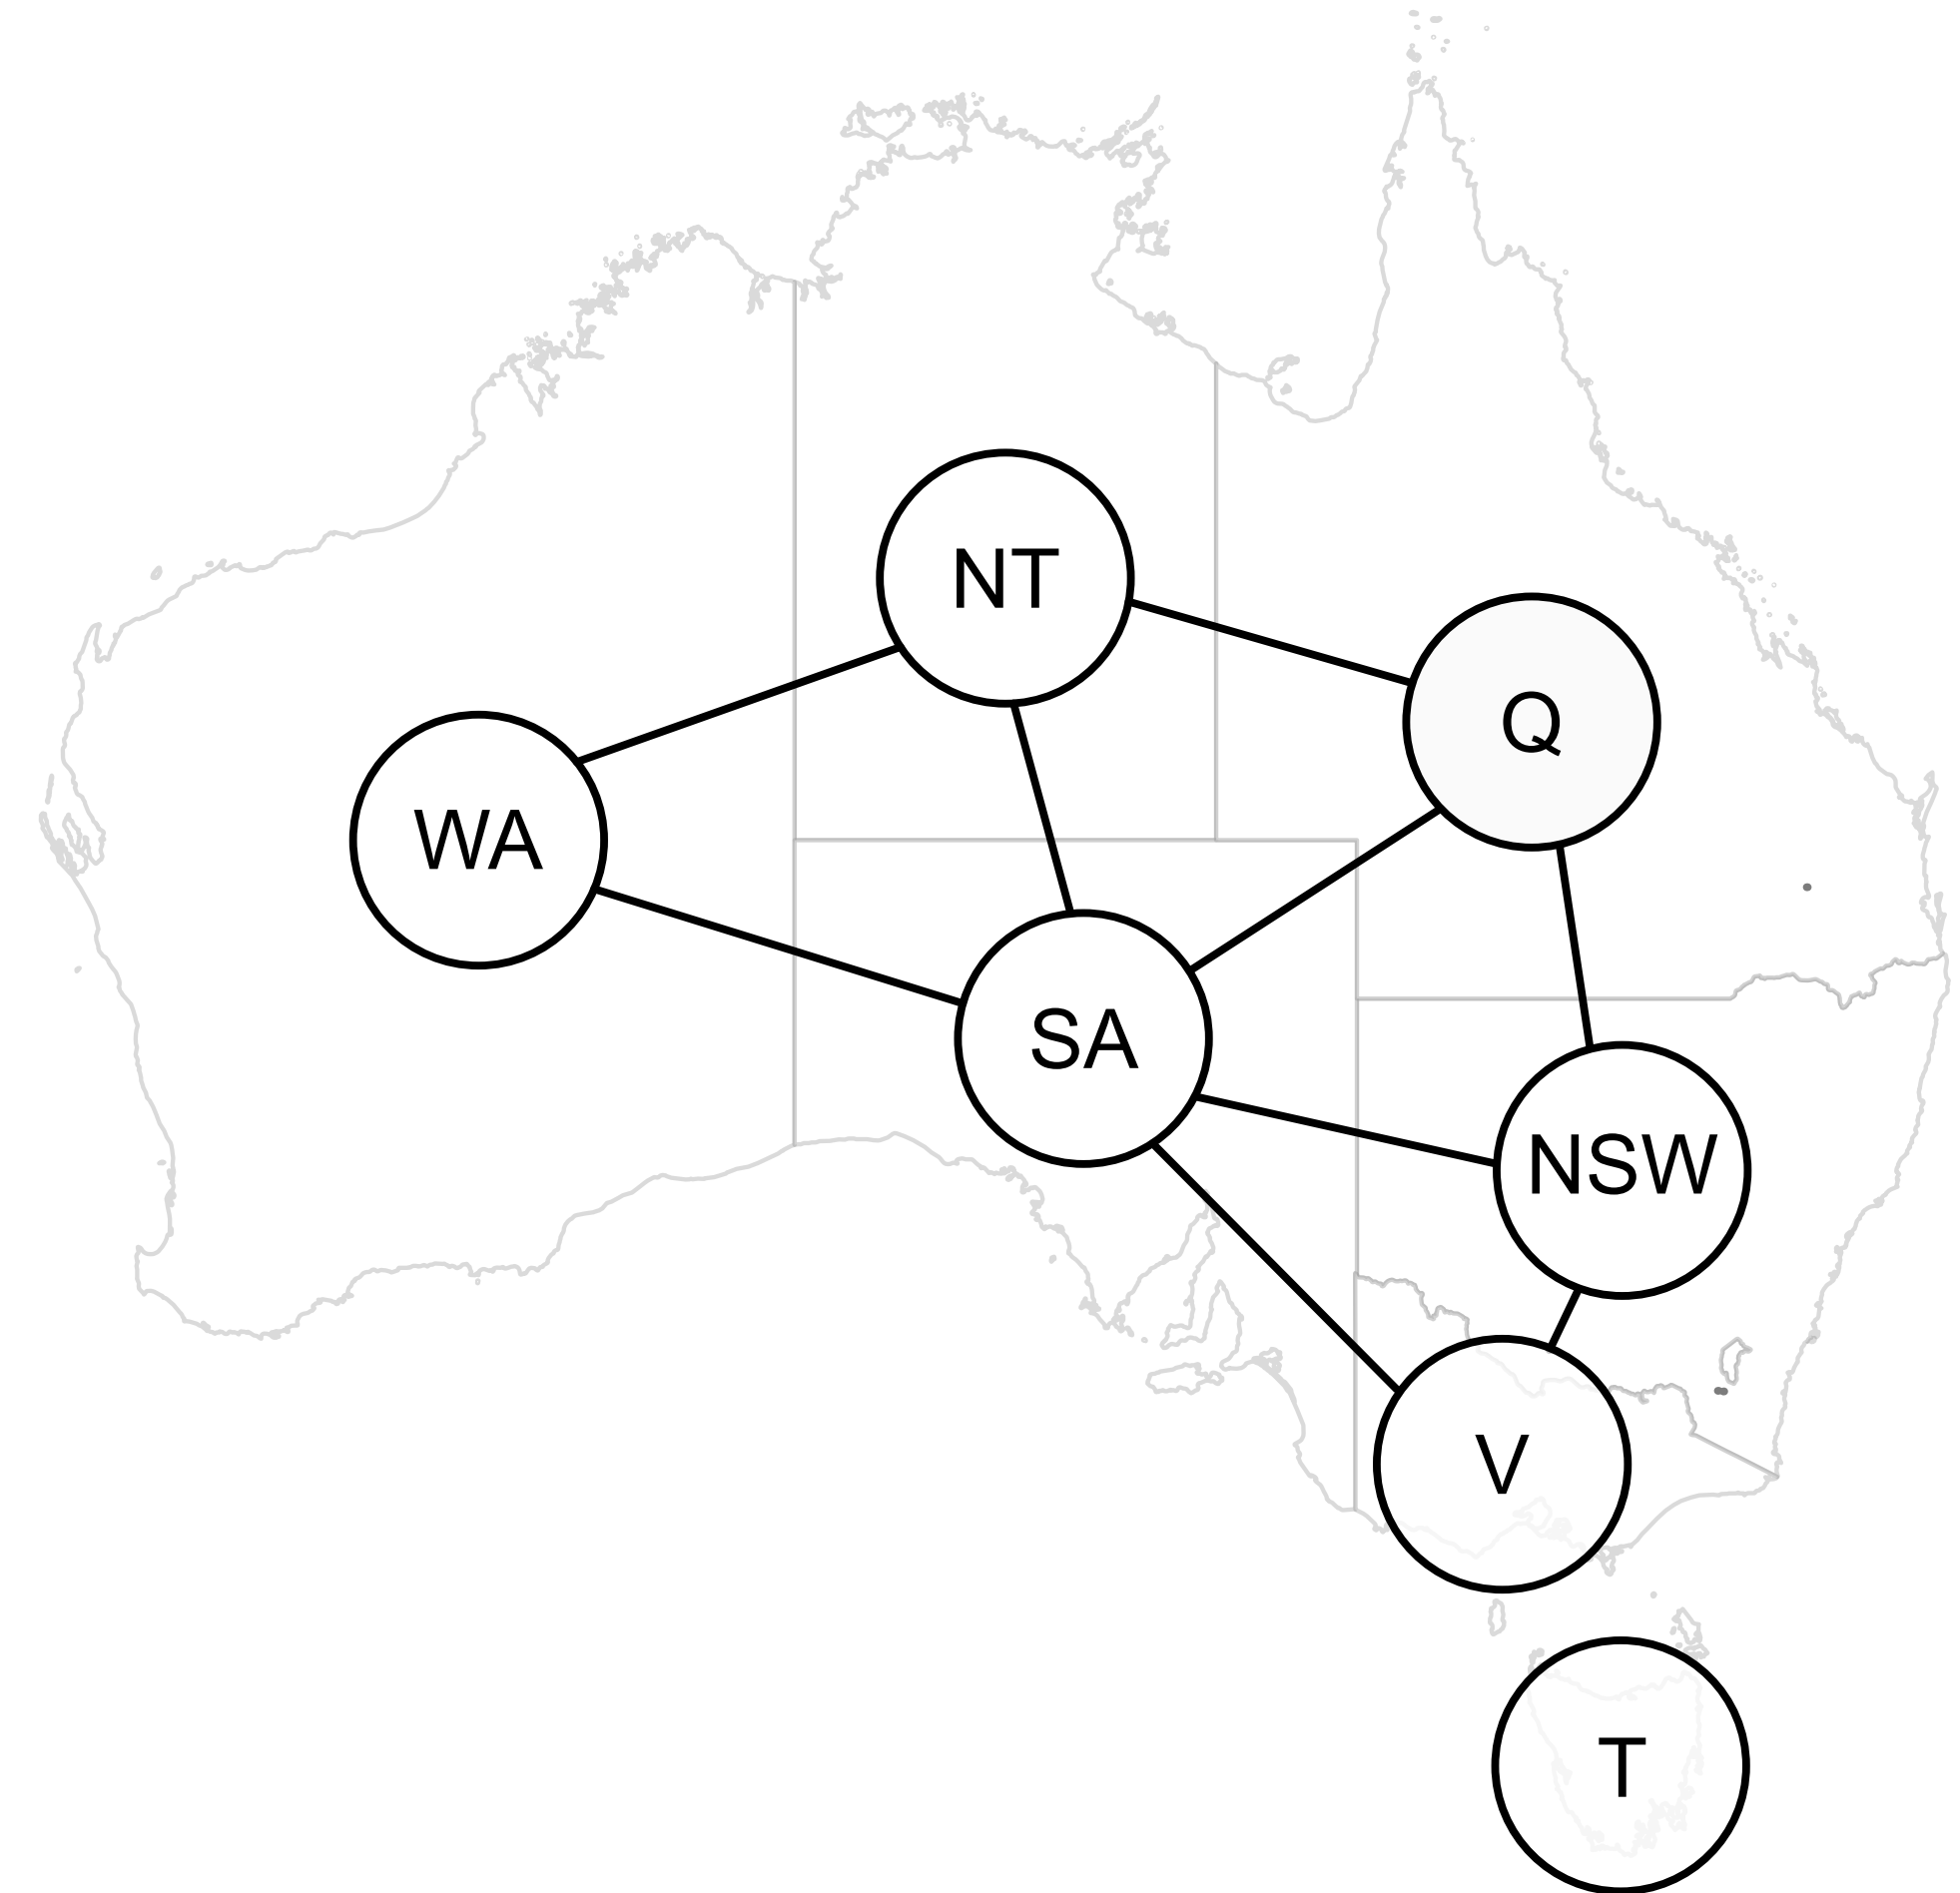
- Given that WA is green, NT is red, apply backpacking
 - With arc consistency checking
 - Using Minimum Remaining Value
 - Degree heuristic
 - Using least constrained value
- ... Select the correct outcome:
 - a. No solution
 - b. NSW = r, Q = g, SA = b, WA = g, T = r
 - c. NSW = r, Q = b, SA = r, WA = g, T = r
 - c. NSW = r, Q = b, SA = r, WA = b, T = r

VARIABLE ELIMINATION

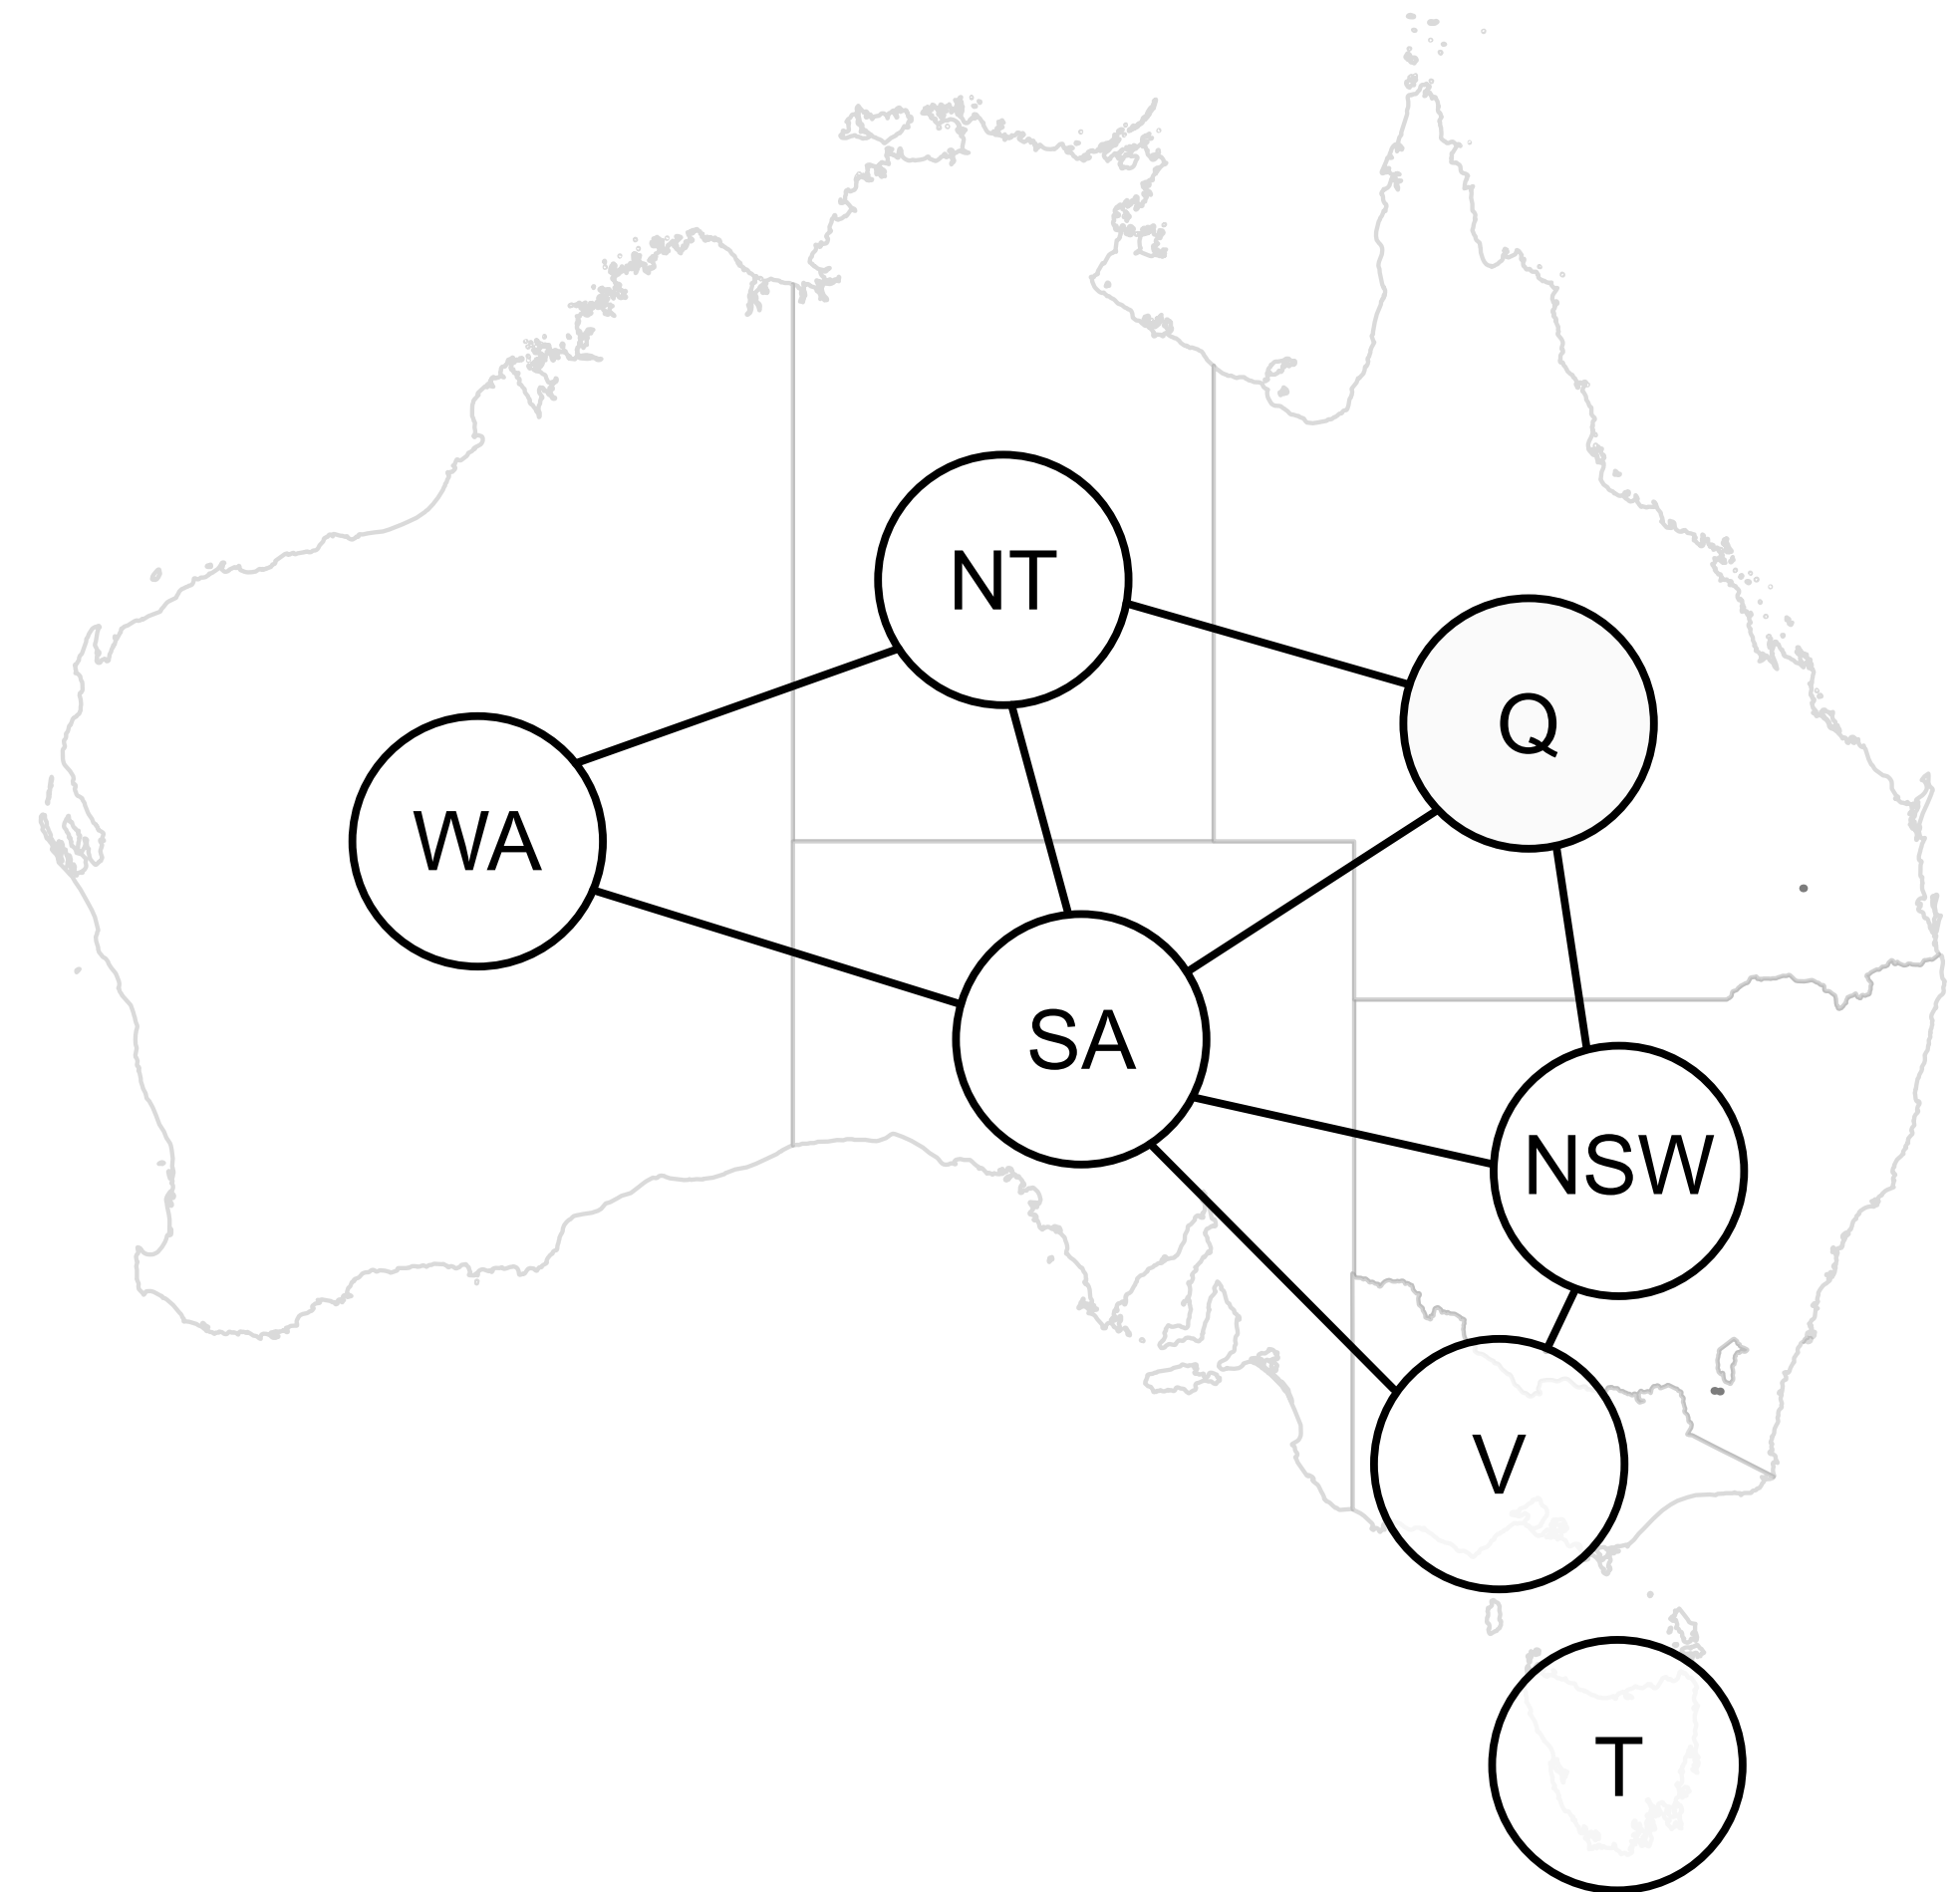
- If there is only one variable, return the intersection of its (unary) constraints
- Otherwise
 - Select a variable X
 - Join the constraints in which X appears, forming constraint $R1$
 - Project $R1$ onto its variables other than X , forming $R2$
 - Replace all of the constraints in which X appears by $R2$
 - Recursively solve the simplified problem, forming $R3$
 - Return $R1$ joined with $R3$

EXAMPLE: VARIABLE ELIMINATION ★

- Eliminate variable V , and then join the new constraint with the constraint $SA \neq NSW$
- $V \neq NSW$, $V \neq SA$



EXAMPLE: VARIABLE ELIMINATION ★



LOCAL SEARCH?

- In Systematic Search a solution can be a path or a state!
- In Local Search a state is a solution!
 - Since it doesn't maintain a path ("only a current state") it can't return a path
 - Can use a goal test or partial goal state
 - (but not a complete goal state – why?)



LOCAL SEARCH ALGORITHMS

Hill Climbing – chose best child

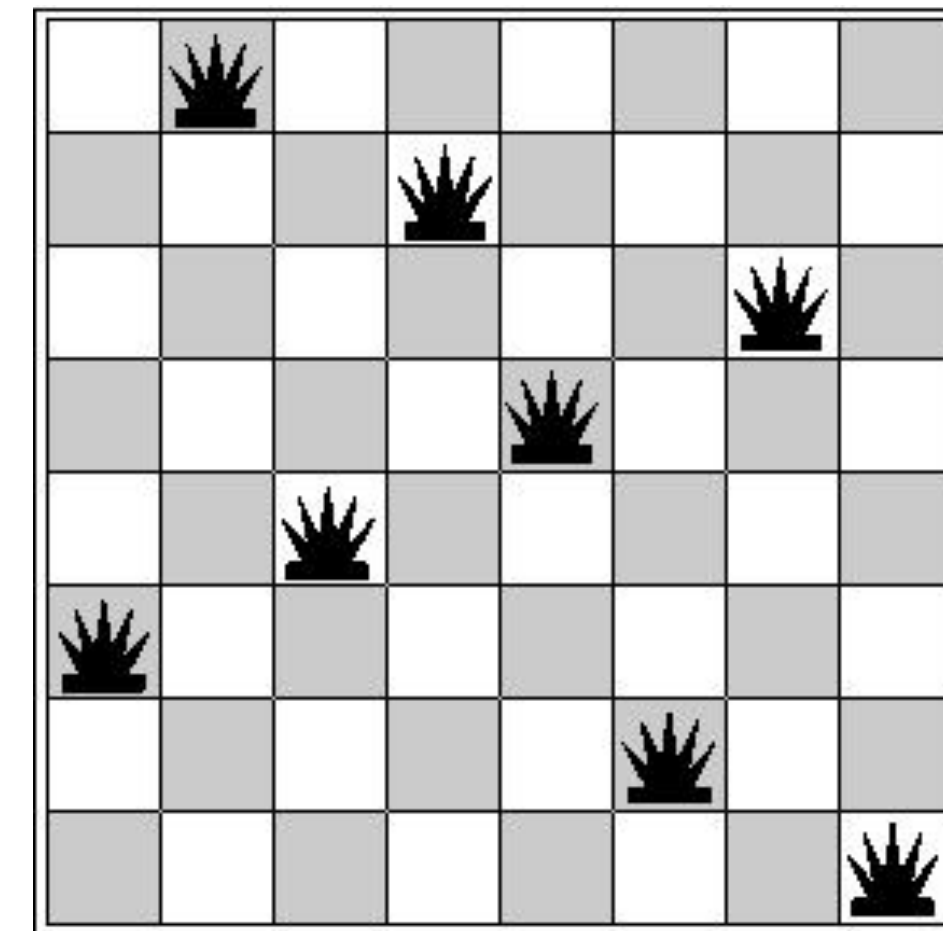
Hill Climbing with Sideways Moves – chose best child or equal child

Tabu Search – keep list of nodes you have been to and don't go back

Enforced Hill Climbing – use breadth first search until you find a “better” node

EXAMPLE: HILL CLIMBING ★

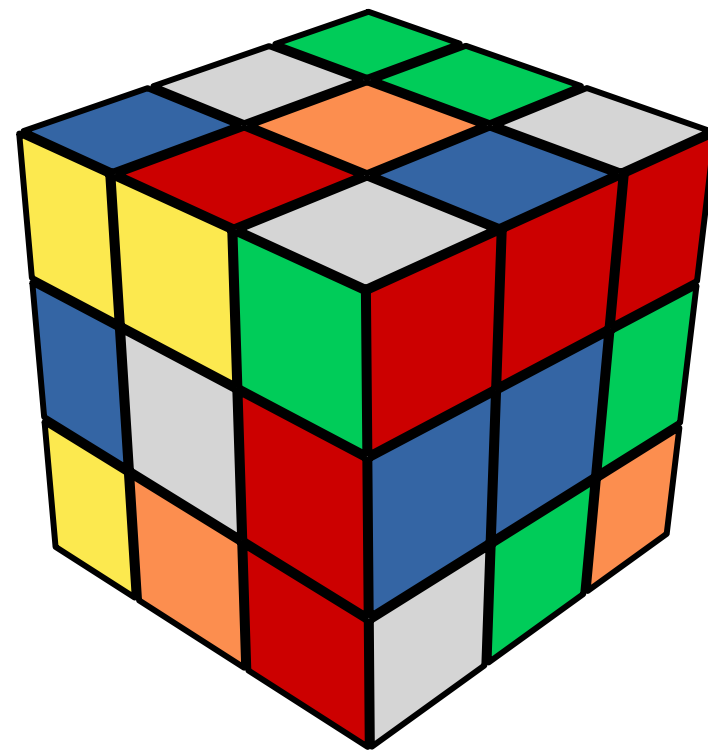
- How many violations are there?
- Label the queens according to the column number. Run Hill-climbing with side way moves and Min Conflicts and in case of a tie prioritise the queen on the left. With respect to the solution, select the position for the 2nd queen:
 - No solution
 - Row 1
 - Row 4
 - Row 7



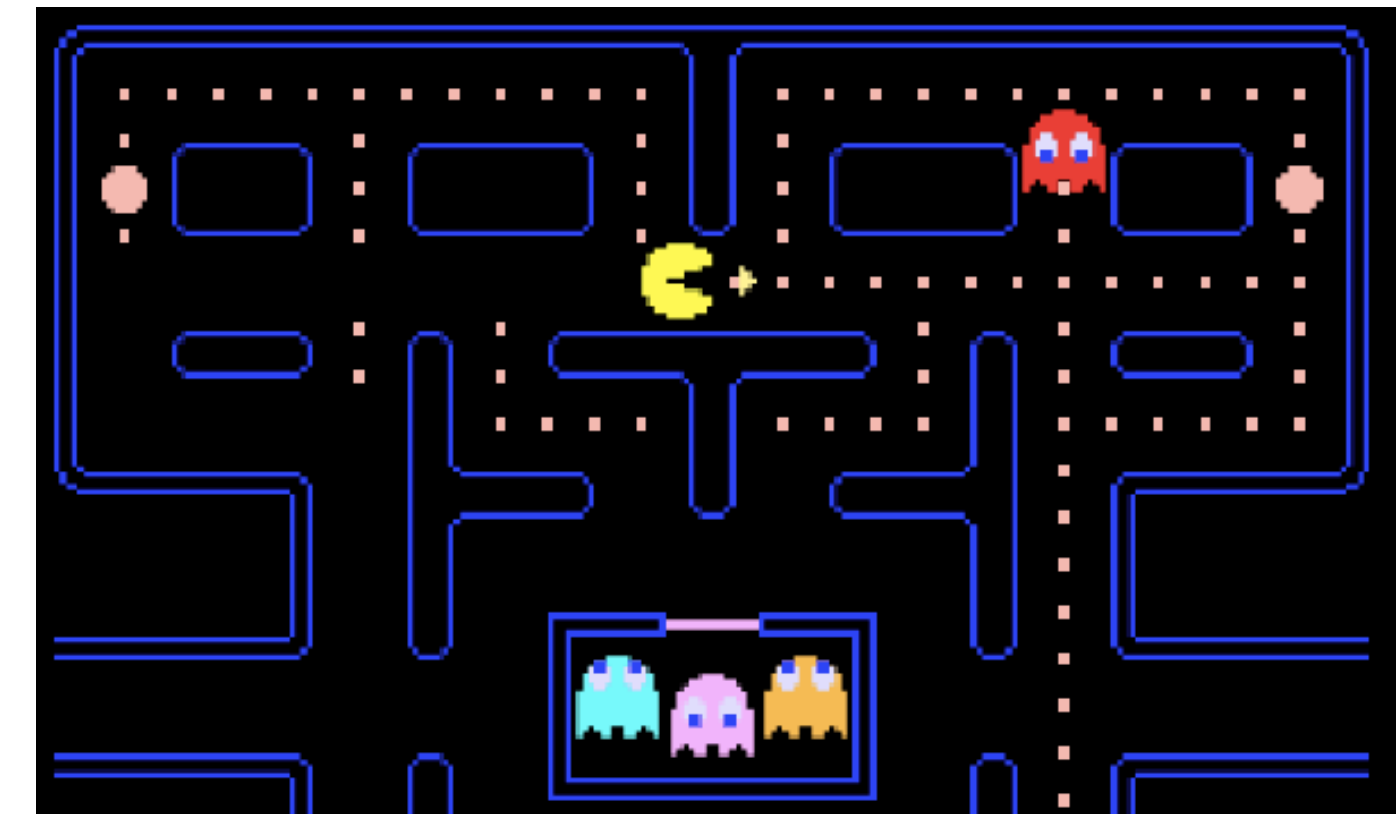
STOCHASTIC SEARCH ALGORITHMS

- Stochastic Hill Climbing – chose probabilistically from among better children based on their fitness
- First Choice Hill-Climbing - chose the first better random child
- Random Walking Hill Climbing – chose probabilistically between the best child and a random child
- Random Restart Hill Climbing – do any hill climbing method with randomly restarts to try and get a better result
- Simulated Annealing – first choice hill climbing with a stochastic chance of choosing a worse node. The chance of choosing a worse node reduces over time with the temperature schedule

SEARCH PROBLEM VS CSP VS GAME



5	3			7				
6			1	9	5			
	9	8					6	
8				6				3
4			8		3			1
7				2				6
	6					2	8	
			4	1	9			5
				8			7	9



- "Unpredictable" opponent → specifying a move for every possible opponent reply
- Time limits → unlikely to find optimal solution, must approximate

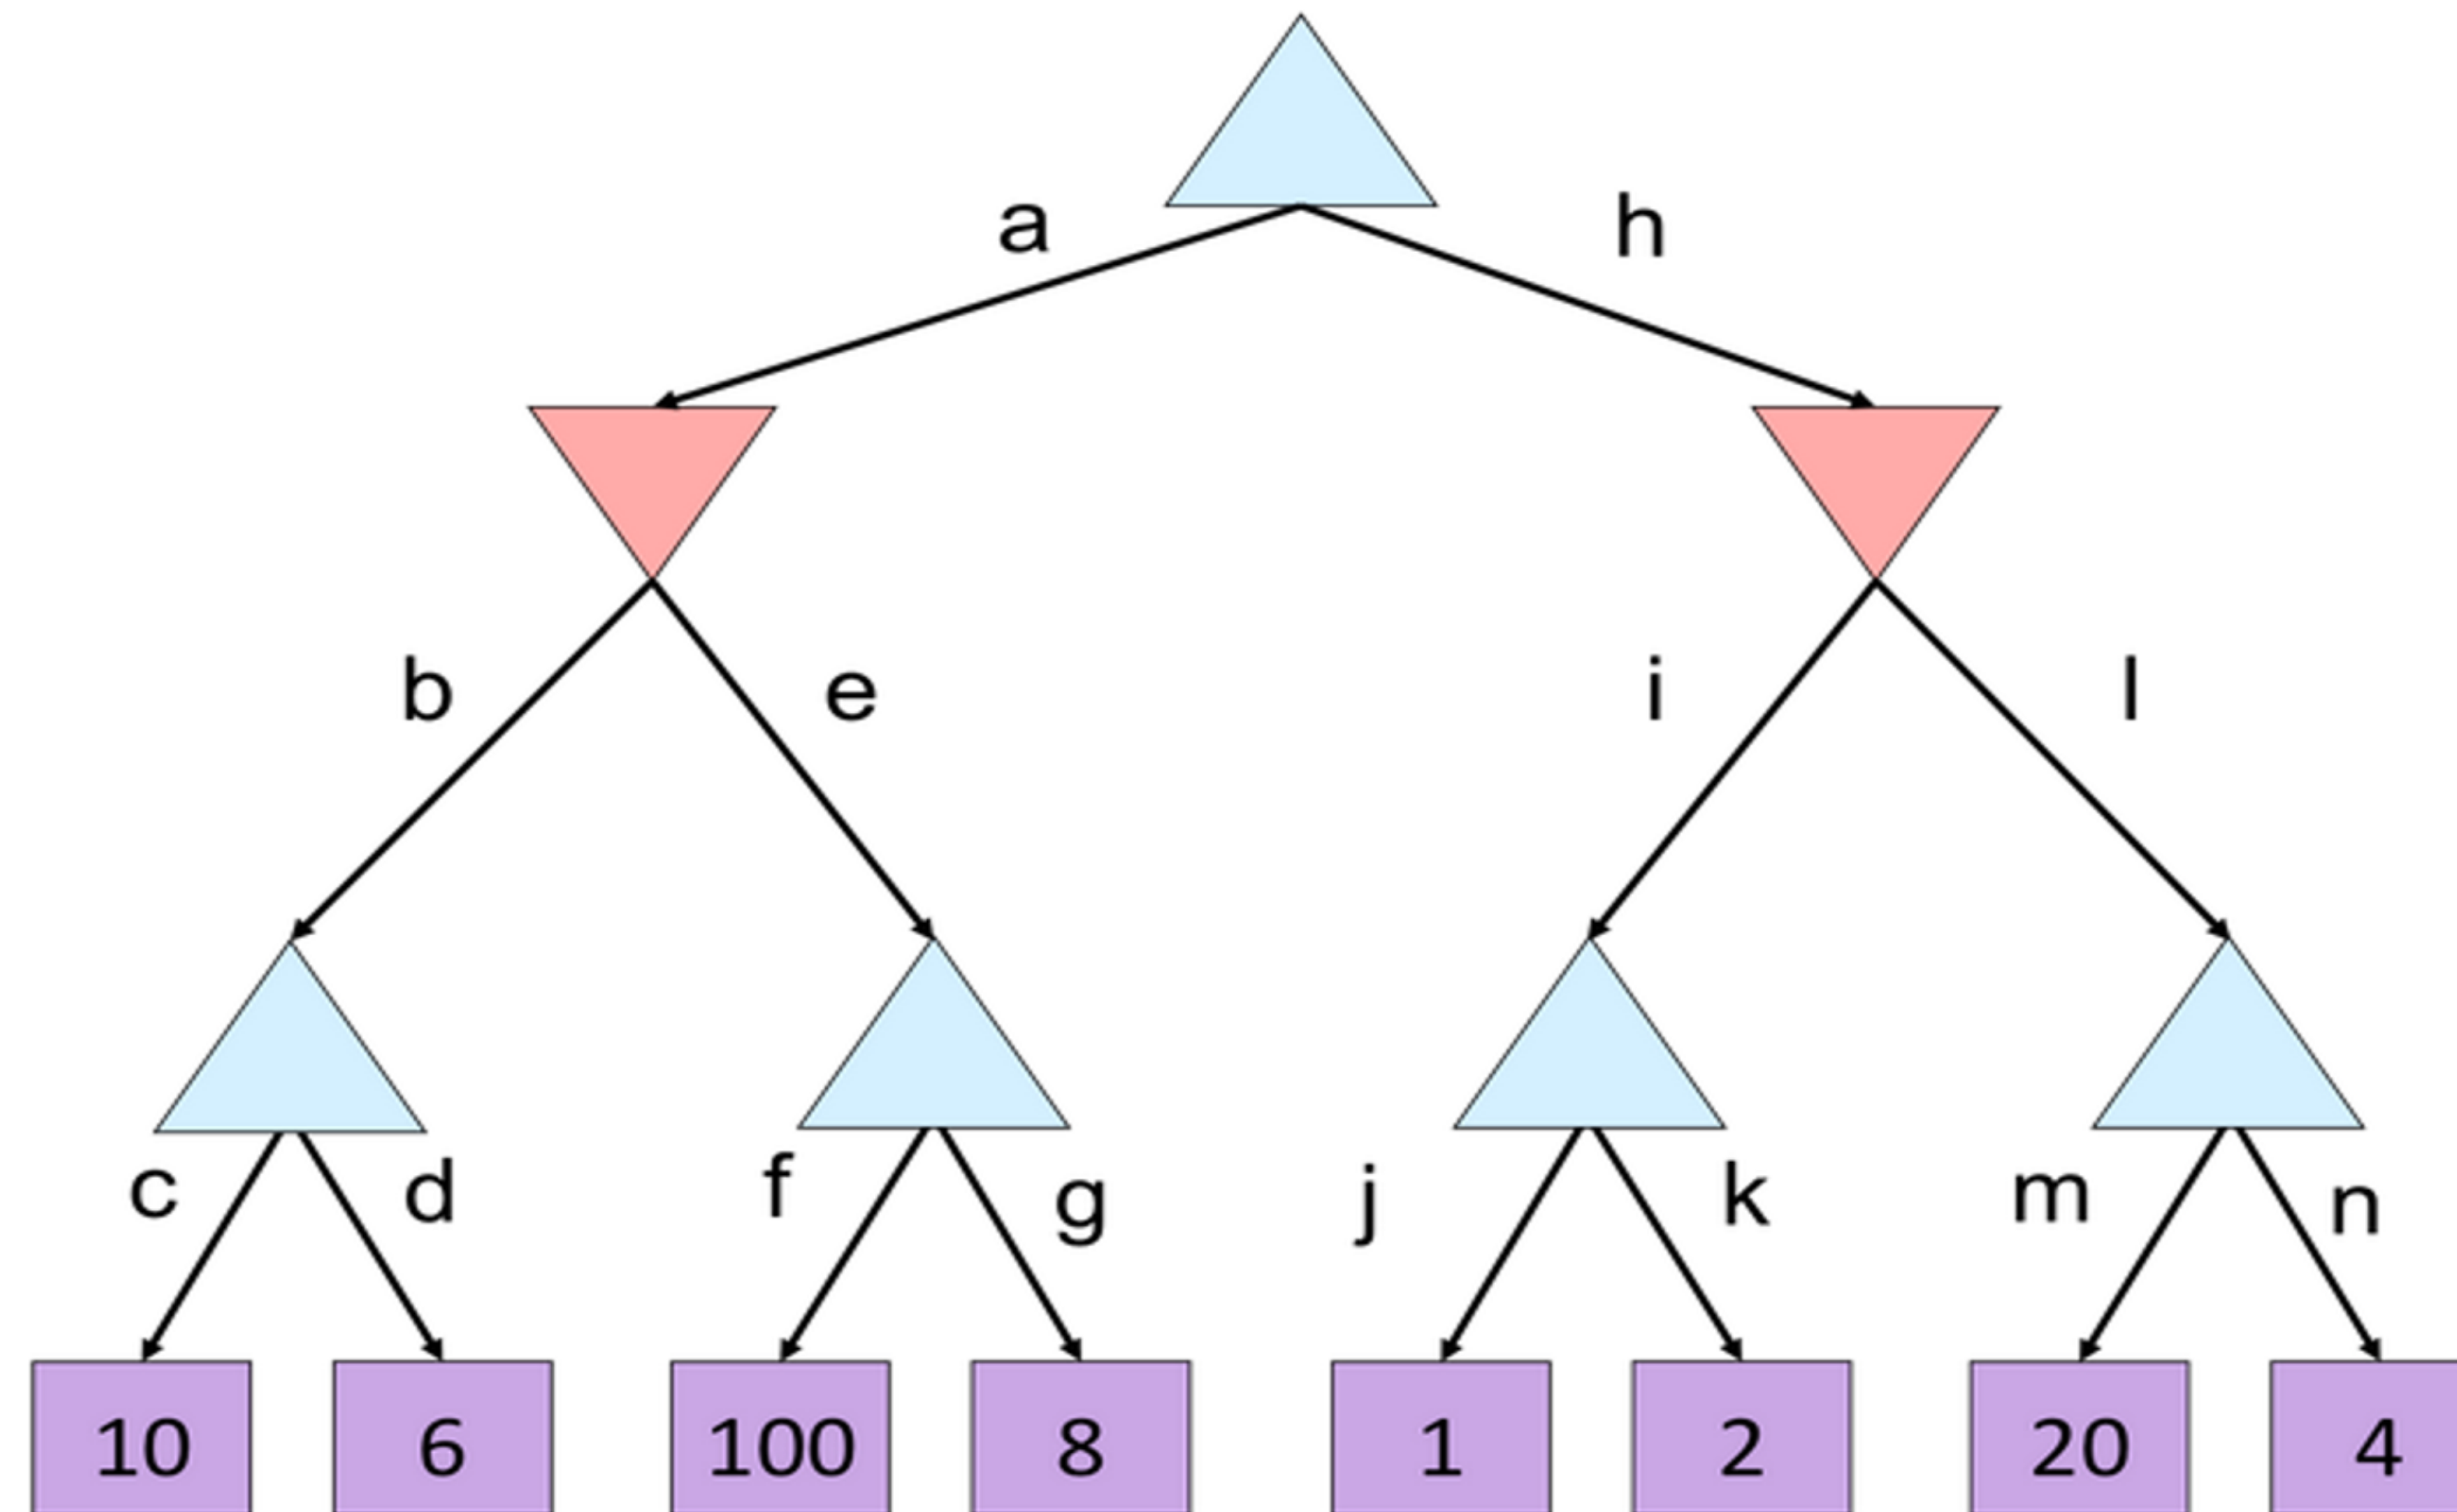
MINIMAX SEARCH

```
function minimax-decision(s) returns an action
  return the action a in Actions(s) with the highest
    minimax_value(Result(s,a))
```



```
function minimax_value(s) returns a value
  if Terminal-Test(s) then return Utility(s)
  if Player(s) = MAX then return maxa in Actions(s) minimax_value(Result(s,a))
  if Player(s) = MIN then return mina in Actions(s) minimax_value(Result(s,a))
```

EXAMPLE: MINIMAX ★



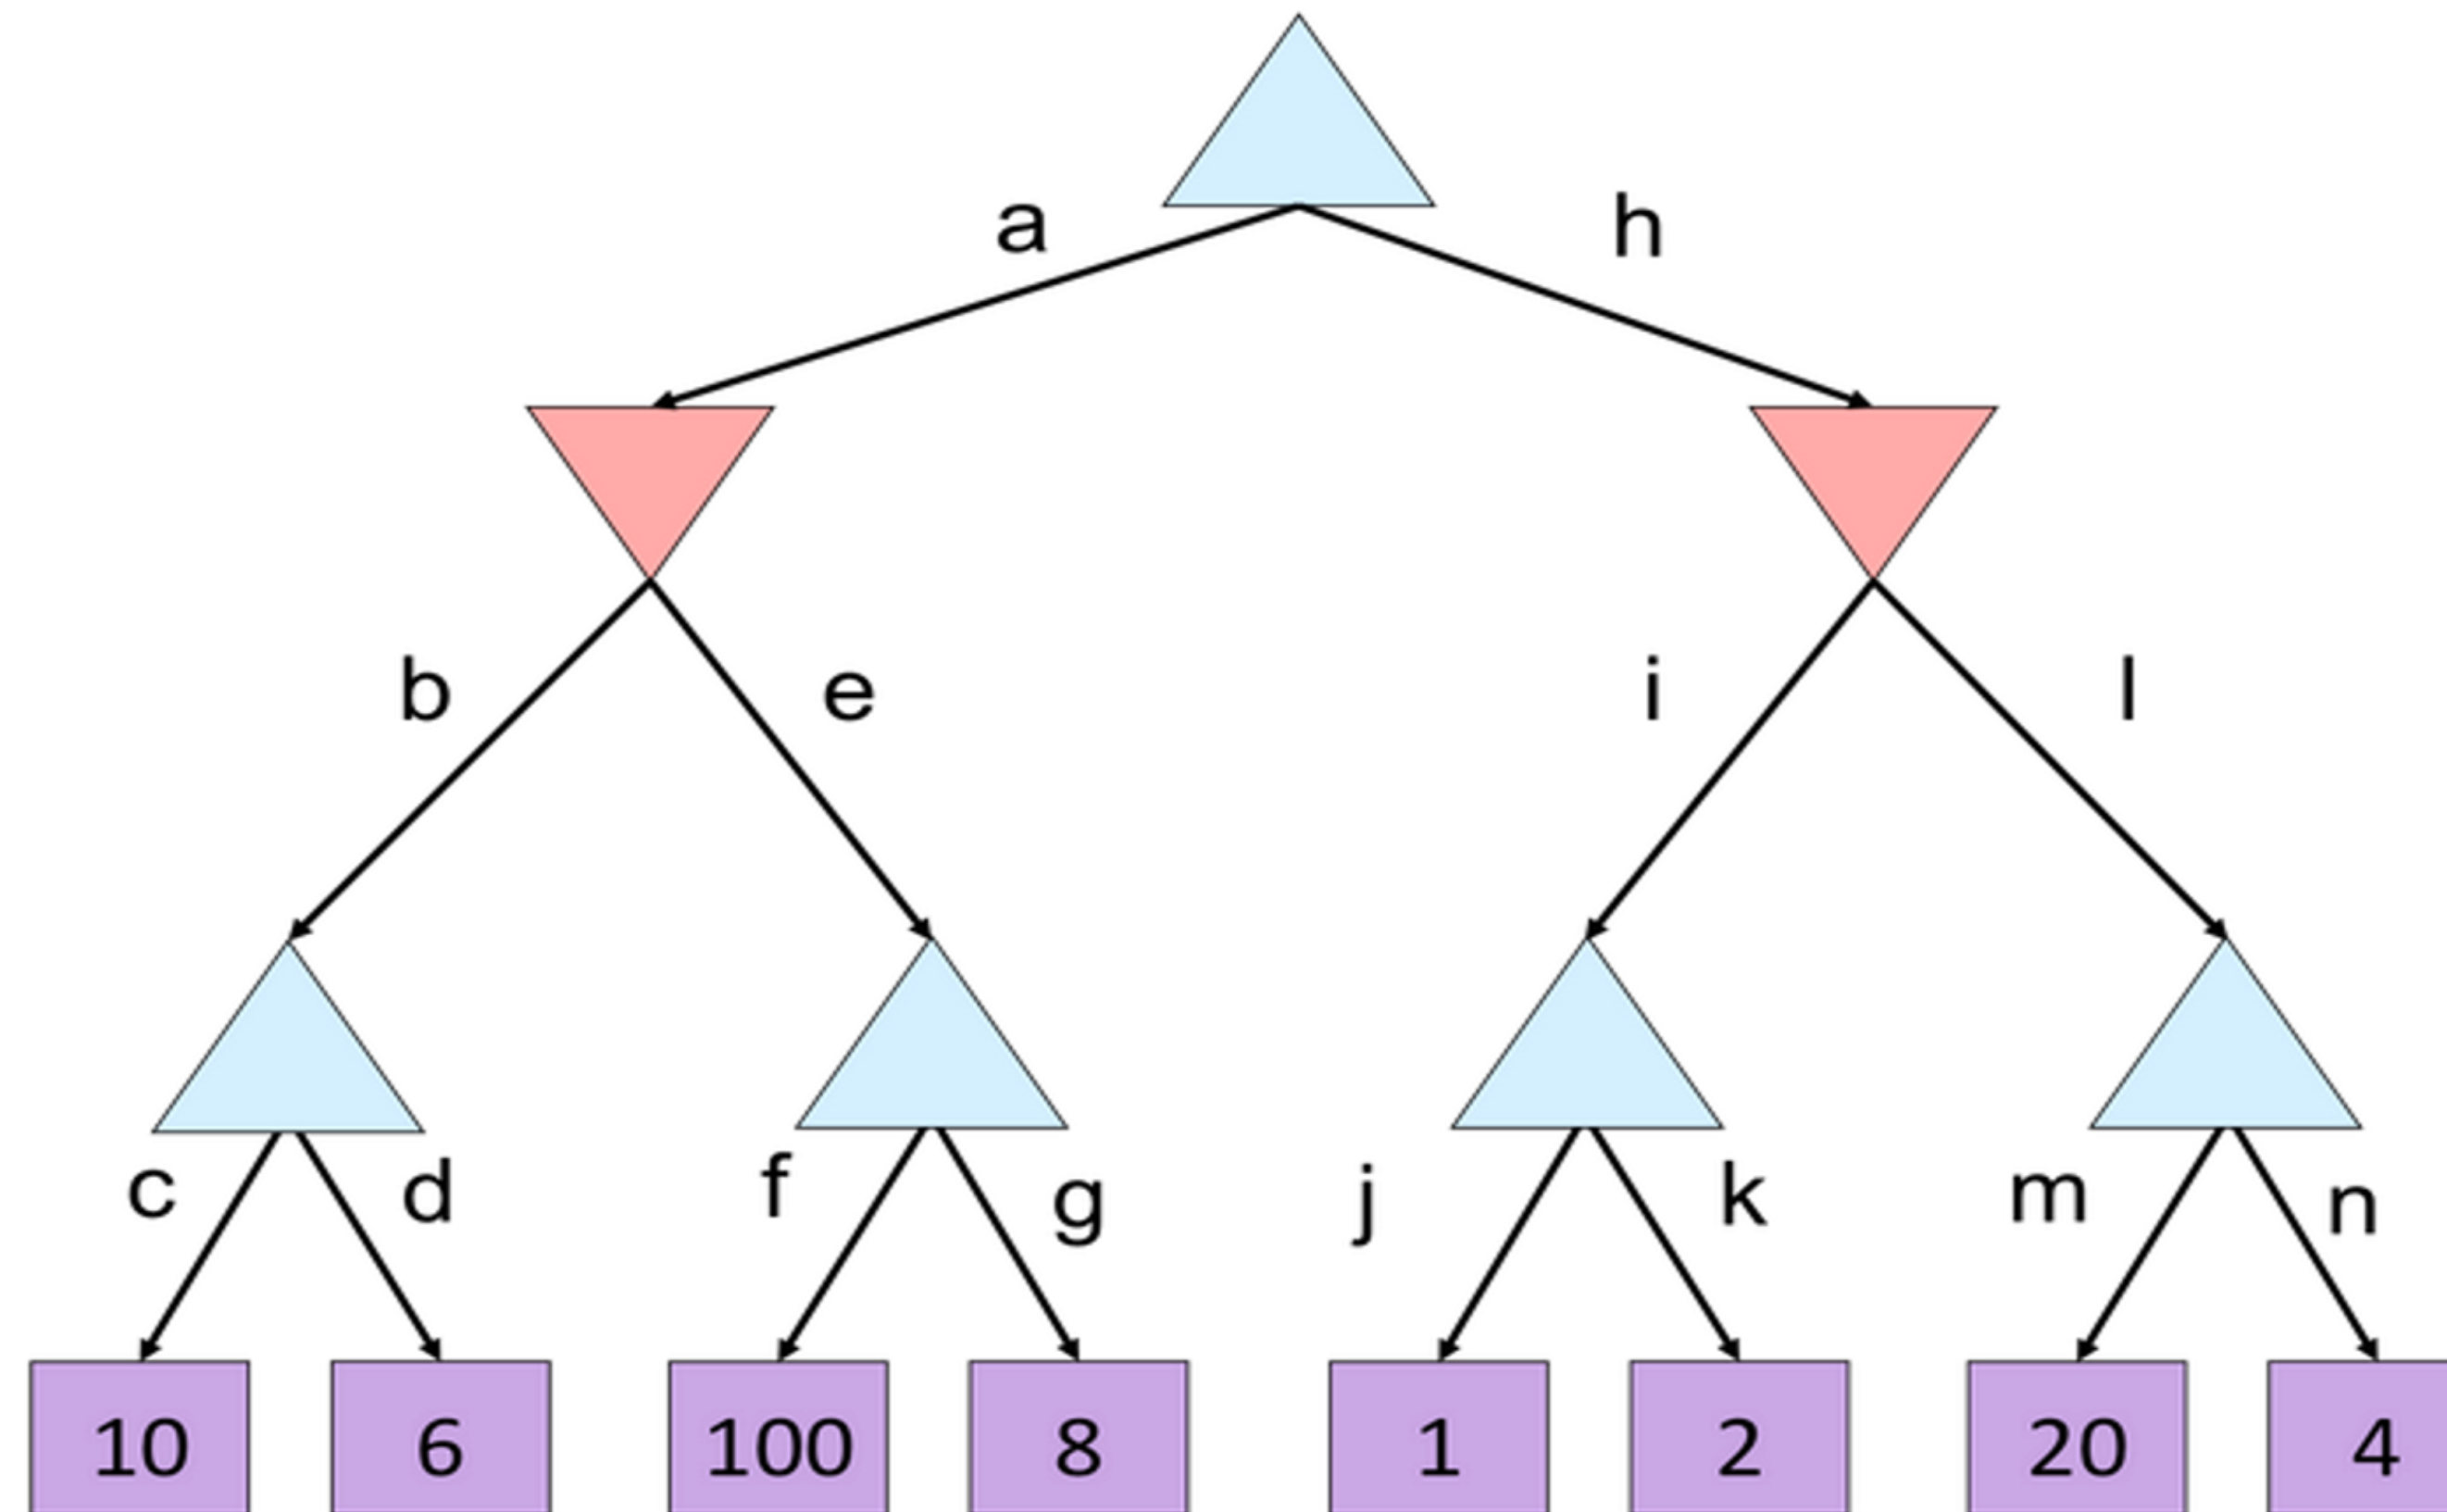
α - β PRUNING ALGORITHM

α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$   
            return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$   
            return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```

EXAMPLE: α - β PRUNING ★



EXPECTIMAX SEARCH

function **decision(s)** returns an action
return the action **a** in **Actions(s)** with the highest
value(Result(s,a))



function **value(s)** returns a value
if **Terminal-Test(s)** then return **Utility(s)**
if **Player(s) = MAX** then return **max**_{a in Actions(s)} **value(Result(s,a))**
if **Player(s) = MIN** then return **min**_{a in Actions(s)} **value(Result(s,a))**
if **Player(s) = CHANCE** then return **sum**_{r in chanceEvent(s)} **Pr(r) * value(Result(s,r))**

EXAMPLE: EXPECTIMAX ★

