

Support Vector Machines I



COMPCSI 762

Instructor: Thomas Lacombe


Based on slides from Meng-Fen Chiang

WEEK 9

RECAP: Machine Learning Systems

- Instance-based Learning


- Compare new data points to known data points
- Non-parametric approaches
- Memory-based approaches
- Prediction can be expensive



use the entire dataset as a model (e.g., k-NN)

- Model-based Learning

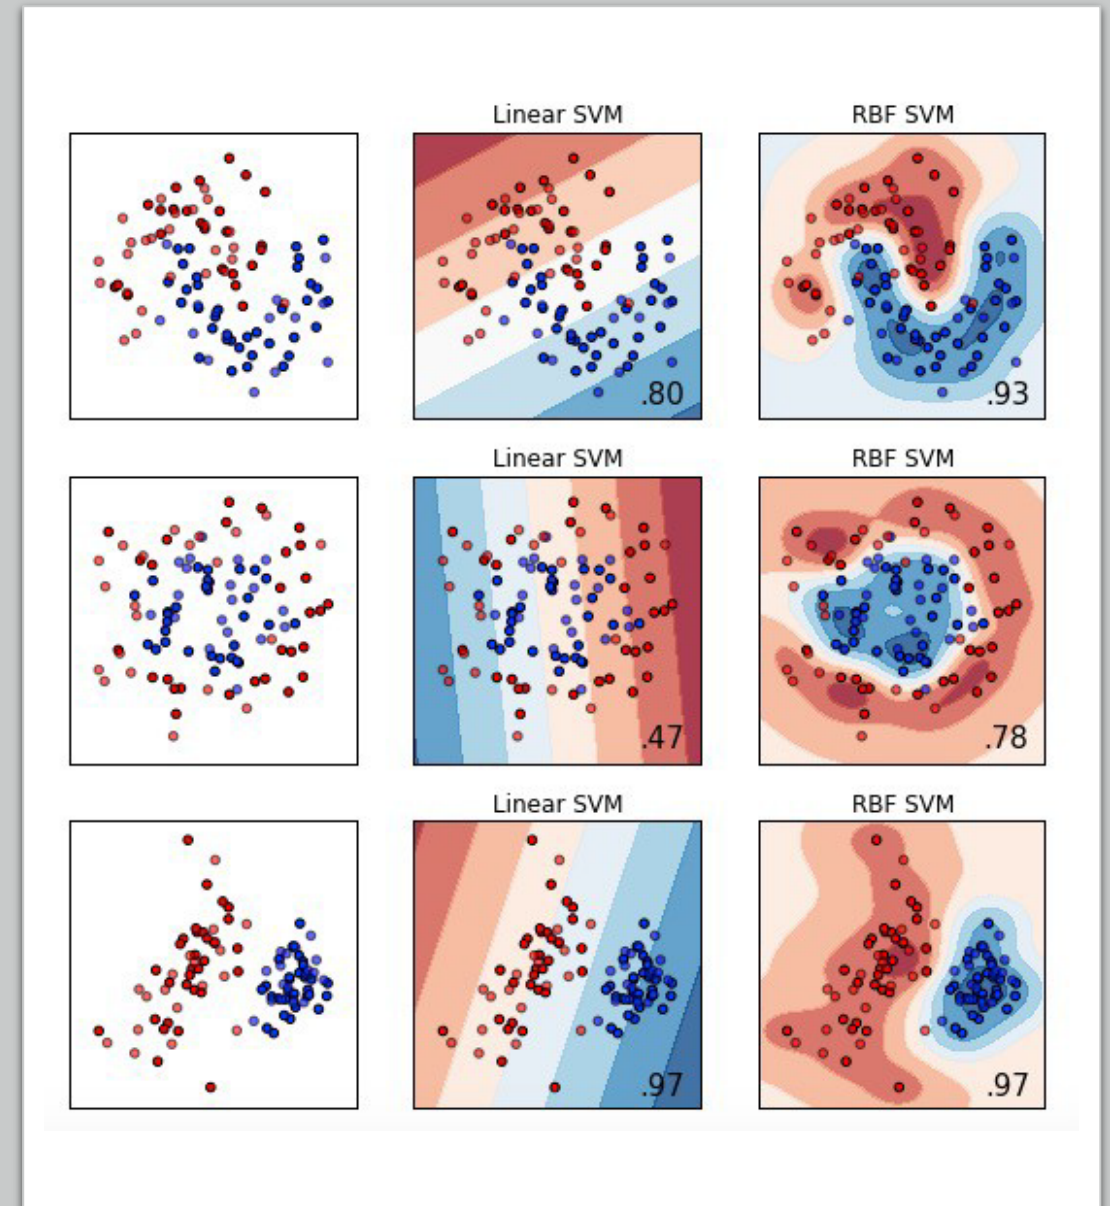
- Detect a pattern in the training data
- Build a predictive model
- Prediction is extremely fast



use the training data to create a model that has parameters learned from the training datasets (e.g., SVM)

OUTLINE


- Data Characteristics
 - Linearly Separable Data
 - Non-Linearly separable Data
- SVM (9.1,9.2,9.3)
 - Linearly Separable: Hard-margin SVMs
 - Non-Linearly Separable: Soft-margin SVMs
 - Non-Linearly Separable: Kernelized SVMs
- Summary




Data Characteristics

- A classification method for both **linear** and **nonlinear** data


Linearly Separable Data




Non-Linearly Separable Data



Data Characteristics



Linearly Separable Data




Non-Linearly Separable Data

Types of Support Vector Machines (SVMs)

- SVM selects the maximum margin linear classifier


Linear SVMs: Hard-margin SVMs



- SVM selects the maximum margin linear classifier with partial misclassifications allowed


Linear SVMs: Soft-margin SVMs

Non-Linear SVMs: Kernelized SVMs





OUTLINE

- Data Characteristics
 - Linearly Separable Data
 - Non-Linearly separable Data
- SVM
 - Linearly Separable Data: Hard-margin SVMs (9.1)
 - Non-Linearly Separable Data: Soft-margin SVMs (9.2)
 - Non-Linearly Separable Data: Kernelized SVMs (9.3)
- Summary



Problem Definition: Margin Maximization



- Given a set of linearly separable training data $S = ((x_1, y_1), \dots, (x_n, y_n))$, $y_i \in \{+1, -1\}$
- We want to find a linear decision boundary (hyperplane) to separate the 2 classes.
- There is an infinite number of lines (hyperplanes) separating the two classes!



- Goal: The **hard-margin** SVM algorithm aims to find a linear classifier that **maximizes** (γ) the margin on S .

Why Margin Maximization?


- Any linear classifier that separates S correctly will have margin $\gamma > 0$
- We want to find the best one (the one that minimizes classification error on unseen data)
- Assumption: the hyperplane with the largest margin will generalise best on unseen data
- SVM searches for the hyperplane with the largest margin, i.e., Maximum Marginal Hyperplane (MMH)




Example: Generalization Performance

We want a classifier which:

- Works well on training data
- Works well on the unseen Data



Margin Maximization Hyperplane (MMH)



Decision rule:
$$f(\vec{x}) = \begin{cases} +1, & \vec{w} \cdot \vec{x} + b \geq +1 \\ -1, & \vec{w} \cdot \vec{x} + b \leq -1 \end{cases}$$

- Distance of closet data \vec{x}_i from the hyperplane H
$$\frac{|\vec{w} \cdot \vec{x}_i + b|}{\|\vec{w}\|}$$

- $$\frac{|\vec{w} \cdot \vec{x}_i + b|}{\|\vec{w}\|} = \frac{y_i(\vec{w} \cdot \vec{x}_i + b)}{\|\vec{w}\|} = \frac{1}{\|\vec{w}\|}$$

- Margin:
$$\gamma = \frac{2}{\|\vec{w}\|}$$

- Maximize γ is equivalent to minimize $\|\vec{w}\|$


$$\min_{w,b} \frac{\|\vec{w}\|}{2}$$

s.t. $y_i(\vec{w} \cdot \vec{x}_i + b) \geq 1$

(discrimination boundary is respected) 11

Margin Maximization Hyperplane (MMH)

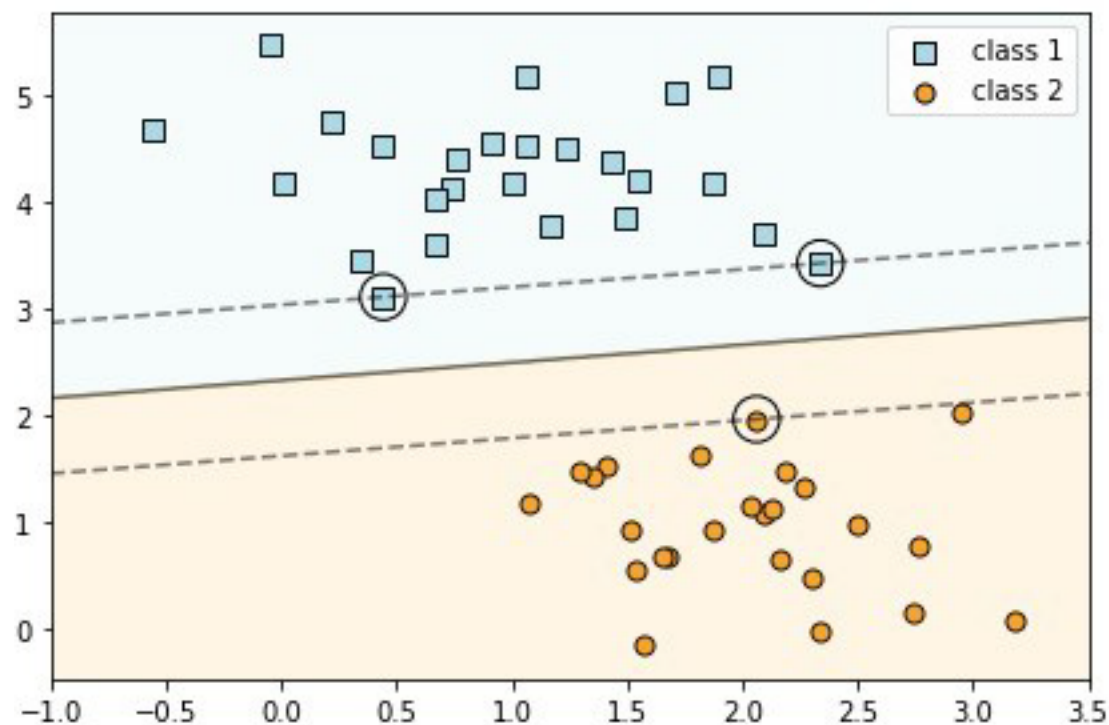
- A **separating hyperplane (H)** can be formally defined as $\vec{w} \cdot \vec{x} + b = 0$
 - $\vec{w} = \{w_1, w_2, \dots, w_n\}$ is a weight vector and b a scalar (bias)
- For 2-D it can be written as: $w_1 \cdot x_{i,1} + w_2 \cdot x_{i,2} + b = 0$
- The hyperplanes defining the sides of the margin:
 - $H_1: w_1 \cdot x_{i,1} + w_2 \cdot x_{i,2} + b \geq 1$, for $y_i = +1$, and
 - $H_2: w_1 \cdot x_{i,1} + w_2 \cdot x_{i,2} + b \leq 1$, for $y_i = -1$
- Any training tuples that fall on margins H_1 or H_2 (i.e., the hyperplanes defining the margin) are **support vectors**



Example: Support Vectors

Three Support Vectors:

1. $[0.44359863 \ 3.11530945]$
2. $[2.33812285 \ 3.43116792]$
3. $[2.06156753 \ 1.96918596]$




Margin Maximization Hyperplane (MMH)

Linear model: $f(\vec{x}) = \begin{cases} +1, & \vec{w} \cdot \vec{x} + b \geq +1 \\ -1, & \vec{w} \cdot \vec{x} + b \leq -1 \end{cases}$

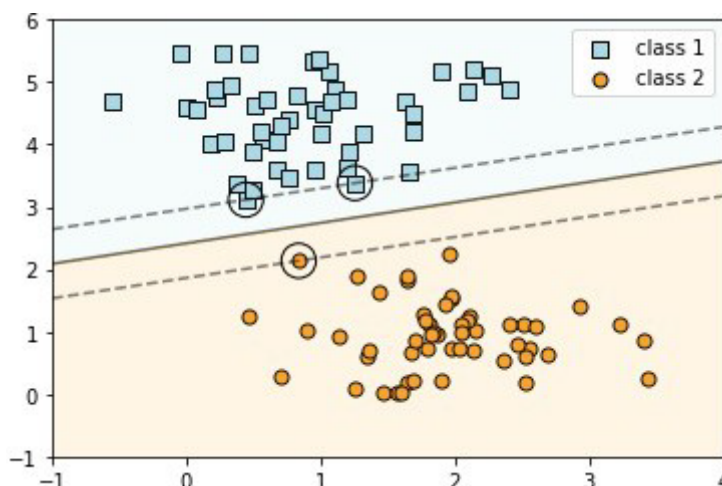
1. **Training Stage:** Learning the model is equivalent to determining the values of \vec{w} and b
 - How to find \vec{w} and b from training data S ?
 2. **Testing Stage:** Once \vec{w} and b are found, given a test data (\vec{x}) , use $f(\cdot)$ to determine the class label
- Decision boundary depends only on support vectors
 - If we have data set with same **support vectors**, decision boundary will not change

Example: Support Vector Matters


- Only the positions of the support vectors matter to decision boundary
- Other points further from the margin which are on the correct side do not modify the decision boundaries



$N=50$



$N=100$



$N=200$


Training Stage

- Objective is to maximize: $\gamma = \frac{2}{\|\vec{w}\|}$
- Equivalently, the objective is to minimize : $\min_{w,b} \frac{\|\vec{w}\|}{2}$
- Subject to the following constraints:


$$y_i = \begin{cases} +1, & \vec{w} \cdot \vec{x}_i + b \geq +1 \\ -1, & \vec{w} \cdot \vec{x}_i + b \leq -1 \end{cases}$$

- Or $y_i(\vec{w} \cdot \vec{x}_i + b) \geq 1, \quad i = 1, 2, \dots, n$
 \rightarrow m inequality constraints

minimization



maximization




This becomes a **constrained (convex) quadratic optimization** problem:

- Quadratic objective function with linear constraints \rightarrow Quadratic Programming (QP)




Quadratic Programming (QP)

- QP is a well-studied solution algorithm
- Lagrange Multipliers and Constrained Optimization
 - Finding the **local minima** and **maxima** of a differentiable function subject to equality or inequality constraints
 - The **point** at which the function and constraint touch each other is the solution to the optimization problem



Lagrange Multipliers and Constrained Optimization


- Optimization function: $f = x_1^2 + x_2^2$
- Subject to the constraint: $g = 2x + 6y = c$, with $c = 5$



- At the minimum of f s.t. the constraint: $\nabla f = \alpha \nabla g$
- Finding the minimum is then equivalent to solving: $\nabla f - \alpha \nabla g = 0$

Lagrange Multipliers and Constrained Optimization

- Optimization function: $f = x_1^2 + y^2$
- Subject to the constraint: $g = 2x + 6y - 5$
- Lagrange function (Lagrangian multiplier α):
 - $\mathcal{L}(x,y,\lambda) = f(x,y) - \alpha g(x,y)$
 $= x_1^2 + y^2 - \alpha(2x + 6y - 5)$
 - Solution for the constrained problem is obtained by solving for the points where the partial derivatives of \mathcal{L} are zero:
 - $\frac{d\mathcal{L}}{dx} = 2x - 2\alpha = 0$
 - $\frac{d\mathcal{L}}{dy} = 2y - 6\alpha = 0$
 - $\frac{d\mathcal{L}}{d\alpha} = -2x - 2y + 5 = 0$



Training Stage: Duality

- SVM **primal problem** form:

$$\min_{w,b} \frac{\|\vec{w}\|}{2} \quad \text{s.t. } y_i(\vec{w} \cdot \vec{x}_i + b) \geq 1, \quad i = 1, 2, \dots, n$$

- We can convert it to the **dual problem** of SVM by introducing Lagrange multipliers (α_i)

$$\mathcal{L}(\vec{w}, b, \alpha) = \frac{\|\vec{w}\|}{2} - \sum_{i=1}^n \alpha_i [y_i(\vec{w} \cdot \vec{x}_i + b) - 1], \quad \alpha_i > 0$$

- Solving the primal problem is equivalent to solving the dual problem:

$$\min_{w,b} \frac{\|\vec{w}\|}{2} \equiv \max_{\alpha} \min_{w,b} \mathcal{L}(\vec{w}, b, \alpha)$$

Training Stage: Duality

- SVM **primal problem** form:

$$\min_{w,b} \frac{\|\vec{w}\|}{2} \quad \text{s.t. } y_i(\vec{w} \cdot \vec{x}_i + b) \geq 1, \quad i = 1, 2, \dots, n$$

- We can convert it to the **dual problem** of SVM by introducing Lagrange multipliers (α_i)

- Set the derivatives of SVM Lagrangian function w.r.t. \vec{w} and b to be zero:

$$\mathcal{L}(\vec{w}, b, \alpha) = \frac{\|\vec{w}\|}{2} - \sum_{i=1}^n \alpha_i [y_i(\vec{w} \cdot \vec{x}_i + b) - 1], \quad \alpha_i > 0$$

- $\frac{d\mathcal{L}}{d\vec{w}} = 0 \Rightarrow \vec{w} = \sum_{i=1}^n \alpha_i y_i \vec{x}_i$
- $\frac{d\mathcal{L}}{db} = 0 \Rightarrow \sum_{i=1}^n \alpha_i y_i = 0$

- Substituting them in the Lagrangian function \mathcal{L} , we obtain the final dual optimization function:

$$\max_{\alpha} \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \vec{x}_i \cdot \vec{x}_j$$

Training Stage: Solving the Dual Problem

- Dual Problem Optimization

$$\begin{aligned} \max_{\alpha} \quad & \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \vec{x}_i \vec{x}_j \\ \text{s.t. } & \alpha_i \geq 0 \text{ and } \sum_{i=1}^n \alpha_i y_i = 0, \quad i = 1, 2, \dots, n \end{aligned}$$


- This can be solved efficiently using numerical optimization
 - $\alpha_i > 0$ for **support vectors** \vec{x}_i that lie on the margin H_1 and H_2
 - $\alpha_i = 0$ for other training points
- Thus, the solution for \vec{w} corresponding to the maximal margin classifier can be written as a linear combination of just the **support vectors**:

$$\vec{w} = \sum_{x_i \in SV} \alpha_i y_i \vec{x}_i$$

Training Stage: Solving the Dual Problem

- For support vectors, we have $y_i - (\vec{w} \cdot \vec{x} + b) = 0$
 - Blue support vector ($y_i = +1$)
 - Green support vector ($y_i = -1$)
- Thus, the solution for b from any of the support vectors

$$b = \frac{1}{|SV|} \sum_{x_i \in SV} y_i - (\vec{w} \cdot \vec{x}_i)$$



Testing Stage


- Given a new data point \vec{x} , we use the learned SVM classifier (\vec{w} and b) to derive the class label as follow:

$$f(\vec{x}) = \begin{cases} +1, & \vec{w} \cdot \vec{x} + b \geq 0 \\ -1, & \vec{w} \cdot \vec{x} + b \leq 0 \end{cases}$$



(Primal Form)

$$= \begin{cases} +1, & \sum_{x_i \in SV} \alpha_i y_i (\vec{x}_i \cdot \vec{x}) + b > 0 \\ -1, & \sum_{x_i \in SV} \alpha_i y_i (\vec{x}_i \cdot \vec{x}) + b < 0 \end{cases}$$

(Dual Form)



Example: Training Stage



X_1	X_2	y	α_i
0.3858	0.4687	1	65.5261
0.4871	0.611	-1	65.5261
0.9218	0.4103	-1	0
0.7382	0.8936	-1	0
0.1763	0.0579	1	0
0.4057	0.3529	1	0
0.9355	0.8132	-1	0
0.2146	0.0099	1	0

**Support
Vectors**

Example: Testing Stage

- Given a new data point $\vec{x} = [0.5, 0.9]$, what is the class label of \vec{x} using the trained SVM?

	X''	$X_{\#}$	y	α_i
Support Vectors	0.3858	0.4687	1	65.5261
	0.4871	0.611	-1	65.5261
	0.9218	0.4103	-1	0
	0.7382	0.8936	-1	0
	0.1763	0.0579	1	0
	0.4057	0.3529	1	0
	0.9355	0.8132	-1	0
	0.2146	0.0099	1	0

$$f(\vec{x}) = \begin{cases} +1, & \sum_{x_i \in SV} \alpha_i y_i (\vec{x}_i \cdot \vec{x}) + b > 0 \\ -1, & \sum_{x_i \in SV} \alpha_i y_i (\vec{x}_i \cdot \vec{x}) + b < 0 \end{cases}$$

$$\begin{aligned} & \text{sign} \left(\sum_{x_i \in SV} \alpha_i y_i (\vec{x}_i \cdot \vec{x}) + b \right) \\ &= \text{sign} \left(65.5261 \cdot 1 \cdot \begin{bmatrix} 0.3858 \\ 0.4687 \end{bmatrix}^T [0.5, 0.9] + 65.5261 \cdot (-1) \cdot \begin{bmatrix} 0.4871 \\ 0.611 \end{bmatrix}^T [0.5, 0.9] \right) \end{aligned}$$


Quiz: Linear SVM

- Question: Suppose we want to build a hard-margin SVM classifier for two-class classification in one dimension space ($d = 1, x_i \in \mathbb{R}$) contains three sample points:
 - point $x_1 = 3$ with label $y_1 = 1$
 - point $x_2 = 1$ with label $y_2 = 1$
 - point $x_3 = -1$ with label $y_3 = -1$

What are the values of \vec{w} and \vec{b} given by our hard-margin SVM?

- Solve the optimization problem for w and b with the following constraints

$$\min_{w, b} \frac{w^2}{2} \quad \text{s.t.} \quad \begin{cases} w * x_1 + b \geq 1 \\ w * x_2 + b \geq 1 \\ w * x_3 + b \leq -1 \end{cases}$$



A: $w = 1, b = 1$

B: $w = 1, b = 0$

C: $w = 0, b = 1$

D: $w = \infty, b = 0$

Advantages v.s. Disadvantages

Advantages

- SVMs depends on relatively few support vectors
 - SVMs are very compact models, and take up very little memory
 - SVMs work well with high-dimensional data, even with more dimensions than samples ($d > |S|$)
- Once the model is trained, the prediction phase is very fast

Disadvantages


- For large numbers of training samples, the computational cost can be prohibitive
- The results do not have a direct probabilistic interpretation

Jupyter Notebook

Hard-margin SVM Coding Example

SUMMARY

- Data Characteristics
 - Linearly Separable Data
 - Non-Linearly separable Data
- Linearly Separable Data: Hard-margin SVMs
 - Margin Maximization Hyperplane (MMH)
 - Primal Form Optimization
 - Duality Form Optimization
 - Training Phase
 - Testing Phase
 - Advantages v.s. Disadvantages



Resources

- SVM Website: <http://www.kernel-machines.org/>
- Representative Implementation
 - **LIBSVM**: an efficient implementation of SVM, multi-class classifications, nu-SVM, one-class SVM, including also various interfaces with java, python, etc.
 - **SVM-light**: simpler but performance is not better than LIBSVM, support only binary classification and only in C
 - **SVM-torch**: another recent implementation also written in C
 - **Scikit-Learn**: a set of supervised learning methods used for classification, regression and outliers detection. [\[link\]](#)

Resources (Contd.)

- Book Chapters: Christopher Bishop, “Pattern Recognition and Machine Learning” ([PDF](#))
 - Sec 7.1.1-7.1.2
 - Sec 4.1.1
 - Sec 6.1, 6.2
 - Appendix E
- Literatures
 - C.J.C. Burges, Chris J.C. Burges "A Tutorial on Support Vector Machines for Pattern Recognition." Data Mining and Knowledge Discovery, 1998